

SUPERDOCS • ROUND 2

THE SHARED OPEN LIST

THE SHARED OPEN LIST, AND HOW TO READ IT

This is one list for every candidate in both roles, and it is deliberately huge: we kept it that way so that everyone finds something they would genuinely enjoy building. Everything on it is voluntary and earns extra marks. Nothing on it is claimed or reserved, nobody is blocked by anybody, and if several people build the same thing, all of them count. Pick anything that pulls you, whichever role you applied for; the lanes are not fences. If your assigned build leaves you cold, finish it honestly at its minimum, then spend yourself here. Whatever you pick, build ON SuperDocs, never a copy of it.

Every entry reads the same way: what to build, who it serves, what a strong build looks like, a difficulty band, the product surfaces it touches, and where the finished work goes. The difficulty band describes the build, not you - judging compares within a band, so a simple build done immaculately beats a hard one left half-finished.

One note on the bands: they are my own sizing estimate, made while writing the cards - reality may differ once you start. Pick by curiosity and interest; a band is information for planning, never a signal about what we want you to choose.

The bands, spelled out. S1, starter: one clear surface path, no external integration, no long-lived state; finish the core early and spend the rest on polish, tests and documentation. S2, solid, the default: two to four surfaces, or one surface against a real-world data shape. S3, stretch: five or more surfaces, a long-running stateful loop, an external integration, or a fidelity round-trip; this is where strong candidates separate themselves. The band is a default, never a gate.

The list is grouped into nine families, so you can jump to your kind of work.

Office and Workspace add-ins. Build inside Word, Google Docs, Outlook, Excel or Sheets rather than beside them.

Vertical mini-apps. Industry-by-document-type apps on the API, across some forty industries.

Integrations and pipelines. Flows that start in another tool and end in an edited, exported document.

Extensions, dev tools and SDK surfaces. Browser extensions, editor plugins, command-line tools, SDK wrappers and editor bridges.

Agent-loop demos and living documents. An agent that maintains a document over time through MCP, rather than editing it once.

RAG, templates and benchmark. Clause libraries, style guides, precedent finders, template packs and benchmark cases.

Growth networking. Map perfect-fit companies, name the use case, design the warm-intro path.

The content wave. Real technical content published where AI engines retrieve from.

Growth experiments and placements. Traction variants, directory and listicle placement research, community launch plans.

The surface words, in plain terms. Chat: the build works through the chat, the way a person does. Templates: it creates or reuses document templates. Review: proposed changes wait for a human yes before they land. Export: it produces finished files people download or receive. API: it drives SuperDocs through the REST API. MCP: it connects through the MCP server, the way coding agents do. Search: it finds things across a document or a library. Memory: it uses what SuperDocs remembers across sessions. Attachments: it works with uploaded source files. Multi-document: it works across several documents at once. Images: images or diagrams inside documents.

WHERE FINISHED WORK GOES

Every entry names its route. The use-cases and extensions routes mean a pull request into the public repository, into that folder. The submission-form route means your files go into Drive through the form. The own-channel route means the work is published on an account you own, with the link submitted through the form. The documentation route means a pull request into our public developer docs. Your task document carries the full submission mechanics; this legend only says where each build lands.

OFFICE AND WORKSPACE ADD-INS

Word add-in: contract redline desk for law firms

WHO IT SERVES

Legal. Associate or contract attorney at a law firm or in-house legal team. Host app: Microsoft Word (Windows, Mac, Word on the web).

WHAT TO BUILD

A task pane opens beside the open contract. The lawyer picks a playbook position (buyer-friendly, seller-friendly, balanced) and asks for a pass. The add-in sends the document, gets clause-level proposals back, and shows each one in the pane as a red/green before-and-after card. Approve and it is written into Word as a tracked change with the reason attached as a Word comment. Reject and nothing moves. Nothing lands unapproved. The finished file is a DOCX the counterparty opens in plain Word with the redlines and comments intact.

WHAT STRONG LOOKS LIKE

survives a real 60-page MSA. Defined terms, clause numbering, and cross-references outside the edited clauses are untouched. Tracked changes open cleanly in desktop Word with correct authorship and dates. Rejecting every proposal leaves the file byte-identical to the original. Sideloads from a manifest in under two minutes and a partner can run it without reading instructions.

BAND S3 · chat, API, Review, export, templates, search · a pull request into the public repository, extensions folder

Word add-in: clause library and fallback ladder pane

WHO IT SERVES

Legal ops / knowledge management lead at a firm or legal department. Host app: Microsoft Word.

WHAT TO BUILD

The firm's approved clauses live in an uploaded library. The pane lets the drafter search it in plain language ("limitation of liability, capped at fees paid"), see the preferred position plus the fallback ladder beneath it, and insert the chosen one directly into the right section of the open document. Insertion adapts party names, defined terms, and numbering to the host contract instead of pasting raw text. A second command audits the whole document and flags every clause that has drifted from the library.

WHAT STRONG LOOKS LIKE

semantic search finds the right clause when the query wording and the library wording share no words. Inserted clauses inherit the document's numbering scheme and styles rather than the library's. The drift audit produces a reviewable list, not a rewrite. Library updates take effect without reinstalling the add-in.

BAND S2 · search, templates, chat, Review, export · a pull request into the public repository, extensions folder

Word add-in: compare and merge two contract versions

WHO IT SERVES

Legal / deal desk. Anyone reconciling counterparty turns. Host app: Microsoft Word.

WHAT TO BUILD

Pick the current document and a second file (the counterparty's turn). The add-in loads both, explains what changed clause by clause in plain language, and proposes a single reconciled version: accept theirs, keep ours, or a merged wording where the two are compatible. Each decision is a card in the pane. The output is one clean document plus a tracked-changes version showing what was conceded, so the deal owner can see the position that was given up.

WHAT STRONG LOOKS LIKE

the comparison is semantic, not a character diff, so a reordered clause reads as "moved", not "deleted and added". Tables and numbered schedules survive the merge. The concession summary is accurate enough that a partner would send it to a client.

BAND S3 · Multi-document, chat, Review, export, search · a pull request into the public repository, extensions folder

Word add-in: investment and credit memo editor

WHO IT SERVES

Analyst or associate at a fund, bank, or credit team. Host app: Microsoft Word.

WHAT TO BUILD

The memo already exists on the firm template with branded tables, a fixed section order, and figures in prose. The pane takes an update instruction ("refresh the Q3 numbers from this attachment and rewrite the risk section to reflect the covenant change") and returns section-scoped proposals. Numbers inside sentences update alongside the table cells that carry them. The reviewer approves each section. Export returns a DOCX where the branded table shading, column widths, and running headers are the same bytes the firm's template shipped with.

WHAT STRONG LOOKS LIKE

a figure changed in three places changes in all three, and nowhere else. Merged cells and alternating row shading survive. The reviewer can see, per section, exactly what will change before it changes.

BAND S2 · chat, Review, templates, export, search · a pull request into the public repository, extensions folder

Word add-in: clinical letter and referral writer with sign-off

WHO IT SERVES

Clinician, medical secretary, or practice manager. Host app: Microsoft Word on a managed clinic desktop.

WHAT TO BUILD

The clinician pastes or attaches consultation notes. The pane drafts the referral, discharge summary, or patient letter onto the practice letterhead template, with the right salutation, clinical structure, and closing. Every clinical statement in the draft appears in the review pane linked to the note text it came from, so nothing is asserted that the notes do not support. The clinician approves section by section, signs, and exports a PDF and a DOCX. An unapproved draft can never be exported.

WHAT STRONG LOOKS LIKE

the letterhead, footer, and practice registration details survive untouched. Nothing clinical appears in the letter without a visible source in the notes. The sign-off gate is real, not a checkbox, and the audit record of who approved what is exportable.

BAND S2 · chat, Review, templates, export, search · a pull request into the public repository, extensions folder

Word add-in: academic manuscript desk

WHO IT SERVES

PhD student, postdoc, or research group admin. Host app: Microsoft Word.

WHAT TO BUILD

The pane handles the three things that break manuscripts: equations, references, and the journal template. It edits prose surgically ("tighten the discussion to 800 words, keep the hedging"), keeps every equation as real Word math that stays editable, keeps footnotes as real footnotes, and reflows the document into a target journal's template on request. A reformat command switches citation style across the whole document without touching the argument.

WHAT STRONG LOOKS LIKE

a 90-page thesis edits in seconds per pass without the numbering, figure captions, or cross-references drifting. Equations round-trip as native Word math, not pictures. Switching journal template does not orphan a single figure or reference.

BAND S3 · chat, export, search, images, templates · a pull request into the public repository, extensions folder

Word add-in: redaction pane with human confirmation

WHO IT SERVES

Paralegal, compliance analyst, or records officer. Host app: Microsoft Word.

WHAT TO BUILD

Scan the open document for personal data, health data, account identifiers, and privileged passages. Every candidate appears in the pane grouped by type with its surrounding sentence, so the reviewer can confirm or dismiss each one. Confirmed items are replaced in place with a redaction marker that keeps the sentence readable and the layout intact. Export produces a clean file with the removed text genuinely gone from the stored document, plus a redaction log listing what was removed and by whom.

WHAT STRONG LOOKS LIKE

nothing is removed without a human click. The exported file carries no residue of the redacted text in comments, tracked changes, or document properties. Page and paragraph structure is preserved so the document still reads as the same document.

BAND S2 · search, Review, chat, export · a pull request into the public repository, extensions folder

Honesty note: document properties and file metadata are outside the pipeline; your pane's job covers the visible document text, and your write-up should say plainly what your redaction does and does not touch. A version that states its boundary honestly is the strong version.

Word add-in: fill this template from a record

WHO IT SERVES

Operations, HR, or back-office coordinator at any templated-document shop. Host app: Microsoft Word.

WHAT TO BUILD

Open the house template. The pane accepts a record (pasted JSON, a CSV row, a form submission, or a pasted email) and fills the template from it: names, dates, amounts, addresses, jurisdiction-specific paragraphs, conditional sections that appear or disappear based on the record. Fields that the record does not answer are listed as gaps rather than invented. The filled document keeps the template's exact styling, so the output is indistinguishable from a hand-filled one.

WHAT STRONG LOOKS LIKE

conditional sections resolve correctly on both branches. Missing data becomes a visible gap list, never a plausible invention. Ten different records produce ten documents whose formatting is identical to the template and to each other.

BAND S1 · templates, chat, API, Review, export · a pull request into the public repository, extensions folder

Word add-in: track-changes triage

WHO IT SERVES

Editor, legal reviewer, or supervisor handling a marked-up return. Host app: Microsoft Word.

WHAT TO BUILD

Open a document that came back full of tracked changes from several reviewers. The pane groups every change by reviewer and by kind (substantive, stylistic, typo) and lets the user act by instruction: "accept all of the proofreader's typo fixes", "reject the deletions in section 2", "show me every change that alters an obligation". Each batch is previewed before it is applied. Reviewer comments can be answered and resolved from the same pane.

WHAT STRONG LOOKS LIKE

reviewer names and timestamps survive the round-trip. A substantive-versus-stylistic split that a human agrees with. Applying a batch leaves every other reviewer's marks untouched and still attributed.

BAND S2 · chat, Review, export, search · a pull request into the public repository, extensions folder

Honesty note: multi-reviewer attribution across a full file round-trip is not something the product promises today, so do not build on it. If you round-trip a file and attribution survives, tell us, because we want to know.

Word add-in: enterprise centrally-deployed build

WHO IT SERVES

IT administrator deploying to 500 to 50,000 seats. Host app: Microsoft Word via central deployment.

WHAT TO BUILD

The same editing pane, packaged the way enterprise IT requires: admin-pushed to users or groups with no per-user install, single sign-on so nobody pastes a key, an organization-level key held by the tenant rather than the individual, per-department feature toggles (redlining on for legal, off for everyone else), and an admin console showing usage by department against the organization's allowance. Includes the deployment runbook and the security questionnaire answers IT will ask for.

WHAT STRONG LOOKS LIKE

a user opens Word and the pane is simply there, already signed in. No key is ever visible to an end user. Feature toggles take effect without redeployment. The admin can answer "what did legal spend last month" without a support ticket.

BAND S3 · API, Review, export, templates · a pull request into the public repository, extensions folder

Word add-in: page-limit-disciplined proposal editor

WHO IT SERVES

Capture manager or proposal writer in government contracting or consulting. Host app: Microsoft Word.

WHAT TO BUILD

Government and consulting proposals have hard page limits per section. The pane shows a live per-section page budget against the actual rendered pagination, and edits to a target: "cut section 3.2 to two pages without losing the past-performance references", "expand the technical approach by half a page using the attached win themes". Because edits are section-scoped, tightening one section does not reflow the rest of the document out of compliance.

WHAT STRONG LOOKS LIKE

the page count shown matches what the exported PDF actually paginates to. Cutting a section does not shift the following section's start page. Required boilerplate and compliance language is never cut in the name of length.

BAND S3 · chat, search, export, templates, Review · a pull request into the public repository, extensions folder

Word add-in: house-style enforcer

WHO IT SERVES

Editorial lead, brand manager, or knowledge manager at an agency, publisher, or professional-services firm. Host app: Microsoft Word.

WHAT TO BUILD

Upload the style guide once (the real PDF or DOCX, not a rules file someone has to maintain). The pane then audits any open document against it and proposes fixes in two lanes: language (terminology, tone, banned phrasing, capitalisation of product names) and formatting (heading levels, list style, spacing, table treatment). Each proposal cites the style-guide passage it comes from. Approve individually or accept a whole lane.

WHAT STRONG LOOKS LIKE

the guide is read, not transcribed into hardcoded rules, so swapping in a different client's guide changes the behaviour with no code change. Citations point at real passages. Formatting fixes do not disturb content, and language fixes do not disturb formatting.

BAND S2 · templates, search, chat, Review, export · a pull request into the public repository, extensions folder

Word add-in: comment resolver

WHO IT SERVES

Anyone who receives a document back with twenty reviewer comments. Host app: Microsoft Word.

WHAT TO BUILD

The pane lists every comment in the document and, for each, proposes the actual edit that would satisfy it, anchored to the commented text. "This paragraph is too vague" becomes a concrete rewrite; "add the 2025 figure here" becomes the inserted sentence. The user approves, the edit lands, and the comment thread is replied to and resolved automatically with a note saying what was done. A summary shows which comments were addressed, which were deferred, and which need the author.

WHAT STRONG LOOKS LIKE

comments stay real Word comments through the round-trip, with the original author preserved. An ambiguous comment produces a question, not a guess. The end state is a document where the resolved threads genuinely correspond to applied edits.

BAND S2 · chat, Review, export, search · a pull request into the public repository, extensions folder

Word add-in: numbering and cross-reference repair

WHO IT SERVES

Legal, pharma regulatory, or standards author working on documents where the numbering is load-bearing. Host app: Microsoft Word.

WHAT TO BUILD

After a document has been edited by several hands, clause numbering, exhibit letters, figure and table numbers, and internal cross-references drift out of alignment. The pane audits all four, shows every break ("clause 7.3 is referenced twice but no longer exists", "Exhibit C is missing between B and D"), and proposes a repair that renumbers and rewrites the references consistently. Repairs are previewed as a list before anything is written.

WHAT STRONG LOOKS LIKE

cross-reference targets survive the round-trip and still jump correctly in Word after export. Renumbering never silently changes the meaning of a reference. The audit finds breaks a careful human reader would find, and does not invent ones they would not.

BAND S3 · search, chat, Review, export · a pull request into the public repository, extensions folder

Scope: page, page-count, citation and table-of-contents fields are the field codes that survive our export; internal REF cross-references are not in that set. Build the repair against the surviving set, and say so in your write-up.

Word add-in: translate in place, layout preserved

WHO IT SERVES

Localisation coordinator, international legal or HR team. Host app: Microsoft Word.

WHAT TO BUILD

Translate the open document into a target language while keeping the document itself: same tables, same headers and footers, same numbering, same styles, with text expansion or contraction handled inside the existing layout. Right-to-left targets flip paragraph direction and alignment correctly without inverting embedded numbers or code. A bilingual mode produces a two-column source-and-target document instead. Terminology can be pinned by attaching a glossary.

WHAT STRONG LOOKS LIKE

the translated file opens as the same document, not a text dump. Right-to-left output is genuinely correct, including mixed-direction runs. Glossary terms are honoured everywhere, including inside table cells and footnotes.

BAND S2 · chat, templates, export, Multi-document · a pull request into the public repository, extensions folder

Word add-in: offer letter and policy composer from an HR record

WHO IT SERVES

HR generalist or people-ops coordinator. Host app: Microsoft Word.

WHAT TO BUILD

The pane pulls a candidate or employee record (pasted from the ATS, or via a simple connector) and composes the offer letter, contract addendum, or policy acknowledgement on the company template: role, compensation, start date, jurisdiction clauses, benefit summary. Jurisdiction-specific paragraphs swap automatically. The recruiter reviews, approves, and exports a signature-ready PDF and DOCX. A batch mode does a whole cohort with the same review gate applied once per document.

WHAT STRONG LOOKS LIKE

compensation and dates are never wrong or invented. The template's branding, signature block, and footer survive exactly. The jurisdiction swap is driven by the record, not by a hardcoded country list.

BAND S1 · templates, chat, Review, export, API · a pull request into the public repository, extensions folder

Word add-in: insurance policy and endorsement builder

WHO IT SERVES

Underwriter, MGA operations, or policy-servicing team. Host app: Microsoft Word.

WHAT TO BUILD

Build a policy document or an endorsement from the carrier's wording library plus the risk's variables: named insured, limits, deductibles, territory, selected coverage forms, and the endorsements that attach. The pane assembles the document in the correct order with the correct schedule pages, then lets the underwriter review each attached form. A second command drafts the endorsement that amends an existing in-force policy, showing exactly which policy wording it changes.

WHAT STRONG LOOKS LIKE

form order and schedule pagination are correct. The endorsement quotes the base policy wording accurately. Limits and deductibles appear identically in the schedule and in the prose that refers to them.

BAND S2 · templates, Multi-document, chat, Review, export · a pull request into the public repository, extensions folder

Word add-in: visual content pane

WHO IT SERVES

Technical writer, consultant, marketer, trainer. Host app: Microsoft Word.

WHAT TO BUILD

Word is where the document lives but not where visuals get made, so people leave and come back with a screenshot. This pane keeps them in Word: describe an image and it is generated and inserted at the cursor; describe a process and a flowchart, sequence diagram, or org chart is rendered and inserted; sketch on a small canvas and the drawing lands as an image. Every inserted visual can be re-edited by instruction later ("add a verification step after background check") rather than regenerated from scratch.

WHAT STRONG LOOKS LIKE

inserted visuals land at the cursor, not at the end of the document. Diagrams re-edit in place and keep their position and size. Everything survives export to DOCX and PDF at print resolution.

BAND S2 · images, chat, export · a pull request into the public repository, extensions folder

Word add-in: equation authoring desk

WHO IT SERVES

Engineer, quant, physicist, or maths teacher. Host app: Microsoft Word.

WHAT TO BUILD

Type or dictate maths in familiar LaTeX notation, or describe it in words ("the quadratic formula", "a summation from i equals one to n"), and get real, editable Word equations inserted at the cursor. A conversion command turns a document full of pasted equation images back into native maths. A check command finds equations whose notation is inconsistent with the rest of the document.

WHAT STRONG LOOKS LIKE

equations arrive as native Word maths that a colleague can click and edit, not as pictures. Round-tripping a document with fifty equations changes none of them. Notation stays consistent across the document.

BAND S2 · chat, images, export · a pull request into the public repository, extensions folder

Honesty note: math conversion degrades per construct and reports partial fidelity when it does. A version that reports honestly which of the fifty equations it could not carry perfectly is stronger than one that claims all fifty and is wrong about three.

Word add-in: headers, footers, footnotes and comments editor

WHO IT SERVES

Formal-document producer: legal secretary, publisher, compliance author. Host app: Microsoft Word.

WHAT TO BUILD

The parts of a document that live outside the body are the ones generic AI tools flatten or drop. This pane edits them explicitly and by instruction: change the running header on first pages only, add a confidentiality footer to every section, fix the typo in footnote 12, convert three inline parentheticals into real footnotes, add a reviewer comment on the payment clause. Each part edit appears as its own before-and-after card so a header change and a body change in the same request are approved separately.

WHAT STRONG LOOKS LIKE

footnotes stay real footnotes and headers stay real header parts through edit and export, including per-section and first-page variants. Editing a long header cannot disturb the body. Comments keep their author.

BAND S2 · chat, Review, export · a pull request into the public repository, extensions folder

Word add-in: precedent side-by-side

WHO IT SERVES

Drafter in legal, procurement, or grants who always starts from a prior document. Host app: Microsoft Word.

WHAT TO BUILD

The pane holds a shelf of reference documents (prior deals, winning proposals, gold-standard contracts) alongside the open draft. Ask "how did we handle assignment on the Acme deal" and the matching passage surfaces with context. Ask to bring it across and it is adapted to the current document's parties, defined terms, and numbering rather than pasted. A coverage view shows which sections of the draft still have no precedent behind them.

WHAT STRONG LOOKS LIKE

the shelf survives across sessions so yesterday's references are still there today. Ported passages read as if drafted for this document. Retrieval works on meaning, not on shared wording.

BAND S2 · Multi-document, search, templates, chat, Review · a pull request into the public repository, extensions folder

Word add-in: version timeline and rewind

WHO IT SERVES

Anyone who has ever wanted the document back the way it was three edits ago. Host app: Microsoft Word.

WHAT TO BUILD

A timeline pane lists every AI turn on this document with a one-line description of what changed. Click any point and both the conversation and the document rewind to that state together, with a confirmation first. Undo the rewind if it was the wrong call. Each entry expands into the exact before-and-after for the sections it touched, so the timeline doubles as a change log the user can export.

WHAT STRONG LOOKS LIKE

rewinding restores the document exactly, including tables, images, and footnotes, not an approximation. The undo of a rewind works. The exported change log is readable by someone who was not there.

BAND S2 · chat, Review, export, API · a pull request into the public repository, extensions folder

Word add-in: one build across Windows, Mac, web and iPad

WHO IT SERVES

Any team where half the users are on Word for the web and the rest are on desktop. Host app: Microsoft Word, all four hosts.

WHAT TO BUILD

Most add-ins are built against desktop Word and quietly break elsewhere. This build treats host parity as the product: the same pane, the same features, and honest degradation where a host genuinely cannot do something, with a clear in-pane explanation instead of a silent failure. Ships with a published host-capability matrix and an automated check that runs the same scenario suite against every host before release.

WHAT STRONG LOOKS LIKE

the same document, edited from Word on the web and then desktop Word, ends in the same state. Every feature either works or explains itself. The capability matrix matches reality, verified by the suite rather than by hand.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Word add-in: regulated audit-trail pane

WHO IT SERVES

Quality, compliance, or validation lead in pharma, medical devices, aviation, or financial services. Host app: Microsoft Word.

WHAT TO BUILD

In controlled document work the record of who changed what and why is part of the deliverable. This pane logs every proposed change, its reason, the approver, and the timestamp, and refuses to write anything that has not been approved by a named user. The log exports as an evidence appendix attached to the document. A review mode replays the document's history for an auditor without needing the original editors present.

WHAT STRONG LOOKS LIKE

the log is complete, with no path that changes the document without an entry. The evidence appendix is legible to an auditor who has never seen the tool. Replay reconstructs each intermediate state faithfully.

BAND S3 · Review, chat, export, API · a pull request into the public repository, extensions folder

Word add-in: batch runner over a folder

WHO IT SERVES

Operations lead who has 400 documents and one instruction. Host app: Microsoft Word (the pane launches the job, the job runs server-side).

WHAT TO BUILD

Point the pane at a folder (local, OneDrive, or SharePoint), give one instruction ("update the company name and address everywhere, keep everything else identical"), and preview it on three sample documents first. Approve the samples and the rest run as a background job with a live progress list, a per-document result, and a downloadable report of what changed in each. Failures are listed honestly rather than skipped silently.

WHAT STRONG LOOKS LIKE

the sample preview genuinely predicts the batch. Four hundred documents keep four hundred different layouts. The report is trustworthy enough that nobody opens the files to check.

BAND S3 · API, chat, templates, export, Review · a pull request into the public repository, extensions folder

Word add-in: organisation building-block library

WHO IT SERVES

Brand or knowledge lead responsible for what everyone else pastes into documents. Host app: Microsoft Word.

WHAT TO BUILD

Cover pages, bios, disclaimers, standard methodology sections, approved boilerplate, and signature blocks live in one managed library. Users insert them from the pane and they arrive styled for the document they land in. When the library owner updates a block, the pane finds every document that still carries the old version and proposes the update. This is the enterprise template-governance shape without the enterprise template-governance price.

WHAT STRONG LOOKS LIKE

inserted blocks inherit the host document's fonts and spacing instead of importing their own. Stale-block detection finds real matches across differently formatted documents. Library edits propagate without reinstalling anything.

BAND S2 · templates, search, chat, Review, export · a pull request into the public repository, extensions folder

Outlook add-in: inbound attachment redline responder

WHO IT SERVES

In-house counsel, deal desk, or procurement lead whose contracts arrive by email. Host app: Microsoft Outlook (desktop and web).

WHAT TO BUILD

A contract arrives as an attachment. From the reading pane, one click loads it, runs the review against the company playbook, and shows every proposed change as an approve-or-reject card without leaving Outlook. Approve the set and the add-in composes the reply with the redlined DOCX attached and a plain-language cover note listing the positions taken. The user still presses send.

WHAT STRONG LOOKS LIKE

the returned file opens in the sender's Word with clean tracked changes. The cover note matches what was actually changed. Nothing sends automatically. Works in Outlook desktop and Outlook on the web from one manifest.

BAND S3 · API, chat, Review, templates, export · a pull request into the public repository, extensions folder

Outlook add-in: thread to letterhead client letter

WHO IT SERVES

Professional-services fee earner: lawyer, accountant, consultant, broker. Host app: Microsoft Outlook.

WHAT TO BUILD

A long email thread contains the advice, scattered across replies. The add-in turns it into a formal client letter on firm letterhead: correct salutation, structured advice sections, caveats, scope wording, and signature block. The draft appears for review with each paragraph traceable to the thread messages it came from. Export produces the DOCX and PDF and attaches them to a reply, or files them.

WHAT STRONG LOOKS LIKE

letterhead, margins, and footer are exactly the firm's. The letter says only what the thread supports. Quoted client wording is accurate. A partner would sign it after a read, not a rewrite.

BAND S2 · chat, templates, Review, export · a pull request into the public repository, extensions folder

Outlook add-in: quote and proposal from an enquiry

WHO IT SERVES

SMB owner or inside-sales rep whose pipeline is their inbox. Host app: Microsoft Outlook.

WHAT TO BUILD

An enquiry email arrives. The add-in reads it, pulls the matching line items from the price list the business uploaded, and produces a branded quote or short proposal on the company template: scope, pricing table, terms, validity date. The rep adjusts in the pane ("add installation, apply the ten percent partner discount") and sends it attached to the reply. Sent quotes are kept, so the next similar enquiry starts from the closest prior one.

WHAT STRONG LOOKS LIKE

pricing maths is correct and the table formatting matches the template exactly. Discounts apply where the rep said and nowhere else. The second quote of the day takes under a minute.

BAND S1 · chat, templates, export, memory, Review · a pull request into the public repository, extensions folder

Outlook add-in: pre-send safety gate

WHO IT SERVES

Anyone who has ever sent a document with the previous client's name still in it. Host app: Microsoft Outlook.

WHAT TO BUILD

Before an email with a document attachment sends, the add-in inspects the attachment and warns about what should not leave the building: leftover tracked changes and hidden comments, author and company metadata, the wrong client's name in a footer, an internal-only section, a pricing annex that should not go to this recipient. Each finding is a one-click fix producing a clean copy, and the user chooses to send the cleaned version or the original.

WHAT STRONG LOOKS LIKE

it catches the real embarrassments, including text hidden in comments and rejected tracked changes. The cleaned copy is genuinely clean. False alarms are rare enough that people stop skipping the dialog.

BAND S2 · search, chat, Review, export · a pull request into the public repository, extensions folder

Google Docs add-on: suggesting-mode redline

WHO IT SERVES

In-house legal, HR, or editorial team that lives in Google Docs. Host app: Google Docs.

WHAT TO BUILD

Google Docs already has a review culture built on suggestions. This add-on uses it: proposed edits are written as native Docs suggestions attributed to the add-on, with the reasoning in an attached comment, so the existing accept-and-reject workflow is the approval interface and reviewers learn nothing new. A pane lists the proposals so a long document can be worked in order, and an export command produces a Word file where the suggestions become tracked changes.

WHAT STRONG LOOKS LIKE

suggestions are real Docs suggestions, not coloured text. Accepting them all yields the intended document. The Word export carries the same changes as genuine track changes, so a Docs team and a Word counterparty can work the same document.

BAND S3 · chat, Review, export, search · a pull request into the public repository, extensions folder

Marketplace note: if your build wants a marketplace listing, build it and sideload it; we handle actual listings.

Google Docs add-on: clause and snippet sidebar

WHO IT SERVES

Startup in-house counsel or operations lead running contracts in Docs. Host app: Google Docs.

WHAT TO BUILD

The approved clause set, standard terms, and reusable paragraphs sit in a sidebar with plain-language search. Insert one and it adapts to the document's party names and defined terms and takes the document's styling. An audit command reads the open agreement and reports which standard positions are missing, weakened, or replaced with counterparty language, with a jump link to each.

WHAT STRONG LOOKS LIKE

search works on intent, so "cap our liability" finds the limitation clause. The audit's misses and false alarms are both rare. Inserted text never carries foreign formatting into the document.

BAND S2 · search, templates, chat, Review · a pull request into the public repository, extensions folder

Google Docs add-on: brand finisher

WHO IT SERVES

Consultant, agency producer, or marketer who drafts in Docs and delivers in Word or PDF. Host app: Google Docs.

WHAT TO BUILD

Docs is a good place to write and a poor place to finish. The add-on takes the finished draft and produces the deliverable: the client's or firm's template applied, a cover page, per-section running headers and footers with working page numbers, real footnotes, a generated table of contents, and consistent table and figure treatment. Pick the brand from a set of uploaded templates. Output is a DOCX and a print-ready PDF.

WHAT STRONG LOOKS LIKE

the exported Word file has genuine header and footer parts and real footnotes, not text pretending to be them. Page numbering and section geometry are correct. Switching brand template reflows cleanly with no orphaned content.

BAND S2 · templates, export, chat, images · a pull request into the public repository, extensions folder

Google Docs add-on: export-fidelity and accessibility preflight

WHO IT SERVES

Public-sector communications, university, or enterprise team with accessibility obligations. Host app: Google Docs.

WHAT TO BUILD

Before a Doc is published, this add-on checks what will actually survive and what will actually be readable: heading levels that are real headings rather than big bold text, tables with header rows, images without alternative text, colour-only meaning, reading order, and every element that would degrade on export. Findings are a list with one-click fixes, then it exports a tagged, accessible PDF and a matching Word file.

WHAT STRONG LOOKS LIKE

the fixes are real structural fixes, not cosmetic ones. The exported PDF passes a standard accessibility checker. The report is specific enough to defend in a procurement review.

BAND S2 · chat, Review, export, images · a pull request into the public repository, extensions folder

One Office add-in across Word, Outlook, Excel and PowerPoint

WHO IT SERVES

IT buyer who wants one thing to approve, not four. Host apps: the Office suite.

WHAT TO BUILD

A single manifest that surfaces the right document capability in each host: redline and edit in Word, attachment handling and letter drafting in Outlook, data-to-narrative reporting in Excel, deck-to-leave-behind document in PowerPoint. Shared sign-in, shared templates, shared usage view. One AppSource listing, one security review, one deployment for the admin.

WHAT STRONG LOOKS LIKE

the manifest validates and installs to all four hosts, and each host's pane feels designed for that host rather than a copy of the Word one. Templates and sign-in are genuinely shared. One deployment reaches every user.

BAND S3 · API, chat, templates, export, Review · a pull request into the public repository, extensions folder

Scope: Excel and PowerPoint are read-only sources on this card. You read the sheet or the deck and produce a document from it. Nothing writes back into a spreadsheet or a slide deck, and nothing produces an .xlsx or a .pptx file. That is why this card does not collide with the note in your document about spreadsheet editing.

Word "reformat to our template" ribbon command

WHO IT SERVES

Ops / brand owner

WHAT TO BUILD

One ribbon button maps the open document onto the organisation's style set: headings, numbering, spacing, cover page, running heads. The mapping is a JSON style contract the user can edit.

WHAT STRONG LOOKS LIKE

the content is untouched and only presentation moves. Runs twice with an identical result. Handles a document that was pasted in from three sources.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

Word comment-driven agent

WHO IT SERVES

Editorial / reviewer workflows

WHAT TO BUILD

Reads every unresolved Word comment, turns each into an edit instruction scoped to the commented range, proposes the batch, and resolves the comment thread once the user approves.

WHAT STRONG LOOKS LIKE

comment to range mapping is exact, a comment that says "no change needed" is honoured, and nothing outside the anchor moves.

BAND S2 · API, chat, Review · a pull request into the public repository, extensions folder

Word footnote, endnote and citation pane

WHO IT SERVES

Academic and legal

WHAT TO BUILD

Edit and renumber footnotes and endnotes by instruction, insert a citation in a chosen style, and repair cross references after the body changes.

WHAT STRONG LOOKS LIKE

numbering stays continuous after an insertion in the middle, and a footnote deleted in Word does not orphan its marker.

BAND S2 · chat, API, export · a pull request into the public repository, extensions folder

Word whole-document agent run with a progress panel

WHO IT SERVES

Cross-industry

WHAT TO BUILD

Kicks off an async job over the entire open document, streams milestones into a pane, and applies the result in one reviewable batch. Built on chat_async plus job polling so a long run survives the pane being closed.

WHAT STRONG LOOKS LIKE

close the pane, reopen it, and the run is still there with its state intact. Cancel works. A failed turn reports honestly instead of claiming success.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Word selection rewrite with a live token counter

WHO IT SERVES

Positioning demo

WHAT TO BUILD

Rewrite the selected paragraph and show, side by side, the tokens a whole-document regeneration would have cost versus what the surgical edit actually cost.

WHAT STRONG LOOKS LIKE

the counter is measured, not asserted, and the same document at 3 pages and 300 pages shows the gap widening.

BAND S1 · chat, API · a pull request into the public repository, extensions folder

Word offline queue add-in

WHO IT SERVES

Field and travel users

WHAT TO BUILD

Instructions typed offline are queued locally, and on reconnect they replay in order against the current document version with a conflict prompt if the base moved.

WHAT STRONG LOOKS LIKE

no lost instruction, no double application, and a visible conflict resolution rather than a silent overwrite.

BAND S2 · API, Review · a pull request into the public repository, extensions folder

Word two-document compare pane

WHO IT SERVES

Legal and procurement

WHAT TO BUILD

Opens the current document plus a second file in one SuperDocs session, asks a cross-document question, and applies the answer as edits to whichever document the user names.

WHAT STRONG LOOKS LIKE

the answer cites which document and which section it came from, and edits land in the intended document only.

BAND S2 · Multi-document, search, chat, Review · a pull request into the public repository, extensions folder

Word mail merge with judgment

WHO IT SERVES

Sales and HR

WHAT TO BUILD

Classic merge fields, except each recipient's letter also gets a short reasoned edit from the row data rather than a blind field substitution.

WHAT STRONG LOOKS LIKE

200 letters where every one is individually correct and the template's formatting is identical in all of them. A dry-run mode shows three samples before the batch.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

Word accessibility repair pane

WHO IT SERVES

Government and enterprise

WHAT TO BUILD

Fixes heading hierarchy, list semantics, table headers and image alt text by instruction, then exports a properly structured document.

WHAT STRONG LOOKS LIKE

a before and after run of Word's own Accessibility Checker with the issue count falling, plus no visual change to the document.

BAND S2 · chat, API, export · a pull request into the public repository, extensions folder

Word two-person approval routing

WHO IT SERVES

Regulated teams

WHAT TO BUILD

The drafter proposes, a named reviewer receives the pending batch, and only the reviewer can approve. Uses the durable approval state so the reviewer can be on a different machine.

WHAT STRONG LOOKS LIKE

the drafter genuinely cannot self-approve, the pending set survives a restart, and the audit trail names who decided what.

BAND S3 · Review, API · a pull request into the public repository, extensions folder

Locked-down enterprise macro client

WHO IT SERVES

Enterprise IT

WHAT TO BUILD

For environments where add-ins are blocked: a VBA or PowerShell client that posts the document to the API, retrieves the edit and writes it back, distributed as a signed template.

WHAT STRONG LOOKS LIKE

works with no add-in install, no admin consent, and documents its own security posture honestly.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Word template gallery pane

WHO IT SERVES

Cross-industry

WHAT TO BUILD

Browse the organisation's templates from inside Word, open one, and populate it from a brief. Clone-and-update rather than start-from-blank.

WHAT STRONG LOOKS LIKE

the clone keeps every style and placeholder, and last quarter's document becomes this quarter's in one turn.

BAND S2 · templates, chat, export · a pull request into the public repository, extensions folder

Word write-back fidelity harness

WHO IT SERVES

Developer tooling

WHAT TO BUILD

A test rig that runs N documents through the add-in, re-reads the result, and reports what changed beyond the intended edit.

WHAT STRONG LOOKS LIKE

a published fidelity score with a reproducible corpus, and at least one real defect found and reported.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

AppSource submission kit

WHO IT SERVES

Distribution

WHAT TO BUILD

A complete, validated add-in manifest, privacy and support pages, store copy, screenshots and the validation checklist, so any of the panes above can actually be listed.

WHAT STRONG LOOKS LIKE

passes Microsoft's validation rules on a dry run, and the copy leads with the editing capability rather than file conversion.

BAND S1 · API · a pull request into the public repository, extensions folder

Docs comment-thread resolver

WHO IT SERVES

Editorial

WHAT TO BUILD

Walks open comment threads, drafts the edit each one implies, and closes the thread on acceptance.

WHAT STRONG LOOKS LIKE

threads with disagreement are surfaced for a human rather than resolved arbitrarily.

BAND S2 · API, chat, Review · a pull request into the public repository, extensions folder

Apps Script brand-normaliser menu

WHO IT SERVES

Ops

WHAT TO BUILD

A custom menu item that pushes any Doc onto the organisation's brand system, deployable across a domain without a marketplace listing.

WHAT STRONG LOOKS LIKE

installs by copying one script file, and the admin can pin the style contract centrally.

BAND S1 · API, templates · a pull request into the public repository, extensions folder

Sheets narrative report add-on

WHO IT SERVES

Finance

WHAT TO BUILD

Turns a selected range into a written report with the numbers embedded in prose and a chart, on the firm template.

WHAT STRONG LOOKS LIKE

the numbers in the prose match the sheet exactly, and re-running after a data change updates only what moved.

BAND S2 · API, images, templates, export · a pull request into the public repository, extensions folder

Gmail add-on: thread to formal letter

WHO IT SERVES

Cross-industry

WHAT TO BUILD

Turns a messy thread into a properly structured letter or memo, attached back to a reply draft.

WHAT STRONG LOOKS LIKE

extracts the actual commitments from the thread rather than summarising politely.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Google Chat document app

WHO IT SERVES

Collaboration

WHAT TO BUILD

A Chat app that holds a per-space session so a team edits one document together through conversation.

WHAT STRONG LOOKS LIKE

the space's session persists across days and members, and approvals are attributed.

BAND S2 · MCP, chat, Review · a pull request into the public repository, extensions folder

Slides speaker notes to a leave-behind

WHO IT SERVES

Consulting

WHAT TO BUILD

Converts a deck plus its speaker notes into a written document that stands on its own.

WHAT STRONG LOOKS LIKE

reads the notes rather than only the slide text, and the result is genuinely readable prose.

BAND S2 · API, images, export · a pull request into the public repository, extensions folder

Reusable Apps Script client library

WHO IT SERVES

Developer

WHAT TO BUILD

A published, clasp-managed Apps Script library wrapping the API, so any of the Workspace builds above start from one dependency.

WHAT STRONG LOOKS LIKE

typed JSDoc, quota-aware retries, and a working example for Docs, Sheets and Gmail.

BAND S1 · API · a pull request into the public repository, extensions folder

Workspace Marketplace publication kit

WHO IT SERVES

Distribution

WHAT TO BUILD

OAuth scopes justified and minimised, verification material, listing copy and screenshots for a Workspace add-on.

WHAT STRONG LOOKS LIKE

the scope list is genuinely minimal and each scope is justified in writing. Passes a dry review.

BAND S1 · API · a pull request into the public repository, extensions folder

VERTICAL MINI-APPS

WordPress AI content block

WHO IT SERVES

Marketing / a content marketer or WordPress site owner who publishes blog posts regularly and wants AI drafting without leaving the block editor.

WHAT TO BUILD

Marketing teams that publish on WordPress draft their copy in a separate AI tool, then lose headings and images pasting it back into the block editor. Build a custom Gutenberg content block that calls the platform API directly from inside the WordPress post editor, so a writer can draft or revise a section, insert a generated image, and keep the structure intact without leaving WordPress. This lives inside the WordPress editor itself as a native block, not a headless-CMS export bridge sitting outside it.

WHAT STRONG LOOKS LIKE

A reviewer drafts a section with an inline image inside a real WordPress post using the block, then exports the same content as a document and confirms headings and images survived unchanged in both directions.

BAND S2 · API, images · a pull request into the public repository, extensions folder

Figma-to-spec doc plugin

WHO IT SERVES

Design / a product designer or design-systems lead who hands off frames to engineering and currently writes explanatory specs by hand.

WHAT TO BUILD

Designers hand off frames to engineering with comments scattered across a Figma file, and someone still writes the functional spec by hand afterward. Build a Figma plugin that reads the selected frames and their attached notes, then drafts a structured spec document (states, interactions, edge cases) through the platform API, so the handoff artifact is a real document instead of a comment thread. This produces a functional specification document from frames and notes, not a branded visual document template.

WHAT STRONG LOOKS LIKE

A reviewer selects a multi-state Figma frame with notes, runs the plugin, and checks the generated spec document lists every visible state and interaction called out in the notes, not just a restated visual description.

BAND S2 · API, images, export · a pull request into the public repository, extensions folder

NDA/MSA self-serve generator

WHO IT SERVES

Freelancer, consultant, or small-agency owner who signs a handful of NDAs and MSAs a year and cannot justify counsel for each

WHAT TO BUILD

Build a public self-serve web app: the user answers a short question flow (parties, term, jurisdiction preferences, mutual or one-way, payment and IP posture) and receives a redline-ready NDA or MSA as a clean DOCX, drafted through the API from a maintained template pair. Every generated document opens with a plain-language summary of what was chosen and why.

WHAT STRONG LOOKS LIKE

two different answer sets provably change only the clauses those answers govern, with the rest byte-identical; the generated document is redline-ready in a real word processor; unanswered questions produce bracketed placeholders, never invented terms.

BAND S2 · chat, templates, Review, export · a pull request into the public repository, use-cases folder

Curriculum/lesson-plan editor

WHO IT SERVES

EdTech curriculum leads and teachers who already have plan drafts and need them cleaned, restructured, and exported fast.

WHAT TO BUILD

Teachers and curriculum leads spend hours reformatting lesson plans and assessments every time a standard changes or a co-teacher hands off a unit. Build a chat-driven editor that takes an existing lesson plan or assessment and restructures, reformats, and polishes it in place, never generating original lesson content from a blank page. Load a plan from a template, edit sections by chat instruction, export a clean copy.

WHAT STRONG LOOKS LIKE

Reviewer opens a rough lesson plan, asks for a reformat pass (objectives/materials/steps/assessment sections) plus one wording edit, and gets both changes applied to the same document, checkable diff against the original.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Loan-document packet assembler

WHO IT SERVES

Banking operations teams assembling disclosure and closing packets from loan origination data.

WHAT TO BUILD

Closing a loan means collecting disclosures, schedules, and closing documents into one consistent packet, and re-assembling that packet by hand every time a term changes is slow and error-prone. Build an API-driven assembler that takes structured loan terms plus a set of disclosure/closing templates and produces a complete, consistently formatted packet as a multi-document set.

WHAT STRONG LOOKS LIKE

Reviewer posts one loan-terms payload and gets back a multi-document packet where every document reflects the same borrower, rate, and term fields with zero mismatches between documents.

BAND S2 · API, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Real-estate listing one-pager maker

WHO IT SERVES

Real estate

WHAT TO BUILD

Agents need a polished one-pager for every new listing, combining property details with photos, and doing that by hand in a general word processor is slow. Build a chat-driven generator that takes listing data (address, price, features) plus images and produces a styled one-pager ready to export and share.

WHAT STRONG LOOKS LIKE

Reviewer supplies listing fields plus two photos and gets one export-ready one-pager with images placed correctly and all listing facts present, no manual layout work.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Claude Code "edit-docx" skill

WHO IT SERVES

DevTools/Agents

WHAT TO BUILD

Claude Code can already reach SuperDocs' MCP tools once a user registers the server, but there is no packaged Skill that tells it when and how to use them mid-task, so most sessions never discover the tools or call them wrong. Build a Claude Code Agent Skill, a SKILL.md with YAML frontmatter inside `.claude/skills/edit-docx/`, following the open Agent Skills spec, that Claude Code loads automatically whenever a task touches a `.docx` file. The skill should walk the agent through opening the file via SuperDocs MCP, making the requested edit, and exporting the result back to disk, all without the user leaving their coding session.

WHAT STRONG LOOKS LIKE

install the skill in a repo containing a sample docx, ask Claude Code mid-session to make a real edit to it, and the transcript shows the edit completed end to end through the skill with zero manual MCP invocation from the user.

BAND S3 · MCP, export · a pull request into the public repository, use-cases folder

Cursor document-editing skill

WHO IT SERVES

DevTools/Agents

WHAT TO BUILD

Wiring SuperDocs into Cursor's `.cursor/mcp.json` gets the tools connected, but nothing tells Cursor's agent to reach for them in the middle of a task instead of trying to edit the binary file directly. Build a Cursor Project Rule, a `.cursor/rules/edit-docx.mdc` file with frontmatter scoping it to `*.docx` so Cursor auto-attaches it on relevance, paired with the SuperDocs `.cursor/mcp.json` entry itself. The rule should govern behavior during an edit already in progress, not read like a one-time setup checklist.

WHAT STRONG LOOKS LIKE

in a repo with a tracked docx, a mid-conversation edit request makes Cursor's agent call the SuperDocs MCP tool on the strength of the rule alone, with no extra prompting from the reviewer.

BAND S3 · MCP, export · a pull request into the public repository, use-cases folder

Windsurf/Cline docs skill

WHO IT SERVES

DevTools/Agents

WHAT TO BUILD

Windsurf and Cline are two more MCP-capable coding agents with no packaged SuperDocs skill for either, and each has its own file layout for rules and MCP config, so a Claude Code or Cursor package cannot just be copied over. Build a dual-host package: for Windsurf, a workspace rule under `.windsurf/rules/` plus an entry in its `mcp_config.json`; for Cline, an entry in its `.clinerules/` folder plus its own `cline_mcp_settings.json`. Both halves should teach the same behavior, open the target document through SuperDocs' MCP tools before touching a docx, pdf, or html file directly, using each host's native mechanism rather than a shared shim.

WHAT STRONG LOOKS LIKE

installing the Windsurf half in a Windsurf workspace and the Cline half in VS Code, a document-revision request in each routes through SuperDocs MCP using only that host's own rule and config files.

BAND S3 · MCP · a pull request into the public repository, use-cases folder

LangChain/LangGraph document-tool

WHO IT SERVES

Agents

WHAT TO BUILD

LangChain and LangGraph builders have no drop-in way to give a graph document-editing ability; today it means hand-rolling raw HTTP calls. Build a small adapter package that binds SuperDocs' document operations as native LangChain tools and wires them into a runnable LangGraph graph. Ship a runnable demo graph alongside the adapter, not just the tool definitions in isolation, so the wiring is provably working.

WHAT STRONG LOOKS LIKE

running the demo graph with a real prompt produces an actually-edited document, and the loaded tool list is visibly the live SuperDocs set, not a hand-written re-implementation.

BAND S3 · API, MCP · a pull request into the public repository, use-cases folder

CrewAI document-agent role

WHO IT SERVES

Agents

WHAT TO BUILD

CrewAI builders assign every crew member a role, goal, and tool list by hand, and nobody ships a ready-made document-editor role backed by SuperDocs. Build a packaged CrewAI tool wrapping SuperDocs' document operations with typed inputs, plus a pre-configured document-editor crew member with a goal and backstory pinned to real document work, ready for another crew to import and assign work to. Ship a minimal demo crew, a writer or researcher role handing off to this one, so the tool is proven inside an actual crew run.

WHAT STRONG LOOKS LIKE

a reviewer drops the packaged agent into their own crew, assigns it a task against a real document, and the run produces an actually-edited document, and CrewAI's own execution log shows the document operation being called, not a hardcoded shortcut.

BAND S3 · API · a pull request into the public repository, use-cases folder

AutoGen doc-writer agent

WHO IT SERVES

Agents

WHAT TO BUILD

AutoGen pipelines pass messages between conversable agents, and any new capability has to be registered on at least two agents, one to call it and one to execute it, a step nobody has packaged for document work. Build a document-writer conversable agent plus a companion executor agent, with SuperDocs' operations registered between them the way AutoGen expects, each operation described so the assistant knows when to call it inside a normal conversation. Ship a runnable group chat, at least two agents, that proves a document actually gets produced from a natural-language request.

WHAT STRONG LOOKS LIKE

running the sample group chat with a plain-language drafting request ends with a real exported document, and the transcript shows the document operation invoked through the agents' registered tooling, not a hardcoded shortcut.

BAND S3 · API · a pull request into the public repository, use-cases folder

MCP pipeline: transcript → prescription → PDF

WHO IT SERVES

Healthcare/Agents

WHAT TO BUILD

Clinics running any transcription tool already get raw text from a visit, but turning that into a properly formatted prescription still means someone copy-pastes into a template by hand. Build an agent pipeline chaining an external transcription MCP tool with SuperDocs: the transcript is parsed into structured prescription fields, dropped into a SuperDocs template, and held for a clinician's explicit review and approval before it is exported as a PDF. Frame the output honestly throughout as a drafting aid for the clinician to approve, never an auto-finalized or certified clinical record.

WHAT STRONG LOOKS LIKE

feeding the pipeline a sample dictation transcript produces a correctly formatted prescription PDF that stopped at a visible approval step before export, with any clinician edits reflected in the final document.

BAND S3 · MCP, chat, export · a pull request into the public repository, use-cases folder

MCP pipeline: repo → release-notes PDF

WHO IT SERVES

DevTools/Agents

WHAT TO BUILD

Engineering teams already have a full commit history in git, but turning raw commits into a release-notes document worth sending to a customer is still manual curation. Build an agent pipeline that pulls commits between two tags from GitHub's own MCP server (list_commits/get_commit), groups them into categories such as features, fixes, and breaking changes with an LLM pass, and drops the result into a styled SuperDocs release-notes document, exported as a PDF. Source the pipeline from commits themselves, never from an issue tracker or an already-written changelog, so the chain is genuinely repo-to-document.

WHAT STRONG LOOKS LIKE

pointing the pipeline at two real tags in a repository produces a release-notes PDF whose sections trace back to actual commit messages in that range, chaining two live MCP servers rather than one script hand-rolling both.

BAND S3 · MCP, export, images · a pull request into the public repository, use-cases folder

Agent "report writer" skill

WHO IT SERVES

Cross-industry. Any team running a research, analyst, or support agent whose output today ends as a wall of chat text instead of a document.

WHAT TO BUILD

Most research and analyst agents finish a task with a good answer buried in chat text, and nothing generic turns that into an actual document. Build an agent skill, packaged to the open Agent Skills spec so it is not tied to one coding-agent host, that takes whatever research the agent already produced in the conversation and turns it into a properly structured report through SuperDocs: headings, a summary section, and export in the requested format. Keep the skill itself topic-agnostic; it supplies the report-writing behavior, not any one industry's template.

WHAT STRONG LOOKS LIKE

handing the skill a messy multi-paragraph research answer with no formatting instructions produces a properly headed, exportable report document, with no structure specified by the user.

BAND S3 · MCP, chat, export · a pull request into the public repository, use-cases folder

Agent "fill this template" skill

WHO IT SERVES

Cross-industry. Ops and back-office teams who fill the same template with new data on a recurring basis and want the agent to do the filling.

WHAT TO BUILD

Filling a template with new data by hand is a recurring, low-value task across a lot of back-office work, and there is no generic agent skill that does it. Build a skill that takes any SuperDocs template plus a structured data source (a spreadsheet row, a JSON object, or a form submission), discovers the template's own placeholders rather than assuming fixed field names, and produces a filled document exported in the requested format. The skill should work against a genuinely new template on first use, with no per-template code written ahead of time.

WHAT STRONG LOOKS LIKE

handing the skill a template with five distinct placeholders and a data record covering all five produces an exported document with every placeholder correctly replaced and no manual field mapping written by the user.

BAND S3 · MCP, templates, export · a pull request into the public repository, use-cases folder

HTML → styled DOCX API demo

WHO IT SERVES

DevTools / a developer evaluating the API in their first hour who wants a minimal, copyable code sample rather than a full app.

WHAT TO BUILD

Developers evaluating the API want to see, in a few lines of code, that arbitrary app-generated HTML becomes a properly styled Word document, not a formatting dump. Build a small, single-purpose demo: one script that takes an HTML page (say, an invoice or report generated inside your own app) and calls the export API to return a branded DOCX, with the source HTML and the resulting document shown side by side.

WHAT STRONG LOOKS LIKE

A reviewer runs the demo against a sample HTML file with headings, a table, and an image, and confirms the DOCX output preserves all three without manual cleanup.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

Markdown → PDF/DOCX service

WHO IT SERVES

DevTools / a developer or technical writer who wants a standing markdown-to-document conversion service their own pipeline can call repeatedly, not a one-off script.

WHAT TO BUILD

Developers writing documentation in markdown want styled, shareable output, including diagrams and math, without hand-building a rendering pipeline. Build a small deployable service, not just a script, that accepts markdown (with diagram blocks and math notation) over an API and returns styled PDF or DOCX, so another tool or job can call it repeatedly as a standing dependency.

WHAT STRONG LOOKS LIKE

A reviewer sends a markdown file containing a diagram and a math expression to the running service and confirms both render correctly in the returned PDF and DOCX, not just the prose.

BAND S2 · API, export, images · a pull request into the public repository, use-cases folder

Scanned-PDF → editable doc app

WHO IT SERVES

Cross-industry / anyone with a paper trail, a consumer, student, or ops professional, turning a scanned, screenshotted, or photographed document into something editable.

WHAT TO BUILD

People with a scanned PDF, a screenshot, or a photo of a printed page have no clean way to get from a flat image to an editable document. Build an app with three capture paths into the same pipeline: upload a scanned PDF, upload a screenshot, or upload a photo taken on a phone, all OCR'd and ingested into an editable document the user can then revise through chat and export cleanly. This is a consumer-facing capture-and-edit tool, not back-office scan-intake routing.

WHAT STRONG LOOKS LIKE

A reviewer runs all three capture paths (scanned PDF, screenshot, phone photo) on legible source material and confirms each lands in the same editable document state, ready to edit, not just a text dump.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

Bilingual/side-by-side doc generator

WHO IT SERVES

Legal/Localization / a company or law firm preparing one contract or notice that must read correctly in two languages side by side, such as a cross-border sales agreement.

WHAT TO BUILD

Companies producing a bilingual contract or notice need one document where both language versions sit side by side and read as a single authoritative file, not two separate documents to keep in sync. Build an app that takes source content and a target language, drafts the translation, and lays both out in a two-column document with matching structure end to end.

WHAT STRONG LOOKS LIKE

A reviewer generates a two-column contract from a single source, then checks that a heading or clause added on one side lines up with its counterpart on the other at the same position, in the same exported document.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Whitepaper/ebook builder

WHO IT SERVES

Marketing / a content marketer or founder producing a long-form flagship whitepaper or ebook.

WHAT TO BUILD

Marketers asked for a flagship whitepaper or ebook end up stitching together a table of contents, diagrams, and a hero image from three different tools before it ever reaches a document. Build a single long-form app: chat-driven drafting of a multi-section paper with an auto-generated table of contents, inline diagrams, and a generated hero image, exported as one styled PDF. This is one illustrated document, not a multi-artefact launch kit.

WHAT STRONG LOOKS LIKE

A reviewer drafts a five-section whitepaper through chat and confirms the exported PDF has a working table of contents, at least one rendered diagram, and a hero image, all in one file.

BAND S2 · chat, images, export · a pull request into the public repository, use-cases folder

Case-study generator

WHO IT SERVES

Marketing / a customer marketing or content lead turning a customer's results into a structured, prospect-facing case study.

WHAT TO BUILD

Customer marketers turning a happy customer's results into a case study spend more time formatting than writing. Build a case-study app that takes inputs like a customer quote, metrics, and a few images and drafts a structured, prospect-facing case study, headline, challenge, solution, results, with the supplied images placed correctly, ready to export. This targets prospects evaluating a purchase, not a job applicant's portfolio.

WHAT STRONG LOOKS LIKE

A reviewer feeds in a customer quote and three metrics and confirms the generated case study states the results as a distinct, skimmable section, not buried in narrative prose.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Diagram-doc generator (Mermaid)

WHO IT SERVES

DevTools / an engineer or technical writer documenting a system who wants prose and diagrams generated together instead of drawn and written separately.

WHAT TO BUILD

Engineering docs that mix prose and architecture diagrams tend to drift apart because the diagram gets edited separately from the words around it. Build a chat-driven app where a request like "explain the request flow" produces prose and an auto-rendered diagram together in one turn, staying in the same document, not a separate build or rebuild pipeline for existing files.

WHAT STRONG LOOKS LIKE

A reviewer asks the app to explain a multi-step process and confirms the response contains both explanatory prose and a rendered diagram of those same steps in one generated section.

BAND S2 · chat, images, export · a pull request into the public repository, use-cases folder

Pitch one-pager maker

WHO IT SERVES

Startups / a founder preparing a single branded one-pager for investor or partner outreach.

WHAT TO BUILD

Founders reaching out to investors or partners need one clean, branded page fast, not a full deck or a multi-document kit. Build a small app that takes a short pitch input (what the company does, traction, the ask) and drafts a single branded one-pager with a generated hero image, exported and ready to send.

WHAT STRONG LOOKS LIKE

A reviewer enters a three-line pitch and confirms the output is exactly one page with a hero image and the ask clearly stated, not a multi-page deck.

BAND S2 · chat, images, export · a pull request into the public repository, use-cases folder

Infographic-in-doc app

WHO IT SERVES

Marketing / a marketing or content professional who needs a simple, legible infographic embedded directly in a document.

WHAT TO BUILD

Marketers who want a simple infographic, a few stats with labels laid out cleanly, currently have to leave the document for a separate design tool. Build an app that generates a legible, text-bearing infographic image from a short data or stat input and embeds it directly inside the document being written, no external design tool involved.

WHAT STRONG LOOKS LIKE

A reviewer requests an infographic with four labeled stats and confirms every label and number in the generated image is legible and matches the input, not garbled or approximated text.

BAND S2 · images, chat · a pull request into the public repository, use-cases folder

Job-description standardizer

WHO IT SERVES

HR/Recruitment / a talent acquisition or people-ops lead standardizing every open job description onto one template and voice.

WHAT TO BUILD

Companies with job descriptions written by a dozen different hiring managers end up with a dozen different formats and tones. Build an app that takes a batch of existing job descriptions and rewrites each to a single structural template and a consistent brand voice, leaving the actual role content and requirements untouched. This normalizes format and voice, it does not audit language for bias or restructure resumes.

WHAT STRONG LOOKS LIKE

A reviewer runs three job descriptions written in different styles through the app and confirms all three come back in the same section order and voice, with the original responsibilities and requirements still present.

BAND S2 · chat, templates · a pull request into the public repository, use-cases folder

Study-guide generator

WHO IT SERVES

EdTech / a student or tutor turning raw class notes into a structured, equation-bearing study guide.

WHAT TO BUILD

Students with a pile of raw class notes have to build their own structured study guide before they can revise from it. Build an app that takes messy notes, including ones with equations, and drafts a structured study guide: organized sections, key terms, and properly rendered equations, exported for printing or review. This builds a new guide from notes, it does not adapt an existing course document.

WHAT STRONG LOOKS LIKE

A reviewer feeds in a page of notes containing at least one equation and confirms the generated guide renders the equation correctly and organizes the material under clear section headings, not a flat restatement.

BAND S2 · chat, images, export · a pull request into the public repository, use-cases folder

Rubric + assessment builder

WHO IT SERVES

EdTech / a teacher or instructional designer generating new assessments and their grading rubrics from a template.

WHAT TO BUILD

Teachers building a new quiz or assignment often write the assessment and its grading rubric as two separate, unlinked documents. Build an app that generates both together from a topic and a template: a new assessment and a matching rubric, aligned question by question, exported as one document or two linked ones. This generates new assessments, it does not grade or comment on a student's existing submission.

WHAT STRONG LOOKS LIKE

A reviewer requests a five-question assessment on a topic and confirms the generated rubric has one scoring row per question, not a generic rubric detached from the actual questions.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

SOP-onboarding kit builder

WHO IT SERVES

Ops / an HR or operations lead assembling a new-hire onboarding document set for a specific role.

WHAT TO BUILD

Assembling a new-hire onboarding packet for one role, a welcome note, a first-week checklist, a role SOP, a point-of-contact sheet, means writing several linked documents from scratch every time. Build an app that takes a role name and its existing procedures and assembles a linked onboarding document set for that new hire, exported together. As an optional agent-loop variant of the same build, a scheduled agent assembles the packet when a new hire is added and refreshes already-assembled packets when a source procedure changes. This is internal new-hire onboarding for a role, not customer onboarding and not a retraining update to an existing SOP.

WHAT STRONG LOOKS LIKE

A reviewer requests an onboarding kit for one named role and confirms the output is a linked set of at least three distinct documents, not one long file, each referencing the same role correctly.

BAND S2 · Multi-document, templates, export · a pull request into the public repository, use-cases folder

Document-assembly microservice

WHO IT SERVES

Legal/Finance / an engineering team at a legal-tech or fintech company that needs to assemble compliant multi-section documents from a clause library, called from their own system.

WHAT TO BUILD

Legal and finance teams that assemble documents from a library of standard clauses and sections currently do it by hand, copying blocks between files. Build a callable service: given an ordered list of clause or section identifiers, it assembles them into one coherent multi-section document through the API and returns it, so another system, a CLM tool, an internal portal, can request an assembled document without a human doing the copy-paste. This composes documents from blocks, it does not build or cache a project's own source files.

WHAT STRONG LOOKS LIKE

A reviewer calls the service with an ordered list of five clause blocks and confirms the returned document has all five in the requested order with consistent formatting, not a copy-paste seam between them.

BAND S2 · API, Multi-document, templates · a pull request into the public repository, use-cases folder

Async large-export service

WHO IT SERVES

Cross-industry / a developer or product team whose app needs to hand end users large generated documents without blocking on a slow export.

WHAT TO BUILD

An app that needs to hand a user a 100-plus-page generated document can't make them wait on a blocking request, and building the async job, polling, and delivery logic from scratch is its own project. Build a small reference client around the existing async large-export job: kick off generation of an oversize document, poll or wait for completion, and deliver the finished file by a signed link or email, without the caller's app blocking on the request.

WHAT STRONG LOOKS LIKE

A reviewer generates a document long enough to trigger the async path and confirms they receive a working download link or email, not a timed-out request or a truncated file.

BAND S3 · API, export · a pull request into the public repository, use-cases folder

Framer/Webflow "generate doc" widget

WHO IT SERVES

Marketing / a growth or marketing-ops lead running a public site who wants an embeddable widget that turns a short form into a personalized PDF lead magnet.

WHAT TO BUILD

Marketing sites that want a lead-gen offer, fill in a few fields and get a personalized PDF, currently have to build that generation logic themselves. Build an embeddable widget shell for Framer or Webflow: a short public-facing form that calls the platform API to generate a personalized PDF, a custom report or checklist, and delivers it to the visitor, with the actual generation logic living behind the widget, not inside the site builder.

WHAT STRONG LOOKS LIKE

A reviewer fills out the embedded form on a live Framer or Webflow test page and confirms a personalized PDF reflecting the entered fields downloads or arrives by email within the same session.

BAND S1 · API, export · a pull request into the public repository, use-cases folder

"Recruitment AI that writes the offer letter" reference app

WHO IT SERVES

HR and recruiting teams evaluating an agent that reasons and writes in the same pass.

WHAT TO BUILD

Hiring teams want one flow that scores a candidate against a role and produces the resulting paperwork, instead of switching tools between a scoring step and a drafting step. Build a reference app where the agent reads a candidate profile against a role brief, produces a scored recommendation, and then drafts the matching offer letter from a template when the recommendation is a hire.

WHAT STRONG LOOKS LIKE

Reviewer feeds one fictional candidate profile plus a role brief, sees a structured score and recommendation, and receives an offer letter whose terms (title, comp, start) trace directly back to that recommendation.

BAND S2 · MCP, templates, export · a pull request into the public repository, use-cases folder

Email-to-contract agent

WHO IT SERVES

Sales/Legal

WHAT TO BUILD

Sales and legal teams negotiate deal terms over loose email threads, then someone has to manually pull those terms into a first contract draft. Build a chat agent that reads a pasted, fictional negotiation thread, extracts the agreed terms, and drafts a first-pass contract from a template built on those terms.

WHAT STRONG LOOKS LIKE

Reviewer pastes a short back-and-forth thread with three negotiated terms and gets one contract draft where all three terms appear correctly and nothing outside the thread is invented.

BAND S3 · chat, templates, export · a pull request into the public repository, use-cases folder

Research-agent → whitepaper

WHO IT SERVES

Marketing/Research

WHAT TO BUILD

Marketing and research teams want a polished, styled whitepaper built from a topic brief, not just a wall of unformatted research notes. Build a chat agent that researches a given topic, structures findings into a whitepaper outline, and assembles a styled, exportable document with section headers and supporting imagery.

WHAT STRONG LOOKS LIKE

Reviewer gives one topic brief and receives a styled multi-section whitepaper with a coherent structure and at least one supporting image, exportable in one step.

BAND S3 · chat, images, export · a pull request into the public repository, use-cases folder

Codegen-agent → user-manual

WHO IT SERVES

DevTools

WHAT TO BUILD

Developer tools that ship a feature rarely ship its documentation at the same time, leaving manuals stale or missing. Build an agent reference app that, given a described feature or small code sample, produces both an explanation of what was built and a matching end-user manual section, kept consistent with each other.

WHAT STRONG LOOKS LIKE

Reviewer gives one small feature description and gets a manual section whose steps and terminology match the feature description exactly, plus one illustrative diagram.

BAND S3 · MCP, images, export · a pull request into the public repository, use-cases folder

Proposal-agent with e-sign handoff

WHO IT SERVES

Sales teams and technical evaluators who want proposal generation chained directly into signing without a manual handoff.

WHAT TO BUILD

Sales proposals still dead-end in a document that someone has to manually shepherd into a signing tool. Build an MCP-driven agent that generates a sales proposal from deal terms and templates, then hands the finished document off to an external e-signature MCP for signing, never performing the signature step itself.

WHAT STRONG LOOKS LIKE

Reviewer supplies deal terms, gets a generated proposal document plus a clearly logged handoff call to an external signing MCP, with no in-house signature capability implied anywhere in the flow.

BAND S3 · MCP, templates, export · a pull request into the public repository, use-cases folder

Comparison intent pages (vs TipTap AI / CKEditor AI)

WHO IT SERVES

DevTools. An engineer who reads a competitor's documentation as carefully as their own and writes without hype.

WHAT TO BUILD

Engineers picking a rich text editor's AI story want to know exactly what they are signing up for: hosting a model, holding an API key, writing the prompt layer, or none of that. Build a short, factual comparison of what wiring AI editing through the SuperDocs public API requires versus what TipTap's Content AI extension and CKEditor's AI Assistant require, with every claim cited to that product's own current public documentation. Pair the writeup with a small runnable demo that wires the SuperDocs side into a plain editor, so a reader can see the integration rather than take the comparison on faith. Keep the tone strictly factual and technical; this is a developer comparison, not a sales pitch.

WHAT STRONG LOOKS LIKE

Every comparative claim links to the specific doc page it came from, a cold reader would call the tone neutral, and the demo installs and runs from a fresh clone in under ten minutes.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

Agent-skill gallery entries

WHO IT SERVES

Agents. A developer who already knows the API and MCP surface and wants reach into agent ecosystems, not just another repo.

WHAT TO BUILD

Agent platforms are starting to support installable skills, small packaged capabilities described in a manifest with worked examples, that an agent can discover and load. Package a focused set of document-editing skills built on the SuperDocs API and MCP server, such as draft-from-outline or export-to-pdf, as real loadable skill packages, then publish them into the public skill directories where such packages are discovered. This is packaging and publishing, not building an integration pack for another agent framework's own tooling, and not just adding a template into an existing gallery: the deliverable is a live, installable skill listing.

WHAT STRONG LOOKS LIKE

A skill installs from the published listing and completes a real edit through the public API on the first try, and the listing itself, not only the underlying code, is live and publicly visible.

BAND S3 · MCP, export · a pull request into the public repository, use-cases folder

Interactive API playground app

WHO IT SERVES

DevTools. A first-time evaluator who wants to feel the product before writing any code.

WHAT TO BUILD

New evaluators should not have to write code or manage an API key just to feel what chat-driven editing is like. Build a hosted, in-browser playground: paste or upload a short document, type an edit instruction, watch the proposed change render live, and export the result, all without leaving the browser or opening a terminal. This is distinct from a command-line REPL or a downloadable sandbox: nothing to install, the whole loop runs in a browser tab.

WHAT STRONG LOOKS LIKE

A first-time visitor with no account gets from landing on the page to a rendered, exported document in under two minutes, and the playground's own API key never shows up in browser network traffic a visitor could inspect.

BAND S2 · API, chat, export · a pull request into the public repository, use-cases folder

Metadata / redline-history scrubber

WHO IT SERVES

Legal/Finance / a professional preparing a document for external send who needs author names and edit history gone before it leaves the building.

WHAT TO BUILD

Before a document goes to someone outside the company, it often still carries author names in its file properties and a full tracked-change history nobody meant to send. Build a pre-send cleanup app: accept every tracked change, strip the visible author and editing-history metadata from the file properties, and export a clean copy, so what goes out shows only the final content. This is mechanical file-property hygiene, not an audited or certified redaction tool, and it does not claim to find every possible trace of prior authorship.

WHAT STRONG LOOKS LIKE

A reviewer opens a document with several tracked changes and a non-default author name in its properties, runs it through the app, and confirms the exported file shows no pending tracked changes and no longer carries the original author name in its visible file properties.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

"Convert any file" web app

WHO IT SERVES

Cross-industry / a consumer or small-business user who needs a fast utility to turn one file format into another without opening a full document editor first.

WHAT TO BUILD

People who just need a file in a different format land on generic converter sites that mangle formatting. Build a named consumer entry-point web app: drop in a DOCX, PDF, RTF, HTML, markdown, or image file, pick any of the five supported export formats, docx, pdf, html, markdown, or txt, and get a converted file back, with an optional step to open and edit the content before exporting. As the largest of these builds, this is a full public-facing app with upload, format selection, optional in-editor review, and download, not a single API call.

WHAT STRONG LOOKS LIKE

A reviewer uploads a PDF, picks docx as the target, and confirms the downloaded file opens with structure and text intact, then repeats with a DOCX input to a markdown target and gets equally clean output.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

Document merge/combine tool

WHO IT SERVES

Operations or admin generalist assembling one deliverable from many contributors' files

WHAT TO BUILD

Build the cross-industry merge tool that owns the "merge files into one document" job: take several documents in mixed formats, order them, produce one clean combined document with a generated cover and table of contents, consistent heading levels and numbering, and per-source attribution. This is the general entry point, not an industry binder.

WHAT STRONG LOOKS LIKE

source ordering is the user's and survives re-export; each source's content is provably present and untruncated; heading hierarchy in the combined file is consistent even when the sources disagree about levels.

BAND S2 · Multi-document, export · a pull request into the public repository, use-cases folder

"Edit this PDF with AI" entry point**WHO IT SERVES**

Cross-industry / someone who needs to edit a PDF they don't have the source file for, using a focused single-purpose tool.

WHAT TO BUILD

Someone who wants to change a few lines in a PDF, and doesn't have the source file, has no clean single-purpose tool for that. Build a named entry-point app: upload a PDF, ingest it through OCR, make targeted edits by describing them in chat, and export a clean PDF or DOCX with just those changes applied.

WHAT STRONG LOOKS LIKE

A reviewer uploads a text-based PDF, asks for one specific sentence to be changed, and confirms the exported file has that exact change with the rest of the document unchanged.

BAND S2 · chat, export · a pull request into the public repository, use-cases folder

Structured field extractor**WHO IT SERVES**

Finance/Legal/Ops / an operations or finance team that receives invoices, contracts, or forms and needs key fields pulled into structured data without hand-retyping.

WHAT TO BUILD

Teams that receive invoices, contracts, or forms need the key fields pulled out as structured data, not retyped by hand, while keeping the original file untouched for the record. Build an app that ingests a document, asks for its key fields through a structured query, and has the wrapper parse the response into labeled JSON fields for the user to copy or use elsewhere, leaving the source document unmodified.

WHAT STRONG LOOKS LIKE

A reviewer uploads a sample invoice, confirms the returned fields (vendor, amount, date, invoice number) are each correctly labeled and correct, and confirms the original uploaded document is unchanged afterward.

BAND S2 · search, API · a pull request into the public repository, use-cases folder

Clip-and-compose workspace

WHO IT SERVES

Research/Content / a researcher, analyst, or writer who gathers material from many source documents over time and assembles the best pieces into one deliverable.

WHAT TO BUILD

Researchers and writers gathering material from many source documents over weeks end up with scattered notes and no single place to pull the best pieces together. Build a workspace where a user saves snippets from any uploaded document over time, then, when ready, assembles the saved snippets into one new, styled document, reordering and editing as needed before export.

WHAT STRONG LOOKS LIKE

A reviewer clips passages from three different uploaded documents across separate sessions, then composes them into one new document and confirms all three clipped passages appear correctly attributed to source, in the chosen order.

BAND S2 · Multi-document, memory, export · a pull request into the public repository, use-cases folder

"Search my whole library" app

WHO IT SERVES

Cross-industry / anyone with a large personal or team library of uploaded documents who needs to find the right passage fast and fix it on the spot.

WHAT TO BUILD

Someone with dozens of uploaded documents can't remember which one contains the passage they need, and opening files one at a time to check is slow. Build a search-first app: semantic and full-text search across every document a user has uploaded, jump straight to the matching passage in its source document, and edit it in place from the search result. This starts from a search across the whole library, not from editing a single document with related material pulled in behind it.

WHAT STRONG LOOKS LIKE

A reviewer uploads five unrelated documents, searches for a phrase that only appears in one of them, and confirms the result jumps directly to that passage in the correct document, ready to edit.

BAND S2 · search, Multi-document · a pull request into the public repository, use-cases folder

Searchable-PDF maker

WHO IT SERVES

Cross-industry / someone with an old scanned PDF who wants to search and select its text, not necessarily edit its content.

WHAT TO BUILD

An old scanned PDF, a book chapter, a form, an archived record, has no selectable text, so searching or copying from it is impossible. Build a named entry-point app: upload a scanned PDF, run it through OCR ingest, and export it back out as a PDF with real, selectable text underneath, searchable and copyable, not a repair or heavy-edit tool.

WHAT STRONG LOOKS LIKE

A reviewer uploads a scanned, image-only PDF, exports it through the app, and confirms they can select and copy text directly from the output PDF in a standard PDF viewer.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

Approval-workflow document app

WHO IT SERVES

Enterprise/HR / an operations or HR team that needs a document to move through several named approval stages before it's finalized.

WHAT TO BUILD

Teams that need a document to pass through several named approval stages, drafter, manager, compliance, before it's final currently chase that routing over email. Build a workflow app: start from a preapproved template, route the draft through defined review stages with human approval at each one, and export the final version once every stage has signed off. This manages the approval lifecycle of one document at a time, it is not a document repository.

WHAT STRONG LOOKS LIKE

A reviewer pushes a drafted document through three named stages, approving at each, and confirms the final export only becomes available after all three approvals are recorded, not before.

BAND S2 · templates, Review, export · a pull request into the public repository, use-cases folder

Find-and-replace with judgment

WHO IT SERVES

Cross-industry / anyone maintaining a long document, a policy, a contract template, a handbook, who needs a term or fact updated everywhere it appears in context.

WHAT TO BUILD

A literal find-and-replace misses every variant of a term or fact that isn't spelled exactly the same way, and doing it by hand across a long document is slow and error-prone. Build a chat-driven replace tool: describe what needs to change in plain language, the app finds every place that instance appears in context across the whole document, proposes each replacement, and applies only the ones a human approves.

WHAT STRONG LOOKS LIKE

A reviewer asks the tool to update a person's title everywhere it's mentioned in a long document, including sentences that phrase it differently, and confirms every genuine instance is proposed for review, with no unrelated text touched.

BAND S2 · chat, Review · a pull request into the public repository, use-cases folder

Handwriting / notes digitizer

WHO IT SERVES

Education/Field-ops / a student, teacher, or field worker who captures handwritten notes on paper and wants them typed up without retyping by hand.

WHAT TO BUILD

Handwritten notes taken on paper, in a classroom or out in the field, stay stuck on paper unless someone retypes them. Build an app that takes a photo or scan of handwritten notes and turns it into a typed, structured, editable document through the same OCR ingest path used for printed scans. Results depend on how legible the handwriting is, this is not a guarantee of perfect transcription.

WHAT STRONG LOOKS LIKE

A reviewer photographs a page of clearly written notes, uploads it, and confirms the resulting document is typed, organized into readable sections, and editable, not a raw wall of unstructured text.

BAND S2 · images, export · a pull request into the public repository, use-cases folder

Document classifier + router

WHO IT SERVES

Ops/Back-office / an operations team that receives a steady stream of mixed document types and currently sorts them by hand.

WHAT TO BUILD

A back-office team that receives a steady stream of mixed document types, invoices, forms, contracts, sorts them by hand before deciding what to do with each one. Build an app that classifies an uploaded document's type through a chat classification step, then automatically kicks off the right next flow for that type, generate, edit, or extract, instead of routing everything manually.

WHAT STRONG LOOKS LIKE

A reviewer uploads three differently-typed documents in one batch and confirms each is classified correctly and routed to a distinct matching flow, not all three landing on the same default path.

BAND S2 · API, search · a pull request into the public repository, use-cases folder

"Repair my broken Word doc" tool

WHO IT SERVES

Cross-industry / someone whose Word document has become corrupted or structurally broken and needs a usable file back.

WHAT TO BUILD

A Word document that's become corrupted or structurally broken, won't open cleanly, formatting has collapsed, leaves its owner with no easy way back to a usable file. Build a consumer-facing repair tool: re-ingest the damaged file, normalize as much of its structure as can be recovered, and export a clean, styled file. The tool aims for graceful, best-effort recovery and says so plainly, it never promises to fully repair every corrupted file, and it should tell the user what was recovered. This is an end-user product, not an engineering QA harness.

WHAT STRONG LOOKS LIKE

A reviewer runs a deliberately messy or partially broken DOCX through the tool and confirms the output is a valid, styled file with a clear summary of what was recovered, with no silent claim of a full or perfect fix.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

Technical-service-bulletin (TSB) authoring app

WHO IT SERVES

Automotive OEM and dealer-network technical writers maintaining service bulletins.

WHAT TO BUILD

Automotive service teams publish technical service bulletins that combine step-by-step repair text with diagrams, and keeping bulletins current as procedures change means constant rewriting. Build a chat-driven authoring app for drafting and updating TSBs, with inline diagram generation or insertion alongside the procedural text.

WHAT STRONG LOOKS LIKE

Reviewer asks for a new bulletin covering one repair procedure with one diagram, then asks for a single-step revision, and the export reflects both the diagram and the revised step correctly.

BAND S2 · chat, images, export · a pull request into the public repository, use-cases folder

Supplier-quality doc drafter (PPAP/8D/FMEA)

WHO IT SERVES

Supplier-quality engineer (automotive/manufacturing)

WHAT TO BUILD

Build the quality-document drafter whose centerpiece is the FMEA nobody else covers: draft failure-mode tables (function, failure mode, effect, severity, occurrence, detection, actions) plus the narrative sections of PPAP and 8D submissions, on the customer's own template, from the engineer's structured inputs. Numbers the engineer did not supply are asked for, never invented.

WHAT STRONG LOOKS LIKE

the severity-occurrence-detection table is arithmetically consistent (priority values derive from their factors) and every action row traces to a named failure mode; the same inputs re-drafted onto a second customer template change presentation only.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Telecom SLA / service-agreement generator

WHO IT SERVES

Telecom sales operations and legal teams generating SLAs at deal-closing volume.

WHAT TO BUILD

Telecom account teams close deals with negotiated terms that then need to become a formal SLA or service agreement, and manually drafting that document for every deal is slow at volume. Build an API-driven generator that takes structured deal terms and produces a complete SLA or service-agreement document from a template.

WHAT STRONG LOOKS LIKE

Reviewer posts one deal-terms payload (uptime target, response times, term length) and gets back one document where every field appears correctly in the right clause, no manual editing needed.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Method-of-procedure (MOP) doc builder

WHO IT SERVES

Telecom network operations teams authoring change-management runbooks.

WHAT TO BUILD

Network-change work in telecom runs on method-of-procedure documents that combine ordered steps with topology diagrams, and authoring a new MOP from scratch for each change is slow. Build a chat-driven builder that drafts network-change MOPs and runbooks, including inline diagrams of the change topology or rollback path.

WHAT STRONG LOOKS LIKE

Reviewer describes one network change and gets a MOP with numbered steps, a rollback section, and at least one diagram showing the change, exportable as one document.

BAND S2 · chat, images, export · a pull request into the public repository, use-cases folder

Game design document (GDD) editor

WHO IT SERVES

Gaming

WHAT TO BUILD

Game teams maintain long, living design documents that keep growing across systems, mechanics, and art direction, and a plain word processor struggles once a GDD splits across linked sections. Build a chat-driven editor for authoring and iterating long GDDs, with image support for concept art or diagrams and multi-document support for splitting a GDD into linked sub-documents such as mechanics, narrative, and art bible.

WHAT STRONG LOOKS LIKE

Reviewer asks for a new GDD section plus one linked sub-document, then requests a revision to one mechanic, and only the correct section changes across both documents.

BAND S2 · chat, images, Multi-document, export · a pull request into the public repository, use-cases folder

Branded menu builder

WHO IT SERVES

Restaurant or cafe owner with a dish list and no designer

WHAT TO BUILD

Build the create-from-zero menu maker: from a plain dish list with prices and descriptions, generate a branded, print-ready menu document - layout, sections, typography, and generated dish imagery where the owner has none - with the brand's colors and voice applied throughout. Existing-menu amendment tools already exist; this one starts from nothing.

WHAT STRONG LOOKS LIKE

the same dish list rendered under two different brand settings changes presentation completely and content not at all; prices survive every layout change byte-exact; the export prints clean on one pass.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

HACCP / food-safety-plan drafter

WHO IT SERVES

Food & Bev

WHAT TO BUILD

Food and beverage operators must maintain a HACCP-methodology food-safety plan and keep it versioned as processes change, and building one from a blank page is a specialist task. Build a chat-driven drafter that walks a HACCP plan onto a structured template (hazard analysis, critical control points, monitoring, corrective actions) and versions it as the process evolves. The output is a draft ready for professional review, never a certified or compliance-verified plan.

WHAT STRONG LOOKS LIKE

Reviewer describes one process step with one hazard, and the draft correctly places it under the matching HACCP section (hazard analysis vs. critical control point) with a clear "for professional review" note retained in the export.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Gym waiver & membership-agreement generator

WHO IT SERVES

Fitness businesses and franchise back-offices generating new-member paperwork at volume.

WHAT TO BUILD

Fitness businesses sign up members at volume and each signup needs a fresh waiver and membership agreement generated from the same template, ready to route to signature. Build an API-driven generator that takes member and plan details and produces a signable-ready waiver and membership agreement from a template, leaving the actual signature step to an external e-sign provider.

WHAT STRONG LOOKS LIKE

Reviewer posts one member and plan payload and gets back a waiver and a membership agreement, both correctly populated from the same input, with no signature step implied.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

OneNote → styled doc exporter

WHO IT SERVES

Cross-industry

WHAT TO BUILD

OneNote pages are not linear documents. The Graph API returns each block on a page as its own absolutely positioned div, image, or object (data-absolute-enabled plus inline top/left coordinates), because the page is really a freeform canvas where anyone can drop a text box or image anywhere on it. A naive HTML-to-document conversion scrambles reading order the moment two boxes sit side by side or stack out of sequence. Build a pipeline that pulls a notebook's pages through the OneNote API, re-linearizes each page by clustering blocks into reading-order rows from their coordinates, then hands the normalized result to SuperDocs to produce a branded, exportable document.

WHAT STRONG LOOKS LIKE

Feed it a real page with two side-by-side text boxes and an inline image; the exported document reads top-to-bottom in a sensible order, not the raw element order from the API response, and opens as a clean docx or pdf with headings and paragraphs intact.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

Shopify product-doc generator

WHO IT SERVES

Retail

WHAT TO BUILD

A retailer's spec sheet, care instructions, and warranty terms for a product usually live scattered across the product description and its metafields, then get rebuilt by hand for every SKU. Shopify's GraphQL Admin API, now the standard path since the REST Admin API moved to legacy status, exposes a product's core fields plus its metafields, the structured extra data merchants attach for exactly this kind of spec, size, and warranty content. Build a generator that queries a product or a whole collection through the GraphQL Admin API, maps its fields and metafields onto a spec/care/warranty template set, and batch-produces one branded document per product through SuperDocs' templates and export surfaces.

WHAT STRONG LOOKS LIKE

Point it at a collection of ten or more products; it produces one correctly filled spec/care/warranty document per product from that product's own metafield values, and a reviewer can pick any two products and see the content genuinely differ, not a single template with the name swapped.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Google Classroom feedback add-on

WHO IT SERVES

EdTech

WHAT TO BUILD

A teacher grading a student's uploaded document today has to leave Classroom, open the file, mark it up separately, and reattach it by hand. Google Classroom's API lets a tool read a student submission's attached Drive file, use `studentSubmissions.modifyAttachments` to attach a new file back onto that same submission, then `studentSubmissions.return` to hand it back, which also transfers ownership of the attached file to the student. Build an add-on that pulls a submission's document into SuperDocs, runs it through a track-changes review pass with a human-in-the-loop approval step, and pushes the marked-up file back onto the submission through the Classroom API, so feedback never leaves the assignment.

WHAT STRONG LOOKS LIKE

Take one real submitted document through the full loop, track-changes suggestions approved through the HITL step, then confirm in Classroom that the returned submission actually carries the marked-up file back to the student, not a copy left orphaned in the tool.

BAND S2 · API, Review, export · a pull request into the public repository, use-cases folder

Canvas/Moodle LMS doc-assignment tool

WHO IT SERVES

EdTech

WHAT TO BUILD

Canvas and Moodle don't share Google Classroom's REST API at all; a tool that wants to live inside either one has to speak LTI 1.3, the launch protocol both platforms implement instead. That means an OIDC third-party-initiated login handshake and a signed JWT launch request just to open the tool, Deep Linking so an instructor can place a "review this document" assignment into the course, and the Assignment and Grade Services endpoint to post a completion grade back to the gradebook once review is done. Build an LTI 1.3 tool that launches into a SuperDocs track-changes review flow for an assigned document and posts the grade back through AGS on completion, an entirely different integration shape from a Classroom REST client even though the student-facing feature looks similar.

WHAT STRONG LOOKS LIKE

The tool completes a real LTI 1.3 launch where the JWT signature verifies against the platform's public keys rather than just opening a link, places itself as an actual assignment through Deep Linking, and a completed review posts a real score back through AGS that shows up in the gradebook.

BAND S2 · API, Review, export · a pull request into the public repository, use-cases folder

Microsoft Loop component exporter

WHO IT SERVES

Cross-industry

WHAT TO BUILD

Microsoft Loop has no dedicated export endpoint; there is no "export workspace" call and no documented API that hands back a finished page. Loop content actually lives in SharePoint Embedded containers as Fluid-Framework-backed component data, so the real path in is Microsoft Graph's file access into that container, followed by parsing each block's component payload (paragraph, table, task list) out of that JSON by hand. Build an exporter that pulls a Loop page through Graph's file endpoints, decodes its component blocks into normal document structure, and hands the result to SuperDocs to produce a finalized, styled document, a materially different mechanic from reading a normal SharePoint file or a Copilot response.

WHAT STRONG LOOKS LIKE

Given a real Loop page with at least three different component types (say a paragraph, a table, and a task list), the exported document renders all three correctly structured, and a reviewer can confirm the exporter is reading actual component data, not just the page's plain-text fallback.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

ServiceNow doc-automation app

WHO IT SERVES

Enterprise/ITSM

WHAT TO BUILD

IT teams running ServiceNow generate change records, policy updates, and knowledge articles that all start as raw ServiceNow ticket data and end up rewritten by hand into readable documents. Build an API-driven app that pulls a ServiceNow change request, problem record, or knowledge-base draft, using fields like short description, affected configuration item, work notes, and resolution code, and turns it into a formatted change doc, policy update, or knowledge article.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional ServiceNow change-request payload with a configuration-item field and work notes, and gets back one formatted document where those exact fields appear in the right sections, not paraphrased away.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Clio legal-matter document app

WHO IT SERVES

Legal

WHAT TO BUILD

Law firms running Clio track matters, contacts, and tasks in the practice-management system, but the letters and documents tied to a matter still get drafted separately by hand. Build an API-driven app that pulls matter data from Clio, such as matter name, client or contact, responsible attorney, and court dates, and drafts the matching matter letters and documents from a template.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional Clio matter payload with a client contact and a court date, and gets back a letter where those exact fields are correctly placed and addressed to the right contact.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Guidewire/Duck Creek policy-doc app

WHO IT SERVES

Insurance carrier policy administration and forms teams generating documents straight from PAS data.

WHAT TO BUILD

Insurance carriers running Guidewire or Duck Creek generate policy and endorsement documents from data that lives in the policy administration system, and today that means exporting data and formatting documents by hand. Build an API-driven app that pulls policy data from a Guidewire or Duck Creek PAS record, such as policy number, named insured, coverage lines, and endorsement effective date, and drafts the matching policy or endorsement document from a wording template.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional PAS policy payload with two coverage lines and an endorsement date, and the generated document shows both coverage lines and the correct effective date in the right clauses.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Procore/Autodesk-Construction RFI app

WHO IT SERVES

Construction

WHAT TO BUILD

project teams log RFIs, submittals, and change orders inside Procore or Autodesk Construction Cloud, but the formal document version of each still gets drafted separately. Build an API-driven app that pulls project data, such as RFI number, spec section, drawing reference, submittal item, and change-order cost or schedule impact, from Procore or Autodesk Construction Cloud and drafts the matching RFI, submittal, or change-order document.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional RFI payload with a spec section and a drawing reference, and the generated document carries both references correctly into the formal RFI format.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Yardi/AppFolio lease-doc app

WHO IT SERVES

Real estate

WHAT TO BUILD

Property managers running Yardi or AppFolio hold unit, tenant, and lease-term data in the property-management system, but new leases and tenant notices still get typed up separately. Build an API-driven app that pulls property-management data, such as unit number, tenant name, lease term, and rent schedule, from Yardi or AppFolio and drafts leases and tenant notices from a template.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional unit, tenant, and lease-term payload and gets back a lease document with the correct unit, tenant, and rent schedule in the right clauses, plus a matching notice document from the same data.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Amazon/marketplace listing-doc pipeline

WHO IT SERVES

Retail and marketplace operations teams generating listing documentation at catalog scale.

WHAT TO BUILD

Retailers selling at scale on Amazon or other marketplaces need listing and spec documents generated in bulk from marketplace catalog data, not written one at a time. Build an API-driven pipeline that pulls marketplace listing data, such as ASIN or SKU, title, bullet points, images, and category attributes, and generates a batch of formatted listing-spec documents with images placed, exported as docx or html, never a spreadsheet format.

WHAT STRONG LOOKS LIKE

Reviewer supplies a small batch of fictional marketplace listing records with images, and gets back one formatted document per listing with the correct title, bullets, and image for that listing, none crossed between listings.

BAND S3 · API, images, templates, export · a pull request into the public repository, use-cases folder

Calendly/Cal.com follow-up-doc automation

WHO IT SERVES

Sales/Ops

WHAT TO BUILD

Sales and ops teams want a recap or proposal ready right after a booked call ends, triggered by the calendar event itself rather than typed up afterward from memory. Build an API-driven automation that listens for a Calendly or Cal.com booking-completed event, such as attendee, meeting type, and scheduled time, and generates a follow-up recap or proposal shell from a template, with no meeting-content input required.

WHAT STRONG LOOKS LIKE

Reviewer fires one fictional booking-completed event and gets back one recap document correctly addressed to the attendee and referencing the right meeting type, generated from the event alone.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

Twilio SMS "text → document" bridge

WHO IT SERVES

Global-south field teams and small operators who need document help over basic SMS, no smartphone data required.

WHAT TO BUILD

In low-smartphone-data regions, a plain text message reaches people that a web app or app-based chat session cannot. Build an API-driven bridge that reads Twilio's inbound SMS webhook fields (From, Body, MessageSid) to pull the sender's instruction and document reference straight off the message body, runs the requested edit or generation, and texts back a link to the resulting file via Twilio's outbound send, no app or data session required.

WHAT STRONG LOOKS LIKE

Reviewer texts one fictional instruction plus a document reference to a test number and receives one SMS reply containing a working link to the correctly edited document, with the inbound From/Body pairing traceable in the logs.

BAND S2 · API, chat, export · a pull request into the public repository, use-cases folder

MCP pipeline: email thread → contract → e-sign**WHO IT SERVES**

Sales/Legal

WHAT TO BUILD

A full deal-to-signature flow needs three separate capabilities chained together: reading an informal email thread, drafting the contract it implies, and routing the result to signature, and no single tool does all three today. Build an MCP pipeline that calls an external parse MCP to extract terms from a fictional negotiated email thread, drafts the contract from those terms using our templates, and hands the finished document to an external e-sign MCP, never performing parsing or signing itself.

WHAT STRONG LOOKS LIKE

Reviewer runs one fictional thread through the full pipeline and can point to three distinct calls in the trace: parse, our own draft, and external e-sign handoff, each with its own artifact.

BAND S3 · MCP, templates, export · a pull request into the public repository, use-cases folder

MCP pipeline: spec → diagram-rich runbook**WHO IT SERVES**

DevTools

WHAT TO BUILD

Engineering specs describe a system in prose, but the operational runbook someone eventually needs has to translate that prose into ordered steps with diagrams, and that translation is usually done by hand. Build a pipeline that takes a technical spec as input and produces an operational runbook with auto-generated Mermaid diagrams illustrating the described flow, using our confirmed diagram-rendering capability end to end.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional spec describing a three-step process and gets back a runbook with three ordered steps and at least one diagram that correctly reflects the process flow.

BAND S3 · MCP, images, export · a pull request into the public repository, use-cases folder

Agent evaluation harness for doc edits

WHO IT SERVES

DevTools/Agents

WHAT TO BUILD

Teams building their own document-editing agents need a repeatable way to check whether a new agent version is actually getting better, not just different. Build a golden-document test harness that runs a third-party agent's proposed edits through our diff and HITL approval surfaces and records what changed, leaving the actual scoring rubric as wrapper-side logic outside SuperDocs itself.

WHAT STRONG LOOKS LIKE

Reviewer runs one golden document through two different edit sets and gets back two diff records plus HITL approval outcomes that can be compared side by side, no manual diffing required.

BAND S3 · API, Review · a pull request into the public repository, use-cases folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Patent-drafting assistant

WHO IT SERVES

Legal/IP

WHAT TO BUILD

Patent claims are numbered and cross-reference each other by dependency, so editing one claim risks silently breaking every claim that depends on it. Build a chat-driven assistant for drafting and amending numbered patent claims that tracks dependency cross-references and keeps them consistent through an edit, producing a draft for attorney review rather than a filed or certified document.

WHAT STRONG LOOKS LIKE

Reviewer amends one independent claim's scope and the assistant flags or updates every dependent claim that references it, with none silently left inconsistent.

BAND S1 · chat, export · a pull request into the public repository, use-cases folder

Deposition/brief summarizer + editor

WHO IT SERVES

Legal

WHAT TO BUILD

Litigation teams read through long deposition transcripts and discovery material to find the passages a brief needs, then draft the brief section around them, and losing the exact source citation along the way is a real risk. Build a search-driven tool that summarizes discovery material, lets the user pull a specific passage into a brief section draft, and keeps the source citation attached through the edit, producing an attorney-review-ready draft, never a filed brief.

WHAT STRONG LOOKS LIKE

Reviewer searches a fictional transcript for one topic, pulls the matching passage into a draft brief section, and the exported section still carries a citation back to that exact passage.

BAND S2 · search, chat, export · a pull request into the public repository, use-cases folder

10-K / filing summarizer to exec brief

WHO IT SERVES

Finance and investor-relations teams turning long filings into executive-ready briefs.

WHAT TO BUILD

Executives do not have time to read a full regulatory filing, but they need the substance, not a generic summary. Build a search-driven tool that takes a long filing, surfaces the sections that matter, and produces a short styled executive brief that preserves the key figures and their source section.

WHAT STRONG LOOKS LIKE

Reviewer feeds one fictional filing excerpt with three key figures and the resulting brief contains all three figures correctly, each traceable back to its source section.

BAND S1 · search, export · a pull request into the public repository, use-cases folder

KYC/AML onboarding-packet assembler

WHO IT SERVES

Fintech compliance and onboarding teams assembling customer-facing KYC/AML bundles.

WHAT TO BUILD

Fintechs onboarding a new customer need to assemble a full KYC/AML documentation bundle, and doing that by hand for every customer does not scale. Build an API-driven assembler that takes customer-provided onboarding data and compiles it into a complete, consistently formatted KYC/AML documentation bundle across multiple linked documents, producing a compliance-review-ready bundle: it assembles the paperwork only and never performs or claims identity verification itself.

WHAT STRONG LOOKS LIKE

Reviewer posts one fictional customer onboarding payload and gets back a multi-document bundle where every document reflects the same customer facts, with no identity-verification claim anywhere in the output.

BAND S2 · API, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Claims-correspondence generator

WHO IT SERVES

Insurance claims teams drafting coverage and decision correspondence under review.

WHAT TO BUILD

Insurance claims correspondence has to reflect real judgment about the claim facts, not just fill in a template, so a coverage or decision letter needs a human to review the reasoning before it goes out. Build an API-driven generator that drafts coverage and decision letters from claim data, routes the draft through HITL approval before anything is finalized, and never auto-sends a decision without review.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional claim payload with a coverage question, and the drafted letter reflects the correct reasoning path, sitting in HITL approval until explicitly approved.

BAND S2 · API, Review, export · a pull request into the public repository, use-cases folder

Underwriting-file assembler

WHO IT SERVES

Insurance underwriters maintaining an ongoing internal file per risk.

WHAT TO BUILD

Underwriters build up a file for each risk over time, adding documents as new information comes in, and keeping that file complete and current is ongoing work, not a one-time packet. Build a multi-document tool that compiles and updates an underwriter's own working file as new information arrives, distinct from a one-time broker submission packet.

WHAT STRONG LOOKS LIKE

Reviewer adds one new piece of information to an existing fictional file and only the affected document updates, with the rest of the file untouched and still exportable as a set.

BAND S2 · Multi-document, templates, export · a pull request into the public repository, use-cases folder

Payer appeals-letter app

WHO IT SERVES

Healthcare revenue-cycle and clinical-documentation teams drafting payer appeals.

WHAT TO BUILD

When a health insurer denies a claim, someone on the provider side has to draft an appeal letter that combines the denial reason with the relevant clinical facts, and getting that pairing right by hand is slow. Build an API-driven app that drafts insurance appeal letters from denial data and clinical notes, routed through HITL approval before anything is finalized.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional denial reason plus one clinical fact, and the drafted appeal letter correctly addresses that denial reason using that fact, held for HITL approval before export.

BAND S2 · API, Review, export · a pull request into the public repository, use-cases folder

Clinical-protocol / SOP editor

WHO IT SERVES

Pharma/Healthcare

WHAT TO BUILD

Pharma and healthcare teams maintain clinical protocols and SOPs that change often and need every change reviewed before it takes effect. Build a chat-driven editor for protocol and SOP content that routes every edit through HITL approval, keeping the document itself as the single authored source rather than a separate approval log.

WHAT STRONG LOOKS LIKE

Reviewer asks for one procedural change to a fictional SOP and the change sits in HITL approval, only entering the document once explicitly approved, with the prior version still recoverable.

BAND S2 · chat, Review, export · a pull request into the public repository, use-cases folder

ISO 9001 quality-manual manager

WHO IT SERVES

Manufacturing quality teams authoring and maintaining an ISO 9001 quality manual for professional review.

WHAT TO BUILD

Manufacturers maintaining ISO 9001 certification keep a quality manual that has to stay current with process diagrams as procedures change, and manually keeping text and diagrams in sync is tedious. Build a chat-driven manager for authoring and versioning the quality manual, including inline process diagrams, producing a manual ready for professional review rather than a certified one.

WHAT STRONG LOOKS LIKE

Reviewer asks for one process-section update with a diagram change, and the exported manual shows both the updated text and the updated diagram consistently, with prior versions still recoverable.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Dangerous-goods / shipping-declaration app

WHO IT SERVES

Logistics and freight-forwarding teams generating dangerous-goods shipping paperwork for professional review.

WHAT TO BUILD

Shipping dangerous goods requires a declaration document and matching shipping instructions built from the item's hazard classification, and getting the classification-to-document mapping right by hand is a specialist task. Build an API-driven app that takes shipment and hazard-classification data and generates the dangerous-goods declaration and shipping instructions from a template, producing a draft ready for professional review, not a certified regulatory filing.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional shipment with a hazard class and UN number, and the generated declaration shows both correctly in the right fields, with a clear review note retained in the export.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Land/lease agreement editor

WHO IT SERVES

Energy/Oil & Gas

WHAT TO BUILD

Oil and gas land teams edit lease and mineral-rights agreements where a single misplaced clause can change who owns what, so formatting fidelity through an edit matters as much as the wording. Build a chat-driven editor for land and lease agreements that preserves document structure and clause numbering exactly through every edit.

WHAT STRONG LOOKS LIKE

Reviewer edits one clause in a fictional lease agreement and every other clause, including numbering and cross-references, is byte-identical in the export.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Script/screenplay formatter

WHO IT SERVES

Media

WHAT TO BUILD

Screenplays follow a strict industry format (scene headings, action lines, character cues, dialogue) that a general word processor does not enforce, so writers spend real time hand-formatting instead of writing. Build a chat-driven formatter that takes a script and reformats it to industry screenplay standard in place, preserving the actual writing while fixing the format.

WHAT STRONG LOOKS LIKE

Reviewer pastes one loosely formatted scene and the export shows correct scene heading, action, and dialogue formatting with none of the original wording changed.

BAND S2 · chat, export · a pull request into the public repository, use-cases folder

Rights/talent-contract generator

WHO IT SERVES

Media/Entertainment

WHAT TO BUILD

Media and entertainment deals need rights and talent agreements drafted from a template for every engagement, and doing that from scratch each time is slower than working off a consistent base. Build a chat-driven generator that drafts rights and talent agreements from a template given the specific deal terms.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional talent deal's terms and gets back one agreement where the term, scope of rights, and compensation are correctly reflected in the right clauses.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

API-reference & guide authoring app

WHO IT SERVES

DevTools

WHAT TO BUILD

Developer-facing teams write API reference docs and getting-started guides that both need to move between plain markdown and a styled, exportable document without losing code blocks or diagrams. Build a chat-driven authoring app with two modes: reference-and-guide authoring that converts between markdown and styled docs while preserving code blocks and diagrams, and a getting-started-guide mode that assembles a new-user onboarding guide from a described API or feature set.

WHAT STRONG LOOKS LIKE

Reviewer pastes one markdown page with a code block and a diagram reference, converts it to styled output and back, and both the code block and the diagram survive the round trip unchanged; the getting-started mode produces a numbered onboarding guide from one feature description.

BAND S2 · chat, images, export · a pull request into the public repository, use-cases folder

RFC/ADR/design-doc workspace

WHO IT SERVES

DevTools

WHAT TO BUILD

Engineering teams draft RFCs and architecture-decision records that need real human review before anything is decided, and pasting draft after draft into chat loses the review thread. Build a chat-driven workspace for drafting design docs where reviewers leave HITL comments directly against the draft and the author resolves them in place, keeping draft and review history together.

WHAT STRONG LOOKS LIKE

Reviewer leaves one comment on a fictional draft section, the author responds with an edit addressing it, and the exported document plus its comment history both reflect the resolution correctly.

BAND S2 · chat, Review, export · a pull request into the public repository, use-cases folder

Vendor-onboarding packet assembler

WHO IT SERVES

Procurement teams assembling new-vendor onboarding and compliance bundles.

WHAT TO BUILD

Procurement teams onboarding a new vendor need a bundle of forms and compliance documents assembled together, and chasing that bundle together by hand across many vendors is repetitive work. Build a multi-document assembler that compiles a vendor's onboarding data into a complete, consistently formatted onboarding and compliance bundle.

WHAT STRONG LOOKS LIKE

Reviewer posts one fictional vendor's onboarding data and gets back a multi-document bundle where every document reflects the same vendor facts, with no document left half-populated.

BAND S2 · Multi-document, templates, export · a pull request into the public repository, use-cases folder

Capability-statement / past-performance app

WHO IT SERVES

GovCon

WHAT TO BUILD

Government contractors need a standardized capability statement that follows a strict one-page, format-compliant layout, and building that layout correctly by hand is fiddly. Build a chat-driven app that assembles a format-compliant capability statement and past-performance summary from a company's core facts, distinct from a solicitation-triggered proposal shell.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional company's core facts and past-performance record and gets back a single-page, correctly sectioned capability statement with nothing overflowing the format.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Architecture spec + design-narrative app

WHO IT SERVES

Architecture firms authoring linked technical specs and design narratives for a project.

WHAT TO BUILD

Architecture firms present a project through both a technical spec and a design narrative with embedded renders, and keeping those two documents aligned as a design evolves is real work. Build a chat-driven app for authoring architecture specs and design narratives with embedded renders or diagrams, using multi-document support to keep the spec and the narrative linked as separate but consistent documents.

WHAT STRONG LOOKS LIKE

Reviewer requests one design change and both the spec and the narrative document reflect the same updated detail, with the correct render or diagram attached to the narrative.

BAND S2 · chat, images, Multi-document, export · a pull request into the public repository, use-cases folder

Split-a-document tool

WHO IT SERVES

Legal/Ops

WHAT TO BUILD

A long document sometimes needs to become several clean standalone files, such as splitting a master agreement into individually exportable exhibits, and there is no dedicated split feature to do that in one step. Build a multi-document utility that takes a long document plus a set of split points and produces one clean, independently exported document per section using our existing multi-document creation primitive.

WHAT STRONG LOOKS LIKE

Reviewer marks three split points in one fictional long document and gets back three separate exported documents, each starting and ending exactly at the marked boundaries with nothing duplicated or dropped.

BAND S2 · Multi-document, export · a pull request into the public repository, use-cases folder

Accessibility-tagged export tool

WHO IT SERVES

Government/Enterprise

WHAT TO BUILD

Government and enterprise teams need documents restructured for accessibility, and exporting a genuinely tagged, PDF/UA-compliant file is not something we can honestly claim today. Build an API-driven tool that restructures a document for accessibility (heading order, alt text, reading order, list and table semantics) and exports it as accessibility-structured HTML or DOCX plus a plain-language accessibility report, never claiming tagged-PDF or PDF/UA output.

WHAT STRONG LOOKS LIKE

Reviewer feeds one fictional document with a heading-order problem and a missing alt text, and the export fixes both, with the accompanying report listing exactly what was checked and fixed, no PDF/UA claim anywhere in the output.

BAND S2 · API, export · a pull request into the public repository, use-cases folder

Page-organizer / reorder-and-export tool

WHO IT SERVES

Anyone restructuring a long document without wanting to touch its prose

WHAT TO BUILD

Build the interactive organizer: load a long document, present its sections as a draggable outline, reorder and remove at the section level, then export a clean result. The prose of every kept section is provably untouched; the only thing that changes is arrangement.

WHAT STRONG LOOKS LIKE

a reorder-only session produces kept sections byte-identical to their originals; removals never orphan cross-references silently (they are listed at export); the second run on the exported file shows a stable outline.

BAND S2 · Multi-document, export · a pull request into the public repository, use-cases folder

MCP tool-explorer web app

WHO IT SERVES

Agents. A DevTools engineer comfortable reading JSON schemas and generating a form from one.

WHAT TO BUILD

There is no live, browsable way to see what the SuperDocs MCP server can do without first wiring up a local MCP client. Build a hosted web app that lists the server's MCP tools, renders each tool's real input schema as a form, and lets a visitor actually invoke it against a live session, showing the real request and the real response. This is distinct from a terminal REPL, a downloadable sandbox, or a static cookbook: it runs in a browser with zero local setup and covers the full current tool surface rather than a curated subset.

WHAT STRONG LOOKS LIKE

Picking any tool, filling its generated form, and invoking it returns a real response from a live call rather than a canned example, and the tool list on the page is read from the server at load time rather than hardcoded.

BAND S3 · MCP · a pull request into the public repository, use-cases folder

llms-full.txt + AGENTS.md generator for our docs

WHO IT SERVES

Agents. A DevTools engineer or technical writer who cares about machine-readable discovery, comfortable with docs crawling and parsing.

WHAT TO BUILD

Agents increasingly discover how to use a product by fetching a couple of predictable files instead of crawling an entire docs site: an llms-full.txt with the full content in one place, and an AGENTS.md with agent-facing setup and usage notes. Build a generic generator, not tied to any single product, that takes a public documentation site's URL and produces both artifacts. Demonstrate it against at least two different real public docs sites, so the tool is provably generic rather than a one-off script written against a single site.

WHAT STRONG LOOKS LIKE

Run against a docs site nobody tested during development, the generator still produces a coherent llms-full.txt and AGENTS.md on the first pass, usable enough that an agent could act on them without a human editing pass first.

BAND S3 · MCP · a pull request into the public repository, use-cases folder

"Build a Contract AI in 45 minutes" starter app

WHO IT SERVES

DevTools. Someone who can build a real app and narrate it accurately against a clock, not just write reference docs.

WHAT TO BUILD

New builders convert fastest when they watch something real get built end to end in one sitting, on a clock, rather than reading reference docs alone. Build a timed, narrated starter kit: a cloneable repo paired with a step-by-step live-coding companion that takes a builder from an empty folder to a working vertical AI document app on the public API. Ship three variants under one kit: a Contract AI in forty-five minutes, an Invoice AI in thirty minutes, and a Compliance AI for SOC2 and ISO policy documents in fifty minutes, each with its own repo and script. This targets a named vertical outcome end to end on the raw API, distinct from a general tech-stack starter, a recipe cookbook of unrelated snippets, or an editor-integration package.

WHAT STRONG LOOKS LIKE

A builder with no prior context follows the companion script against a fresh clone and reaches the working demo within the stated time, for each of the three variants, timed for real rather than rounded down.

BAND S2 · API, templates · a pull request into the public repository, use-cases folder

Token-cost calculator ("why re-read a 50-page doc?")

WHO IT SERVES

Positioning. A builder who can state assumptions plainly rather than dressing an estimate up as a measurement.

WHAT TO BUILD

Engineers evaluating an AI document API want a quick gut check on the real cost gap between resending a whole document on every edit and sending only the changed section. Build a small interactive calculator: a visitor enters a document length and an edit frequency, and it shows the estimated token cost of whole-document regeneration next to section-level surgical editing, with its assumptions stated plainly on the page. Label this clearly as an illustrative estimate for marketing use, never as a measured benchmark result.

WHAT STRONG LOOKS LIKE

Every number the calculator shows traces to a stated, visible assumption on the same page, and changing an input updates the estimate live without a page reload.

BAND S2 · chat · a pull request into the public repository, use-cases folder

Build-vs-buy / ROI calculator app

WHO IT SERVES

Positioning. A builder equally comfortable with an honest calculator and a real document export.

WHAT TO BUILD

A prospect deciding whether to build an in-house AI document pipeline or use SuperDocs wants a number, not a call. Build a self-serve TCO calculator: the visitor enters their own document volume, an engineering-hours estimate, and a loaded hourly cost, and the tool computes a build-versus-buy comparison live. Let them download a real branded PDF of their specific result, generated through the actual export API rather than a template screenshot. This is a repeatable, self-serve tool anyone can run with their own numbers, distinct from a one-time published build log of someone else's DIY attempt.

WHAT STRONG LOOKS LIKE

Changing any input recomputes the comparison immediately, the downloaded PDF's numbers exactly match what was on screen when it generated, and the PDF is a real exported document rather than an image standing in for one.

BAND S2 · chat, export · a pull request into the public repository, use-cases folder

Sales battlecard generator

WHO IT SERVES

Sales. A builder who can structure competitive positioning into a consistent, reusable template using placeholder competitors.

WHAT TO BUILD

A rep switching between deals needs a fast, consistent answer to "why not them," not a scattered doc per competitor. Build a battlecard generator: feed it a competitor name and a few structured inputs, their likely pitch, where we win, where we do not, and it produces a consistent, exportable one-page battlecard on a fixed template. Build and demo it entirely with fictional competitor names. This is a live, rep-facing quick reference, distinct from a quarterly retrospective that synthesizes win-loss patterns from past deal transcripts.

WHAT STRONG LOOKS LIKE

Two different fictional competitor inputs produce battlecards that read as genuinely different documents rather than one paragraph with a name swapped, and every card exports cleanly from the same template.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Objection-handling cheatsheet builder

WHO IT SERVES

Sales. A builder who turns messy notes into something scannable in a live call, without inventing content that was not there.

WHAT TO BUILD

A sales playbook full of prose paragraphs is useless mid-call; a rep needs an objection and a crisp answer, scannable in seconds. Build a tool that takes a raw, unstructured set of objection-and-response notes, a fictional example set is fine for the demo, and reformats it into a clean, consistently formatted, exportable cheatsheet: one objection per block, its response underneath. This is a reformatting tool, not a research tool: the objections arrive as input rather than getting generated from public research.

WHAT STRONG LOOKS LIKE

Handed a genuinely messy input, mixed bullet styles, an objection missing its response, the output cheatsheet is uniformly formatted throughout and no objection silently disappears.

BAND S2 · chat, export · a pull request into the public repository, use-cases folder

Competitor comparison-table doc generator

WHO IT SERVES

Positioning. A builder who can lay out a comparison cleanly and keep it factual-reading even with placeholder competitors.

WHAT TO BUILD

A prospect who asks how we are different deserves something they can forward internally, not an answer that evaporates after the call. Build a generator that produces a shareable, branded feature-comparison table: capability rows, product columns, on a template clean enough to send externally. Build and demo it against fictional competitor names only. This artefact is meant to leave the building, distinct from an internal, confidentiality-stripped competitive brief that never gets shared outside the company.

WHAT STRONG LOOKS LIKE

The exported table reads as a polished, externally shareable document on first look rather than an internal spreadsheet screenshot, and adding a third fictional competitor column does not break the layout.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Industry pitch-deck narrative generator

WHO IT SERVES

Sales. A builder who writes genuinely vertical-specific narrative content rather than one script reskinned four times.

WHAT TO BUILD

A generic pitch never lands as well as one that speaks a vertical's own language. Build a generator that produces a pitch narrative for at least three named verticals, for example legal, fintech, healthcare, or edtech, as a written deck-script document: one section per slide-equivalent, each with its heading, its talking point, full speaker notes, and a generated supporting image where a section calls for one. The output is always a real exported document with structure and speaker notes; SuperDocs does not produce presentation files, so nothing here should imply a slide deck comes out the other end.

WHAT STRONG LOOKS LIKE

Reading two vertical variants back to back, the actual talking points and examples differ in substance rather than a swapped industry noun, and the exported document is unmistakably a script-with-notes format, never confusable with an actual slide file.

BAND S2 · chat, templates, images, export · a pull request into the public repository, use-cases folder

AI-document-editing glossary microsite generator

WHO IT SERVES

SEO. A content-minded builder who writes real definitions and wires a real generator rather than hand-authoring pages.

WHAT TO BUILD

Long-tail searches for terms like surgical document editing or chunk-based AI editing currently have nowhere useful to land. Build a generator that takes a list of AI-document-editing terms and produces a public glossary microsite: one clean, styled page per term, each with a real definition, a worked example, and links to related terms. The generator, not a one-off hand-written page set, is the deliverable: adding a term to the input list should produce a new page in the same style with no hand-edited HTML. This is distinct from a translation do-not-translate term list built for localization control: this glossary is public SEO content explaining our own product concepts.

WHAT STRONG LOOKS LIKE

Every entry reads as a genuine, specific explanation a search visitor would find useful rather than a reworded dictionary definition, and adding a new term to the source list regenerates the site with that term included, with no manual page authoring.

BAND S2 · chat, export · a pull request into the public repository, use-cases folder

Brand-book / style-guide document generator

WHO IT SERVES

Marketing. A brand-minded builder who can turn raw brand inputs into something a designer would actually open.

WHAT TO BUILD

A new collaborator without a real style guide will improvise a slightly wrong version of the brand every time. Build a generator that takes a set of brand inputs, colors, typography choices, logo files, a few voice-and-tone notes, and produces a complete, branded style-guide document: swatches, type specimens, logo usage examples, and written voice guidelines, laid out as something a designer would actually want to open. This generator's output is the style guide itself, distinct from a machine-readable brand kit built to audit and fix other documents against a rule set.

WHAT STRONG LOOKS LIKE

The generated document renders the real colors as swatches and the real fonts in the type specimens rather than text labels describing them, and swapping in a different brand's inputs produces a visibly, correctly updated guide.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Quarterly business review (QBR) doc generator

WHO IT SERVES

Customer-success or account manager producing QBR decks-as-documents per account each quarter

WHAT TO BUILD

Build the QBR compiler: from an account's metrics payload (usage, tickets, milestones, spend), generate the quarterly business review document on the team's template - narrative sections written from the numbers, charts generated from the data, risks and asks surfaced explicitly. One command per account, one consistent document family per quarter.

WHAT STRONG LOOKS LIKE

every number in the narrative traces to the payload (no invented metrics); two accounts through the same template differ only where their data differs; a missing metric becomes a visible gap, never a plausible guess.

BAND S2 · API, templates, images, export · a pull request into the public repository, use-cases folder

Trade-show lead → tailored recap batch

WHO IT SERVES

Sales or field-marketing owner following up a conference with dozens of leads

WHAT TO BUILD

Build the post-event recap machine: from a lead list with notes (who, conversation, interest, next step), generate one tailored recap document per lead - what was discussed, what was promised, what happens next - on a consistent branded template, as a batch. The batch is the point: dozens of personalized documents, one run.

WHAT STRONG LOOKS LIKE

no cross-contamination - lead A's details provably never appear in lead B's document; blank note fields degrade to honest generic sections, never invented specifics; the batch of N inputs yields exactly N documents.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Pitch-competition / accelerator application assembler

WHO IT SERVES

Startups. A builder who understands hard per-field limits and treats tightening as compression, never invention.

WHAT TO BUILD

Accelerator and pitch-competition applications are not a template filled from a record; they are a fixed set of open questions with hard word limits, where every answer has to earn its place. Build an application assembler: load a structured set of fields, the real questions and their word or character limits from a target application format, and give a per-field co-pilot that helps draft and tighten each answer to fit the limit without dropping the founder's actual facts. Assemble the finished set into one exportable document. The hard constraint here is the per-field limit, distinct from filling a template out of a data record or keeping several derived documents in sync with one company narrative.

WHAT STRONG LOOKS LIKE

Every field's drafted answer is under its stated limit on the first pass, and tightening an answer to fit never drops a concrete fact the founder entered, checkable by comparing the draft against the original input line by line.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Interactive prompt-chip demo app

WHO IT SERVES

PLG

WHAT TO BUILD

Visitors who land on a document tool rarely grasp what "AI editing" means until they see it fire. Build a single-page interactive demo where a sample document sits next to a row of clickable chips like "fix grammar," "cut jargon," "shorten this," and each click sends a real chat-driven edit through the product and shows the before/after inline. No signup wall before the first click; the loop is see-a-paragraph, click-a-chip, watch-it-change.

WHAT STRONG LOOKS LIKE

A cold visitor can click at least three different chips against three different sample paragraphs and see three genuinely different, correctly-applied edits, with the diff visibly rendered each time.

BAND S2 · chat · a pull request into the public repository, use-cases folder

Multilingual demo-doc pack generator

WHO IT SERVES

Global

WHAT TO BUILD

Teams selling document AI into non-English markets need localized sample content, not just a translated UI. Build a generator that produces a matched pack of demo documents (a contract, a letter, a report) in each of the platform's top-11 supported user languages from one shared brief, then exports the whole pack cleanly. The pack should read as native-authored in each language, not machine-translated boilerplate.

WHAT STRONG LOOKS LIKE

Generate the same three-document pack in at least five of the eleven languages and confirm each export opens correctly with proper script rendering (no tofu or mojibake) and locale-appropriate formatting.

BAND S2 · chat, export · a pull request into the public repository, use-cases folder

Terms-of-Service & Privacy-Policy generator

WHO IT SERVES

Legal/SaaS

WHAT TO BUILD

Early-stage SaaS and e-commerce teams put off writing their policy pages until launch week, then copy someone else's. Build a guided Q&A intake that produces a matched four-document set (Terms of Service, Privacy Policy, Warranty Policy, Returns Policy) for a fictional company, each editable in the document editor and exportable, with a persistent "drafted for attorney review, not legal advice" notice on every page. Keep the four documents cross-consistent (same company name, contact details, and effective terms), since a policy set that looks assembled from four different eras is the classic tell.

WHAT STRONG LOOKS LIKE

Complete one intake for a fictional online store and produce all four documents with matching company details throughout, each carrying the review notice, each exporting cleanly to docx and pdf.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

EULA / SaaS-agreement generator

WHO IT SERVES

Legal/SaaS

WHAT TO BUILD

Software vendors negotiate the same handful of license clauses (grant of license, usage restrictions, IP ownership, uptime commitments, liability limits) over and over with each new customer. Build a template-driven generator that produces a EULA or SaaS agreement from vendor inputs, then supports redlining a counterparty's requested changes as tracked edits inside the document rather than a fresh rewrite each time. Every draft carries a "for review by qualified counsel, not legal advice" notice, and finished agreements route to an external e-signature provider rather than claiming signature inside the product.

WHAT STRONG LOOKS LIKE

Generate a SaaS agreement from vendor inputs, then apply a simulated customer redline (for example a shortened liability cap) as visible tracked changes the vendor can accept or reject clause by clause.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

DPA (data-processing-agreement) generator

WHO IT SERVES

Compliance/Legal

WHAT TO BUILD

Every SaaS vendor handling customer personal data eventually gets asked for a DPA, and assembling one from scratch means getting the processor-controller mechanics right. Build a questionnaire that walks through the real DPA structure (subject matter and duration of processing, categories of data and data subjects, sub-processor authorization, security measures, international transfer terms, breach-notification timelines) and assembles a complete GDPR/CCPA-aligned draft. The output carries a "for review by qualified counsel, not legal advice" notice and is distinct in shape from a EULA: this is processor obligations, not a license grant.

WHAT STRONG LOOKS LIKE

Complete the questionnaire for a fictional vendor processing EU customer data and confirm the output includes named sections for sub-processors, security measures, and breach notification, not just boilerplate confidentiality language.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

BAA (business-associate-agreement) generator

WHO IT SERVES

Healthcare/Compliance

WHAT TO BUILD

HIPAA requires a signed BAA before a covered entity can share protected health information with a vendor, and the agreement carries its own required clause set separate from a generic contract. Build an API-driven generator that merges party data into a BAA template covering the real HIPAA structure (permitted uses and disclosures of PHI, administrative/physical/technical safeguard commitments, breach-notification obligations, subcontractor flow-down, and PHI return-or-destruction on termination). Every draft is marked "for review by qualified counsel, not legal advice," and final execution hands off to an external e-signature provider.

WHAT STRONG LOOKS LIKE

Call the generator with two fictional parties' details and confirm the draft names all three safeguard categories and a breach-notification timeline, structurally distinct from the DPA's processor-transfer clauses.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Power-of-attorney / POA generator

WHO IT SERVES

Legal

WHAT TO BUILD

People assembling a power of attorney need the powers-granted language and the general/limited/durable distinction to be right, then still have to get it witnessed and notarized in person, which the product never touches. Build a guided intake that asks who the principal and agent are, which POA type applies, and which specific powers are granted (financial, healthcare, real property), then assembles a clean draft with a successor-agent clause. The draft is explicitly marked "for review by qualified counsel, not legal advice," and the document itself states that signing and notarization happen outside the product.

WHAT STRONG LOOKS LIKE

Run the intake for a durable financial POA and a separate healthcare POA and confirm the two drafts differ correctly in granted powers, and both carry the review notice and the external-notarization statement.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Background-check consent-form generator

WHO IT SERVES

HR and recruiting-ops teams issuing pre-employment background-check paperwork.

WHAT TO BUILD

Employers issuing background checks need an FCRA-style standalone disclosure and a separate signed authorization, not one form doing both jobs badly. Build an API-driven generator that merges a company's template and a candidate's details into a disclosure-and-authorization pair structured the way FCRA requires (standalone disclosure, separate signed authorization), plus a summary-of-rights notice, ready for the candidate to sign through an external e-signature step. Keep the disclosure genuinely standalone with no extra clauses folded in, since that separation is the part employers get wrong.

WHAT STRONG LOOKS LIKE

Generate the set for a fictional candidate and confirm the disclosure document contains no unrelated clauses (waivers, liability language) mixed in, with the authorization and rights notice produced as separate documents.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

SOC2 evidence-index builder

WHO IT SERVES

Compliance lead preparing for a certification audit who must hand the auditor one navigable index of all evidence

WHAT TO BUILD

Build the one-time evidence-index builder: from a folder of evidence documents, assemble the audit-ready index - each control mapped to its evidence documents with titles, dates, owners, and gaps named where evidence is missing. The continuous evidence-maintainer agent is a different, existing card; this is the point-in-time builder an audit kickoff needs.

WHAT STRONG LOOKS LIKE

every index row resolves to a real document in the set; missing-evidence rows say so explicitly instead of vanishing; re-running on an updated folder produces a targeted diff of the index, not a reshuffle.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Corrective-action-plan (CAPA) drafter

WHO IT SERVES

Manufacturing/Compliance

WHAT TO BUILD

Manufacturing quality teams responding to a defect need to document root cause, not just describe the symptom and move on. Build a chat-driven drafter that walks an incident through the real 8D structure (team and problem description, containment, root-cause analysis, corrective action, implementation and validation, prevention, closure) and produces a formatted CAPA report. This is a root-cause methodology, not a checklist against a code, which is what keeps it distinct from a health-code-driven corrective plan.

WHAT STRONG LOOKS LIKE

Feed in a sample defect description and confirm the output has all eight labeled sections in order, with the root-cause section containing actual causal analysis, not a restatement of the symptom.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Franchise-disclosure-document (FDD) editor

WHO IT SERVES

Franchise/Legal

WHAT TO BUILD

An FDD is a long, rigidly structured filing (the 23 items required under the FTC Franchise Rule, from initial fees through financial performance representations) that franchise legal teams need to update section by section without breaking the rest of the document. Build a multi-doc editor that loads a full FDD, lets a reviewer target a single numbered item for a scoped edit, and exports the updated filing with every other item byte-for-byte untouched. Every draft carries a "for review by qualified counsel, not legal advice" notice; this is section-scoped editing at length, not a from-scratch generator.

WHAT STRONG LOOKS LIKE

Load a sample FDD, edit only Item 7 (initial investment figures), export, and confirm every other numbered item is unchanged from the original while Item 7 reflects the edit.

BAND S2 · chat, Multi-document, export · a pull request into the public repository, use-cases folder

HOA notice & minutes generator

WHO IT SERVES

Real estate/Community

WHAT TO BUILD

HOA boards send the same handful of document types on repeat: meeting notices, approved minutes, violation notices, all from a standard template with different fill-in details each time. Build an API-driven generator that merges board and meeting data into a clean notice-or-minutes draft, ready to export and distribute to homeowners. Keep the scope to document generation, not a full community-management platform with resident accounts or dues tracking.

WHAT STRONG LOOKS LIKE

Generate a meeting notice and a matching minutes document from the same meeting data and confirm both are internally consistent (same meeting details, same attendee list) and export cleanly.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Property-condition / move-in report app

WHO IT SERVES

Real estate

WHAT TO BUILD

Landlords and property managers documenting a unit's condition at move-in need photos tied to specific rooms and items, not a loose pile of images attached to a paragraph. Build an app that walks through a room-by-room checklist, attaches photos to each item as it's recorded, and assembles a complete condition report that exports to PDF for both landlord and tenant records. The photo-to-item binding is the point: a reviewer should be able to tell exactly which photo documents which line.

WHAT STRONG LOOKS LIKE

Complete a sample walkthrough of three rooms with at least one photo per room and confirm the exported PDF shows each photo next to its specific checklist item, not grouped separately at the end.

BAND S2 · API, images, export · a pull request into the public repository, use-cases folder

Home-inspection report builder

WHO IT SERVES

Real estate/Construction

WHAT TO BUILD

Home inspectors walk a property system by system (roof, electrical, plumbing, HVAC, structure) and need a report that pairs their findings with photo evidence per system, formatted for a buyer who's never read one before. Build a structured report builder organized by inspection system, where each finding carries a severity label, notes, and attached photos, assembled into a formatted field report on export. Keep the language observational ("found," "recommend further evaluation"), never framed as a certification of the property's condition.

WHAT STRONG LOOKS LIKE

Build a sample report covering at least four inspection systems, each with a finding, a severity label, and a photo, and confirm the exported report groups everything correctly by system.

BAND S2 · API, images, templates, export · a pull request into the public repository, use-cases folder

Comparative-market-analysis (CMA) doc app

WHO IT SERVES

real estate agents and brokers preparing listing pitches for sellers.

WHAT TO BUILD

Listing agents pitching a seller need a branded CMA, or broker opinion of value, built from comparable-property data, not a bare spreadsheet of numbers. Build an app that takes a subject property plus a handful of comparable listings and assembles a branded document with photos, a comparables table, and a suggested price range, styled for a homeowner audience rather than an analyst. The visual, branded presentation is the differentiator from a plain data export.

WHAT STRONG LOOKS LIKE

Feed in one subject property and four comparables and confirm the export shows a photo-backed comparables table and a clearly stated suggested price range on branded letterhead.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Contractor estimate & quote app

WHO IT SERVES

Trades/SMB

WHAT TO BUILD

Contractors pricing a job need the estimate to reflect what they actually saw on site, not a generic line-item form filled in from memory back at the office. Build an app that starts from site-visit data (measurements, notes, and photos captured during the walkthrough) and turns it into a branded estimate with itemized labor and materials, keeping each line traceable back to the site photo that justified it. That site-data-to-line-item link is what separates this from a generic CRM quote tool built on typed-in numbers alone.

WHAT STRONG LOOKS LIKE

Enter site notes and at least two photos for a sample job, generate the estimate, and confirm at least one line item visibly references or displays the site photo it's based on.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Restaurant-inspection corrective-plan app

WHO IT SERVES

Food & Bev

WHAT TO BUILD

A restaurant that fails a health inspection needs a corrective-action document organized around the specific violation codes cited, not a generic incident report. Build a chat-driven app that takes a list of cited violations (by code and description) and drafts a corrective plan with a fix and a responsible party per violation, formatted for resubmission to the health department. This is a code-by-code compliance response, not an 8D root-cause investigation, which keeps it distinct from a manufacturing CAPA.

WHAT STRONG LOOKS LIKE

Enter three sample cited violations and confirm the draft addresses each one individually by its code, with a specific corrective action and owner per item, not one paragraph covering all three.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Insurance loss-run summary app

WHO IT SERVES

insurance brokers and underwriters evaluating a prospective client's claims history.

WHAT TO BUILD

Underwriters and brokers reviewing a prospective client's insurance history get handed years of raw loss-run data across multiple policies and claims, and need it turned into a single readable summary before a coverage decision. Build a search-driven app that pulls claims history across multiple loss runs and produces a review document summarizing claim counts, total incurred amounts, and any open claims by policy year. This is a multi-claim, multi-year summary, distinct from drafting the narrative for one complex claim.

WHAT STRONG LOOKS LIKE

Load a sample multi-year, multi-claim loss-run set and confirm the summary correctly groups and totals claims by policy year, flagging any still-open claims separately.

BAND S2 · search, export · a pull request into the public repository, use-cases folder

Personal-statement / SOP editor

WHO IT SERVES

Education/B2C

WHAT TO BUILD

An applicant's personal statement is the one document in an application packet where the writer's own voice matters most, so a heavy AI rewrite can quietly ruin it. Build an editing app that takes a draft personal statement or statement of purpose and makes surgical, targeted edits (tightening a paragraph, fixing a weak opening line, cutting redundancy) while visibly preserving the applicant's original phrasing everywhere it isn't the specific target of an edit. Every edit should be traceable to what changed and why, not a silent full rewrite.

WHAT STRONG LOOKS LIKE

Submit a sample draft, request one targeted edit such as "tighten the opening paragraph," and confirm only that paragraph changed while the rest of the essay is untouched.

BAND S2 · chat, export · a pull request into the public repository, use-cases folder

IEP (individualized education plan) drafter

WHO IT SERVES

Education/Special-ed

WHAT TO BUILD

An IEP has a specific required shape (present levels of performance, measurable annual goals, accommodations and modifications, related services, and a least-restrictive-environment justification), and drafting one well takes real teacher time per student. Build a chat-driven drafter that takes case-manager inputs about a student's needs and produces a structured IEP draft in that shape, then routes it through a human-in-the-loop approval step so the case manager, and eventually the parent, reviews and signs off on the goals before anything is finalized. The HITL gate is the point: nothing in the plan is final until a person approves it.

WHAT STRONG LOOKS LIKE

Draft a sample IEP from student-need inputs and confirm the output has clearly labeled present-levels, goals, accommodations, and services sections, held in a pending state until the HITL approval step is completed.

BAND S2 · chat, Review, templates, export · a pull request into the public repository, use-cases folder

Syllabus + accessibility-statement generator

WHO IT SERVES

Education

WHAT TO BUILD

Course instructors rebuild the same syllabus skeleton every term (schedule, grading policy, required accessibility and accommodations statement), and the required statement text is exactly the kind of boilerplate that gets stale or copy-pasted wrong. Build a generator that assembles a syllabus from course inputs on an institution's template, including a current accessibility-statement clause (how to request accommodations, who to contact) as an editable text section. This produces a written accessibility statement inside the document, not a claim about the exported file itself being tagged for screen readers.

WHAT STRONG LOOKS LIKE

Generate a sample syllabus and confirm it includes a clearly labeled, editable accessibility-statement section with concrete contact and process language, not a placeholder.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

LinkedIn-profile → resume doc app

WHO IT SERVES

Recruitment/B2C

WHAT TO BUILD

Job seekers who've kept their LinkedIn profile current still need a polished, formatted resume document for applications that don't take a profile link. Build an app where a user pastes or uploads their exported profile data (work history, education, skills) and the app maps it into a formatted resume template, letting the user pick a layout and export. The profile import is the user's own data, pasted text or an exported file, since the product doesn't offer a native LinkedIn account connection.

WHAT STRONG LOOKS LIKE

Paste a sample exported profile's work history and education into the intake and confirm the generated resume correctly separates and formats each job entry and degree, not one run-on block.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Eulogy / tribute editor

WHO IT SERVES

B2C-personal

WHAT TO BUILD

Someone asked to speak at a memorial service is usually writing under grief, not with a clear head, and the last thing they need is a tool that flattens their own words into generic phrasing. Build a gentle editing app that takes a draft eulogy or tribute and offers light, supportive edits (pacing for reading aloud, smoothing a rough transition, tightening a section that runs long) while keeping every edit visibly optional and the writer's own voice and specific memories fully intact. No tone should ever read as clinical or transactional; the app is there to help, quietly.

WHAT STRONG LOOKS LIKE

Submit a sample draft tribute, request a pacing edit for reading aloud, and confirm the suggested changes preserve every specific personal detail and name from the original while only adjusting rhythm and flow.

BAND S2 · chat, export · a pull request into the public repository, use-cases folder

Wedding/event program & itinerary maker

WHO IT SERVES

Events/B2C

WHAT TO BUILD

Couples and event planners need a printable, branded program and a day-of itinerary that actually looks designed, not a bulleted list in a generic template. Build an app that takes event details (schedule, wedding party names, venue notes) plus a couple of chosen photos and produces a branded program and a matching itinerary document, leaning on styled templates and image placement as the visual core of the build. The template-and-image pairing is what makes it feel designed rather than typed.

WHAT STRONG LOOKS LIKE

Generate a program and itinerary from one set of event details and at least two photos, and confirm the photos are placed intentionally in the layout, not just appended at the end, with consistent branding across both documents.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

TTRPG / campaign-doc builder

WHO IT SERVES

Gaming/Creative

WHAT TO BUILD

Tabletop RPG game masters accumulate a sprawl of campaign material (world lore, session notes, NPC sheets, player-facing character sheets) that usually lives across a dozen disconnected text files. Build a multi-doc app that keeps a campaign's linked documents (a campaign bible, per-character sheets, a session log) together in one workspace, generates illustrative art for key NPCs or locations, and exports the set for sharing with players. Keep it to document generation and organization, not a live multiplayer session-play tool.

WHAT STRONG LOOKS LIKE

Build a sample campaign with a campaign-bible document, two linked character sheets, and at least one generated illustration, and confirm all three documents export as a coherent, cross-referenced set.

BAND S2 · chat, images, Multi-document, export · a pull request into the public repository, use-cases folder

Museum/gallery catalog & label maker

WHO IT SERVES

Creative/Culture

WHAT TO BUILD

Small museums and galleries preparing an exhibit need a visitor-facing catalog and a matching set of object labels, both pulling from the same object data (title, artist, date, description) so the two never drift out of sync. Build an app that takes a list of exhibit objects with images and produces a branded catalog document alongside a set of short-form printable labels from the same source data. This is visitor-facing exhibit material, distinct from an internal archival finding-aid.

WHAT STRONG LOOKS LIKE

Enter five sample exhibit objects with images and confirm the catalog and the label set both reflect the same five objects with consistent titles and details, generated from one shared object list.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Court-filing format compliance app

WHO IT SERVES

Legal / litigation paralegal or filing clerk at a 20-200 attorney litigation firm

WHAT TO BUILD

A rule pack per court defines caption block, line numbering, font and margin rules, page or word limits, signature block, certificate of service and exhibit naming. The user uploads a draft motion or brief; the app reformats it into the target court's shape, fills the caption from matter data, generates the certificate of service from the service list, and reports every rule as passed, fixed, or needs a human call. Body text is edited in place so the argument, footnotes and citations survive untouched. Output is a filing-ready DOCX plus PDF, with a one-page compliance sheet listing each change and the rule that forced it.

WHAT STRONG LOOKS LIKE

Rule packs are data files, so a new court is a config addition and not a code change; page and word limits are measured after reformatting rather than before; the compliance sheet cites the specific rule number behind each fix; footnotes, cross-references and quoted block formatting survive the reformat intact.

BAND S2 · chat, templates, Review, export · a pull request into the public repository, use-cases folder

Bates-stamped exhibit and privilege-log packet builder

WHO IT SERVES

Legal / e-discovery or litigation-support manager at a firm or discovery vendor

WHAT TO BUILD

Drop a folder of exhibits in mixed formats (DOCX, PDF, scans, photos). The app reads scanned pages, assigns sequential Bates ranges, stamps each page, adds a generated exhibit cover sheet, and builds an exhibit index document whose descriptions come from the content rather than the filename. Reviewers mark documents privileged with a reason and the privilege log assembles itself. Candidate redactions (names, account numbers, medical terms) are proposed and applied only after human approval. Final output is one bundled packet, per-exhibit clean files, and a manifest recording every Bates range and every redaction with its approver.

WHAT STRONG LOOKS LIKE

Bates numbering is deterministic and safely restartable after a partial run; nothing is ever redacted without a recorded approval; the manifest reconciles exactly to the packet page count; scanned exhibits come back genuinely searchable rather than as images.

BAND S3 · API, search, Review, Multi-document, export · a pull request into the public repository, use-cases folder

Clause-deviation memo generator

WHO IT SERVES

Legal / in-house counsel or contract manager reviewing inbound counterparty paper

WHAT TO BUILD

Upload the counterparty draft, your standard form and your playbook. The output is not a redline but a deviations memo: a table of every clause that differs from standard, the business risk stated in plain language, the playbook's fallback position, and a recommended action (accept, negotiate, escalate). Each row quotes the exact source clause. The reviewer can flip any row to redline it and the app produces that one tracked change inside the counterparty document, leaving every other clause and all of the counterparty's formatting untouched.

WHAT STRONG LOOKS LIKE

Every memo claim is grounded in quoted source text; risk language is calibrated rather than alarmist about boilerplate; the redline path edits only the clause the reviewer chose, and a diff of the untouched sections shows zero change.

BAND S2 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

M&A closing-binder assembler

WHO IT SERVES

Legal / corporate associate or transactional paralegal at a deal firm

WHAT TO BUILD

A closing checklist has 60 to 200 items. The app takes the executed documents (mixed DOCX, PDF and scans), matches each to a checklist line by content rather than filename, orders them into the binder structure, builds a hyperlinked index and cover page, inserts tab dividers and consistent running headers, and flags items that are missing or unsigned. Output is one navigable PDF binder, a per-document folder, and an outstanding-items document the deal team can send to the client unedited.

WHAT STRONG LOOKS LIKE

Document-to-checklist matching is semantic and survives inconsistent naming; unsigned-document detection has few false positives; index links resolve inside the final PDF; the outstanding list is accurate enough to be client-ready without editing.

BAND S3 · Multi-document, search, templates, export · a pull request into the public repository, use-cases folder

Contract playbook authoring and governance workspace

WHO IT SERVES

Legal / head of legal operations owning the negotiation playbook

WHAT TO BUILD

Here the playbook itself is the document. Users author positions per clause (preferred, acceptable, walk-away) with rationale and approved alternative language. The app then ingests a batch of recently signed contracts, finds where real outcomes drifted from the stated positions, and proposes playbook edits with the evidence attached. Approved changes apply in place, preserving the playbook's numbering and formatting, and a change log records who approved what and when. Published output is a styled playbook for humans plus a clean structured export that the negotiating team's own tools can consume.

WHAT STRONG LOOKS LIKE

Drift findings quote real signed-contract text rather than asserting a trend; playbook edits are surgical, so clause 7.2 changes while 7.1 does not move or renumber; the human-readable and machine-readable outputs can never diverge because both derive from one source.

BAND S3 · chat, search, Multi-document, Review, export, memory · a pull request into the public repository, use-cases folder

Research memo to client advisory letter converter

WHO IT SERVES

Legal / associate at a law firm converting internal research into client work product

WHAT TO BUILD

Take an internal research memo (authority-heavy, hedged, jargon-dense) and produce a client-ready advisory letter on firm letterhead: restructured to lead with the answer, risk framed in plain language, jurisdiction and date carried forward, authority moved into footnotes, and the firm's standard scope and reliance paragraph inserted. The associate reviews each proposed rewrite as a before and after and approves section by section. The original memo is never modified; the letter opens as a second document in the same session so both stay side by side.

WHAT STRONG LOOKS LIKE

The letter says exactly what the memo said legally, with no accidental widening or narrowing of an opinion; footnote authorities export as real footnotes; the firm's letterhead, margins and signature block come back pixel-faithful.

BAND S1 · chat, Multi-document, templates, Review, export · a pull request into the public repository, use-cases folder

Table of authorities and pin-cite builder

WHO IT SERVES

Legal / brief-writing attorney or appellate paralegal

WHAT TO BUILD

Scan a finished brief, identify every citation (cases, statutes, regulations, secondary sources), normalize each to the required citation style, verify that short-form cites resolve to a prior full cite, add missing pin cites where the quoted passage makes the page recoverable from an attached record, and generate the Table of Authorities with accurate page references grouped by authority type. Anything needing a human (an unresolvable short cite, a quote that does not match the cited source in the record) lands on a review list rather than being silently fixed. Output is the brief with the table inserted, ready to file.

WHAT STRONG LOOKS LIKE

Page references are recomputed after the inserted table shifts pagination; short-cite chains resolve correctly across sections and headings; nothing is ever invented, so an unresolvable cite is flagged instead of fabricated.

BAND S3 · chat, search, Review, export · a pull request into the public repository, use-cases folder

Earnings release and IR pack builder

WHO IT SERVES

Financial services / investor-relations manager at a listed company

WHAT TO BUILD

One session holds three linked documents: the earnings press release, the prepared remarks script, and the analyst Q&A briefing. Feed in the quarter's financial tables plus last quarter's pack. The app rolls the pack forward, replaces the figures, rewrites narrative only where the numbers moved, keeps the release template exactly (headline table, safe-harbour block, boilerplate), and drafts the Q&A document from the variances it detected. Any number appearing in more than one document is cross-checked and mismatches are raised before export. Output: release DOCX and PDF, remarks, and Q&A.

WHAT STRONG LOOKS LIKE

Cross-document number consistency is enforced and visibly demonstrated; the safe-harbour and boilerplate blocks are provably untouched; narrative changes are confined to the sections whose underlying numbers changed.

BAND S3 · Multi-document, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Fund document annual-update engine

WHO IT SERVES

Financial services / product or compliance manager at an asset manager

WHAT TO BUILD

Prospectuses, factsheets and key-information documents refresh on a cycle: performance figures, fees, risk indicators, holdings tables and dates change while the legal wording mostly does not. Upload last year's document set plus the new data. The app updates only the fields and tables that changed, keeps the typeset layout and page geometry (these documents live inside hard page budgets), produces a marked version showing every change for the compliance reviewer, and exports the clean final. A change-log document lists each edit beside its source value so every number ties back to the data file.

WHAT STRONG LOOKS LIKE

Layout survives on documents with tight page budgets, including tables that must not reflow onto a new page; the marked version is genuinely reviewable rather than a raw diff dump; unchanged legal wording is provably byte-identical.

BAND S2 · API, Multi-document, templates, Review, export · a pull request into the public repository, use-cases folder

Investment policy statement and client review pack

WHO IT SERVES

Financial services / financial adviser or wealth-management operations lead

WHAT TO BUILD

Build the IPS from a client intake (objectives, constraints, risk tolerance, allocation targets) on the firm template, then generate the periodic review pack: performance narrative, allocation-drift table, a rebalancing recommendation, and the meeting agenda. At each review the previous IPS is edited in place rather than regenerated, so the client sees precisely what changed since last time and the adviser can defend it. The adviser approves each proposed change. Exports as a branded client PDF plus a DOCX working copy for the file.

WHAT STRONG LOOKS LIKE

The "what changed since last review" section is derived from the document history rather than typed by the adviser; suitability language is never silently altered; the pack assembles for 200 clients as reliably as for one.

BAND S2 · chat, templates, Multi-document, Review, memory, export · a pull request into the public repository, use-cases folder

Covenant compliance certificate worker

WHO IT SERVES

Financial services / treasury or FP&A analyst at a borrower, or a lender's portfolio-monitoring team

WHAT TO BUILD

A server-side job that, each reporting period, takes the financial data plus the credit agreement and produces the compliance certificate in exactly the form annexed to that agreement (schedules, ratio calculations, officer signature block), shows the calculation working, and flags any covenant inside a configurable headroom of breach. Ratio definitions and defined terms are read out of the actual agreement and its amendments rather than hardcoded, so an amended definition changes the certificate automatically. Output is a signable DOCX plus PDF and a lender-ready summary attachment.

WHAT STRONG LOOKS LIKE

Definitions are genuinely sourced from the agreement including later amendments; the annexed form is reproduced faithfully enough that a lender accepts it without comment; the near-breach warning is calibrated and explains its own arithmetic.

BAND S3 · API, search, templates, export · a pull request into the public repository, use-cases folder

Jurisdiction disclosure-insert engine

WHO IT SERVES

Fintech / compliance-minded product owner at a cross-border financial platform

WHAT TO BUILD

One master client document (factsheet, proposal or account letter) must ship into a dozen markets, each demanding its own risk warnings, regulator statements, footnotes and sometimes a different section order. The app holds a disclosure library keyed by market and product, then generates each market's version from the master by inserting, replacing or reordering only the regulated blocks. A reviewer sees each market version as a diff against the master. When a regulator changes wording, one library edit regenerates every affected document and reports exactly which ones moved.

WHAT STRONG LOOKS LIKE

Inserts land in the correct section rather than being appended at the end; the master is never mutated by a market run; a single library change produces a precise, auditable list of regenerated documents with their diffs.

BAND S3 · API, templates, Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Sell-side CIM assembler

WHO IT SERVES

Financial services / M&A associate at a boutique investment bank or corporate-finance advisory

WHAT TO BUILD

From a data tape, management interview notes and the client's own materials, assemble a confidential information memorandum: business overview, market section, financial summary with recast tables, customer-concentration exhibits and appendix. Charts are generated as diagrams inside the document. Each section is drafted then edited as the client comments, section by section, without reflowing the rest. A teaser and a one-page executive summary derive from the same session, and a buyer-safe redacted version strips names and identifying details once approved.

WHAT STRONG LOOKS LIKE

The redacted version is verifiably clean, with a test proving no identifying string survives anywhere including headers and metadata; financial tables keep their number formatting and alignment; a late edit to section 3 leaves section 9's pagination alone.

BAND S3 · chat, Multi-document, images, templates, Review, export · a pull request into the public repository, use-cases folder

Regulator exam response binder

WHO IT SERVES

Fintech / compliance lead at a licensed fintech or bank during an examination

WHAT TO BUILD

Examiners send request lists, then amend them. The app turns each request into a response document: the narrative answer, the evidence list, and a cross-reference to the governing policy or control. Evidence files are indexed and renamed into the examiner's numbering. When a request is amended, only the affected response is edited and the binder index updates around it. A running status document tracks every request, owner, due date and state. Final output is a bookmarked PDF binder plus the working DOCX set.

WHAT STRONG LOOKS LIKE

An amended request updates one response without renumbering the entire binder; evidence cross-references resolve in the exported PDF; the status document is accurate enough to be the team's single source of truth during a live exam.

BAND S2 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Renewal coverage-comparison document

WHO IT SERVES

Insurance / commercial broker or account manager at renewal

WHAT TO BUILD

Upload the expiring policy and the proposed renewal wording. The app produces a client-facing comparison document: a coverage grid (limits, deductibles, sublimits), a plain-language list of what changed, newly added exclusions called out explicitly, and a recommendation section the broker edits. Wording differences are located semantically so a moved or reworded exclusion is still caught. The broker approves each finding before it reaches the client version. Exports as a branded client PDF and a DOCX working copy for the file.

WHAT STRONG LOOKS LIKE

It catches a reworded or relocated exclusion that a text diff would miss, and proves it with a test case; it never reports a pure formatting change as a coverage change; the grid's numbers tie exactly to the policy schedules.

BAND S2 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Broker submission packet builder

WHO IT SERVES

Insurance / commercial-lines broker placing a risk into the market

WHAT TO BUILD

Assemble the market submission: a narrative risk summary, exposure schedules, loss-history summary, the applications, and a cover letter tailored per carrier that names their appetite and the specific reasons this risk fits it. One session holds the shared underlying documents; each carrier version is generated as its own document from that core, so an update to the risk narrative propagates everywhere while each bespoke cover letter stays intact. Output is a per-carrier PDF pack ready to send.

WHAT STRONG LOOKS LIKE

Propagation is surgical, so the shared narrative updates across carriers while bespoke paragraphs are never overwritten; exposure schedules keep their table formatting through export; carrier framing is substantive rather than a name swap.

BAND S2 · Multi-document, templates, chat, export · a pull request into the public repository, use-cases folder

Reinsurance treaty wording workspace

WHO IT SERVES

Insurance / treaty underwriter or reinsurance broker

WHAT TO BUILD

Treaty wordings are long, bespoke and negotiated line by line. Upload the expiring wording plus the market's proposed slip. The app maps slip terms onto the wording's clauses, drafts the amended wording as tracked changes, keeps clause numbering and cross-references consistent when clauses are inserted or removed, and produces a negotiation summary listing every open point with each party's stated position. Amendments are approved individually. Exports a tracked-changes DOCX for the market plus a clean version for signing.

WHAT STRONG LOOKS LIKE

Cross-references survive clause insertion and deletion, which is where naive tools break; the tracked changes open correctly in Word with sane author attribution; the open-points summary reconciles one for one with the redline.

BAND S3 · chat, Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Claim chronology and reserve rationale memo

WHO IT SERVES

Insurance / claims adjuster or claims manager handling a complex loss

WHAT TO BUILD

Drop the claim file (correspondence, reports, invoices, medical or engineering records, photos). The app builds a dated chronology document with a one-line summary per event and a pointer to the source page, extracts the quantified elements into a damages table, and drafts the reserve rationale memo in the carrier's format using the adjuster's reasoning captured through chat. When new documents arrive weeks later, the chronology updates in place and only the memo paragraphs the new facts affect are proposed for edit.

WHAT STRONG LOOKS LIKE

Every date and amount traces to a source page; adding a document a month later produces a targeted update instead of a regenerated document; photographic evidence is referenced and embedded in the right place.

BAND S3 · Multi-document, search, images, chat, Review, export · a pull request into the public repository, use-cases folder

Policy-form filing packet builder

WHO IT SERVES

Insurance / product filing analyst at an insurer

WHAT TO BUILD

When a policy form changes, the regulator wants the new form, a redline against the approved version, a memorandum explaining each change and its rationale, and often a side-by-side. The app generates all four from the old and new forms plus the product manager's notes, preserves the form's exact typography and numbering (filing reviewers reject reformatted forms outright), and produces the state-specific cover letter from a template library. A checklist confirms every required item is present before export.

WHAT STRONG LOOKS LIKE

The form's original typography survives the round trip unchanged; the explanatory memo maps one for one onto the redline entries; state variations live in data rather than in forked code paths.

BAND S2 · Multi-document, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Commercial quote comparison document

WHO IT SERVES

Insurance / broker or corporate risk manager evaluating carrier quotes

WHAT TO BUILD

Upload several carrier quotes in whatever shape they arrived (PDF, DOCX, scan). The app normalizes them into one comparison document: coverage, limits, deductibles, premium, key conditions and subjectivities, with an explicit "not quoted" marker wherever a carrier is silent rather than guessing an answer. Differences that matter (a sublimit, a warranty, a claims-made trigger) get called out in a narrative section. The risk manager asks follow-up questions in chat and has the answer inserted as a note inside the document.

WHAT STRONG LOOKS LIKE

Silence is represented as silence and never inferred; scanned quotes are read reliably enough to trust the numbers; the wide comparison table exports to a client-ready PDF without column collapse.

BAND S2 · Multi-document, search, chat, export · a pull request into the public repository, use-cases folder

Subrogation demand and evidence packet

WHO IT SERVES

Insurance / subrogation specialist at a carrier or a recovery vendor

WHAT TO BUILD

From the claim file, generate the demand letter to the responsible party or their insurer: liability theory, itemized damages breakdown with supporting references, the statutory or contractual basis, and a response deadline. Supporting evidence (repair invoices, expert reports, photos) is indexed, labelled and bundled behind the letter. Second and third demands escalate tone and cite the real dates of prior correspondence. Every dollar claimed traces to an indexed exhibit, and the specialist approves the liability theory before any letter is produced.

WHAT STRONG LOOKS LIKE

The damages total reconciles exactly to the exhibit index; escalation letters reference genuine prior dates rather than placeholders; no letter asserts a liability theory the specialist did not explicitly approve.

BAND S2 · Multi-document, search, images, templates, Review, export · a pull request into the public repository, use-cases folder

Marketplace note: if your build wants a marketplace listing, build it and sideload it; we handle actual listings.

Patient-education leaflet engine (reading level plus language)

WHO IT SERVES

Healthcare / patient-experience lead or clinical content manager at a hospital network

WHAT TO BUILD

Clinicians submit the clinical source (discharge instruction set, medication guide, procedure prep). The app rewrites it to a target reading level, inserts a medication schedule table and a call-us-immediately-if box, generates simple illustrative images where they help, and produces the leaflet in each patient language with the layout preserved. A clinician approves every rewritten paragraph against the source before publication, and the app verifies that clinical facts (doses, timings, red-flag symptoms) still appear unchanged after simplification. Output is a printable PDF per language plus one editable master.

WHAT STRONG LOOKS LIKE

A dose or red-flag symptom can never be dropped by the simplification pass, and there is a test proving it; reading level is measured and reported rather than asserted; non-Latin and right-to-left scripts keep the layout on export.

BAND S2 · chat, images, templates, Review, export · a pull request into the public repository, use-cases folder

Multi-site informed-consent adaptation engine

WHO IT SERVES

Healthcare / clinical-trial start-up manager at a sponsor or CRO

WHAT TO BUILD

One master informed-consent form becomes thirty site-specific versions: local ethics-committee boilerplate, site contacts, local injury-compensation and privacy wording, plus translations. Each site version is produced by editing only the blocks that differ, yielding a tracked-changes copy for the committee submission and a clean copy for use. When the master is amended mid-study, the app regenerates each site version by re-applying that site's accepted local changes on top, reports conflicts where a site had customized the same paragraph, and emits the amendment redline per site.

WHAT STRONG LOOKS LIKE

Mid-study amendment handling is the whole test, so site-local customizations must survive the re-base and conflicts must be surfaced rather than silently resolved; the committee redline is a real tracked-changes DOCX that opens correctly in Word.

BAND S3 · Multi-document, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Accreditation evidence binder builder

WHO IT SERVES

Healthcare / quality and accreditation manager at a hospital

WHAT TO BUILD

A standards framework (chapters of standards, each with measurable elements) drives the binder. The app maps existing policies, committee minutes, training records and audit results to each element, drafts a compliance narrative per element citing its evidence, flags elements with no evidence at all, and assembles a surveyor-ready binder with an index and per-standard tabs. Gaps become a work-list document with owners and dates. Between surveys the binder is maintained rather than rebuilt: when a policy is revised, only the affected narratives are proposed for update.

WHAT STRONG LOOKS LIKE

Evidence-to-standard mapping is semantic and honest, so "no evidence" is reported rather than padded with a plausible-sounding narrative; a single policy revision produces a small and correct set of narrative updates rather than a full rebuild.

BAND S3 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Operative and procedure note generator

WHO IT SERVES

Healthcare / surgeon or procedural specialist, or a clinical-documentation product team

WHAT TO BUILD

From dictation or structured intraoperative capture plus the surgeon's own preferred template, produce the operative note: indication, findings, procedure detail, specimens, estimated blood loss, implants with catalogue numbers, complications and disposition. Implant and specimen fields are required and blocked from being guessed. The surgeon reviews before signing and edits by chat, with each edit changing only the paragraph it touches. The signed note exports in the record format the hospital requires, and a patient-facing summary generates from the same session.

WHAT STRONG LOOKS LIKE

Required fields are genuinely enforced, so a missing implant number halts the note instead of being invented; the surgeon's personal template style is preserved exactly; the patient summary can never contradict the clinical note.

BAND S2 · chat, templates, Review, export · a pull request into the public repository, use-cases folder

Nursing handoff and transfer packet assembler

WHO IT SERVES

Healthcare / charge nurse or transfer coordinator

WHAT TO BUILD

Transfers between units or facilities are assembled by hand today. The app builds the packet: a structured handoff summary, the current medication list with reconciliation flags, active orders, pending results, isolation and mobility status, and the source documents appended in a fixed order the receiving unit expects. Fields that must be human-verified before transfer (allergies, code status, high-risk medications) are highlighted for explicit confirmation and block export until confirmed. Output is one PDF packet plus the structured handoff document.

WHAT STRONG LOOKS LIKE

High-risk fields genuinely gate the export; medication reconciliation is conservative, flagging a possible duplicate rather than silently resolving it; packet ordering is fixed and predictable so the receiving unit can rely on it.

BAND S2 · Multi-document, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Medico-legal record chronology builder

WHO IT SERVES

Healthcare / medico-legal analyst at a law firm, insurer or independent medical review company

WHAT TO BUILD

Ingest a few hundred pages of mixed medical records (scans, faxes, portal exports). The app orders them by encounter date, builds a chronology table of date, provider, setting, findings and a source page reference, extracts a medications timeline and a diagnostics timeline, and drafts a summary highlighting treatment gaps, inconsistencies between providers, and the first documented mention of the claimed condition. Every row links back to its page. The analyst annotates rows in place and those annotations flow into the final report.

WHAT STRONG LOOKS LIKE

Page-level provenance on every single row; duplicated and out-of-order faxed records are handled without double counting; "first documented mention" claims are correct and independently verifiable from the links.

BAND S3 · API, Multi-document, search, chat, export · a pull request into the public repository, use-cases folder

Provider credentialing packet builder

WHO IT SERVES

Healthcare / medical staff services or credentialing coordinator

WHAT TO BUILD

Per provider, assemble the credentialing file: the CV reformatted to the committee template with employment-history gaps detected and listed, a licence and certification summary with expiry dates, attestations, peer references, and the primary-source verification index. A committee packet compiles many providers into one agenda-ordered document with a per-provider summary sheet. Recredentialing reuses the prior file and edits only what changed, producing a "changes since last cycle" page that the committee will actually read.

WHAT STRONG LOOKS LIKE

employment-gap detection works on messy real CVs in inconsistent formats; expiry dates are extracted reliably enough to drive reminders; the recredentialing delta page is derived from the documents rather than retyped by a coordinator.

BAND S2 · API, Multi-document, templates, search, export · a pull request into the public repository, use-cases folder

Synoptic report re-templater for pathology and radiology

WHO IT SERVES

Healthcare / pathologist, radiologist, or a diagnostics-AI product manager

WHAT TO BUILD

Free-text diagnostic reports get restructured into the synoptic format registries and tumour boards require: one field per required data element, standard headings, required negatives stated explicitly, and a rendered summary table. The app reads the narrative report, maps content into the synoptic template, marks any element it cannot ground as "not stated" instead of inferring it, and shows the reporting clinician a side-by-side of narrative and synoptic for approval. Historic report batches convert server-side for registry submission.

WHAT STRONG LOOKS LIKE

"not stated" is used honestly and frequently rather than filled with a confident guess; the original narrative is preserved untouched next to the synoptic version; a thousand-report batch conversion is idempotent and safely re-runnable.

BAND S3 · chat, API, templates, Review, export · a pull request into the public repository, use-cases folder

LOI to lease and purchase-contract converter

WHO IT SERVES

Real estate / commercial broker or in-house transaction manager

WHAT TO BUILD

A signed letter of intent carries the deal economics in shorthand. The app extracts the business terms (rent schedule, escalations, term, options, tenant-improvement allowance, free rent, deposit, permitted use), then drafts the full lease or purchase agreement from the landlord's or seller's standard form, placing each term in its correct clause and generating the rent-schedule exhibit as a real table. A term-to-clause map shows where every LOI point landed, and any LOI term with no home in the form is flagged for a lawyer rather than quietly dropped.

WHAT STRONG LOOKS LIKE

the term-to-clause map is complete and every unplaced term surfaces; the rent schedule computes correctly including escalations and free-rent months; the standard form's formatting and defined terms survive the fill intact.

BAND S3 · Multi-document, search, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Jurisdiction notice engine for property managers

WHO IT SERVES

Real estate / property manager or lettings operations lead running units across several jurisdictions

WHAT TO BUILD

Notices (rent increase, entry, cure or quit, renewal, deposit disposition) carry jurisdiction-specific content, timing and delivery rules. The app holds a notice library keyed by jurisdiction, generates the correct notice for a tenant and unit from the property record, computes and prints the required dates, and refuses to produce a notice whose timing is invalid while explaining exactly why. Batches generate per property with a delivery-log document. Updating a jurisdiction's rules regenerates only the affected pending notices.

WHAT STRONG LOOKS LIKE

date arithmetic is right including weekends, holidays and mailing add-on days; an invalid notice is blocked with a readable explanation instead of being produced; jurisdiction rules are data with a recorded provenance note per rule.

BAND S2 · API, templates, search, Review, export · a pull request into the public repository, use-cases folder

Commercial offering memorandum builder

WHO IT SERVES

Real estate / investment-sales broker at a commercial brokerage

WHAT TO BUILD

From a rent roll, operating statements, photographs and the property record, assemble the offering memorandum: executive summary, investment highlights, tenant and lease-abstract tables, financial summary with the pro forma, market and demographic section, and photo pages. Diagrams such as tenancy mix and site callouts render inside the document. The broker refines section by section in chat while the design template stays fixed. A confidential version and a broadly marketed version generate from the same session with different disclosure blocks.

WHAT STRONG LOOKS LIKE

rent-roll and financial tables tie exactly to the source files; photo and diagram pages survive PDF export at print quality; the two disclosure variants never leak content into each other, proven with a test.

BAND S3 · chat, Multi-document, images, templates, export · a pull request into the public repository, use-cases folder

Estoppel and SNDA batch generator

WHO IT SERVES

Real estate / asset manager or lender's counsel during a portfolio financing or sale

WHAT TO BUILD

Financing requires an estoppel certificate from every tenant and often a subordination agreement. The app reads each lease, fills the estoppel with that lease's actual facts (commencement, expiry, rent, deposits, options, defaults), generates the tenant-specific cover letter, and tracks returns. When a tenant returns a marked-up estoppel, the changes are summarized against the requested form so counsel reviews exceptions instead of re-reading every certificate. A portfolio status document shows outstanding, returned and exception items.

WHAT STRONG LOOKS LIKE

lease facts extract correctly across inconsistent lease forms including later amendments; returned mark-ups diff accurately against the requested form; the status document stays trustworthy at two hundred tenants.

BAND S3 · API, Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Condo and HOA resale disclosure packet assembler

WHO IT SERVES

Real estate / HOA manager or community-association management company

WHAT TO BUILD

A resale requires a packet: governing documents, current budget and reserve-study summary, insurance certificate, minutes for a defined period, the unit's assessment status, violations and litigation disclosure, plus a jurisdiction-required cover form. The app assembles it per unit, generates the unit-specific status and disclosure pages from the association record, builds the index, and produces a completeness certificate stating what is included and what is unavailable with the reason. Repeat orders reuse the association core and regenerate only the unit pages.

WHAT STRONG LOOKS LIKE

unavailable items are disclosed honestly rather than omitted; association-wide content and unit-specific content are cleanly separated so a budget update propagates once; packets regenerate reproducibly for audit.

BAND S2 · Multi-document, templates, API, export · a pull request into the public repository, use-cases folder

Short-term rental house manual and listing pack

WHO IT SERVES

Real estate / short-term-rental operator or a property-management startup

WHAT TO BUILD

One property record produces a branded guest house manual (arrival, wifi, appliances, house rules, local recommendations, emergency info) with photographs and a simple map diagram, plus the listing description written to the platform's tone and length limits and a printable in-unit welcome card. Manuals generate in the guest's language with layout preserved. When a rule or appliance changes, the operator edits once and every property variant sharing that block updates while property-specific text stays put.

WHAT STRONG LOOKS LIKE

shared blocks and property-specific content are modelled cleanly so an update cannot clobber bespoke text; multi-language versions keep both layout and photos; the export looks right printed, not only on screen.

BAND S1 · chat, images, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Appraisal and valuation report builder

WHO IT SERVES

Real estate / appraiser or valuation analyst

WHAT TO BUILD

A valuation report is long, standardized and evidence-heavy. The app assembles it from inspection notes, comparables data, photographs and market research: property description, highest and best use, the three valuation approaches with their calculation tables, reconciliation, assumptions and limiting conditions, and the certification page. Comparable adjustments render as a real table with the arithmetic shown. The appraiser edits narrative sections without disturbing the standard boilerplate, and a completeness checklist confirms every required element before export. Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S3 · chat, images, templates, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Submittal register and transmittal packet builder

WHO IT SERVES

Construction / project engineer or document controller at a general contractor

WHAT TO BUILD

From the specification sections, generate the submittal register (what must be submitted, by whom, in what form). As submittals arrive, the app builds each transmittal packet: cover sheet, the submittal content, the applicable specification extract for the reviewer, and a comparison highlighting where the submitted product deviates from the specified one. Reviewer stamps and comments are captured, the register updates, and an overdue report generates weekly. The deviation analysis is the differentiator because it reads both documents rather than ticking a box.

WHAT STRONG LOOKS LIKE

deviation detection compares real product data against real specification language and quotes both sides; the register stays consistent when specification sections are revised mid-project; the weekly overdue report is accurate enough to send to the owner unedited.

BAND S3 · Multi-document, search, chat, templates, export · a pull request into the public repository, use-cases folder

Change order and pricing narrative generator

WHO IT SERVES

Construction / project manager or contract administrator at a contractor or owner's representative

WHAT TO BUILD

A change starts as an RFI answer, a field directive or a design revision. The app drafts the change-order document: scope narrative, the entitlement paragraph citing the specific contract clause that permits the change, a cost breakdown table, the schedule-impact statement and the signature block. Entitlement language is pulled from the actual executed contract open in the session, not a generic library. Rejected or revised changes are edited in place with a version history the owner can follow across rounds. Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S2 · Multi-document, search, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Site safety plan and toolbox-talk generator

WHO IT SERVES

Construction / safety manager at a contractor

WHAT TO BUILD

From the project scope and the activity schedule, generate the site-specific safety plan (hazard analysis per activity, controls, emergency procedures, site map) and the weekly toolbox-talk documents matching the coming week's activities. Site photographs annotate the hazards. Talks generate in the crew's languages with layout preserved for printing and posting. After an incident, the app drafts the incident report and proposes the specific plan sections to revise, so the plan behaves as a living document rather than a shelf binder.

WHAT STRONG LOOKS LIKE

hazard analyses are activity-specific rather than generic boilerplate; multilingual talks keep formatting suitable for printing and posting on site; the incident-to-revision loop proposes narrow, correct edits instead of rewriting the plan.

BAND S2 · chat, images, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Daily field report pipeline

WHO IT SERVES

Construction / superintendent or field engineer; built phone-first

WHAT TO BUILD

The superintendent dictates or types rough notes and uploads photos from a phone. The app produces the branded daily report: manpower and equipment tables, work performed by area, weather, delays with their causes, safety observations, and photo pages with captions generated from the images. Reports batch across projects nightly. A weekly rollup document summarizes progress and flags recurring delay causes, with every claim in the rollup linking back to the daily report that supports it.

WHAT STRONG LOOKS LIKE

photo captions are specific and useful rather than generic; delay causes are categorized consistently enough to roll up meaningfully; the nightly batch is resilient to partial input and never emits a half-empty report silently.

BAND S2 · API, chat, images, templates, export · a pull request into the public repository, use-cases folder

Owner-contractor agreement redline workspace

WHO IT SERVES

Construction / contracts manager at a contractor or owner's representative

WHAT TO BUILD

Standard-form construction agreements arrive with owner amendments attached, often as a separate supplementary-conditions document that silently rewrites the base form. The app reconciles the base form, the supplementary conditions and any exhibits into one effective document showing what each clause actually says after every modification, then redlines it against the contractor's risk playbook (indemnity, consequential-damages waiver, notice periods, payment terms, termination for convenience). Each finding carries a risk rating and proposed alternative language the manager approves one by one.

WHAT STRONG LOOKS LIKE

the reconciliation across base form plus supplementary conditions is correct, which is the genuinely hard part; notice periods and payment terms are extracted exactly; approved alternatives insert without disturbing the surrounding document.

BAND S3 · Multi-document, search, chat, Review, templates, export · a pull request into the public repository, use-cases folder

Public-comment response matrix builder

WHO IT SERVES

Government / rulemaking analyst or policy officer at a regulatory agency

WHAT TO BUILD

A proposed rule draws hundreds or thousands of comments. The app ingests the comment set (PDF letters, form emails, portal exports), clusters them by issue, isolates the distinct substantive points, drafts the agency's answer per point in the required comment-and-response format, and links each response to the commenters who raised it. Form-letter campaigns are detected and counted rather than repeated. The response-to-comments document publishes alongside the final rule, and every change made to the rule text because of a comment is cross-referenced to the response explaining it.

WHAT STRONG LOOKS LIKE

clustering is honest, so a minority substantive point is never swallowed into a majority cluster; form-letter detection counts accurately without discarding the individualized paragraphs some senders add; every rule change traces to the comment that caused it.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Bill and amendment markup tool

WHO IT SERVES

Government / legislative counsel or committee staffer

WHAT TO BUILD

Amendments arrive as instructions ("in section 4(b), strike 'thirty' and insert 'sixty'"). The app applies each instruction to the bill text as tracked changes, rennumbers affected sections, repairs internal cross-references, and produces both the amended bill and a clean comparison print. It also runs the reverse: given a marked-up bill, generate the formal amendment instructions. Conflicting amendments touching the same line are detected and reported to staff rather than resolved automatically.

WHAT STRONG LOOKS LIKE

instruction-to-edit application is exact, including quoted-string matching and whitespace; the reverse direction produces instructions that reapply cleanly to the original text; conflicts are always reported and never silently merged.

BAND S3 · chat, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Plain-language notice rewriter with reading-level report

WHO IT SERVES

Government / communications or service-design lead at a public agency

WHAT TO BUILD

Take existing public notices, benefit letters and forms and rewrite them into plain language while keeping every legally required statement verbatim. Required statements are marked as protected blocks that cannot be reworded. The output is the rewritten document, a before-and-after reading-level report, and a checklist proving each protected statement is present. Heading and list structure is rebuilt so the document is navigable with assistive technology, and translated versions are generated with the protected blocks handled correctly per language.

WHAT STRONG LOOKS LIKE

protected blocks are provably untouched; the reading-level improvement is measured on the real output rather than claimed; the rebuilt heading and list structure survives export instead of only looking right on screen.

BAND S2 · chat, search, templates, Review, export · a pull request into the public repository, use-cases folder

Council and board agenda packet assembler

WHO IT SERVES

Government / city clerk or board secretary

WHAT TO BUILD

Each meeting produces a packet: agenda, staff reports per item, resolutions and ordinances, attachments and exhibits. The app builds it from submitted staff reports, normalizes each into the standard format, numbers items and generates the hyperlinked agenda, then produces the public PDF plus a redacted version where confidential attachments are removed and listed as withheld. After the meeting, minutes draft from the agenda and the vote record with each item's outcome and motion text filled in.

WHAT STRONG LOOKS LIKE

staff reports from different departments normalize without losing content; the withheld list is accurate and the redacted packet is genuinely clean; minutes item numbering matches the agenda exactly.

BAND S2 · Multi-document, templates, search, Review, export · a pull request into the public repository, use-cases folder

Funding opportunity and applicant guideline pack

WHO IT SERVES

Government / grants program manager at a public agency or funder

WHAT TO BUILD

A funding program needs the opportunity notice, applicant guidelines, evaluation rubric, budget template and FAQ, all consistent with one another. The app drafts the set from the program parameters, then enforces consistency: a deadline or eligibility rule changed in one document is proposed for update everywhere it appears, with the reviewer approving each. Applicant questions during the open period convert into FAQ entries and, where they expose ambiguity, into proposed guideline clarifications with a properly dated amendment notice.

WHAT STRONG LOOKS LIKE

cross-document consistency enforcement is real and catches a date changed in only one place; the amendment notice describes exactly what changed and when; the rubric maps one for one onto the stated criteria.

BAND S2 · Multi-document, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Permit and licence decision-letter engine

WHO IT SERVES

Government / permitting officer at a regulatory or municipal authority

WHAT TO BUILD

A decision letter must state the decision, the statutory basis, the conditions imposed, the appeal route and the deadline. The app generates it from the case record plus a conditions library, inserting the exact statutory citation behind each condition and the correct appeal wording for the decision type. A refusal requires the officer to select and justify each ground before the letter can be produced at all. Bulk cases generate overnight with an exception queue for anything the rules could not resolve.

WHAT STRONG LOOKS LIKE

statutory citations are correct and current because they are sourced from the attached legislation rather than hardcoded; a refusal cannot issue without recorded grounds; the exception queue is small and every entry carries a readable reason.

BAND S2 · API, templates, search, Review, export · a pull request into the public repository, use-cases folder

Emergency public-information kit generator

WHO IT SERVES

Government / emergency management or public information officer

WHAT TO BUILD

During an incident the same facts must ship within minutes as a press release, a public notice, a social copy set, a shelter or evacuation instruction sheet with a map diagram, and an accessible plain-language version, in every community language. The app takes one incident record and generates the whole kit; when the situation changes it edits every artefact in place so version numbers, timestamps and facts stay consistent. A change bulletin lists exactly what differs from the previous release.

WHAT STRONG LOOKS LIKE

an update propagates to every artefact with no stale fact surviving anywhere, proven by a test; multilingual and accessible versions are produced in the same pass rather than bolted on later; the change bulletin is derived rather than hand-written.

BAND S2 · chat, Multi-document, images, templates, export · a pull request into the public repository, use-cases folder

Accreditation self-study assembler

WHO IT SERVES

Education / assessment or accreditation director at a university or college program

WHAT TO BUILD

A self-study maps program evidence to accreditation standards across hundreds of pages written by many faculty. The app holds the standards framework, collects section drafts from contributors, normalizes voice and formatting into the institution's template, maps evidence documents to each criterion with an appendix index, drafts narrative where evidence exists, and flags criteria whose evidence is thin or missing. Reviewer comments are captured in the document and resolved through the same review flow. Output is the submission PDF plus the evidence appendix.

WHAT STRONG LOOKS LIKE

many faculty voices normalize into one without losing substance or attribution; the evidence index resolves correctly in the exported PDF; thin-evidence flags are specific and honest instead of a blanket warning.

BAND S3 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Reading-level and curriculum adaptation engine for course materials

WHO IT SERVES

Education / instructional designer or open-educational-resources publisher

WHAT TO BUILD

Take an existing chapter or unit and produce adapted versions: a lower reading level, a language-learner version with a glossary sidebar, a reduced-load version, and a version aligned to a different curriculum's terminology. Figures, equations and captions are preserved exactly across every version and only the prose adapts. A comparison view shows the adaptations side by side so a teacher can see what changed. Adaptations regenerate when the source chapter is revised.

WHAT STRONG LOOKS LIKE

mathematical content and figures survive every adaptation byte-identical; the glossary is drawn from the actual text rather than a generic word list; a source revision regenerates adaptations without discarding teacher edits already made to them.

BAND S2 · chat, Multi-document, images, templates, export · a pull request into the public repository, use-cases folder

Multilingual family communication engine for schools

WHO IT SERVES

Education / district communications lead or school administrator

WHAT TO BUILD

Report-card comment sheets, permission slips, behaviour notices and district announcements must reach families in their home languages, on the district template, with the child's and teacher's details merged in. The app generates each family's document in their recorded language, preserves layout and district branding across scripts, and produces both a print run and per-family PDFs. Teacher comment banks are edited once and the change flows to every generated letter still pending.

WHAT STRONG LOOKS LIKE

layout survives right-to-left and CJK scripts on the same template; merged fields are correct at scale with a reconciliation report proving it; a comment-bank edit affects only pending letters and never already-sent ones.

BAND S2 · API, templates, Multi-document, export · a pull request into the public repository, use-cases folder

School improvement plan drafter

WHO IT SERVES

Education / principal or district improvement officer

WHAT TO BUILD

The plan follows a mandated structure: needs assessment with data tables, goals with measurable targets, strategies with evidence citations, resource allocation and a monitoring schedule. The app drafts it from assessment data plus last year's plan, keeps the mandated structure exactly, generates the data tables and trend diagrams, and carries forward in-progress goals with their status attached. Mid-year monitoring edits the plan in place and produces the board progress report from the same document.

WHAT STRONG LOOKS LIKE

the mandated structure is never reorganized by the drafting pass; carried-forward goals keep their identity and history rather than being rewritten; the mid-year report derives from the plan so the two can never drift apart.

BAND S2 · chat, images, templates, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Teacher observation and evaluation report builder

WHO IT SERVES

Education / instructional coach, department head or school leader

WHAT TO BUILD

During an observation the leader captures evidence notes against a rubric. The app assembles the report: evidence quoted under each rubric indicator, a rating with justification, strengths, growth areas and agreed next steps. Prior observations are referenced so the report can state what actually improved. The teacher adds a written response that is appended verbatim rather than merged. Output is a signed PDF for the personnel file plus an editable copy for the coaching conversation.

WHAT STRONG LOOKS LIKE

every rating is backed by quoted evidence from the notes and an unsupported rating is blocked; the teacher's response is appended verbatim and never edited; statements about improvement cite the specific earlier observation.

BAND S1 · chat, templates, Multi-document, memory, Review, export · a pull request into the public repository, use-cases folder

Standards alignment mapper for existing course documents

WHO IT SERVES

Education / curriculum director or an EdTech content team

WHAT TO BUILD

A district or publisher holds hundreds of existing lesson and assessment documents plus a standards framework they must align to. The app reads each document, proposes the standards it addresses with the supporting passage quoted, inserts an alignment table into the document itself, and builds a coverage map showing which standards are over-covered, under-covered or absent. Reviewers accept or correct each mapping and those corrections shape later suggestions. Output is the annotated document set plus the coverage report.

WHAT STRONG LOOKS LIKE

every proposal quotes the passage that justifies it rather than asserting a match; the coverage map is honest about gaps instead of smoothing them; accepted corrections measurably change subsequent proposals rather than being logged and ignored.

BAND S3 · API, Multi-document, search, Review, memory, export · a pull request into the public repository, use-cases folder

Global employment contract generator with jurisdiction clause packs

WHO IT SERVES

HR / people-operations lead, or a product manager at an employer-of-record platform

WHAT TO BUILD

Hiring in a new country means a different employment agreement: probation limits, notice periods, working time, leave entitlements, restrictive-covenant enforceability and mandatory local language. The app holds a clause pack per country, assembles the agreement from the company's global core plus the local pack, produces a bilingual version where local language is mandatory, and reports which clauses were replaced and why. When a country rule changes, the affected templates and live agreements are identified and updated behind a review step.

WHAT STRONG LOOKS LIKE

the global core and local packs compose without contradiction, and contradictions are detected rather than shipped; bilingual documents keep aligned layout across scripts; the change-impact report names exactly which templates and which signed agreements are affected.

BAND S3 · Multi-document, templates, search, Review, export · a pull request into the public repository, use-cases folder

Workplace investigation report writer

WHO IT SERVES

HR / employee-relations investigator or HR business partner

WHAT TO BUILD

An investigation produces interview notes, documents and a chronology. The app builds the report in the standard structure: allegations, scope and methodology, evidence per allegation with witness attribution, credibility assessment prompts the investigator completes, findings against the stated standard of proof, and recommendations. Every finding must cite the evidence supporting it and a finding with no citation is blocked. A redacted version for the complainant and the respondent is generated from the same document with witness identities protected.

WHAT STRONG LOOKS LIKE

uncited findings are genuinely blocked rather than warned about; the redacted version is verifiably free of protected identities including in headers, footers and file metadata; the report reads as the investigator's own work rather than generic HR prose.

BAND S3 · chat, Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Total-rewards statement generator

WHO IT SERVES

HR / compensation and benefits manager at a mid-size or large employer

WHAT TO BUILD

Each employee receives a personalized statement: base pay, bonus, equity with a vesting schedule table, retirement contributions, insurance values, leave value and perks, plus a chart of total package composition. The app generates them on the company template at scale, in each employee's language and currency, including or omitting country-specific benefit sections correctly. A preview batch lets HR check a sample before the full run, and a reconciliation report ties every statement back to its source data row Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S2 · API, images, templates, export · a pull request into the public repository, use-cases folder

Reduction-in-force and separation packet assembler

WHO IT SERVES

HR / HR director or employment counsel running a restructuring

WHAT TO BUILD

Each affected employee needs a coordinated set: notice letter, separation agreement with the correct consideration and revocation periods, the disclosure schedule of ages and job titles where required, benefit continuation information, and any jurisdiction-specific notice. The app assembles per employee from the roster and approved templates, computes and prints every date, and produces a manager talking-points document. Counsel approves the template set once, then every packet is generated and checked against a rules checklist before release.

WHAT STRONG LOOKS LIKE

date computations and the disclosure schedule are exactly right, because those two items are what void an agreement; jurisdiction rules select correctly per employee; nothing releases without passing the checklist.

BAND S3 · API, Multi-document, templates, Review, export · a pull request into the public repository, use-cases folder

Job architecture and levelling framework builder

WHO IT SERVES

HR / compensation or people-strategy lead building a career framework

WHAT TO BUILD

A job architecture is a document set: level definitions, track descriptions, per-family competency matrices and one profile per role. The app drafts the framework from existing job descriptions by clustering them into families and levels, generates consistent profiles on the template, and reports which existing roles do not fit the proposed structure. Changing a level definition proposes edits to every profile that references it. Output is a published framework document plus per-role profiles employees can actually read.

WHAT STRONG LOOKS LIKE

clustering produces a structure a compensation professional recognizes as defensible rather than arbitrary groupings; a level-definition change updates dependent profiles surgically; misfit roles are reported honestly instead of being forced into a level.

BAND S3 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Open-enrolment communication kit

WHO IT SERVES

HR / benefits manager at a mid-size or large employer

WHAT TO BUILD

Benefits change every year and everyone must be told. The app produces the kit from the plan data: an enrolment guide with plan comparison tables, a summary of what changed since last year, per-population variants by location, union status and employment type, personalized enrolment letters carrying each employee's current elections and deadline, and a manager FAQ. The comparison tables and the "what changed" section are derived from a diff of this year's plan data against last year's rather than written by hand.

WHAT STRONG LOOKS LIKE

the "what changed" section is genuinely derived and complete; population variants are correct so no employee receives the wrong plan set; wide comparison tables export legibly to both print and PDF.

BAND S2 · API, Multi-document, templates, images, export · a pull request into the public repository, use-cases folder

Statutory financial statements builder with a note library

WHO IT SERVES

Accounting / financial reporting manager, or a reporting specialist at an accounting firm

WHAT TO BUILD

Statutory accounts are the primary statements plus twenty to sixty notes, each with required disclosures under the applicable standard. The app rolls last year's statements forward, updates the primary statements and comparatives from the trial balance, regenerates each note's figures, flags notes whose required disclosures are missing given this year's transactions (a new lease, a business combination), and keeps the typeset layout and note numbering intact. A disclosure checklist maps every required item to the note that satisfies it.

WHAT STRONG LOOKS LIKE

comparatives and the cross-references between notes and primary statements tie exactly; a newly required note is detected from the data rather than remembered by a human; note numbering and the checklist stay consistent after an insertion mid-document.

BAND S3 · Multi-document, templates, search, Review, export · a pull request into the public repository, use-cases folder

Engagement letter and client acceptance pack

WHO IT SERVES

Accounting / partner or practice manager at an accounting or advisory firm

WHAT TO BUILD

Each client and service line needs an engagement letter with the right scope, fee basis, limitation of liability and independence wording, plus acceptance and anti-money-laundering documentation. The app generates the pack from the client record and service selection, assembles multi-service engagements into one letter with a correct combined scope, and runs the annual re-issue cycle where last year's letter is updated only where fees, scope or standards changed, with a change summary for the partner to approve.

WHAT STRONG LOOKS LIKE

multi-service scope composition never produces contradictory or duplicated clauses; the annual re-issue is a genuine diff rather than a regeneration; independence and liability wording is never altered by the drafting pass.

BAND S2 · templates, Multi-document, chat, Review, export · a pull request into the public repository, use-cases folder

Tax technical memorandum builder

WHO IT SERVES

Accounting / tax manager at a firm or an in-house tax lead

WHAT TO BUILD

A position memo states facts, question, authorities, analysis and conclusion with a confidence level. The app drafts it from the facts provided plus the authorities the user attaches (legislation, rulings, guidance), quoting each authority exactly with its citation, building an authorities table, and marking any point where the analysis leans on an authority that is not in the attachments so it can be added or the point dropped. Client-ready and file-ready versions generate from the same session Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S2 · chat, Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Audit findings register and management letter generator

WHO IT SERVES

Accounting / audit senior at a firm or an internal audit manager

WHAT TO BUILD

Fieldwork scatters findings across workpapers. The app compiles the findings register (condition, criteria, cause, effect, recommendation, rating, owner, due date), drafts the management letter grouping findings by process and severity, and generates the management-response template for the client to complete. Prior-year findings carry forward with their status so the letter can state what was remediated and what recurs. A one-page audit committee summary derives from the same register.

WHAT STRONG LOOKS LIKE

each finding keeps its five-element structure with nothing collapsed or missing; prior-year carry-forward is accurate including partially remediated items; the committee summary never states a conclusion the register does not support.

BAND S2 · Multi-document, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Client tax organizer and year-end letter batch

WHO IT SERVES

Accounting / tax practice manager at a small or mid-size firm

WHAT TO BUILD

Before filing season every client gets an organizer pre-filled from last year's return (the schedules that applied, the deductions claimed, the entities involved) plus a personalized letter naming the law changes that affect them specifically. The app generates the batch, produces a fillable and a print version, and once documents come back drafts the missing-items follow-up per client from what is still outstanding. Client segments such as business, rental and expatriate receive the right sections and nothing more.

WHAT STRONG LOOKS LIKE

pre-fill genuinely comes from the prior return rather than a generic checklist; the law-change paragraph is client-specific and correct; follow-up letters are accurate about what has and has not been received.

BAND S2 · API, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Month-end close narrative and variance commentary pack

WHO IT SERVES

Accounting / client-accounting-services lead at a firm, or an in-house controller

WHAT TO BUILD

After close, the numbers need a story. The app takes the trial balance, prior period and budget, generates the variance table, drafts commentary for every variance beyond a threshold with a proposed explanation drawn from the ledger detail, and assembles the client or board reporting pack on the branded template with charts. The controller edits commentary in place, and next month the pack carries forward the narrative structure and prior explanations so recurring items are described consistently rather than freshly invented. Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S2 · API, chat, images, templates, memory, export · a pull request into the public repository, use-cases folder

Pursuit proposal engine with credentials library

WHO IT SERVES

Consulting / proposal manager or partner at a consulting or professional-services firm

WHAT TO BUILD

A pursuit proposal assembles from reusable parts: methodology sections, case studies, team biographies, credentials, rate cards and firm boilerplate, plus bespoke understanding-of-need and approach sections. The app searches the credentials library for the case studies most relevant to the client's sector and problem, assembles the document on the firm template, tailors each reused section into the client's own language, and reports which content is reused versus bespoke so the partner knows where to spend review time.

WHAT STRONG LOOKS LIKE

case-study selection is genuinely relevant and its reasoning is shown; tailoring changes reused sections meaningfully without breaking their factual claims; the reuse-versus-bespoke report is accurate enough to triage a partner's review.

BAND S3 · search, Multi-document, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Consultant biography and CV pack builder

WHO IT SERVES

Consulting / resourcing or proposal coordinator at a consulting firm

WHAT TO BUILD

Every pursuit needs tailored biographies, because the same consultant reads differently in a healthcare bid than in a public-sector one. The app holds master consultant records (projects, sectors, qualifications, publications) and generates per-pursuit biographies at the required length, selecting the most relevant experience, on the firm's biography template with the photograph placed correctly. A team pack assembles the selected biographies behind a capability summary. Master-record updates propagate to future pursuits without touching already-submitted packs.

WHAT STRONG LOOKS LIKE

relevance selection visibly and correctly changes between sectors; length targets are hit exactly, since bid biographies carry hard word limits; already-submitted packs are immutable and reproducible.

BAND S2 · API, templates, search, Multi-document, export · a pull request into the public repository, use-cases folder

Engagement status and steering-committee report automation

WHO IT SERVES

Consulting / engagement manager on a live project

WHAT TO BUILD

Weekly status and monthly steering documents are assembled by hand from the plan, risk log, workstream updates and financials. The app generates them: progress against milestones with a timeline diagram, status per workstream with the reasoning behind each rating, risks and issues with their movement since last period, decisions needed, and budget burn. Workstream leads submit short updates and the app normalizes them. Every "since last time" statement derives from a comparison with the previous report rather than being retyped.

WHAT STRONG LOOKS LIKE

movement-since-last-period statements are genuinely derived and correct; a status change always carries an explanation; the timeline diagram regenerates cleanly week after week without manual repair.

BAND S2 · Multi-document, images, templates, chat, memory, export · a pull request into the public repository, use-cases folder

Post-merger integration playbook document set

WHO IT SERVES

Consulting / integration management office lead at an acquirer or an advisory firm

WHAT TO BUILD

Integration runs on a document set: the integration charter, workstream charters per function, day-one readiness checklists, a synergy tracking summary and the communications pack. The app generates the set from the deal parameters and the target's profile, keeps workstream charters consistent with the master charter, and updates readiness status from workstream inputs. Where two workstreams claim the same synergy or the same dependency, the conflict is flagged in both documents rather than silently duplicated.

WHAT STRONG LOOKS LIKE

cross-workstream conflicts are genuinely detected and surfaced in both places; master-to-workstream consistency holds after edits made on either side; the day-one checklist is specific to this deal rather than a generic list.

BAND S3 · Multi-document, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Multi-author deliverable harmoniser

WHO IT SERVES

Consulting / engagement manager assembling a final report written by six people

WHAT TO BUILD

Six authors produce sections in six voices, six heading styles and three template versions. The app merges them into one deliverable: consistent voice and tense, one heading hierarchy, unified terminology so the client is called one thing throughout, consistent number and date formatting, and a single template. Every change is presented as a before and after so authors can see what happened to their words, and substantive content is never altered in the name of style. A terminology report lists every term unified and what it was unified to.

WHAT STRONG LOOKS LIKE

substance is provably preserved while style converges, with a test showing no factual claim changed; the terminology report is complete; authors can accept or reject changes to their own sections individually.

BAND S2 · Multi-document, chat, templates, Review, search, export · a pull request into the public repository, use-cases folder

Quote and order form generator with approval routing

WHO IT SERVES

Sales / deal desk or revenue operations lead at a B2B software company

WHAT TO BUILD

A quote becomes a document: a pricing table with the configured line items, the discount and its approval status, term and renewal mechanics, and a non-standard terms schedule. The app generates it from the deal record, computes the approval level required from the discount and term, blocks export until the approvals are recorded, and produces the final order form with the approved terms merged in. Non-standard terms are extracted into a short summary the approver reads instead of the entire document. Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S2 · API, templates, Review, export · a pull request into the public repository, use-cases folder

Mutual action plan document

WHO IT SERVES

Sales / enterprise account executive running a complex deal

WHAT TO BUILD

The mutual action plan is a shared document listing every step to go live, with owners on both sides, dates, dependencies and a timeline diagram. The app drafts it from the deal context, updates it after each meeting from the account executive's notes (moving dates, marking steps complete, adding newly discovered steps), and regenerates the timeline. A customer-facing version omits internal notes. Each update produces a short "what changed" note the account executive can paste straight into the follow-up email.

WHAT STRONG LOOKS LIKE

internal and customer-facing versions can never leak into each other; the timeline diagram regenerates correctly as dates move; the change note is derived from the document rather than typed by hand.

BAND S1 · chat, images, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Per-stakeholder business case memo builder

WHO IT SERVES

Sales / account executive or solutions consultant arming an internal champion

WHAT TO BUILD

One deal needs three documents: a finance memo about payback and risk, a technical memo about architecture and security, and a one-page executive summary. The app builds all three from a single deal fact base so the numbers cannot diverge, tailors framing and depth per audience, embeds the value calculation as a table with visible assumptions, and produces a version the champion can circulate internally without vendor branding dominating. Changing an assumption regenerates every affected number across all three documents.

WHAT STRONG LOOKS LIKE

one fact base drives all three documents so an assumption change is reflected everywhere, with a test proving it; the technical memo is genuinely technical rather than marketing prose; assumptions are always visible and never buried in a footnote.

BAND S2 · Multi-document, chat, templates, images, export · a pull request into the public repository, use-cases folder

Renewal and true-up documentation engine

WHO IT SERVES

Sales / renewals manager or customer-success operations lead

WHAT TO BUILD

Ninety days before renewal the app assembles the pack: a usage and value summary from product data, the renewal quote, any true-up for overage with the calculation shown, the amendment document if terms change, and a talk track for the customer-success manager. Where usage exceeded the contracted entitlement, the true-up calculation is shown transparently with the governing contract clause quoted from the signed agreement. Batches run for the whole renewal cohort with genuine exceptions flagged for a human.

WHAT STRONG LOOKS LIKE

the governing clause is quoted from the actual signed contract rather than assumed from a template; true-up arithmetic is transparent and reproducible by the customer; the cohort batch surfaces real exceptions instead of dumping everything into review.

BAND S2 · API, Multi-document, search, templates, images, export · a pull request into the public repository, use-cases folder

Partner and reseller enablement kit generator

WHO IT SERVES

Sales / channel or partnerships manager

WHAT TO BUILD

Onboarding a partner means a document set: the partner agreement on the right tier template, a co-branded solution overview, a pricing and margin sheet, a demo script and an objection-handling sheet, all carrying the partner's logo and the correct regional terms. The app generates the set per partner, keeps co-branding consistent across every asset, regenerates the whole channel's collateral when product messaging or pricing changes, and reports which partners are still holding stale material.

WHAT STRONG LOOKS LIKE

co-branding applies correctly across every asset including the exported files; a messaging change regenerates precisely the affected assets and nothing else; the stale-material report is accurate per partner.

BAND S2 · Multi-document, templates, images, API, export · a pull request into the public repository, use-cases folder

Sales to customer-success handoff document

WHO IT SERVES

Sales / onboarding lead or customer-success manager receiving a closed deal

WHAT TO BUILD

What was promised during the sale is scattered across call notes, the proposal, the contract and email. The app builds the handoff document: what the customer bought, the commitments made including anything promised verbally that appears in notes, success criteria, stakeholders and their concerns, technical environment, and risks. Every commitment cites where it came from so the customer-success manager can verify it, and anything appearing in notes but not in the contract is flagged explicitly as an unpapered commitment.

WHAT STRONG LOOKS LIKE

the unpapered-commitment flag is the differentiator and has to be accurate in both directions; every commitment carries a source citation; the document stays short enough that a customer-success manager actually reads it.

BAND S2 · Multi-document, search, chat, export · a pull request into the public repository, use-cases folder

Win-loss debrief and quarterly competitive brief

WHO IT SERVES

Sales / product marketing or competitive intelligence lead

WHAT TO BUILD

Each closed deal produces a debrief from the interview or call transcript: why we won or lost, competitors present, pricing dynamics, objections raised and the deciding factor, all on a standard template so fields stay comparable. Quarterly, the app synthesizes the debrief set into a competitive brief: patterns by competitor and segment, wording that worked, losses attributable to a specific capability gap, with every claim linked to the debriefs supporting it. Confidential customer details are stripped from the shared version.

WHAT STRONG LOOKS LIKE

synthesis claims link to their supporting debriefs and small-sample patterns are labelled as small-sample; customer identity is verifiably stripped from the shared brief; the standard template keeps debriefs genuinely comparable across quarters.

BAND S2 · Multi-document, search, chat, memory, Review, export · a pull request into the public repository, use-cases folder

Annual and integrated report production line

WHO IT SERVES

Publishing / corporate communications or reporting manager at a listed or large private company

WHAT TO BUILD

The annual report joins a narrative front half (chair statement, strategy, business review, governance) to a financial back half, written by many contributors on a compressed timetable with several review rounds. The app holds the whole report as one linked document set, tracks which section sits at which review stage, enforces house style and the consistent use of figures that also appear in the financial statements, and produces the print-ready PDF plus a plain HTML version. A number-consistency report lists every figure appearing in more than one place with its values side by side.

WHAT STRONG LOOKS LIKE

the number-consistency check genuinely catches a figure that differs between the narrative and the accounts; review-round state survives across weeks and reloads; the print export is production quality with correct pagination and running heads.

BAND S3 · Multi-document, images, templates, search, Review, export · a pull request into the public repository, use-cases folder

Product launch documentation kit

WHO IT SERVES

Marketing / product marketing manager preparing a launch

WHAT TO BUILD

A launch needs a coordinated set: the messaging and positioning document, a customer-facing FAQ, a sales one-pager, the press release, the release notes and an internal enablement brief. The app generates the whole set from one source of truth (the positioning document) so claims cannot drift, tailors each artefact's depth and tone, and when positioning shifts during the run-up it proposes the specific edits needed in each downstream document. A claims register lists every substantive claim, where it appears, and whether it has been approved.

WHAT STRONG LOOKS LIKE

downstream propagation is surgical and the claims register is complete; the press release reads like a press release rather than a reformatted positioning document; unapproved claims can be blocked from the external artefacts entirely.

BAND S2 · Multi-document, chat, templates, images, Review, export · a pull request into the public repository, use-cases folder

Regulated marketing review workflow (claims to references)

WHO IT SERVES

Marketing / marketing operations lead in pharma, medical devices, financial services or any claims-regulated sector

WHAT TO BUILD

Every substantive claim in a piece of collateral must link to an approved reference and pass medical, legal and regulatory review. The app annotates the document, extracting each claim, matching it to the reference library, flagging claims with no support and claims whose cited reference does not actually say what is claimed, then routes the annotated document through the review stages with every reviewer comment and approval captured in the document itself. Approved pieces export with the reference list appended and a review certificate.

WHAT STRONG LOOKS LIKE

the "reference does not support this claim" detection is real and demonstrated rather than a presence check; reviewer comments and approvals persist against the exact claim across rounds; the certificate reconstructs the full review history for an auditor.

BAND S3 · search, Multi-document, Review, chat, templates, export · a pull request into the public repository, use-cases folder

Regulatory change to policy update pipeline

WHO IT SERVES

Compliance / compliance manager at a regulated firm

WHAT TO BUILD

A regulator publishes an amendment. The app ingests it, summarizes what changed, searches the firm's policy and procedure set for every section the change touches, and proposes the specific edits per document with the regulatory citation attached to each. An impact memo lists affected documents, owners and the deadline. Approved edits apply in place, keeping each policy's numbering and formatting intact, and the version history shows the change beside the regulation that caused it. Nothing applies without a named approver.

WHAT STRONG LOOKS LIKE

the affected-section search is complete, since missing a policy is the failure that matters, and precise enough to stay usable; every proposed edit carries the citation that justifies it; the version history reconstructs why a policy says what it says today.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Business continuity and disaster recovery plan builder

WHO IT SERVES

Compliance / resilience or business continuity manager

WHAT TO BUILD

The plan set is a business impact analysis, recovery strategies per process, the continuity plan with roles and call trees, per-system recovery runbooks with diagrams, and the annual test report. The app builds it from the process and system inventory, keeps recovery time objectives consistent between the impact analysis and the runbooks, generates the test scenario, and afterwards drafts the test report with findings plus the plan updates they imply. A distribution pack produces the printable crisis-card version.

WHAT STRONG LOOKS LIKE

recovery objectives stay consistent across documents and an inconsistency is reported rather than shipped; test findings map to the specific plan sections needing update; the crisis card is genuinely usable on paper during an outage with no systems available.

BAND S3 · Multi-document, images, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Same continuity and recovery plan set built from one process and system inventory with diagrams and a test cycle

Incident and breach notification factory

WHO IT SERVES

Compliance / privacy officer or incident response lead

WHAT TO BUILD

A confirmed incident triggers notifications with different content and different deadlines per regulator, per contract and per affected population. The app builds the whole set from the incident record: each regulator's required content in their required form, the individual notification in plain language and in the affected people's languages, the customer and partner notices the contracts require, and a holding statement. A deadline board tracks each obligation. The facts, which must be identical everywhere, derive from one source, and legal approves before anything exports.

WHAT STRONG LOOKS LIKE

obligation-specific required content is correct per regulator and traceable to the rule that demands it; the factual account is provably identical across every notification; the deadline board computes correctly from the discovery timestamp including non-business days.

BAND S3 · Multi-document, templates, search, Review, export · a pull request into the public repository, use-cases folder

Policy attestation pack generator

WHO IT SERVES

Compliance / compliance operations lead running the annual attestation cycle

WHAT TO BUILD

Nobody can meaningfully attest to a 90-page policy set. The app generates a role-specific extract per population, since what a warehouse worker must know differs from what a trader must know, produced from the master policies so an extract can never contradict its source, along with the attestation form, a short quiz where required, and a reading-time estimate. Attestation status is tracked in a per-population document, non-responders receive a personalized reminder, and a mid-cycle policy change identifies exactly which extracts and attestations are affected.

WHAT STRONG LOOKS LIKE

extracts are provably derived from the master so they cannot drift out of sync; role targeting is correct and defensible to an auditor; a mid-cycle policy change produces an accurate re-attestation list instead of resetting the whole population.

BAND S2 · Multi-document, search, templates, API, export · a pull request into the public repository, use-cases folder

Third-party risk assessment report generator

WHO IT SERVES

Compliance / third-party risk or vendor management analyst

WHAT TO BUILD

A vendor submits evidence: policies, certifications, penetration test summaries, questionnaire answers. The app reads it and drafts the assessment report with a finding per control domain (what the evidence shows, what is missing, residual risk), extracts certification expiry dates, compares the vendor's contract against the standard security and privacy clause set, and recommends specific contract language wherever a gap exists. The report closes with a risk rating and conditions of approval that the analyst approves before issue.

WHAT STRONG LOOKS LIKE

findings quote the vendor's own evidence and clearly separate "the evidence shows this" from "the evidence is silent on this"; the contract-gap comparison is accurate against the real agreement; recommended clause language is drop-in usable by counsel.

BAND S3 · Multi-document, search, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Control narrative and process flow generator from walkthroughs

WHO IT SERVES

Compliance / internal controls lead, financial-controls (SOX-style) manager, or an external auditor

WHAT TO BUILD

A walkthrough conversation with a process owner becomes the control documentation: a narrative describing the process end to end, a process flow diagram rendered inside the document, the control matrix (objective, activity, frequency, owner, evidence, assertion), and a list of the questions the walkthrough did not answer. When the process changes the following year, the app compares the new walkthrough against the existing narrative and proposes only the edits needed, with a change memo for the auditor.

WHAT STRONG LOOKS LIKE

the diagram and the narrative always agree because both derive from one description; unanswered questions are listed honestly instead of being filled with assumptions; the year-two update is surgical and the change memo is something an auditor can rely on.

BAND S2 · chat, images, Multi-document, templates, Review, export · a pull request into the public repository, use-cases folder

eCTD Module 2 summary assembler

WHO IT SERVES

Pharma / regulatory writer at a mid-size biotech or a CRO regulatory-affairs team

WHAT TO BUILD

Regulatory writers build Module 2 summaries (clinical overview, clinical summaries) by hand from the Module 5 study reports. Build a session that opens the study reports as attachments and the Module 2 shell as the active document, then drafts each subsection by pulling from and citing the source study, one subsection at a time. Every drafted paragraph carries a pointer back to the source report and section. Section numbering and cross-references survive repeated revision. Export a submission-shaped .docx with real footnotes, running headers and a live table of contents.

WHAT STRONG LOOKS LIKE

a reviewer can click any drafted claim and land on the exact source study section; re-running after one study report is replaced changes only the subsections that depended on it, and the diff proves it.

BAND S3 · Multi-document, search, chat, export, Review · a pull request into the public repository, use-cases folder

Clinical study report narrative builder

WHO IT SERVES

Life sciences / medical writer producing ICH E3 clinical study reports

WHAT TO BUILD

A CSR is a fixed shell with narrative sections that must agree exactly with the statistical output tables. Build an app that loads the CSR shell as the active document, takes the tables, listings and figures as attachments, and writes each narrative section against the numbers it must describe. When a table is re-issued, the app finds only the sentences that quote the changed values and rewrites those, leaving the rest of the 200-page report byte-identical. Numbers quoted in prose are checked against the source table before the edit is proposed.

WHAT STRONG LOOKS LIKE

a re-issued table triggers a review card with three sentence-level edits, not a rewritten section; a deliberately wrong number in the shell is caught and flagged rather than silently copied forward.

BAND S3 · chat, search, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Informed-consent readability and site localization app

WHO IT SERVES

Clinical operations / study start-up lead running a multi-country trial

WHAT TO BUILD

Consent forms must be readable at a set grade level and adapted per site and per language while keeping the sponsor's approved core text intact. Build a tool that takes the master ICF, marks the protected core sections, rewrites only the participant-facing explanatory sections to the target reading level, and then produces one localized version per site with the layout, headings and signature blocks preserved. Each output ships with a change report showing exactly which paragraphs were simplified and which were left untouched.

WHAT STRONG LOOKS LIKE

protected clauses come through the localization run character-for-character; the readability improvement is measured and reported per section, not claimed.

BAND S2 · chat, Multi-document, templates, Review, export · a pull request into the public repository, use-cases folder

Investigator's Brochure annual-update engine

WHO IT SERVES

Pharma / drug-safety and regulatory writer owning the annual IB refresh

WHAT TO BUILD

Each year the IB is re-issued with new nonclinical and clinical safety data, and reviewers want to see what moved. Build an app that opens last year's approved IB, takes the new data summaries as attachments, and produces the new edition as a set of targeted in-place edits rather than a fresh document. It emits a companion change-log document listing every edited section with the before and after text, and exports the new edition with track changes so the medical reviewer sees the redline in Word.

WHAT STRONG LOOKS LIKE

the unchanged 80 percent of the brochure is provably untouched (a byte-diff of the export shows it), and the change log matches the redline exactly.

BAND S2 · chat, Multi-document, Review, export, templates · a pull request into the public repository, use-cases folder

Pharmacovigilance case-narrative batch writer

WHO IT SERVES

Drug safety / PV case processor at a pharma company or safety vendor

WHAT TO BUILD

Every adverse-event case needs a written narrative built from structured case data in a fixed house style. Build a server-side worker that takes a batch of case records, drafts each narrative on the company template, and routes every one through review before it is finalized. Narratives that cite a lab value or a dechallenge outcome must trace back to the record field they came from. Medically significant cases get flagged for a human reviewer with the specific sentences that need attention; routine cases pass with one approval click.

WHAT STRONG LOOKS LIKE

throughput (cases per hour) reported alongside a reviewer-correction rate, and a deliberate data gap produces an honest "not stated in the source" sentence rather than an invented one.

BAND S3 · API, templates, Review, export · a pull request into the public repository, use-cases folder

Core-label change propagation across country packs

WHO IT SERVES

Regulatory labeling manager maintaining a product's label family across markets

WHAT TO BUILD

One safety change to the core company data sheet has to land in every country label (EU SmPC, patient leaflet, US prescribing information) without disturbing the local-only sections. Build an app that holds the label family as a document set, takes the approved core change, and proposes the matching edit in each local label in that label's own wording, section numbering and language. Local deviations are surfaced for a decision rather than overwritten, and each market gets its own export plus a submission-ready summary of what changed.

WHAT STRONG LOOKS LIKE

a market whose local wording legitimately differs is raised as a question, not silently harmonized; the run produces N exports and one consolidated change table.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

SOP-revision training packet and attestation generator

WHO IT SERVES

GxP quality / training coordinator at a manufacturing or clinical site

WHAT TO BUILD

When a controlled SOP is revised, every affected operator needs a training handout, a short assessment and a read-and-understand attestation record. Build a tool that compares the new SOP revision against the superseded one, extracts what actually changed, and generates a role-specific training handout covering only the deltas, plus a quiz and a signature-block attestation form on the site template. Each handout embeds the relevant procedure diagram. The pack exports as one PDF per role with the revision number in the running header.

WHAT STRONG LOOKS LIKE

the handout covers the deltas and nothing else, and the role split is real (an operator handout and a supervisor handout differ in content, not just in the title).

BAND S2 · chat, Multi-document, images, templates, export · a pull request into the public repository, use-cases folder

Photo-driven work-instruction refresher

WHO IT SERVES

Manufacturing engineer maintaining illustrated work instructions on a production line

WHAT TO BUILD

Work instructions go stale when the line changes and the photos no longer match the fixture. Build a tool where a technician attaches the new station photos and names the steps they belong to, and the app swaps the images in the existing work instruction in place, updates only the step text that the new photo contradicts, and bumps the revision block and the change history table. The rest of the document, including its formatting and step numbering, stays exactly as approved. Export the controlled PDF for the station tablet.

WHAT STRONG LOOKS LIKE

image replacement is a true swap into the existing layout (the caption, border and step numbering all survive), and step text changes only where the photo genuinely contradicts it.

BAND S2 · images, chat, Review, export · a pull request into the public repository, use-cases folder

Operator quick-reference card generator

WHO IT SERVES

Plant training lead turning OEM equipment manuals into shop-floor cards

WHAT TO BUILD

A 400-page OEM manual is useless at the machine. Build an app that ingests the manual, lets the user name a task ("clear a jam on the infeed", "weekly lubrication"), searches the manual for every section that bears on that task, and produces a two-page quick-reference card with the steps, the torque and safety values quoted exactly from the source, a simple flow diagram, and a page reference back into the manual for each step. Batch mode produces a card set for a whole machine.

WHAT STRONG LOOKS LIKE

every quoted specification value is verifiably lifted from the manual with a page pointer, and the card fits the page budget without dropping a required safety step.

BAND S2 · search, chat, images, export · a pull request into the public repository, use-cases folder

GHS safety data sheet reformatter and language pack

WHO IT SERVES

EHS / product-stewardship specialist at a chemicals or coatings manufacturer

WHAT TO BUILD

Supplier SDSs arrive in a dozen layouts and have to be normalized into the company's 16-section GHS format, then issued per market language. Build a service that ingests an incoming SDS in any format, maps its content into the house 16-section template without inventing or dropping any hazard statement, flags sections the source did not cover, and produces the localized versions with the section headings and hazard-statement numbering intact. Missing data is reported, never filled in.

WHAT STRONG LOOKS LIKE

a source that omits section 9 data produces an explicit gap report rather than plausible filler, and hazard-statement codes survive the reformat unaltered.

BAND S2 · API, templates, chat, export · a pull request into the public repository, use-cases folder

Factory acceptance test protocol builder

WHO IT SERVES

Commissioning engineer at an equipment builder or systems integrator

WHAT TO BUILD

FAT and commissioning protocols are long test documents with numbered steps, acceptance criteria, result columns and sign-off blocks. Build an app that generates the protocol from the equipment specification and the customer's requirement list, keeping each test step traceable to the requirement it verifies. During execution, the engineer attaches result photos and dictates outcomes, and the app fills the result cells and the deviation log in place. The signed-off protocol exports as the record copy with photos embedded and the requirement traceability table current.

WHAT STRONG LOOKS LIKE

every requirement maps to at least one test step and the traceability table is generated, not hand-maintained; filling a result cell does not disturb the table's formatting.

BAND S3 · chat, templates, images, Review, export · a pull request into the public repository, use-cases folder

Part-change impact sweep across the document set

WHO IT SERVES

Configuration-management or engineering-change owner at a manufacturer

WHAT TO BUILD

When a part number or supplier changes, nobody knows which of 300 procedures, manuals and inspection sheets mention it. Build a tool that holds the controlled document set, searches it for every reference to the changed part (including descriptive references, not just the exact number), produces an impact report listing document, section and context, and then applies the approved replacement as a targeted edit in each affected document. Documents with no hit are never opened or rewritten. Every applied edit lands in a per-document change history.

WHAT STRONG LOOKS LIKE

the sweep finds references phrased differently from the search term, the impact report has no false positives after review, and untouched documents are provably untouched.

BAND S3 · search, Multi-document, Review, API, export · a pull request into the public repository, extensions folder

Calibration and maintenance record-book generator

WHO IT SERVES

Maintenance planner or metrology technician keeping auditable equipment records

WHAT TO BUILD

Auditors want a per-asset record book: calibration certificates, maintenance entries, out-of-tolerance investigations and signature blocks, in date order with a cover index. Build a tool that takes the maintenance log data and the scanned certificates, assembles the record book per asset, writes the out-of-tolerance investigation narrative from the recorded readings, and keeps the running index current as entries are appended. Appending a new entry edits the index and the log table in place instead of rebuilding the book.

WHAT STRONG LOOKS LIKE

appending the 200th entry is as cheap as the second, and the scanned certificates round-trip into the export without being flattened to text.

BAND S2 · API, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Letter-of-credit discrepancy checker and document fixer

WHO IT SERVES

Trade-finance or export-documentation officer at a freight forwarder, trading house or bank

WHAT TO BUILD

A letter of credit is rejected when the presented documents disagree with its terms or with each other, and each discrepancy costs money and days. Build an app that opens the credit alongside the presentation set (commercial invoice, bill of lading, packing list, certificate of origin, insurance certificate), checks every named term (description of goods, consignee, ports, latest shipment date, amounts, weights) across all of them, produces a discrepancy report with the exact conflicting text quoted from each document, and then applies the agreed corrections in place, document by document.

WHAT STRONG LOOKS LIKE

a seeded discrepancy set is caught with zero misses and quoted precisely from each source; corrections change only the conflicting field and never reflow the document.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Freight-claim packet assembler

WHO IT SERVES

Claims coordinator at a shipper, 3PL or freight forwarder

WHAT TO BUILD

A cargo damage claim needs a claim letter that references the proof of delivery exception, the commercial invoice value, the photos and the repair or salvage quote, all bundled with an index. Build a tool where the coordinator attaches the evidence and the app writes the claim letter quoting the exception notation and the amounts exactly as they appear in the source documents, embeds the damage photos with captions, assembles the indexed packet, and produces the carrier-specific cover page. A follow-up chaser letter reuses the same packet with only the dates and tone changed.

WHAT STRONG LOOKS LIKE

amounts and reference numbers in the letter always match the attached evidence, and the follow-up letter is a targeted edit of the original rather than a fresh draft.

BAND S2 · chat, Multi-document, images, templates, export · a pull request into the public repository, use-cases folder

Tariff classification memo builder

WHO IT SERVES

Customs broker or trade-compliance analyst defending an HS classification

WHAT TO BUILD

Classification decisions have to be written up defensibly: product description, composition, essential character reasoning, the rules of interpretation applied, and prior rulings relied on. Build an app where the analyst attaches the product spec sheet, photos and the rulings they are relying on, and the app drafts the classification memo in the house format, quoting the ruling language it uses and embedding the product image. Alternative classifications considered and rejected get their own section. Reused reasoning is drawn from the firm's prior memos rather than rewritten.

WHAT STRONG LOOKS LIKE

the memo distinguishes what the source rulings actually say from the analyst's reasoning, and reused reasoning is cited to the prior memo it came from.

BAND S2 · chat, search, templates, images, export · a pull request into the public repository, use-cases folder

Booking to export-packet generator with cross-document consistency

WHO IT SERVES

Export documentation clerk at a manufacturer or forwarder

WHAT TO BUILD

One booking produces five documents that must agree: shipping instruction, commercial invoice, packing list, certificate of origin and the draft bill of lading. Build a service that generates all five from the booking record on the house templates, then runs a consistency pass across the finished set on consignee, marks and numbers, gross and net weights, package counts, HS codes and incoterm, and reports any field where the documents disagree before anything is sent. A correction applied once propagates to every document that carries that field.

WHAT STRONG LOOKS LIKE

the consistency pass is a genuine cross-document check (seed a mismatched weight and it is caught), and a single correction updates all carriers of that field in one reviewable batch.

BAND S2 · API, Multi-document, templates, export, Review · a pull request into the public repository, use-cases folder

Demurrage and detention dispute letter generator

WHO IT SERVES

Ocean freight cost-control analyst disputing terminal charges

WHAT TO BUILD

Disputing demurrage means writing a letter that reconstructs the timeline from gate-in and gate-out records, the free-time terms in the service contract, and the events that stopped the clock (port closure, customs hold, carrier delay). Build a tool that takes the charge invoice, the equipment interchange records and the contract, builds the timeline as a rendered diagram inside the letter, cites the contract clause that governs free time, and states the calculated amount in dispute. Batch mode produces one letter per container with the shared reasoning stated once.

WHAT STRONG LOOKS LIKE

the timeline diagram is generated from the records rather than hand-drawn, and the disputed amount is arithmetically reproducible from the numbers quoted in the letter.

BAND S2 · chat, Multi-document, images, search, export · a pull request into the public repository, use-cases folder

Marketplace note: if your build wants a marketplace listing, build it and sideload it; we handle actual listings.

Export-licence application and screening memo assembler

WHO IT SERVES

Trade-compliance officer at an exporter of controlled or dual-use goods

WHAT TO BUILD

An export licence application needs a technical description written for a licensing officer, an end-use and end-user statement, and a screening memo recording the checks performed on the parties. Build an app that takes the engineering specification and the customer paperwork, drafts the licensing-officer-facing technical description at the right level of detail (neither marketing copy nor a full spec dump), assembles the end-use statement on the agency template, and records the screening results as a dated memo with the evidence attached. Prior approved applications are reused as precedent for wording.

WHAT STRONG LOOKS LIKE

the technical description is genuinely rewritten for the reader rather than pasted from the spec, and precedent reuse is shown with a pointer to the prior application it came from.

BAND S2 · chat, templates, Multi-document, search, export · a pull request into the public repository, use-cases folder

Environmental incident report in regulator format

WHO IT SERVES

HSE coordinator at a utility, pipeline operator or upstream producer, filing against a deadline

WHAT TO BUILD

A reportable spill or release triggers a filing on the agency's own form-shape within a fixed number of hours, built from field notes, photos and sampling results. Build an app that takes the field notes as dictated or typed text plus the site photos, drafts the narrative into the agency's required section order and headings, places the photos with captions and timestamps, and flags every field the regulator requires that the field notes did not supply. A follow-up amended report is produced as a targeted edit of the original with an amendment log.

WHAT STRONG LOOKS LIKE

missing mandatory fields are named explicitly before submission rather than filled with plausible text, and the amended report shows exactly what changed from the initial filing.

BAND S2 · chat, images, templates, Review, export · a pull request into the public repository, use-cases folder

AFE and well-file packet builder

WHO IT SERVES

Drilling or completions engineer preparing an authorization for expenditure package for partners

WHAT TO BUILD

An AFE package bundles the cost estimate, the drilling programme narrative, the geological prognosis, the location plat and the partner cover letter, and each partner needs their working-interest share reflected. Build a tool that assembles the packet from the estimate data and the technical inputs, writes the programme narrative from the engineer's bullet notes, embeds the location exhibit, and produces per-partner cover letters with the correct interest share and election deadline. A revised estimate updates the cost table and every letter that quotes a number from it.

WHAT STRONG LOOKS LIKE

partner-specific numbers are never crossed, and a revised estimate propagates to every quoting document in one reviewable batch.

BAND S2 · API, Multi-document, templates, chat, export · a pull request into the public repository, use-cases folder

Marketplace note: if your build wants a marketplace listing, build it and sideload it; we handle actual listings.

Rate-case testimony and exhibit assembler

WHO IT SERVES

Regulatory affairs lead at a utility filing a rate case with a state commission

WHAT TO BUILD

Rate-case filings are enormous: prepared testimony per witness, numbered exhibits, schedules, and a strict commission format including line numbering and Q-and-A structure. Build an app that drafts and edits prepared testimony in the commission's format, keeps every exhibit reference in the testimony synchronized with the exhibit index as exhibits are added or renumbered, and produces witness-ready and filing-ready exports. Editing a single answer must not renumber the rest of the testimony or break a cross-reference.

WHAT STRONG LOOKS LIKE

inserting an exhibit in the middle renumbers the index and every in-text reference correctly and touches nothing else; line numbering and Q-and-A structure survive round-trip to Word and PDF.

BAND S3 · chat, Multi-document, search, export, Review · a pull request into the public repository, use-cases folder

Storm restoration after-action report builder

WHO IT SERVES

Emergency-response manager at an electric or water utility

WHAT TO BUILD

After a major outage event, the utility owes the commission and its board a report covering the timeline, crew mobilization, restoration curve, customer communications and lessons learned. Build a tool that takes the event log, crew records and customer-communication archive, renders the restoration timeline as a diagram inside the document, drafts each narrative section against the logged facts, and separates what the records show from what the team concluded. Lessons-learned items are carried forward into the next event's report as open or closed.

WHAT STRONG LOOKS LIKE

every timeline claim in the narrative is traceable to a log entry, and open actions genuinely carry forward across two consecutive event reports.

BAND S2 · chat, images, Multi-document, memory, export · a pull request into the public repository, use-cases folder

Integrity-management plan regulation sweep

WHO IT SERVES

Pipeline or asset-integrity engineer maintaining a regulator-approved plan

WHAT TO BUILD

When a standard or rule is amended, the operator has to find every place the plan relied on the old requirement and update it, then evidence the review. Build a tool that holds the plan and the referenced standards, takes the amendment text, searches the plan for every dependent provision (including ones that paraphrase rather than quote the rule), proposes the specific edits, and produces a review record listing every provision examined with the disposition. Provisions found compliant are recorded as reviewed and left alone.

WHAT STRONG LOOKS LIKE

paraphrased dependencies are found, not just literal citations, and the review record covers examined-and-unchanged provisions as well as edited ones.

BAND S3 · search, Multi-document, chat, Review, export · a pull request into the public repository, use-cases folder

Right-of-way landowner notice batch

WHO IT SERVES

Land agent or right-of-way coordinator on a transmission, pipeline or fibre project

WHAT TO BUILD

Hundreds of landowners each need a legally correct notice carrying their own parcel description, access dates, compensation terms and the statutory language for their jurisdiction, often in more than one language. Build a service that merges the parcel register into the notice template with the legal description reproduced exactly, selects the jurisdictional statutory paragraph, produces each notice in the landowner's preferred language with the layout preserved, and outputs a mailing manifest. Any parcel whose register data is incomplete is held back and reported rather than sent with a gap.

WHAT STRONG LOOKS LIKE

legal descriptions are reproduced character-for-character, incomplete records are quarantined rather than guessed, and the non-English notices keep the same layout as the English ones.

BAND S2 · API, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Job safety analysis and toolbox-talk generator

WHO IT SERVES

Field supervisor or HSE advisor preparing daily crew paperwork

WHAT TO BUILD

Every task needs a JSA listing steps, hazards, controls and required PPE, plus a short toolbox talk the crew signs. Build a tool where the supervisor describes the day's task in plain language and the app builds the JSA from the company hazard library, keeps the control wording exactly as the safety team approved it, adds a task-specific talk with a signature block, and embeds a simple site sketch when the supervisor drops one in. Reused tasks pull last time's JSA and only change what the supervisor says is different today.

WHAT STRONG LOOKS LIKE

approved control wording is never paraphrased, and the reuse path is a real targeted edit of the prior JSA rather than a regeneration.

BAND S1 · chat, templates, images, memory, export · a pull request into the public repository, use-cases folder

Budget justification and compliance narrative builder

WHO IT SERVES

Grants manager at a nonprofit applying for institutional or government funding

WHAT TO BUILD

The budget justification has to explain every line of the budget in prose, and the compliance narrative has to answer the funder's standard assurances. Build an app that takes the budget table and writes a justification paragraph per line, keeping the totals in the prose arithmetically consistent with the table, then assembles the assurances narrative from the organization's approved policy language. When a budget line changes, only the paragraphs that quote it are rewritten and the totals sentence is corrected. Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S2 · chat, templates, search, Review, export · a pull request into the public repository, use-cases folder

Multi-funder report splitter

WHO IT SERVES

Programme director reporting one programme to several funders with different formats

WHAT TO BUILD

One programme's results have to be told six ways: each funder wants its own section order, indicator names, period and page limit. Build a tool that holds the programme's result data and the funder templates, generates each funder's report in that funder's structure and vocabulary, and then cross-checks the numbers across all six so the same indicator never appears with two values. Reports over the funder's page limit are tightened section by section rather than truncated, and each output ships with the indicator mapping used.

WHAT STRONG LOOKS LIKE

the cross-report number check is real (seed a divergence and it is caught), and page-limit compliance is achieved by editing, not by cutting the last section.

BAND S3 · Multi-document, templates, chat, export, Review · a pull request into the public repository, use-cases folder

Theory-of-change and logic-model document builder

WHO IT SERVES

Monitoring and evaluation lead or programme designer

WHAT TO BUILD

Funders ask for a logic model or theory of change as a diagram plus a narrative explaining each link. Build an app where the team describes inputs, activities, outputs, outcomes and assumptions in conversation, and the app renders the model as a diagram inside the document and writes the accompanying narrative with the causal assumptions stated explicitly. Editing one arrow ("add a community-mobilization activity feeding the retention outcome") rewrites the diagram in place and updates only the affected narrative paragraph.

WHAT STRONG LOOKS LIKE

the diagram edit is a genuine in-place rewrite of the diagram source rather than a regenerated picture, and the narrative and diagram never drift apart.

BAND S1 · chat, images, export · a pull request into the public repository, use-cases folder

Field-office policy pack localizer

WHO IT SERVES

Operations or safeguarding lead rolling policies out to country offices

WHAT TO BUILD

Safeguarding, security and volunteer policies must reach every field office in the working language, with the non-negotiable core intact and the local annexes swapped per country (reporting contacts, legal references, local escalation paths). Build a tool that marks the protected core, generates each country pack with its own annex and language, produces the acknowledgement form, and issues a change notice when the core is amended so offices see what moved rather than a whole new document.

WHAT STRONG LOOKS LIKE

protected core text is identical across all packs, annexes are genuinely country-specific, and a core amendment produces a per-country change notice rather than a full reissue.

BAND S2 · Multi-document, templates, chat, export · a pull request into the public repository, use-cases folder

Annual report and public-filing narrative companion

WHO IT SERVES

Communications and finance pair producing the yearly public-facing report

WHAT TO BUILD

The annual report has to tell the story and agree with the audited numbers and the public filing schedules. Build an app that takes the financial schedules and the programme results, drafts the narrative and the programme highlights on the brand template with generated charts and photos placed, and then verifies every figure quoted in the prose against the schedules. Last year's report is used as the layout source so the new edition matches it exactly, with only content and figures changed.

WHAT STRONG LOOKS LIKE

the layout of the previous edition is reproduced faithfully (not approximated), and every prose figure ties to a schedule with a check the reviewer can see.

BAND S2 · chat, images, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Beneficiary story consent and anonymization workflow

WHO IT SERVES

Communications officer publishing field stories under a safeguarding policy

WHAT TO BUILD

A publishable story needs a signed consent form, identifying details removed to the policy's standard, and a version history proving what was removed. Build a workflow that generates the consent form in the participant's language, ingests the raw story and interview notes, proposes every redaction or pseudonymization for a human decision, and produces both the publishable version and a sealed internal version. The redaction review card shows the original text next to the proposed replacement so nothing is removed silently.

WHAT STRONG LOOKS LIKE

every proposed redaction is reviewable and reversible, indirect identifiers (a named school, a rare condition) are caught alongside names, and the two versions never get mixed up.

BAND S2 · Review, chat, search, templates, export · a pull request into the public repository, use-cases folder

Proposal boilerplate freshness auditor

WHO IT SERVES

Bid or development manager reusing organizational boilerplate across many proposals

WHAT TO BUILD

Capability statements, staffing tables and impact figures get copied into every proposal and quietly go stale. Build a tool that holds the boilerplate library and the submitted proposals, finds every instance of a fact that has changed (headcount, country count, audited figures, certifications), reports where each stale claim currently lives, and applies the correction across the library and any live draft in one reviewable batch. Facts are stored once with their source and date.

WHAT STRONG LOOKS LIKE

the audit finds paraphrased versions of a stale fact, not just exact copies, and applying the fix updates the library and drafts without disturbing the surrounding proposal text.

BAND S2 · search, Multi-document, memory, Review, export · a pull request into the public repository, extensions folder

Fact-check annotation workflow

WHO IT SERVES

Fact-checker or research editor at a magazine, broadcaster or documentary team

WHAT TO BUILD

Build an app that reads a draft, extracts every checkable assertion, and anchors each one as a comment on the exact sentence it belongs to, with a status (verified, unsupported, needs source) and the source the checker attached. The writer sees the comments in the document itself and in the exported Word file as real comments. Resolved checks are cleared and logged; unresolved ones block the export of the publish-ready version, and the check log becomes the record for the legal read.

WHAT STRONG LOOKS LIKE

comments are real comments that survive the round-trip to Word, anchors stay attached to the right sentence after the writer edits around them, and the block on unresolved checks is enforced rather than advisory.

BAND S2 · chat, search, Review, Multi-document, export · a pull request into the public repository, use-cases folder

Call sheet and shooting schedule generator

WHO IT SERVES

First assistant director or production coordinator on a film, commercial or series

WHAT TO BUILD

Build a tool that takes a script breakdown (scenes, cast, locations, day or night, special requirements) and produces the daily call sheet and the shooting schedule on the production's template, with per-department call times, the weather and sunrise or sunset block, the map thumbnail, and the cast list numbered as the breakdown numbers them. A schedule change (a location drops, a scene moves) updates only the affected day's call sheet and the strip order, and reissues carry a revision letter and a change summary for the crew.

WHAT STRONG LOOKS LIKE

a mid-week location change reissues one call sheet with a change summary rather than rebuilding the week, and the template's dense table layout survives every reissue.

BAND S2 · chat, templates, images, Multi-document, export · a pull request into the public repository, use-cases folder

Chain-of-title and clearance binder assembler

WHO IT SERVES

Production legal or distribution coordinator delivering a title to a buyer

WHAT TO BUILD

Delivery requires a binder proving the chain of title: option and purchase agreements, writer and talent deals, music licences, clip clearances and errors-and-omissions paperwork, with a cover index and a gap memo naming what is still missing. Build a tool that ingests the executed documents, builds the index with document dates and parties read from the documents themselves, writes the chain narrative, and reports gaps against the buyer's delivery schedule. Adding a late-arriving licence updates the index, the narrative and the gap memo together.

WHAT STRONG LOOKS LIKE

parties and dates in the index are read from the documents rather than typed by the user, and the gap memo is checked against the buyer's actual delivery schedule, not a generic list.

BAND S2 · Multi-document, search, chat, export · a pull request into the public repository, use-cases folder

Blameless incident postmortem builder

WHO IT SERVES

Incident commander or SRE writing the postmortem after a production incident

WHAT TO BUILD

Build a tool that assembles the postmortem from the incident's raw trail: the alert timeline, the response channel transcript, the deploy and change records, and the customer-impact numbers. It renders the timeline as a diagram inside the document, drafts the contributing-factors and remediation sections in the team's template, and keeps a strict separation between what the records show and what the team concluded. Action items become a tracked table that carries forward to the next review, and the customer-facing summary is generated as a separate document from the same facts.

WHAT STRONG LOOKS LIKE

every timeline entry is traceable to a source record, speculation is visibly labelled as such, and the internal and customer-facing documents cannot contradict each other because they are generated from one fact set.

BAND S2 · chat, images, Multi-document, templates, export · a pull request into the public repository, use-cases folder

API spec diff to integration guide updater

WHO IT SERVES

Developer-experience engineer maintaining partner integration guides

WHAT TO BUILD

When the API specification changes, the human-written integration guide silently rots. Build a tool that takes two versions of the specification, computes what actually changed (new fields, removed endpoints, changed defaults, tightened validation), then finds the passages in the guide that describe the changed behaviour and proposes surgical rewrites with the code samples updated to match. Purely additive changes produce a new section placed correctly in the guide's structure. The run outputs a reviewable change set plus a partner-facing "what changed" note.

WHAT STRONG LOOKS LIKE

a renamed field is found in prose that never names it literally, code samples are updated consistently, and untouched sections of the guide are byte-identical after the run.

BAND S3 · search, chat, Review, export, API · a pull request into the public repository, extensions folder

Documentation drift detector

WHO IT SERVES

Staff engineer or docs owner on a codebase where documentation lags behind reality

WHAT TO BUILD

Build a job that runs against a documentation set and a change feed (release notes, changed configuration defaults, deprecated flags, renamed commands) and answers one question: which documentation sections are now wrong? For every hit it reports the document, the section, the sentence, and the evidence of the drift, then proposes a targeted fix for review. Confidence is separated from certainty: sections it cannot verify are listed as unknown rather than claimed correct. Approved fixes are applied in place and exported back to the docs pipeline.

WHAT STRONG LOOKS LIKE

measured precision and recall against a seeded drift corpus, an explicit unknown bucket, and applied fixes that never touch a neighbouring paragraph.

BAND S3 · search, Multi-document, Review, API, export · a pull request into the public repository, extensions folder

Dependency notice and attribution document generator

WHO IT SERVES

Release engineer or open-source compliance owner shipping a product release

WHAT TO BUILD

Every release needs an attribution document listing dependencies, versions, licences and required notice text, formatted for legal review and for the product's about screen. Build a tool that takes the dependency inventory, groups components by licence family, reproduces each required notice text exactly as the licence demands, writes the summary sections legal wants, and produces both the long legal document and the short in-product notice. A new release edits only the entries that changed and lists the deltas for the reviewer.

WHAT STRONG LOOKS LIKE

required notice text is reproduced verbatim with no reflowing or paraphrase, and release-to-release runs produce a delta list rather than a fresh unreadable document.

BAND S1 · API, templates, export, chat · a pull request into the public repository, extensions folder

Technical brief translator for executive readers

WHO IT SERVES

Principal engineer or architect who has to get a decision from a non-engineering board

WHAT TO BUILD

Build a tool that takes a deep technical document (a design doc, an architecture review, a security assessment) and produces the executive brief: the decision being asked for, the options with their trade-offs in business terms, cost and risk, and a simplified version of the architecture diagram. The source document is not modified. Terms of art are replaced with plain language while the numbers stay exact, and every claim in the brief links back to the section of the source it came from so the author can defend it.

WHAT STRONG LOOKS LIKE

the brief survives a hostile read (no claim without a source pointer, no number that drifted from the original), and the simplified diagram is genuinely derived from the source diagram.

BAND S2 · search, chat, images, Multi-document, export · a pull request into the public repository, use-cases folder

Code-safe developer-docs localization pipeline

WHO IT SERVES

Documentation platform owner shipping docs in several languages

WHAT TO BUILD

Translating developer documentation breaks when code blocks, command flags, identifiers, file paths and product terms get translated too. Build a pipeline that localizes a documentation set while treating code fences, inline code, identifiers and a maintained do-not-translate glossary as protected content, keeps the page structure and anchors intact so links keep working, and produces a per-page report of protected spans and any term it was unsure about. Source updates re-translate only the changed paragraphs, not the whole page.

WHAT STRONG LOOKS LIKE

code and identifiers come through byte-identical across every language, anchors and cross-links still resolve, and an incremental source change re-translates only what moved.

BAND S3 · API, Multi-document, templates, export, memory · a pull request into the public repository, extensions folder

Deprecation campaign kit

WHO IT SERVES

Product or platform manager retiring an API version or a feature

WHAT TO BUILD

A deprecation is not one document, it is a wave: a migration guide, a changelog entry, in-page deprecation banners across every affected doc page, customer email variants by usage tier, and an internal support brief. Build a tool that takes the deprecation decision and the affected surface list and produces all of them from one fact set, so the dates, replacement names and migration steps cannot disagree. Extending the sunset date edits every artifact that quotes it in a single reviewable batch.

WHAT STRONG LOOKS LIKE

changing the sunset date once updates every carrier of that date and nothing else, and the banner insertion lands in the right place on each page without reformatting it.

BAND S2 · Multi-document, chat, templates, Review, export · a pull request into the public repository, extensions folder

Executive-search dossier book builder

WHO IT SERVES

Researcher or consultant at an executive-search firm presenting a shortlist

WHAT TO BUILD

A search shortlist is presented as a confidential book: a firm-branded cover, the search brief, then one dossier per candidate with a photograph, career narrative, compensation summary, referencing notes and a fit assessment against the brief's criteria, ending with a comparison matrix. Build an app that assembles the book from candidate materials and the consultant's notes, writes each fit assessment against the actual brief criteria, keeps every candidate's section formatting identical, and produces both the client copy and an internal copy with the referencing notes retained.

WHAT STRONG LOOKS LIKE

the client copy provably contains no internal-only content, per-candidate formatting is genuinely uniform, and the comparison matrix is derived from the individual assessments rather than written separately.

BAND S2 · Multi-document, templates, images, chat, export · a pull request into the public repository, use-cases folder

Placement compliance packet per jurisdiction

WHO IT SERVES

Compliance or operations manager at a staffing agency placing contractors across regions

WHAT TO BUILD

Every placement needs a bundle whose contents depend on the worker's location and engagement type: assignment confirmation, contract or schedule, right-to-work record, statutory notices, timesheet and expense instructions and any local mandatory disclosure. Build a service that selects the correct document set for the jurisdiction, fills each one from the placement record, produces the bundle with a checklist cover page, and refuses to issue a bundle whose mandatory items are incomplete, naming the missing item instead of shipping a gap.

WHAT STRONG LOOKS LIKE

jurisdiction selection is data-driven and auditable, incomplete bundles are blocked with a named reason, and adding a new jurisdiction is a content change rather than a code change.

BAND S2 · API, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Candidate feedback and decline letter desk

WHO IT SERVES

Recruiting coordinator handling high-volume candidate communications

WHAT TO BUILD

Most candidates get a generic decline, and the ones who deserve real feedback do not get it because writing it does not scale. Build a tool that drafts each message from the interview record at the right level of specificity for the stage reached, in the candidate's language, on the employer's template, while enforcing the legal review rules on what may and may not be said. Messages that reference a specific assessment go through review; standard early-stage declines do not. Every sent message is retained as a record.

WHAT STRONG LOOKS LIKE

the routing rule between reviewed and auto-sent messages is explicit and enforced, prohibited phrasing never survives review, and non-English messages read natively rather than as translations.

BAND S2 · API, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Contract-to-hire conversion paperwork bundle

WHO IT SERVES

Account manager at a staffing firm converting a contractor to a client's payroll

WHAT TO BUILD

A conversion triggers a set of interlocking documents: the conversion notice to the client with the fee calculation shown, the amended or terminated contractor agreement, the client-facing confirmation, and the internal record. Build a tool that generates the set from the placement and contract terms, reads the conversion-fee formula from the governing agreement rather than assuming it, shows the calculation in the notice, and keeps dates and party names consistent across all four documents. A negotiated fee change updates every document that quotes it.

WHAT STRONG LOOKS LIKE

the fee formula is genuinely read out of the governing contract and quoted back, and a renegotiated number updates all carriers in one batch with nothing else touched.

BAND S2 · Multi-document, search, templates, chat, export · a pull request into the public repository, use-cases folder

Internal mobility and promotion packet builder

WHO IT SERVES

HR business partner running promotion or internal transfer cycles

WHAT TO BUILD

A promotion case needs a written business case against the level criteria, an updated job description, a compensation letter, and a manager talking-points sheet, all consistent with each other and with the levelling framework. Build a tool that takes the manager's evidence and drafts the case against the framework's actual criteria (flagging criteria with no evidence supplied), updates the existing job description in place rather than rewriting it, and produces the letter and talking points from the approved outcome. Declined cases produce a development-plan document instead.

WHAT STRONG LOOKS LIKE

unevidenced criteria are named rather than papered over, the job description is edited in place with its formatting intact, and the approved and declined paths both produce complete artifacts.

BAND S2 · chat, templates, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Panel debrief to defensible hiring decision memo

WHO IT SERVES

Hiring manager or talent lead in a regulated or high-scrutiny hiring process

WHAT TO BUILD

Build a tool that ingests each interviewer's written notes and scores and produces the decision memo: the criteria, what each panel member observed, where they disagreed, the decision and the reasoning, with every assertion tied to the note it came from. It refuses to smooth over disagreement, surfacing split assessments explicitly. The memo goes through panel review before it is finalized, and the finalized record exports with the panel's approvals attached for the audit file.

WHAT STRONG LOOKS LIKE

disagreement between panel members survives into the memo rather than being averaged away, and no assertion appears without a traceable note behind it.

BAND S2 · Multi-document, chat, Review, export · a pull request into the public repository, use-cases folder

Bid tabulation and award recommendation memo

WHO IT SERVES

Procurement officer or category manager closing a competitive tender

WHAT TO BUILD

After the bids are in, the buyer needs a scored comparison and a defensible award recommendation. Build an app that opens all bid responses at once, extracts the priced and scored elements against the tender's own evaluation criteria, builds the tabulation table, and drafts the award memo explaining why the winner won in the criteria's language, with each scoring statement quoting the bidder's own words. Non-responsive bids are identified with the specific requirement they failed. Unsuccessful-bidder debrief letters are generated from the same analysis.

WHAT STRONG LOOKS LIKE

scores trace to quoted bid text, a non-responsive bid is caught on a requirement no keyword search would find, and the debrief letters cannot contradict the award memo.

BAND S3 · Multi-document, search, chat, templates, export · a pull request into the public repository, use-cases folder

Supplier renewal and price-increase response desk

WHO IT SERVES

Procurement manager handling a wave of supplier price-increase notices

WHAT TO BUILD

Each supplier's notice has to be answered against that supplier's own contract: what the escalation clause allows, the notice period, the index it is tied to and the cap. Build a tool that opens the contract alongside the notice, extracts the governing terms, drafts a response that accepts, disputes or counters with the clause quoted back, and produces an internal summary of exposure across the whole wave. Suppliers whose notice breaches their own contract are flagged first with the breach quoted.

WHAT STRONG LOOKS LIKE

the escalation terms are genuinely read from each contract rather than assumed, non-compliant notices are surfaced with the exact clause, and the portfolio exposure summary reconciles with the individual letters.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Buy-side RFP authoring from requirements intake

WHO IT SERVES

Sourcing lead building the tender document itself

WHAT TO BUILD

Build an app that turns a stakeholder requirements intake into the tender document: scope of work, evaluation criteria with weightings, response templates the bidders must complete, mandatory terms and the submission instructions. Requirements are traced through to the evaluation criteria so nothing important goes unscored, and the response template is generated from the criteria so bids arrive comparable. The app checks the finished tender for contradictions between sections (a deadline stated twice, a criterion with no weighting) before it is issued.

WHAT STRONG LOOKS LIKE

every intake requirement is traceable to a scope clause and an evaluation criterion, and the internal-contradiction check finds seeded conflicts across distant sections.

BAND S2 · chat, templates, Multi-document, search, export · a pull request into the public repository, use-cases folder

Supplier code-of-conduct and ESG questionnaire cycle

WHO IT SERVES

Responsible-sourcing or supplier-compliance manager running an annual supplier attestation

WHAT TO BUILD

Build a tool that issues the supplier code of conduct and the ESG questionnaire per supplier tier and region (localized, with the tier-specific annexes), ingests the returned responses in whatever format suppliers send them, normalizes each into the internal assessment document, and drafts the follow-up letter for suppliers whose answers fall short, quoting their own response back. Aggregate findings become a programme report with the response-rate and risk profile shown as charts.

WHAT STRONG LOOKS LIKE

returns in several formats normalize into one comparable assessment shape, follow-up letters quote the supplier's actual answer, and the aggregate report reconciles with the individual assessments.

BAND S2 · API, Multi-document, templates, images, export · a pull request into the public repository, use-cases folder

Purchase-order terms conflict checker

WHO IT SERVES

Contract manager or procurement counsel dealing with the battle of the forms

WHAT TO BUILD

A purchase order carries the buyer's terms, the supplier's acknowledgement carries theirs, and nobody reads either until something goes wrong. Build a tool that compares the two term sets clause area by clause area (warranty, liability, delivery, title and risk, payment, IP, governing law), produces a conflict memo quoting both sides on each contested point, states which document is likely to govern under the sequence of events described, and drafts the objection or confirmation letter that preserves the buyer's position.

WHAT STRONG LOOKS LIKE

the comparison is by substance rather than by heading (matching terms that use different section names), and silence on a topic in one document is reported as a distinct finding from a conflict.

BAND S3 · Multi-document, search, chat, export · a pull request into the public repository, use-cases folder

Adverse-action and decision letter engine

WHO IT SERVES

Consumer-lending operations or compliance owner sending decline and counter-offer notices

WHAT TO BUILD

Decline letters have to state the actual principal reasons in approved language, include the required disclosures, and never say more than the regulation permits. Build a service that generates each letter from the decision record, maps reason codes to the approved wording without paraphrasing it, assembles the required notices for the product and jurisdiction, and produces the customer's language version with the layout preserved. Letters whose decision record contains a reason with no approved wording are held and reported rather than improvised.

WHAT STRONG LOOKS LIKE

approved wording is reproduced exactly, unmapped reasons are blocked rather than free-written, and the multilingual versions keep the disclosure block layout identical.

BAND S2 · API, templates, Review, export · a pull request into the public repository, use-cases folder

Credit memo builder with covenant schedule

WHO IT SERVES

Commercial credit analyst preparing a memo for a credit committee

WHAT TO BUILD

Build an app that assembles the credit memo from the spread data, the borrower's financial statements and the relationship history: business description, financial analysis with the ratios tied to the spreads, risk rating rationale, structure and the covenant schedule with each covenant's test, level and headroom. Every figure in the narrative is checked against the spread it came from. When the structure changes in committee, the covenant table and the paragraphs that describe it are edited together, and the memo keeps the bank's exact template. Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S3 · chat, Multi-document, templates, search, export · a pull request into the public repository, use-cases folder

Closing document reconciliation memo

WHO IT SERVES

Mortgage closing or post-closing auditor checking a file before funding

WHAT TO BUILD

Build a tool that opens the loan estimate, the closing disclosure, the note, the settlement statement and the underwriting approval together and reconciles the fields that must agree: loan amount, rate, term, escrow, fees by category, cash to close, borrower names and property address. It produces a variance memo quoting each document's value side by side, classifies each variance as tolerable or curable under the file's own rules, and drafts the correction letter or revised disclosure request for the ones that need action.

WHAT STRONG LOOKS LIKE

field extraction is right across documents that label the same value differently, seeded variances are all caught, and the classification is explained rather than asserted.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Loan modification and hardship packet generator

WHO IT SERVES

Loss-mitigation specialist working borrower hardship cases

WHAT TO BUILD

Build a tool that produces the modification packet from the approved workout terms: the borrower letter explaining the new payment in plain language, the modification agreement, the revised amortization summary, and the document checklist the borrower must return. The plain-language letter and the legal agreement are generated from the same term set so they cannot disagree, the payment figures are shown with the arithmetic the borrower can follow, and the packet is produced in the borrower's language with the signature blocks intact.

WHAT STRONG LOOKS LIKE

the plain-language letter and the legal agreement are provably generated from one term set, and translated packets keep the signature and notary blocks positioned correctly.

BAND S2 · chat, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Regulatory exam response binder

WHO IT SERVES

Compliance officer responding to an examination or audit findings letter

WHAT TO BUILD

An exam findings letter requires a written response per finding with the remediation plan, owner, date and the evidence, assembled into a binder with an index the examiner can navigate. Build an app that parses the findings letter into individual items, drafts each response against the institution's actual policies (citing the policy section), assembles the evidence exhibits with an index, and tracks which findings still lack evidence. Follow-up examinations reuse the prior binder and show what changed.

WHAT STRONG LOOKS LIKE

each finding is answered individually with a real policy citation, missing evidence is listed rather than hidden, and the follow-up binder is a delta on the prior one.

BAND S3 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Disclosure library version sweep

WHO IT SERVES

Retail-banking compliance owner maintaining the consumer disclosure set

WHAT TO BUILD

A regulation change or a product change means every affected disclosure, application form, fee schedule and website policy has to move together, with an effective date and a version record. Build a tool that holds the disclosure library, finds every document affected by a stated change (including the ones that only imply the old rule), proposes the edits with the required wording applied exactly, produces the version pack with a change log per document, and keeps the superseded versions retrievable for the retention period.

WHAT STRONG LOOKS LIKE

implied dependencies are found, not just literal quotes of the old rule, and the version pack plus change log is complete enough to hand an examiner without further work.

BAND S3 · search, Multi-document, templates, Review, export · a pull request into the public repository, use-cases folder

Collections and remediation letter factory

WHO IT SERVES

Collections or customer-remediation operations lead running a letter programme

WHAT TO BUILD

Remediation and collections programmes send letters at volume where the wording is legally constrained, varies by jurisdiction and stage, and must reflect each customer's actual figures. Build a service that assembles each letter from the approved wording blocks for its jurisdiction and stage, inserts the customer's calculated figures with the arithmetic reproducible, produces the customer's language variant, and holds back any record whose figures fail an internal sanity check. The run reports counts by variant and every held record with its reason.

WHAT STRONG LOOKS LIKE

approved wording blocks are never paraphrased, held records are reported with named reasons, and the same programme runs across jurisdictions without a code change.

BAND S2 · API, templates, Review, export · a pull request into the public repository, use-cases folder

Reviewer response letter builder

WHO IT SERVES

Corresponding author revising a manuscript after peer review

WHAT TO BUILD

Revision means two artifacts that must stay in lockstep: a point-by-point response letter and a revised manuscript. Build an app that parses the reviewers' comments into numbered points, and for each one lets the author state the intended change in plain language. It then makes the change surgically in the manuscript and writes the response paragraph quoting the new text and its location, so the letter and the manuscript can never disagree. Output is the response letter plus both a clean and a tracked-changes manuscript.

WHAT STRONG LOOKS LIKE

every response paragraph quotes text that genuinely exists at the stated location in the revised manuscript, equations and citations survive the edits, and the tracked-changes export opens correctly in Word.

BAND S3 · chat, Multi-document, Review, export, search · a pull request into the public repository, use-cases folder

Systematic review package assembler

WHO IT SERVES

Research fellow or evidence-synthesis specialist producing a systematic review

WHAT TO BUILD

A systematic review ships with a search-strategy appendix per database, a study-flow diagram with exact counts, evidence tables, a risk-of-bias summary and the reporting checklist mapped to the page where each item is addressed. Build a tool that assembles all of them, renders the flow diagram from the screening counts so the numbers cannot drift from the text, builds the evidence tables from the extraction sheet, and updates the checklist page references automatically when the manuscript is re-paginated.

WHAT STRONG LOOKS LIKE

the flow diagram is generated from the counts (change a count and the diagram and the prose both move), and the checklist page references are still correct after the manuscript is edited.

BAND S3 · chat, images, Multi-document, search, export · a pull request into the public repository, use-cases folder

Research grant packet assembler

WHO IT SERVES

Principal investigator or research-office administrator preparing a funder submission

WHAT TO BUILD

A research grant is a bundle of separately formatted documents: the research narrative under a page limit, the data-management plan, biosketches per investigator in the funder's format, facilities statement, and the budget justification. Build a tool that assembles the bundle, converts each investigator's existing CV into the funder's biosketch format without losing content, enforces the page limits by tightening prose rather than truncating, and checks that every named collaborator has a document in the packet. The final check reports each funder requirement as met or missing.

WHAT STRONG LOOKS LIKE

page limits are met by real editing with the argument intact, biosketch conversion preserves content across very different layouts, and the requirement check catches a deliberately omitted document.

BAND S3 · Multi-document, templates, chat, search, export · a pull request into the public repository, use-cases folder

Methods section drafter from the lab record

WHO IT SERVES

Bench scientist or graduate researcher writing up reproducible methods

WHAT TO BUILD

Methods sections are written from memory weeks later and lose the details that make the work reproducible. Build a tool that takes the lab notebook entries, protocol versions and instrument settings and drafts the methods section in the target journal's conventions, listing every parameter the record contains and explicitly flagging the ones the record does not (catalogue numbers, incubation times, software versions) instead of inventing them. Deviations from the written protocol recorded in the notebook are surfaced for the author to address.

WHAT STRONG LOOKS LIKE

missing parameters are listed as gaps rather than filled with typical values, and recorded protocol deviations are surfaced rather than smoothed over.

BAND S2 · chat, Multi-document, search, export · a pull request into the public repository, use-cases folder

Co-author merge desk

WHO IT SERVES

Corresponding author reconciling several co-authors' edited copies of one manuscript

WHAT TO BUILD

Four co-authors return four edited Word files with overlapping tracked changes and comments. Build a tool that opens them together against the base version, attributes every change to its author, merges the non-conflicting edits automatically, and presents genuine conflicts (two authors rewriting the same sentence differently) as a choice with both versions shown. Comments are preserved with their authors. The output is one master with a merge report showing what was taken from whom and what was rejected.

WHAT STRONG LOOKS LIKE

attribution survives the merge, conflicts are surfaced rather than resolved silently by last-writer-wins, and the master exports with the merged tracked changes and comments intact.

BAND S3 · Multi-document, Review, chat, export · a pull request into the public repository, use-cases folder

Figure caption and accessibility pass

WHO IT SERVES

Journal production editor or an author meeting accessibility requirements

WHAT TO BUILD

Build a tool that walks every figure, table and equation in a manuscript, checks that each has a caption in the required style and a numbered reference in the text, writes alternative text that describes the figure's content rather than restating the caption, and flags figures whose meaning depends on colour alone. Renumbering after a figure is inserted or removed updates the captions and every in-text reference together. The pass reports what it changed and what needs the author's judgment.

WHAT STRONG LOOKS LIKE

alt text describes what the figure shows (verified by a reviewer who cannot see it), and inserting a figure mid-manuscript rennumbers captions and references correctly with nothing else touched.

BAND S2 · images, chat, search, export · a pull request into the public repository, extensions folder

Thesis format compliance checker

WHO IT SERVES

Doctoral candidate or graduate-school reviewer against a university's submission specification

WHAT TO BUILD

Universities reject theses for margins, heading styles, front-matter order, page numbering that must switch from roman to arabic at the right page, table and figure list accuracy, and citation style. Build a tool that holds the university specification, checks the thesis against every rule, reports each violation with the location, and fixes the mechanical ones in place while leaving the judgment calls to the candidate. It regenerates the tables of contents, figures and tables from the actual headings and captions rather than trusting the existing lists.

WHAT STRONG LOOKS LIKE

page-numbering transitions and section-specific page geometry survive the export, generated lists match the document's real content, and the fix pass never disturbs the body text.

BAND S2 · chat, search, export, Review · a pull request into the public repository, use-cases folder

Developmental edit letter with anchored manuscript notes

WHO IT SERVES

Freelance developmental editor or an author self-assessing a draft

WHAT TO BUILD

A developmental edit is a long structural letter plus notes anchored where the problems are. Build a tool that reads the full manuscript and produces both: an editorial letter covering structure, pacing, point of view and character arcs with chapter evidence for each claim, and comments anchored on the specific passages the letter refers to. The letter and the comments cross-reference each other. The author works through the comments in the manuscript, and resolved notes update the letter's progress record.

WHAT STRONG LOOKS LIKE

every structural claim in the letter is backed by a chapter and passage rather than generic craft advice, and the anchored comments stay attached to the right passages as the author revises.

BAND S2 · search, chat, Review, Multi-document, export · a pull request into the public repository, use-cases folder

Print-ready interior formatter

WHO IT SERVES

Self-publishing author or small-press production person preparing a print interior

WHAT TO BUILD

Build a tool that takes a finished manuscript and produces the print interior: the chosen trim size and gutter-aware margins, mirrored running heads (author on verso, title on recto), chapter openers with drop position and no head on the opening page, front and back matter in the correct order, and page numbering that starts where it should. Editing a chapter afterwards reflows only that chapter's pages and updates the contents, and the tool reports orphan and widow problems for the author to resolve.

WHAT STRONG LOOKS LIKE

the exported interior honours trim size, mirrored margins and per-section page geometry, running heads behave correctly on chapter openers, and a late text change does not force a full reformat.

BAND S3 · chat, templates, export · a pull request into the public repository, use-cases folder

Query letter and book proposal package

WHO IT SERVES

Author submitting to literary agents or acquiring editors

WHAT TO BUILD

A submission package is a query letter, a synopsis at the length each agent asks for, an overview and comparable-titles section, a chapter outline, an author bio and sample chapters formatted to standard. Build a tool that generates the set from the manuscript, produces the synopsis at several required lengths from one source so they cannot contradict each other, tailors the query per agent from their stated interests, and assembles the per-agent package with a submission tracker. Sample chapters are formatted to standard manuscript layout without altering the prose.

WHAT STRONG LOOKS LIKE

the multiple synopsis lengths are genuinely derived from one summary rather than written separately, and per-agent tailoring changes the letter's argument, not just the salutation.

BAND S2 · Multi-document, chat, templates, export · a pull request into the public repository, use-cases folder

Voice-preserving line-edit pass with a change ledger

WHO IT SERVES

Author or line editor working through a full-length manuscript

WHAT TO BUILD

Build a tool that line-edits a book one chapter at a time against a stated brief (tighten, remove tics, fix rhythm) while keeping the author's voice, dialogue and deliberate stylistic choices untouched. Every change lands as a reviewable edit with a one-line reason, and the run produces a ledger showing edit density by chapter so the author can see where the prose needed the most work. The author's rejections train the pass, and a chapter can be reverted to its pre-edit state without disturbing the rest of the book.

WHAT STRONG LOOKS LIKE

dialogue and intentional fragments survive; a reader cannot tell the edited chapters from the author's own best pages; reverting chapter nine leaves chapters one to eight and ten onward untouched.

BAND S3 · chat, Review, memory, Multi-document, export · a pull request into the public repository, use-cases folder

Translation-ready manuscript packet

WHO IT SERVES

Rights manager or self-publishing author licensing into other languages

WHAT TO BUILD

Build a tool that prepares a manuscript for translation: a term base of invented names, places and coined words with pronunciation and meaning notes, a do-not-translate list, a style note describing register and voice, and a clean source file with the layout and any formatting-carried meaning documented. Once a translation returns, the tool re-replaces it into the original layout, checks that every protected term was handled per the term base, and reports mismatches chapter by chapter.

WHAT STRONG LOOKS LIKE

the term base is extracted from the manuscript rather than typed by hand, and the return check catches a protected term that was translated when it should not have been.

BAND S2 · search, Multi-document, templates, export, memory · a pull request into the public repository, use-cases folder

Second-edition refresh engine

WHO IT SERVES

Nonfiction author or publisher updating a backlist title

WHAT TO BUILD

A second edition needs updated facts, figures and references, a new preface, removal of dated material, and a change list for the reader and for marketing. Build a tool that sweeps the book for time-sensitive claims (dated statistics, superseded tools, "currently" and "recently" phrasing, dead references), proposes a targeted update for each with the new source attached, drafts the new-edition preface from the accumulated changes, and produces both the updated interior and the marketing change summary. Untouched chapters stay untouched.

WHAT STRONG LOOKS LIKE

the sweep catches implicitly dated claims, not just explicit years, and the change summary is generated from the actual applied edits rather than written separately.

BAND S2 · search, chat, Review, Multi-document, export · a pull request into the public repository, use-cases folder

Support letter drafting desk

WHO IT SERVES

Immigration attorney or paralegal producing expert and employer support letters

WHAT TO BUILD

Petitions rest on letters signed by different people, and letters that read as if one person wrote them all are a liability. Build a tool that drafts each support letter from that signer's own CV, their real relationship to the beneficiary and the specific criterion the letter must address, in a voice consistent with the signer's background and seniority. It refuses to assert anything the attached evidence does not support, produces a signer-facing plain summary of what the letter claims, and tracks which criteria are covered and which are still bare.

WHAT STRONG LOOKS LIKE

the letters are demonstrably distinct in voice and content, unsupported assertions are refused with the gap named, and the criteria coverage map is accurate.

BAND S2 · chat, Multi-document, templates, Review, export · a pull request into the public repository, use-cases folder

Evidence exhibit indexer and cover-page engine

WHO IT SERVES

Immigration paralegal assembling a filing-ready exhibit set

WHAT TO BUILD

Build a tool that takes the evidence pile in mixed formats, proposes an exhibit order that follows the argument of the brief, generates a numbered cover page for each exhibit describing what it proves, builds the exhibit index with page ranges, and inserts the exhibit citations into the brief at the points that rely on them. Adding an exhibit late rennumbers the index, the cover pages and every in-brief citation together. The assembled packet exports as one paginated file.

WHAT STRONG LOOKS LIKE

a late insertion rennumbers the whole set consistently including in-brief citations, and page ranges in the index match the assembled export exactly.

BAND S3 · Multi-document, chat, search, export · a pull request into the public repository, use-cases folder

Request-for-evidence response builder

WHO IT SERVES

Immigration attorney responding to an officer's notice against a deadline

WHAT TO BUILD

Build a tool that parses the officer's notice into each specific request, retrieves what the original petition already argued and evidenced on that point, drafts a response that addresses the officer's stated concern directly (rather than restating the original), identifies the new evidence needed, and assembles the response with a new exhibit set. It produces a checklist mapping every request in the notice to the part of the response that answers it, so nothing in the notice goes unanswered.

WHAT STRONG LOOKS LIKE

every request in the notice maps to an answer, responses engage the officer's specific concern instead of repeating the petition, and prior-petition citations are accurate.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Certified translation packet builder

WHO IT SERVES

Case manager handling foreign-language supporting documents

WHAT TO BUILD

Foreign documents need a translation, a translator certification in the required form, and a presentation that lets an officer match source to translation. Build a tool that ingests the source document (often a scan), produces the translation preserving the original's layout so seals, stamps and table structure remain locatable, generates the certification page with the translator's details, and assembles a side-by-side or facing-page packet. Illegible or ambiguous source content is marked as such in the translation rather than guessed.

WHAT STRONG LOOKS LIKE

the translated document mirrors the source layout closely enough to compare line by line, and illegible regions are explicitly marked instead of invented.

BAND S2 · chat, images, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Case status and document checklist letter desk

WHO IT SERVES

Immigration practice manager keeping clients informed at volume

WHAT TO BUILD

Clients need regular, understandable updates and a precise list of what they still owe. Build a service that generates each client's update letter from the case record: the current stage in plain language, what happens next and when, and a personalized checklist of outstanding documents with the exact form of each one required. Letters go out in the client's language, avoid legal advice beyond the approved language set, and the outstanding checklist is regenerated from the actual case file rather than a stale list.

WHAT STRONG LOOKS LIKE

the checklist reflects the true current state of the file, the plain-language stage description is accurate for each case type, and approved-language boundaries are enforced.

BAND S1 · API, templates, chat, export · a pull request into the public repository, use-cases folder

Employer compliance file assembler

WHO IT SERVES

Business-immigration paralegal or HR compliance owner maintaining employer-side records

WHAT TO BUILD

Employer-sponsored cases require posting notices, wage and position documentation, and a public access file assembled within set deadlines and retained for a set period. Build a tool that generates each required notice with the case's actual position, wage and worksite details, assembles the file with an index and a retention calendar, and audits existing files against the required contents, reporting each gap with the deadline it relates to. Re-use across many cases is template-driven, so a rule change updates the generator once.

WHAT STRONG LOOKS LIKE

the audit finds real gaps in a seeded imperfect file set, deadlines are calculated rather than typed, and a rule change is a content edit rather than a code change.

BAND S2 · API, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Customer-facing incident letter desk

WHO IT SERVES

Customer success or support lead communicating after an outage or a data issue

WHAT TO BUILD

After an incident, different customers need different letters: enterprise accounts with contractual notification duties, regulated customers with specific disclosure requirements, and everyone else with a clear summary. Build a tool that takes the internal postmortem and generates each tier's letter from the same fact set, checks each enterprise letter against that customer's own contract for the notification wording and deadline it requires, and holds any letter whose contract terms it could not verify. All variants stay factually identical where they overlap.

WHAT STRONG LOOKS LIKE

contractual notification requirements are genuinely read from each customer's agreement, the tiered letters never state contradicting facts, and unverifiable cases are held rather than sent.

BAND S3 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Customer onboarding runbook generator

WHO IT SERVES

Implementation or onboarding manager at a B2B software company

WHAT TO BUILD

Every new customer gets a generic onboarding deck and then asks the same questions. Build a tool that generates a customer-specific runbook from the implementation plan and the account's configuration: their environments, their integrations, their named contacts, their milestone dates, with the standard procedures included only where they apply. Diagrams of their specific integration topology are generated inside the document. As the project progresses, the runbook is edited in place and becomes the handover document to support.

WHAT STRONG LOOKS LIKE

the runbook contains no procedure the customer will never use, the topology diagram reflects their actual configuration, and the handover version is the same living document rather than a fresh write-up.

BAND S2 · chat, templates, images, Multi-document, export · a pull request into the public repository, use-cases folder

Macro and canned-response library governance

WHO IT SERVES

Support operations manager owning hundreds of agent response templates

WHAT TO BUILD

Macro libraries drift into inconsistent tone, outdated policy statements and duplicated near-identical entries. Build a tool that audits the library for policy statements that no longer match current policy, tone that violates the voice guide, and near-duplicate macros, then proposes consolidations and rewrites for review. Approved changes are applied with a version record per macro. New macros are checked against the library before being added so duplicates never enter, and the whole library can be regenerated in another language.

WHAT STRONG LOOKS LIKE

near-duplicate detection catches macros that say the same thing in different words, policy drift is detected against the current policy document rather than a hardcoded list, and versioning survives a bulk apply.

BAND S2 · Multi-document, search, templates, Review, API · a pull request into the public repository, extensions folder

Voice-of-customer digest for product

WHO IT SERVES

Support insights analyst reporting to product and leadership each quarter

WHAT TO BUILD

Build a tool that turns a quarter of support conversations into a styled digest for the product team: the themes with their volume trend, representative verbatim quotes with identifying details removed, the issues that grew fastest, and what changed since last quarter's digest. Charts and trend visuals are generated inside the document. Last quarter's digest is the source of the comparison, so the "what changed" section is derived rather than remembered, and every theme links to the evidence behind it.

WHAT STRONG LOOKS LIKE

the quarter-over-quarter comparison is genuinely computed against the prior document, quotes are anonymized before they land in the digest, and each theme's volume claim is traceable.

BAND S2 · search, Multi-document, images, chat, export · a pull request into the public repository, use-cases folder

Service-credit claim response desk

WHO IT SERVES

Customer success or service-delivery manager answering SLA credit claims

WHAT TO BUILD

When a customer claims a service credit, the answer depends on their contract's availability definition, exclusion list, measurement window and credit schedule, none of which are in the ticket. Build a tool that opens the customer's agreement alongside the incident record and the availability data, calculates the entitlement under that contract's own terms, and drafts the formal response showing the calculation and quoting the governing clauses. Claims where the contract genuinely supports a larger credit than the customer asked for are flagged rather than quietly under-paid.

WHAT STRONG LOOKS LIKE

the calculation follows the specific contract rather than a house default, exclusions are applied only when the contract lists them, and the under-claimed case is surfaced honestly.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Round-Terms Propagation Desk

WHO IT SERVES

Startups & VC · founder-ops lead or startup counsel at a seed to Series A company

WHAT TO BUILD

Load the financing set (term sheet, SAFE or SPA, board consent, disclosure schedule, side letters) as one multi-document session. The user changes a single deal variable in chat ("cap moves to 12M, pro rata drops for sub-100K checks"). The app finds every place that number or right appears across all open documents, including inside tables and inside prose, and stages a per-document edit batch for approval. It also writes a change-ledger document listing every touched clause with before and after. Export the updated set plus the ledger.

WHAT STRONG LOOKS LIKE

zero missed occurrences across the set; figures embedded in sentences caught as reliably as figures in tables; denied changes leave the document byte-identical; the ledger is good enough to hand to counsel.

BAND S3 · Multi-document, chat, search, Review, export · a pull request into the public repository, use-cases folder

Investor Update That Differs Last Month

WHO IT SERVES

Startups & VC · founder or chief of staff sending a monthly investor update

WHAT TO BUILD

Two documents open: last month's update and this month's metrics export. The app reads the prior update, works out what actually changed (metrics, hiring, risks, asks), then drafts this month by editing a copy of last month in place, so section order, layout and chart placement stay identical. Every changed sentence is highlighted for approval. With cross-session memory enabled, commitments made in earlier updates ("we said we would close two enterprise pilots") are carried forward and checked.

WHAT STRONG LOOKS LIKE

month two is visibly the same document, not a regeneration; the change list is accurate and specific; a promise made in month one resurfaces in month three unprompted.

BAND S2 · Multi-document, chat, memory, Review, export · a pull request into the public repository, use-cases folder

Diligence Q&A Dossier Builder

WHO IT SERVES

Startups & VC · founder preparing for diligence, or an associate running it

WHAT TO BUILD

Point the app at a data room of 10 to 40 files. It runs a standard diligence question set (revenue quality, churn, cap table, IP assignment, key-person risk, security posture, customer concentration) as semantic searches across every file, then writes a dossier where each question carries a drafted answer, the filename and section the answer came from, and an explicit flag when nothing in the room answers it. Founders then edit answers in place before sharing. Second mode: the investor side runs the same pass and gets a gap list.

WHAT STRONG LOOKS LIKE

every answer cites a real location in a real file; unanswerable questions are flagged honestly rather than invented; quality holds at 40 files, not just 5.

BAND S2 · search, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Investment-Committee Memo Redline Applier

WHO IT SERVES

VC / private credit · investment associate turning partner markup into a final memo

WHAT TO BUILD

Partners return the IC memo as a Word file with track changes and margin comments. Upload it. The app reads the markup with reviewer names and dates attached, groups changes by reviewer and by memo section, and proposes each as an approve-or-deny entry against the working memo. Where two partners asked for opposite things, it surfaces the conflict with both names instead of silently choosing. Export a clean memo plus a one-page resolution note recording who asked for what and how it landed.

WHAT STRONG LOOKS LIKE

reviewer attribution survives ingest; contradictions are raised, never resolved silently; the clean export carries no residual markup or comment anchors.

BAND S2 · chat, Review, search, export · a pull request into the public repository, use-cases folder

Data-Room Readiness Auditor

WHO IT SERVES

Startups & VC · founder or fractional CFO preparing a room before a raise or an exit

WHAT TO BUILD

Run a checklist pass over every document in a data room and produce one scored readiness report: missing signature blocks, unsigned versions sitting next to signed ones, stale dates, inconsistent legal entity name, placeholder text left in, documents referenced by another document but absent from the room. Each finding cites file and section. A second pass fixes the mechanically fixable class (entity name, dates, placeholder tokens, file naming) as reviewable edits and re-scores.

WHAT STRONG LOOKS LIKE

the report is specific enough to act on without opening a single file; the fix pass touches only the paragraph each finding names; the re-score is honest.

BAND S2 · Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Per-Investor Memo Tailoring Batch

WHO IT SERVES

Startups & VC · founder running a structured raise across 25 target funds

WHAT TO BUILD

Input is one master fundraising memo plus a roster of investors with thesis, stage, geography and notable portfolio companies. For each investor the app produces a tailored copy: emphasis reordered, a "why this fits your thesis" section written, relevant portfolio comparables named, everything else left untouched. Runs server side over the roster and returns per-investor exports plus a summary table of exactly what changed for each. Shared sections must stay identical across all outputs.

WHAT STRONG LOOKS LIKE

tailoring is substantive rather than a find-and-replace of the fund name; the shared 80 percent of the memo is byte-identical across every output; 25 runs complete without drift or duplicated work.

BAND S2 · API, templates, chat, export · a pull request into the public repository, use-cases folder

One Narrative, Five Derived Artefacts

WHO IT SERVES

Startups & VC · founder or head of comms maintaining a single company story

WHAT TO BUILD

A master narrative document drives four derivatives in the same session: a one-pager, a two-page memo, an executive summary, and an FAQ. Edit the master and the app proposes the matching edit in each derivative and marks which derivatives are now stale, per section. Each derivative keeps its own length target and register (investor, customer, press, internal). Export all five, or just the stale ones. A staleness badge per derivative makes the sync state visible at a glance.

WHAT STRONG LOOKS LIKE

derivatives stay in sync without being regenerated from scratch; staleness is detected per section not per file; length targets still hold after three rounds of sync.

BAND S3 · Multi-document, chat, Review, export · a pull request into the public repository, use-cases folder

Page-Limit Proposal Compressor

WHO IT SERVES

GovCon · proposal manager or capture lead on a federal or defense bid

WHAT TO BUILD

Government proposals live and die on page limits. Load the volume, set the constraint (for example 30 pages, 12pt, one-inch margins), and the app reports the real page count, ranks sections by evaluated value against the solicitation's evaluation factors, then proposes compression edits section by section until the volume fits. Every cut arrives as a reviewable diff so nothing evaluated disappears silently. Export the compliant volume plus a cut log. Bonus mode: report how many pages each section is over or under its allocation.

WHAT STRONG LOOKS LIKE

the page count is real, not estimated from word count; compression preserves every answered requirement; the reviewer can see precisely what was removed and why.

BAND S3 · chat, search, Review, export · a pull request into the public repository, use-cases folder

Section L and M Outline Compiler

WHO IT SERVES

GovCon · proposal coordinator building the compliant shell on day one of a bid

WHAT TO BUILD

Upload the solicitation. The app extracts the instructions section and the evaluation-factors section, builds a compliant proposal shell where every required heading appears in the required order with the evaluation criterion quoted underneath it, then drafts into each heading from an attached past-performance and boilerplate library. Output is the shell document plus a coverage table showing which instruction each heading answers and which instructions are still unmapped.

WHAT STRONG LOOKS LIKE

heading order matches the solicitation exactly, including sub-levels; no instruction item is left unmapped without being listed; the coverage table survives a re-run on a different solicitation.

BAND S2 · search, templates, chat, Multi-document, export · a pull request into the public repository, use-cases folder

Color-Team Review Workflow

WHO IT SERVES

GovCon · capture manager running pink, red and gold team reviews

WHAT TO BUILD

Reviewers leave feedback in whatever form they like: Word comments, PDF annotations, or a separate comment document. The app converts every comment into a concrete proposed edit against the correct proposal section, deduplicates reviewers who said the same thing in different words, and routes one batch for capture-manager approval. It produces a review disposition log recording each comment, the reviewer, the disposition, and the reason for any rejection.

WHAT STRONG LOOKS LIKE

comments become edits rather than a summary of comments; duplicates merge without losing either author; the disposition log would stand up in a post-award debrief.

BAND S2 · Review, chat, search, export · a pull request into the public repository, use-cases folder

Teaming and Subcontract Packet Assembler

WHO IT SERVES

GovCon · business-development lead standing up a teaming arrangement

WHAT TO BUILD

From a single deal record, assemble the teaming packet: teaming agreement, mutual NDA, subcontractor data call, letter of commitment, and small-business subcontracting plan, each on the prime's template with the partner's details filled. Multi-document session with a cover index. Partner-specific terms such as workshare percentage and exclusivity window propagate to every document that mentions them, so one change updates the whole packet.

WHAT STRONG LOOKS LIKE

one workshare change updates every document that references it; the packet exports as a single clean bundle with a working index; no document carries a trace of the previous partner.

BAND S2 · Multi-document, templates, chat, export · a pull request into the public repository, use-cases folder

Amendment and Q&A Delta Applier

WHO IT SERVES

GovCon · proposal manager mid-bid when the agency changes the rules

WHAT TO BUILD

Solicitations change while you are writing. Feed the app the amendment or the agency's published answers to bidder questions. It diffs them against the original solicitation, identifies which proposal sections are now affected (page limits, due dates, evaluation weightings, new requirements, deleted requirements), and stages the required edits against the in-flight proposal, each annotated with the amendment paragraph that drove it. Produces an amendment-response log for the compliance file.

WHAT STRONG LOOKS LIKE

the affected-section mapping is correct including second-order effects; every edit is traceable to an amendment paragraph; sections the amendment did not touch are provably untouched.

BAND S3 · Multi-document, search, Review, chat, export · a pull request into the public repository, use-cases folder

Brand-Kit Switcher for Client Deliverables

WHO IT SERVES

Agencies · design lead or account director serving several brands from one studio

WHAT TO BUILD

One deliverable, several client brand systems. Each client's brand kit (typography, palette, cover page, running header and footer, tone rules, logo) lives as an attachment set. Pick a client and the app re-skins the open document in place: type styles, table shading, cover, headers, and a voice pass, without touching the substance. Ship the demo rendering the same report in three brands side by side, plus a text diff proving only style changed.

WHAT STRONG LOOKS LIKE

a text diff between two branded outputs shows style-only changes; tables and running headers genuinely re-skin, not just body text; the brand survives export into the file.

BAND S2 · templates, chat, images, memory, export · a pull request into the public repository, use-cases folder

Campaign Performance Recap Builder

WHO IT SERVES

Agencies · account manager doing monthly client reporting

WHAT TO BUILD

Take the metrics export plus the campaign brief and generate the recap on the client's brand: results narrative, channel table, auto-rendered charts, and a recommendations section. The key move is month two: the app updates the same document rather than creating a new one, so the year reads as one living report with a monthly history section, and layout stays byte-stable across months.

WHAT STRONG LOOKS LIKE

numbers stated in prose match the table exactly; the month-two update preserves layout and only touches what changed; recommendations reference last month's recommendations.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Scope Diff and Change-Order Generator

WHO IT SERVES

Agencies / professional services · project lead defending a shrinking margin

WHAT TO BUILD

Compare the signed statement of work against what the project actually delivered (status document, ticket export, or delivery log). The app produces a scope-variance report listing in-scope, out-of-scope and unrequested work, each mapped to the SOW clause it sits under or falls outside of, then drafts the change-order document with pricing lines for the out-of-scope items. Ambiguous items are flagged for a human rather than assumed either way.

WHAT STRONG LOOKS LIKE

every variance cites a real SOW clause; the change order is client-ready without rewriting; ambiguity is surfaced rather than resolved silently.

BAND S2 · Multi-document, search, chat, templates, export · a pull request into the public repository, use-cases folder

Stakeholder Feedback Consolidator

WHO IT SERVES

Agencies · account lead receiving the same deliverable back from four people in four formats

WHAT TO BUILD

Four stakeholders return feedback four ways: a PDF with sticky notes and highlights, a Word file with track changes, an email of bullet points, and inline comments. Load them all into one session. The app normalizes every item into a single proposed-edit batch against the master, groups duplicates, flags direct contradictions with both authors named, and lets the account lead approve once. Export the resolved deliverable plus a feedback disposition sheet showing who asked for what and what happened to it.

WHAT STRONG LOOKS LIKE

annotation authorship and dates survive ingest; contradictions surface with names attached; one approval pass resolves all four sources without losing a single item.

BAND S3 · Multi-document, Review, search, chat, export · a pull request into the public repository, use-cases folder

New-Business Pitch Pack Assembler

WHO IT SERVES

Agencies · new-business director responding to an inbound brief

WHAT TO BUILD

Given a prospect (sector, size, brief), assemble the pitch document from the agency's credentials library: the three most relevant case studies, matching sector experience, the right team bios, and a tailored approach section. Case studies are chosen by semantic match to the brief rather than by tag, and metrics inside them are carried through unmodified. Export on the agency brand with a cover, contents and a credentials appendix.

WHAT STRONG LOOKS LIKE

case-study selection is defensibly relevant and a second brief produces a genuinely different pack; numbers inside case studies are never paraphrased or invented; bios come through unedited.

BAND S2 · search, templates, Multi-document, images, export · a pull request into the public repository, use-cases folder

Glossary and Do-Not-Translate Enforcement Pass

WHO IT SERVES

Localization · localization manager at an LSP or an in-house globalization team

WHAT TO BUILD

A quality pass that runs over an already-translated document. Attach the client term base (approved renderings), the do-not-translate list (product names, trademarks, UI strings) and the style rules, then let the app scan the translation for violations: an approved term rendered three different ways, a product name translated, a forbidden register. Each violation is proposed as an in-place correction with the term-base entry quoted as justification. Output is the corrected document plus a QA report counting violations by category.

WHAT STRONG LOOKS LIKE

inflected and declined forms of a glossary term are caught, not only exact strings; do-not-translate items are never "corrected" into a translation; corrections do not disturb surrounding formatting.

BAND S2 · search, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Back-Translation Drift Check

WHO IT SERVES

Localization · quality lead in regulated or high-stakes localization (medical, legal, financial)

WHAT TO BUILD

Open the source and the translation together. The app back-translates the target into the source language section by section, aligns each back-translation with the original section, and flags meaning drift: a negation lost, a number changed, a qualifier dropped, a legal obligation softened. Findings are ranked by severity and land as proposed corrections on the translation, never on the source. Output is a bilingual drift report with the source, the back-translation and the flagged difference per row.

WHAT STRONG LOOKS LIKE

genuine meaning changes are caught while acceptable stylistic variation is not flagged; numerals, negations and modal obligations get special attention; the source document is provably never edited.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Right-to-Left Layout Adaptation Studio

WHO IT SERVES

Localization · localization engineer taking English collateral into Arabic, Hebrew, Persian or Urdu

WHAT TO BUILD

Take a styled English document and produce a genuinely right-to-left version: text direction flipped, table column order mirrored, list markers and indentation reversed, headers and footers mirrored, page numbering side swapped, and mixed-direction runs (Latin brand names, numerals, code) handled correctly inside RTL paragraphs. Images and captions keep their pairing. Export to both DOCX and PDF and show the two documents side by side.

WHAT STRONG LOOKS LIKE

the exported file is RTL at the file level, not merely right-aligned text; mixed-direction runs read correctly; tables mirror without losing header association; the PDF matches the editor.

BAND S3 · chat, images, export, templates · a pull request into the public repository, use-cases folder

CJK Typography and Layout Adapter

WHO IT SERVES

Localization · Japanese, Chinese or Korean localization specialist handling formal business documents

WHAT TO BUILD

Adapt a Western-typeset document into CJK conventions: appropriate typeface families and sizes, line spacing that suits dense glyphs, punctuation substitution, no orphaned closing punctuation at line start, correct spacing around embedded Latin and numerals, and page geometry that suits the target market's business-document norms. Handles the reverse direction too. Ships with a side-by-side comparison and a checklist report of every convention applied.

WHAT STRONG LOOKS LIKE

the output reads as a native business document to a native speaker rather than a translated Western one; the export preserves the typography choices; the checklist matches what actually changed.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Prior-Translation Reuse Desk

WHO IT SERVES

Localization · freelance translator or small LSP without a full translation-memory stack

WHAT TO BUILD

Before translating a new document, the app searches the owner's previously translated documents for sections that closely match the incoming one, and offers the previously approved rendering for reuse instead of a fresh translation. Matches arrive with the source document name and a similarity note so the translator can accept, adapt or reject each. Reused sections are marked in a leverage report showing how much of the document came from prior work.

WHAT STRONG LOOKS LIKE

matches are semantically right rather than string-similar; every reuse is owner-scoped and never crosses accounts; the leverage report is accurate enough to base a quote on.

BAND S3 · memory, search, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Multilingual Regulatory Insert Set

WHO IT SERVES

Localization · regulatory-affairs or packaging lead shipping one product into many markets

WHAT TO BUILD

One approved source insert (product information leaflet, safety sheet, warranty statement) becomes a market-specific set: each language rendered in that market's required layout, mandatory headings in the mandated order, required warnings verbatim from an attached regulator wording list, and local units and date formats. Any market where the source lacks a mandated section is flagged rather than filled with a guess. Export the full set plus a compliance matrix per market.

WHAT STRONG LOOKS LIKE

mandated wording is reproduced verbatim and never paraphrased; missing mandatory sections are refused loudly; a change to the source propagates into all market versions as reviewable edits.

BAND S3 · Multi-document, templates, search, Review, export · a pull request into the public repository, use-cases folder

Post-Editing Workbench

WHO IT SERVES

Localization · post-editor cleaning up machine output before delivery

WHAT TO BUILD

A machine-translated draft arrives alongside the source. The workbench presents each section as a before-and-after proposal so the linguist approves, edits or rejects section by section, with the source visible beside it. Approved sections lock. The app tracks post-editing effort per section (untouched, light, heavy, rewritten) and delivers three artefacts: the clean target document, a bilingual side-by-side review file, and an effort report that supports post-editing pricing.

WHAT STRONG LOOKS LIKE

section alignment holds across the whole document including tables and lists; approvals are durable across a page reload; the effort classification is consistent enough to invoice against.

BAND S2 · Review, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Supplier Spec-Sheet Normalizer

WHO IT SERVES

Retail · category or sourcing manager receiving specs from 200 suppliers in 200 formats

WHAT TO BUILD

Suppliers send product specs as scanned PDFs, Word files, spreadsheets pasted into documents and photographs of printed sheets. The app ingests the lot, extracts the fields that matter (dimensions, materials, certifications, care instructions, country of origin, packaging), and re-issues each as the retailer's house spec document with a consistent layout. Anything unreadable or missing is listed in an exceptions report per supplier rather than filled with a guess.

WHAT STRONG LOOKS LIKE

scanned and photographed sheets are handled as well as digital ones; the exceptions report is honest and per-field; a 200-file run completes without manual babysitting.

BAND S2 · API, images, search, templates, export · a pull request into the public repository, extensions folder

Seasonal Catalog Re-Issue

WHO IT SERVES

Retail · merchandising lead reissuing the catalog every season

WHAT TO BUILD

Last season's catalog is the template. Feed the new season's product data (new SKUs, discontinued lines, new prices, new photography) and the app updates the existing catalog in place: rows swapped, prices changed, images replaced, discontinued spreads removed, new spreads inserted in the right category, contents and page references repaired. Layout, typography and grid stay exactly as designed. Export print-ready.

WHAT STRONG LOOKS LIKE

the new catalog is visibly the same designed artefact, not a re-layout; image swaps keep their crops and captions; the contents and any cross-references still resolve after insertion and deletion.

BAND S3 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Wholesale Line Sheet Generator

WHO IT SERVES

Retail / wholesale · brand sales lead selling into stockists

WHAT TO BUILD

From a product list with images, wholesale and retail pricing, minimums, case packs and delivery windows, generate the line sheet buyers expect: a branded cover, product grid with photography, size and colour runs, pricing tiers per account type, and terms. Produce account-specific variants where pricing tier, minimums and available assortment differ per stockist, from one master, in a single batch.

WHAT STRONG LOOKS LIKE

per-account variants differ only where they should; imagery is placed at consistent size and crop; a 40-product sheet paginates cleanly without orphaned rows.

BAND S2 · API, images, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Marketplace Compliance Doc Pack

WHO IT SERVES

E-commerce · marketplace operations manager listing across several platforms and regions

WHAT TO BUILD

Each marketplace and region demands its own paperwork: supplier declarations, conformity statements, materials disclosures, responsible-person details, restricted-substance attestations, each in its own required shape. The app takes one product record plus the certificates on file and assembles the correct pack per marketplace and per region, flags which required documents are missing or expired, and produces a submission checklist per destination.

WHAT STRONG LOOKS LIKE

the per-destination differences are real and correct; expiry checking works against the certificate documents themselves; missing items block the pack instead of being fabricated.

BAND S2 · Multi-document, templates, search, export · a pull request into the public repository, use-cases folder

Product Manual and Quick-Start Builder

WHO IT SERVES

Retail / consumer hardware · product manager shipping a physical product

WHAT TO BUILD

From the specification and a parts list, generate the product manual: safety notices, an exploded parts diagram rendered from a described structure, numbered setup steps with a diagram per step, troubleshooting table, warranty and disposal statements. Then generate the folded quick-start card as a separate document from the same source, so the two never disagree. Localize the pair into the launch markets keeping pagination stable.

WHAT STRONG LOOKS LIKE

the quick-start and the manual stay consistent when the manual changes; diagrams regenerate on edit rather than being re-drawn; localized versions keep the same page structure.

BAND S2 · chat, images, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Dispute and Chargeback Evidence Packet

WHO IT SERVES

E-commerce · payments or customer-operations lead fighting chargebacks and marketplace claims

WHAT TO BUILD

For a disputed order, assemble the evidence packet the processor or marketplace expects: order record, delivery confirmation, customer communications, the applicable clause of the published returns policy quoted verbatim with a citation, product listing snapshot, and a cover narrative that walks the reviewer through the timeline. Each processor has its own required order and length, so the packet is shaped per destination. Export as one clean PDF within the size limit.

WHAT STRONG LOOKS LIKE

policy quotes are verbatim with an accurate citation; the timeline narrative matches the evidence; the packet respects the destination's page and size limits without dropping evidence silently.

BAND S2 · Multi-document, search, images, templates, export · a pull request into the public repository, use-cases folder

Consumer Safety Notice Pack

WHO IT SERVES

Retail · quality or compliance lead running a product withdrawal or safety notice

WHAT TO BUILD

One incident becomes a coordinated document set: the regulator notification in the required format, the customer letter (personalized in batch from the purchaser list), in-store signage, the retailer partner instruction sheet, the call-centre script, and a press statement. Every artefact draws the same facts (affected batch codes, hazard description, remedy) from one source of truth, so a correction to the hazard wording updates every document as a reviewable edit.

WHAT STRONG LOOKS LIKE

a single fact change propagates to all six artefacts; the regulator format is followed exactly; batch personalization never leaks one customer's details into another's letter.

BAND S3 · Multi-document, templates, API, Review, export · a pull request into the public repository, use-cases folder

Disruption Re-Issue Engine

WHO IT SERVES

Travel · operations lead at a tour operator, corporate travel desk or DMC when a flight cancels

WHAT TO BUILD

A disruption event (cancelled flight, closed hotel, weather closure) hits a group of travellers. The engine identifies every affected traveller, regenerates each personalized itinerary document with the new arrangements, writes the explanation letter in the traveller's language, and issues the voucher or claim document where entitled. Everything runs in one batch and returns a per-traveller bundle plus an operations summary of who was moved where.

WHAT STRONG LOOKS LIKE

per-traveller correctness at 200 travellers with no cross-contamination; unaffected segments of an itinerary stay identical; letters read as written by a person, in the right language.

BAND S3 · API, templates, chat, Multi-document, export · a pull request into the public repository, use-cases folder

Event Operations Pack

WHO IT SERVES

Events / hospitality · event operations manager turning one event record into everything the floor needs

WHAT TO BUILD

One event record produces the operational document set in a single multi-document session: rooming list, banquet event order, run of show with timings, staffing and station sheet, AV and rigging schedule, and the client-facing summary. Change the guest count or the start time once and every document that depends on it updates as a reviewable edit. Export the whole pack as a single operations bundle plus per-department pull-outs.

WHAT STRONG LOOKS LIKE

a change to one input updates every dependent document and nothing else; per-department exports contain exactly that department's pages; the run of show arithmetic stays consistent after edits.

BAND S2 · Multi-document, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Speaker and Sponsor Kit Generator

WHO IT SERVES

Events · conference producer serving 40 speakers and 20 sponsors

WHAT TO BUILD

From the event database, produce a personalized kit per speaker (session brief, timings, stage layout diagram, AV notes, travel details, headshot and bio proof for approval) and per sponsor (activation brief, booth spec with a rendered floor diagram, deliverable deadlines, brand asset requirements). Each kit is branded to the event and personalized end to end. A change to the schedule reissues only the affected kits.

WHAT STRONG LOOKS LIKE

60 kits generated with zero cross-contamination; diagrams reflect the actual room and booth data; a schedule change reissues exactly the affected subset.

BAND S2 · API, images, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Destination Proposal with Visual Boards

WHO IT SERVES

Hospitality · destination management company or venue sales lead pitching an incentive or a wedding

WHAT TO BUILD

From a client brief (dates, headcount, budget band, style words), produce the proposal document: concept narrative, day-by-day programme, venue options with photography, generated mood-board imagery for the styling concept, a costed table with options priced separately, and terms. The visual concept is generated inside the document rather than assembled in a design tool, so the salesperson can iterate by chat ("make the styling warmer, drop the second venue") Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S2 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Property Welcome Book and Policy Set

WHO IT SERVES

Hospitality · property manager or short-let operator running many units in many languages

WHAT TO BUILD

Per property, generate the guest welcome book (arrival instructions, appliance guides with diagrams, local recommendations, house rules, emergency contacts) and the matching policy documents (rental terms, damage policy, cancellation terms) on the operator's brand. Generate each in the guest's language on demand. Property-specific facts live in one record so a lock-code change or a rule change updates every language version at once.

WHAT STRONG LOOKS LIKE

a single fact change propagates across all language versions; language versions keep the same structure and pagination; diagrams and images survive translation untouched.

BAND S2 · Multi-document, templates, chat, images, export · a pull request into the public repository, use-cases folder

Post-Event Impact Report

WHO IT SERVES

Events · event organizer proving value to sponsors and to the board

WHAT TO BUILD

After the event, pull the numbers (attendance, sessions, satisfaction survey, sponsor leads, spend) and the qualitative material (survey comments, coverage, photographs) into one branded impact report: headline results, charts, per-sponsor delivery against the contracted activation list, verbatim attendee quotes, and a recommendations section. Produce a per-sponsor variant showing that sponsor's own delivery, from the same master.

WHAT STRONG LOOKS LIKE

per-sponsor variants show only that sponsor's numbers; quotes are verbatim; the contracted-versus-delivered comparison is honest about shortfalls.

BAND S2 · chat, images, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Traveller Entry-Requirements Pack

WHO IT SERVES

Travel · corporate travel manager or agency ops handling mixed-nationality groups

WHAT TO BUILD

For each traveller (nationality, residency, route, purpose, dates), assemble a personal entry-requirements document from an attached reference set of the operator's own verified requirement notes: visa or authorization needed, document validity rules, transit requirements, health and insurance documentation, plus a personalized checklist with dates worked back from departure. Every statement cites the reference note it came from, and anything not covered by the reference set is flagged as needing a human check rather than answered.

WHAT STRONG LOOKS LIKE

every claim traces to an attached source; uncovered cases are refused loudly rather than guessed; a 30-traveller group produces 30 genuinely different checklists.

BAND S2 · search, templates, API, Multi-document, export · a pull request into the public repository, use-cases folder

Substitution Request and Response Packet

WHO IT SERVES

Architecture / AEC · project architect handling contractor substitution requests during construction

WHAT TO BUILD

A contractor proposes an alternative product to the one specified. Upload the request with its cut sheets alongside the relevant specification section. The app compares the proposed product against every specified performance requirement line by line, marks each as met, not met or unclear with the spec clause cited, drafts the architect's response letter with the reasoning, and, when the substitution is accepted, stages the edit to the specification section and to the finish schedule so the record stays true.

WHAT STRONG LOOKS LIKE

every performance requirement in the section is checked, not just the headline ones; cut-sheet values are read from the attachment including scanned ones; acceptance updates the spec and the schedule together.

BAND S3 · Multi-document, search, images, Review, export · a pull request into the public repository, use-cases folder

Competition and Award Submission Assembler

WHO IT SERVES

Architecture / design · marketing lead at a practice entering awards and competitions

WHAT TO BUILD

Every award has its own format: word limits per section, image counts and captions, credit lists, anonymity rules. From the practice's project record (narrative, drawings, renders, photography, consultant credits), assemble the submission to the target award's exact shape, enforce the word limits by compressing rather than truncating, place images with their required captions and credits, and strip identifying references when the award requires anonymous entries. Produce a compliance checklist alongside.

WHAT STRONG LOOKS LIKE

word limits are met by real compression that keeps the argument; anonymity stripping catches practice name, staff names and recognizable project titles; a second award format produces a genuinely different submission from the same source.

BAND S2 · templates, chat, images, search, export · a pull request into the public repository, use-cases folder

Code-Compliance Narrative Builder

WHO IT SERVES

Architecture · project architect producing the code analysis for a permit set

WHAT TO BUILD

From project parameters (occupancy classification, construction type, storeys, area, sprinklered or not) and the applicable code attached as a reference document, produce the code-compliance narrative: occupancy and construction analysis, allowed area and height calculations shown as worked steps, egress computation with a rendered diagram, fire-resistance rating schedule, and accessibility statements. Every conclusion cites the code section it rests on, and anything the attached code does not cover is flagged for the reviewer.

WHAT STRONG LOOKS LIKE

citations point to real sections in the attached code; the calculations are shown as steps, not asserted; uncovered conditions are refused rather than invented.

BAND S3 · search, chat, images, templates, export · a pull request into the public repository, use-cases folder

Finish and FF&E Schedule Builder

WHO IT SERVES

Interior design / architecture · designer assembling the schedule that drives procurement

WHAT TO BUILD

From selected products (cut sheets, supplier pages, photographs) build the finish and furniture schedule: a row per item with code, room assignment, manufacturer, model, finish, dimensions, lead time and price, plus a product image, assembled into the practice's schedule format. Then generate the room-by-room specification narrative from the same data so the two never disagree. A product swap updates every row and every narrative mention as one reviewable batch.

WHAT STRONG LOOKS LIKE

data pulled from supplier sheets is accurate and attributed; a swap updates schedule and narrative together; the schedule exports with images intact and columns aligned.

BAND S2 · Multi-document, images, search, templates, export · a pull request into the public repository, use-cases folder

Design-Review Decision Register

WHO IT SERVES

Architecture / design · project lead running client and stakeholder design reviews

WHAT TO BUILD

After each review, feed in the marked-up drawings, the meeting recording or notes and the client's comments. The app produces the decision register: each item numbered with the decision, the person who made it, the date, the drawings or documents affected, and the cost or programme implication where stated. It then stages the corresponding edits to the design narrative and the specification, and carries forward every open item to the next review's agenda automatically.

WHAT STRONG LOOKS LIKE

decisions are attributed correctly; open items genuinely persist across reviews without being re-entered; the staged spec edits match the decisions and nothing else.

BAND S2 · Multi-document, images, search, Review, memory, export · a pull request into the public repository, use-cases folder

Sustainability Credit Documentation Packet

WHO IT SERVES

Architecture / AEC · sustainability lead assembling a green-certification submission

WHAT TO BUILD

Per pursued credit, assemble the documentation packet: the credit narrative written against the rating system's stated requirement, the calculation with its inputs shown, the evidence index listing each supporting document with where it lives, and a status flag (documented, evidence outstanding, at risk). The whole submission assembles as a multi-document bundle with a scorecard cover. Changing a credit's status updates the scorecard and the narrative together.

WHAT STRONG LOOKS LIKE

narratives are written against the actual requirement text attached, not from general knowledge; the evidence index matches what is really in the bundle; the scorecard arithmetic is right after edits.

BAND S2 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Project-Sheet Portfolio Factory

WHO IT SERVES

Architecture / engineering · marketing manager who rebuilds the same project sheets for every pursuit

WHAT TO BUILD

Maintain one project record per job (facts, imagery, credits, metrics, awards) and generate project sheets on demand in whichever shape the pursuit demands: the practice's standard one-pager, a client-requested format, a government pre-qualification form, or a slimmed version for a shortlist submission. Select projects by semantic match to the pursuit brief rather than by tag. A fact corrected once corrects every future sheet.

WHAT STRONG LOOKS LIKE

the same project renders correctly in three different required formats; selection by brief beats selection by tag on a blind test; a corrected fact never reappears in its old form.

BAND S2 · search, templates, images, Multi-document, memory, export · a pull request into the public repository, use-cases folder

Manual Revision-Service Diff Applier

WHO IT SERVES

Aviation · technical publications lead at an airline or an operator maintaining customized manuals

WHAT TO BUILD

The manufacturer or the revision service issues a change package against the base manual. The operator's manual is not the base manual: it carries years of operator-specific customization. The app diffs the incoming revision against the operator's live document, identifies only the paragraphs genuinely affected, and stages each as a reviewable edit that shows the incoming text, the operator's current text, and the merged proposal, flagging every place where the revision collides with a customization so a human decides. Untouched customizations must survive perfectly.

WHAT STRONG LOOKS LIKE

operator customizations survive a revision cycle; collisions are surfaced rather than resolved; paragraphs unrelated to the revision are byte-identical afterwards.

BAND S3 · Multi-document, search, Review, chat, export · a pull request into the public repository, use-cases folder

Regulator Cross-Reference Compliance Matrix

WHO IT SERVES

Aviation · compliance monitoring manager preparing for an authority audit

WHAT TO BUILD

Attach the applicable regulation set and the operator's manual suite as reference documents. The app builds the compliance cross-reference matrix: each regulatory requirement mapped to the manual paragraph that satisfies it, with the paragraph quoted, plus a list of requirements with no mapping and manual paragraphs that claim compliance with a requirement that no longer exists. Then it drafts the missing paragraphs into the right manual and section for review.

WHAT STRONG LOOKS LIKE

mapping accuracy holds across a large regulation set; orphans are found in both directions; drafted paragraphs match the surrounding manual's voice and numbering.

BAND S3 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Revision Bars and Effective-Pages Generator

WHO IT SERVES

Aviation · technical publications specialist issuing a controlled revision

WHAT TO BUILD

Controlled documents ship with machinery around the change: change bars in the margin next to every altered paragraph, a revision record table listing revision number, date, affected pages and a summary of change, a list of effective pages, and a highlights-of-change summary for crew. After an editing session, the app generates all of it automatically from what actually changed, stamps the headers and footers with the revision number and date, and exports the controlled PDF.

WHAT STRONG LOOKS LIKE

change bars align with real changes and nothing else; the effective-pages list matches the exported pagination; headers and footers carry the revision identity on every page.

BAND S2 · chat, export, templates, Review · a pull request into the public repository, use-cases folder

Crew Bulletin and Acknowledgment Pack

WHO IT SERVES

Aviation · flight operations manager pushing a time-critical notice to a multinational crew base

WHAT TO BUILD

From the source change (a regulation update, an operational restriction, a procedure amendment), draft the crew bulletin in the operator's bulletin format, produce the plain-language crew summary, generate the read-and-sign acknowledgment form, and localize the whole set for each crew base's working language while keeping the technical terms and any mandated phraseology untouched. Track which bulletin supersedes which, and produce a distribution index.

WHAT STRONG LOOKS LIKE

mandated phraseology survives localization verbatim; supersession is tracked correctly; every language version carries the same content in the same structure.

BAND S2 · chat, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Occurrence Report Writer

WHO IT SERVES

Aviation · safety manager turning raw reports into filed occurrence reports

WHAT TO BUILD

A crew member's free-text report arrives with attachments (photographs, a technical log page, a screenshot of a system message). The app drafts the formatted occurrence report the authority expects: factual narrative in the required order, classification, contributing factors, immediate action taken, and recommendations, with the personal identifiers removed and replaced by role descriptors for the de-identified copy. Two outputs: the internal report with identities and the filed report without, both from one source.

WHAT STRONG LOOKS LIKE

de-identification is complete including indirect identifiers such as rosters and aircraft registrations where required; the narrative adds nothing the reporter did not say; both outputs stay in sync when the source is corrected.

BAND S2 · chat, images, Review, templates, export · a pull request into the public repository, use-cases folder

Defect Log to Work Package

WHO IT SERVES

Aviation maintenance · planner turning reported defects into actionable work

WHAT TO BUILD

Take the technical log defect entries plus photographs and turn each into a work package document: defect description, applicable manual task references found by searching the attached maintenance manual, parts and tooling list, deferral assessment against the equipment list with the deferral category and time limit stated, and the sign-off block. Where a defect is deferred, the app also drafts the deferral entry and the crew-facing placard note.

WHAT STRONG LOOKS LIKE

manual task references are real and found by search rather than guessed; deferral categories and limits match the attached source; photographs land with the right defect.

BAND S3 · search, Multi-document, images, templates, export · a pull request into the public repository, use-cases folder

Training Syllabus and Assessment Forms

WHO IT SERVES

Aviation · training manager maintaining an approved training footprint

WHAT TO BUILD

Maintain the training syllabus, lesson plans, briefing guides and assessment forms as one linked set. When a procedure changes in the operations manual, the app finds every training document that teaches or assesses that procedure and stages the matching edits, so the training footprint never drifts from the operational one. Assessment forms regenerate with the updated competency wording and keep their grading structure. Produces a change-impact report before anything is edited.

WHAT STRONG LOOKS LIKE

the impact report finds every dependent document including indirect references; grading structures survive the wording update; nothing outside the impacted scope changes.

BAND S3 · Multi-document, search, Review, templates, export · a pull request into the public repository, use-cases folder

Crew Brief Digest

WHO IT SERVES

Aviation · dispatcher or flight-ops officer producing the pre-flight briefing package

WHAT TO BUILD

Raw operational text (notices to air missions, weather reports and forecasts, company notices, route restrictions) arrives as dense abbreviated strings. The app produces the human-readable crew brief: decoded and grouped by relevance to the specific route and time window, the significant items lifted to the top with a plain-language explanation, a route diagram, and the raw source preserved in an appendix so nothing is lost. Export as the one-page brief plus the full package.

WHAT STRONG LOOKS LIKE

decoding is accurate and the significance ranking matches what an experienced dispatcher would choose; the raw source is preserved verbatim; the one-page brief actually fits one page.

BAND S2 · chat, images, search, export · a pull request into the public repository, use-cases folder

Owner Manual Multi-Market Re-Issue

WHO IT SERVES

Automotive · technical documentation lead shipping one vehicle into many markets

WHAT TO BUILD

A feature changes mid-cycle. Rather than reissuing the whole owner manual, the app locates every affected passage across the manual and its market variants, stages the edit in each, and keeps market-specific differences (units, regulatory warnings, feature availability, tyre pressure tables) intact. Localized versions update in the same pass with mandated warning wording preserved verbatim per market. Output is a per-market change report plus the updated manuals.

WHAT STRONG LOOKS LIKE

market-specific content survives a global change; mandated warnings are never paraphrased; the change report is accurate per market and per section.

BAND S3 · Multi-document, search, Review, export, images · a pull request into the public repository, use-cases folder

Dealer and Franchise Agreement Redline Desk

WHO IT SERVES

Automotive · network-development or legal counsel managing hundreds of dealer agreements

WHAT TO BUILD

A dealer returns the standard agreement with their own markup. The app compares it against the approved master and the negotiation playbook (which deviations are pre-approved, which need escalation, which are refusals), classifies every deviation, drafts the counter-position language for each, and produces a negotiation summary for the approver. Approved positions apply back into the agreement as tracked changes ready to return to the dealer.

WHAT STRONG LOOKS LIKE

deviation detection catches meaning changes not just wording changes; playbook classification is correct including partial matches; the returned file carries genuine tracked changes.

BAND S3 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Type-Approval Dossier Assembler

WHO IT SERVES

Automotive · homologation engineer preparing a market approval submission

WHAT TO BUILD

Assemble the approval dossier per market: the information document with its numbered sections, test reports indexed and cross-referenced, technical descriptions, drawings and photographs, and the declaration of conformity. The app checks the assembled set against the target market's required content list, flags what is missing or expired, and maintains cross-reference numbering across the bundle so an inserted test report rennumbers everything that points at it. Produces the submission index as the cover document.

WHAT STRONG LOOKS LIKE

cross-references survive insertion and deletion; the completeness check is per-market and honest; the index matches the assembled bundle exactly.

BAND S3 · Multi-document, search, templates, images, export · a pull request into the public repository, use-cases folder

Repair-Order Narrative Normalizer

WHO IT SERVES

Automotive · dealer service manager or warranty administrator whose claims get denied for bad paperwork

WHAT TO BUILD

Technicians write shorthand. Warranty claims get denied when the repair-order story does not clearly state the complaint, the cause and the correction. The app takes the raw technician note plus the diagnostic codes and turns it into a compliant narrative in the manufacturer's required structure, keeping every technical fact and adding none, flagging where the technician's note is genuinely missing a required element instead of inventing it. Runs over a day's worth of orders in batch with a per-order confidence flag.

WHAT STRONG LOOKS LIKE

no fact is invented and none is lost; missing elements are flagged rather than filled; a batch of 100 orders produces consistent structure without homogenizing genuine differences.

BAND S2 · API, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Recall Regulator and Dealer Instruction Set

WHO IT SERVES

Automotive · product-safety lead running a recall campaign

WHAT TO BUILD

One campaign produces the regulator-facing side and the network-facing side from one source of truth: the defect information report in the authority's required structure, the quarterly completion report template with its data tables, the dealer service instruction with step-by-step procedure and diagrams, the parts ordering note, and the technician sign-off sheet. Correcting the defect description or the affected build range updates every artefact as a reviewable batch.

WHAT STRONG LOOKS LIKE

a single source change propagates everywhere; regulator formats are followed exactly; the dealer instruction is genuinely followable with diagrams in step order.

BAND S3 · Multi-document, templates, images, Review, export · a pull request into the public repository, use-cases folder

Supplier Submission Binder Assembler

WHO IT SERVES

Automotive supply chain · supplier quality engineer assembling a production part approval binder

WHAT TO BUILD

The approval binder is a fixed multi-element bundle (design records, process flow, control plan, measurement studies, capability results, appearance approval, part submission warrant and more). The app assembles the binder from whatever the supplier has, in the required element order, with a cover index and per-element status, flags missing or superseded elements, checks internal consistency (the part number and revision on every element must match), and exports the whole binder as one navigable document plus the individual elements.

WHAT STRONG LOOKS LIKE

cross-element consistency checking actually catches a mismatched revision; the index is accurate; a missing element blocks submission rather than being silently skipped.

BAND S2 · Multi-document, search, templates, export · a pull request into the public repository, use-cases folder

Enterprise Network Bid Pack

WHO IT SERVES

Telecom · bid manager at a carrier responding to a multi-site enterprise tender

WHAT TO BUILD

Enterprise network tenders are structurally different from generic proposals: the answer is mostly annexes. From the customer's site list and requirement matrix, assemble the bid pack: the response narrative, the per-site solution annex (access type, bandwidth, resilience, delivery window) generated per site, the service-level annex assembled from the standard catalogue with the customer's negotiated variations, and a pricing schedule that stays arithmetically consistent with the site list. A site added late updates every annex. Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S3 · Multi-document, templates, search, chat, export · a pull request into the public repository, use-cases folder

Telecom Regulatory Filing Drafter

WHO IT SERVES

Telecom · regulatory affairs manager filing with a national authority

WHAT TO BUILD

Filings (licence applications, spectrum submissions, universal-service reports, tariff notifications) each have a mandated structure and mandated attachments. From the operator's data and the authority's published form requirements attached as reference, the app drafts the filing in the required structure, assembles the annexes, cross-references them from the body, and produces a completeness checklist against the requirement list. Prior-year filings are reused where the answer has not changed, marked as carried forward.

WHAT STRONG LOOKS LIKE

structure matches the authority's requirement exactly; carried-forward answers are marked and verified as still true; cross-references to annexes resolve.

BAND S2 · search, Multi-document, templates, memory, export · a pull request into the public repository, use-cases folder

Rate-Card Amendment Propagation

WHO IT SERVES

Telecom · wholesale or interconnect manager repricing across a partner base

WHAT TO BUILD

Wholesale rates change. Each partner contract embeds its own rate schedule inside its own negotiated document. The app opens the affected contracts, updates only the rate tables while leaving the negotiated terms alone, generates the formal amendment letter per partner referencing the correct clause and effective date, and produces a control sheet listing every partner, the old and new rates, and the notice period each contract requires. Partners whose contract forbids the change are flagged, not amended.

WHAT STRONG LOOKS LIKE

only rate tables change and negotiated terms are provably untouched; contracts with a pricing lock are detected and excluded; notice periods are read from each contract rather than assumed.

BAND S3 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Post-Implementation Evidence Pack

WHO IT SERVES

Telecom · change manager closing out a network change window

WHAT TO BUILD

After the change window, produce the post-implementation report from the method of procedure plus the evidence: step-by-step actual versus planned with timings, command output and screenshots placed against their step, validation test results, deviations and their approvals, and the rollback decision record. Where a step was skipped or altered, the report says so plainly. Export as the closure document that the change board signs, with the raw evidence in an appendix.

WHAT STRONG LOOKS LIKE

actual versus planned is honest about deviations; evidence lands beside the right step; the report is complete enough that the board does not ask for the raw logs.

BAND S2 · Multi-document, images, chat, Review, export · a pull request into the public repository, use-cases folder

Site Acquisition Packet Builder

WHO IT SERVES

Telecom · site acquisition specialist rolling out towers, small cells or fibre

WHAT TO BUILD

Per candidate site, assemble the packet: landlord lease or licence on the operator's template with the site's specifics, access agreement, planning or zoning application in the local authority's required shape, site survey report with photographs and a location diagram, and the structural or environmental annexes. Terms negotiated on one site (rent escalation, term length, access hours) can be applied to a whole batch of similar sites in one pass with a per-site variance report.

WHAT STRONG LOOKS LIKE

local authority formats differ correctly between jurisdictions; batch application of a negotiated term does not overwrite site-specific exceptions; photographs and diagrams place correctly per site.

BAND S2 · Multi-document, templates, images, API, export · a pull request into the public repository, use-cases folder

Provisioning and Service-Order Packet

WHO IT SERVES

Telecom · service delivery manager onboarding a business customer

WHAT TO BUILD

From the signed order, assemble everything delivery needs: the service order form, the site survey request, the installation scope of work with an as-designed diagram, the customer readiness checklist, the acceptance test plan, and the handover certificate. Each document draws from one order record, so a change to the delivery date or the bandwidth updates every dependent document. Produces the customer-facing pack and the internal pack separately from the same source.

WHAT STRONG LOOKS LIKE

internal and customer packs differ only in what should be hidden; a date change updates every document that carries it; the diagram reflects the ordered configuration.

BAND S2 · Multi-document, templates, images, API, export · a pull request into the public repository, use-cases folder

Vendor Response Evaluation Report

WHO IT SERVES

Telecom · procurement or technical evaluation lead comparing six vendor responses

WHAT TO BUILD

Load the requirement set and every vendor response. The app extracts each vendor's answer per requirement, normalizes the wording so answers are comparable, scores against the stated criteria with the evidence quoted, and produces the evaluation report: a comparison matrix, per-vendor strengths and gaps with citations, unanswered requirements per vendor, and a recommendation section the evaluation panel edits before sign-off.

WHAT STRONG LOOKS LIKE

answers are matched to requirements correctly even when vendors answer out of order; every score cites the vendor's own words; non-answers are marked as non-answers rather than scored generously.

BAND S3 · Multi-document, search, chat, Review, export · a pull request into the public repository, use-cases folder

Regulator Outage Report

WHO IT SERVES

Telecom · incident manager filing a mandatory outage report after a major service failure

WHAT TO BUILD

From the incident timeline, monitoring evidence and the internal post-mortem, produce the regulator's mandatory outage report in its required structure: affected services and customer counts, timeline to the minute, root cause, restoration actions, preventive measures with owners and dates, and the emergency-services impact statement. A second output is the customer-facing explanation, and a third is the executive summary, all generated from the same facts so no version contradicts another.

WHAT STRONG LOOKS LIKE

the three audiences get genuinely different documents that never contradict each other; the timeline is reconstructed accurately from evidence; regulator formatting is exact.

BAND S2 · Multi-document, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Milestone Acceptance Document

WHO IT SERVES

Gaming · producer at a studio delivering against a publisher-funded milestone schedule

WHAT TO BUILD

Milestone payments hang on a document nobody enjoys writing. From the milestone definition in the publishing agreement plus the build's actual state (task exports, build notes, screenshots, capture footage stills), produce the milestone acceptance document: each contracted deliverable listed with its evidence, a status of delivered, partial or deferred, screenshots placed against the feature they prove, and a deferral note with the agreed new date for anything not delivered. Export the version the publisher signs.

WHAT STRONG LOOKS LIKE

every contracted deliverable is accounted for including the ones that slipped; evidence sits next to its claim; the document reads as honest rather than defensive.

BAND S2 · Multi-document, images, search, templates, export · a pull request into the public repository, use-cases folder

In-Game Text Localization Kit

WHO IT SERVES

Gaming · localization producer sending a game's text to translation vendors

WHAT TO BUILD

Game text is not prose: it is thousands of short strings with character limits, speaker context, gender and plural variables, and formatting tokens that must survive untouched. The app builds the localization kit per language: a context document explaining the game world and tone, the string tables laid out with the source, the character limit, the speaker and the scene note per row, a do-not-translate list covering tokens and proper nouns, and a style sheet. Returned translations are checked back against limits and tokens with violations flagged per row.

WHAT STRONG LOOKS LIKE

formatting tokens survive the round trip untouched; character-limit checking accounts for the target script's width; context notes are genuinely useful to a translator who has not played the game.

BAND S3 · Multi-document, templates, search, Review, export · a pull request into the public repository, use-cases folder

Age-Rating Submission Pack

WHO IT SERVES

Gaming · compliance or publishing lead submitting for content ratings across territories

WHAT TO BUILD

Every rating board asks similar questions in different shapes and different thresholds. From one content disclosure record (violence, language, themes, monetization, user interaction, gambling mechanics) plus supporting evidence stills and footage notes, the app generates the submission questionnaire response per board, in each board's format and vocabulary, with the evidence index attached. Where two boards would classify the same content differently, it flags the divergence so the publisher can plan for it.

WHAT STRONG LOOKS LIKE

the same content produces correctly different answers per board; divergences are flagged rather than smoothed; nothing is understated relative to the source disclosure.

BAND S2 · Multi-document, templates, images, search, export · a pull request into the public repository, use-cases folder

Game Press Kit Builder

WHO IT SERVES

Gaming · indie developer or PR lead preparing a launch press kit

WHAT TO BUILD

Build the press kit journalists and creators expect: fact sheet (studio, release date, platforms, price, availability), description in three lengths (one line, one paragraph, long form), feature list, history and inspiration section, verbatim quote block, awards and coverage list, and an asset index describing every screenshot, key art and trailer still with its caption and credit. Regenerate the whole kit in another language keeping the asset index and the fact sheet identical.

WHAT STRONG LOOKS LIKE

the three description lengths say the same thing at different resolutions rather than being truncations; the asset index matches the actual assets; the localized kit keeps facts identical.

BAND S1 · chat, images, templates, export · a pull request into the public repository, use-cases folder

Portfolio and Case-Study Builder for Creatives

WHO IT SERVES

Creative portfolio · freelance artist, animator or designer applying for studio work

WHAT TO BUILD

Turn a body of work into a portfolio document that reads as a case-study set rather than an image dump: per project, the brief, the role and contribution, the process shown in stages, the final work, and the outcome, with imagery placed and captioned. Then produce role-targeted variants from the same master (an environment-art variant, a technical-art variant) that reorder and re-emphasize without duplicating the source. Export a print-ready PDF and a lightweight version under a submission size limit.

WHAT STRONG LOOKS LIKE

role variants genuinely re-emphasize rather than reorder captions; the size-limited export drops quality, not content; captions and credits stay attached to the right image.

BAND S2 · chat, images, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Storefront Copy Pack per Platform

WHO IT SERVES

Gaming · marketing lead publishing one title across several storefronts and regions

WHAT TO BUILD

Each storefront has its own copy shape: different character limits per field, different required disclosures, different screenshot counts and aspect ratios, different tone conventions. From one master description, produce the per-store copy pack: every field filled within its limit, required disclosures included verbatim, an asset checklist per store, and the localized variants per region with limits re-checked against the target script. A change to the master reissues only the fields it touches.

WHAT STRONG LOOKS LIKE

limits are respected per script rather than per character count alone; required disclosures appear verbatim; a master edit updates only affected fields across every store and language.

BAND S2 · Multi-document, templates, chat, export · a pull request into the public repository, use-cases folder

Allergen and Nutrition Sheet Generator

WHO IT SERVES

Food & Beverage · food-safety or product lead at a multi-site restaurant group or a packaged-food brand

WHAT TO BUILD

From recipe data (ingredients, quantities, supplier ingredient declarations, preparation method), generate the allergen matrix and nutrition sheet per dish: allergens listed by the target market's declared list, cross-contact statements based on the kitchen's stated handling, per-portion and per-100g values with the calculation basis shown, and the market's mandatory wording. Recipe changes reissue only the affected dishes. Missing supplier declarations block the sheet and appear on an outstanding-information list instead of being estimated.

WHAT STRONG LOOKS LIKE

allergens derive from supplier declarations rather than ingredient names alone; missing data blocks rather than estimates; a recipe change reissues exactly the affected sheets.

BAND S1 · API, templates, search, Review, export · a pull request into the public repository, use-cases folder

Seasonal Menu Re-Issue

WHO IT SERVES

Food & Beverage · operator reprinting the designed menu every season or every price rise

WHAT TO BUILD

The existing designed menu is the template, not a starting point. Feed the changes (six dishes out, four in, eighteen prices up, one section renamed) and the app edits the menu in place: dish entries replaced in position, prices updated inside their own typography, removed dishes closed up without reflowing the whole page, new dishes inserted into the right section with matching styling, and dish photography swapped. Export print-ready plus a proofing sheet listing every change for the chef to sign off.

WHAT STRONG LOOKS LIKE

the reissued menu is visibly the same designed artefact; untouched sections are byte-identical; the proofing sheet matches the changes exactly.

BAND S2 · chat, images, Review, export · a pull request into the public repository, use-cases folder

Recipe Costing and Re-Cost Pass

WHO IT SERVES

Food & Beverage · head chef or operations manager protecting gross margin

WHAT TO BUILD

Build the costed recipe document per dish: ingredient lines with unit cost, yield and waste factors, portion cost, sell price, gross margin, and a plating note. Then the re-cost pass: when supplier prices change, the app updates every affected ingredient line across every recipe, recalculates margins, and produces an exceptions report listing dishes that have fallen below the target margin with a suggested price or a suggested specification change. The recipe documents keep their kitchen-facing layout throughout.

WHAT STRONG LOOKS LIKE

arithmetic stays consistent across recipe, portion and margin after a bulk price change; only affected lines are touched; the exceptions report ranks by money at risk, not by percentage alone.

BAND S2 · Multi-document, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Multi-Language Menu and Delivery Listing Pack

WHO IT SERVES

Food & Beverage · multi-site or franchise operator listing on several delivery platforms in several languages

WHAT TO BUILD

One menu source becomes the pack: the in-restaurant menu, the per-platform listing documents with each platform's field limits and required allergen or calorie disclosure, and the language versions for the local market, with dish names handled by the operator's rule (translate, transliterate or keep). Prices differ per platform because commission differs, so the pack carries platform-specific pricing derived from one base price and a stated uplift rule.

WHAT STRONG LOOKS LIKE

platform-specific pricing derives correctly and is auditable back to the base price; dish-name rules are applied consistently across languages; disclosures appear verbatim where mandated.

BAND S2 · Multi-document, templates, chat, export · a pull request into the public repository, use-cases folder

Audit Evidence Binder for Food Safety

WHO IT SERVES

Food & Beverage · quality manager two weeks out from a third-party food-safety audit

WHAT TO BUILD

Assemble the audit binder against the scheme's clause list: for each clause, the policy or procedure that answers it, the records that evidence it, and a status of complete, partial or missing with what is missing named. Records arrive as photographs of paper logs, spreadsheets and signed forms, so the binder handles mixed sources and indexes them. Output is one navigable binder with a clause-by-clause cover matrix plus a prioritized pre-audit action list.

WHAT STRONG LOOKS LIKE

photographed paper records are read and indexed correctly; partial coverage is called partial rather than complete; the action list is ordered by audit risk.

BAND S2 · Multi-document, images, search, templates, export · a pull request into the public repository, use-cases folder

Franchisee Opening Packet

WHO IT SERVES

Food & Beverage / franchise · franchise operations lead opening a new location

WHAT TO BUILD

From the site record, assemble the opening packet: pre-opening checklist with dates worked back from the opening day, local licence and permit application drafts on the relevant authority's forms, supplier account setup forms, staff training record templates, the opening-week roster and prep plan, and the local marketing pack. Site-specific facts (address, licensee name, opening date, seating count) fill everywhere they appear, and a date slip regenerates the whole schedule.

WHAT STRONG LOOKS LIKE

a date slip cascades correctly through the whole checklist; local authority forms differ correctly between jurisdictions; the packet exports as one bundle the franchisee can work from.

BAND S2 · Multi-document, templates, chat, API, export · a pull request into the public repository, use-cases folder

Event Order to Kitchen Production Plan

WHO IT SERVES

Food & Beverage / catering · executive chef turning confirmed events into kitchen work

WHAT TO BUILD

Take the week's confirmed event orders and produce the kitchen-facing plan: an aggregated pull list across all events, per-station prep sheets with quantities scaled to headcount, a timed production schedule working back from each service time, allergen and dietary alerts flagged per event with the guest count affected, and a packing and transport list for off-site events. One event's headcount change updates every dependent quantity and time.

WHAT STRONG LOOKS LIKE

aggregation across events is arithmetically right; scaling respects yield rather than multiplying naively; a headcount change updates quantities, timings and the pull list together.

BAND S3 · Multi-document, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Client Progress and Re-Assessment Report

WHO IT SERVES

Fitness · personal trainer or studio owner running quarterly client reviews

WHAT TO BUILD

At the re-assessment point, generate the client's progress report: baseline versus current measurements in a comparison table, charted trends, adherence summary, what changed in the programme and why, and the next block's goals. The report is generated as an update to the client's existing plan document so the whole engagement reads as one continuous record rather than a pile of separate PDFs. Branded, exportable, and produced in batch across a coaching roster.

WHAT STRONG LOOKS LIKE

the same client's second and third reports extend one document coherently; charts match the table; a 40-client batch produces genuinely individual narratives.

BAND S1 · API, chat, images, templates, export · a pull request into the public repository, use-cases folder

Corporate Wellness Proposal and Reporting Pack

WHO IT SERVES

Fitness / wellness · B2B sales lead selling a wellness programme into an employer

WHAT TO BUILD

Two linked builds. The proposal: programme design, session schedule, staffing, pricing, expected participation and outcome measures, on the provider's brand and tailored to the employer's stated goals. Then the quarterly report on the same engagement: participation rates, session delivery against plan, anonymized outcome measures, testimonials, and a renewal recommendation, generated as an update to the same document family so the proposal's promises are visibly tracked against delivery.

WHAT STRONG LOOKS LIKE

the report explicitly measures against the proposal's own stated commitments; individual participants are never identifiable; the renewal case is evidenced rather than asserted.

BAND S2 · Multi-document, chat, images, templates, export · a pull request into the public repository, use-cases folder

Certification Course Workbook Builder

WHO IT SERVES

Fitness / wellness education · course provider running trainer certifications

WHAT TO BUILD

From the syllabus, build the course document set: the learner workbook with diagrams and practical checklists, the assessor guide with marking criteria, the practical assessment forms, and the certificate of completion. When the governing body updates a competency, the app finds every place across the set that teaches or assesses it and stages the edits, keeping the workbook's page structure stable so printed stock stays usable where possible. Certificates generate in batch per cohort.

WHAT STRONG LOOKS LIKE

a competency change reaches workbook, assessor guide and assessment form together; page structure is preserved where it can be; certificate batches personalize without cross-contamination.

BAND S2 · Multi-document, search, images, templates, API, export · a pull request into the public repository, use-cases folder

Screening to Clearance Letter Pack

WHO IT SERVES

Fitness / wellness · studio manager handling clients who flag health conditions at intake

WHAT TO BUILD

A completed pre-exercise screening questionnaire comes back with flags. The app produces the pack: the medical clearance request letter addressed to the client's physician summarizing the intended programme and the specific concerns, the client-facing explanation of why clearance is needed, the modified-programme note the coach works from until clearance arrives, and the file note for the record. Every clinical statement stays inside a fitness scope of practice, and anything outside it is escalated rather than written.

WHAT STRONG LOOKS LIKE

the letter is specific to the flagged conditions and the actual programme; scope-of-practice boundaries hold under adversarial inputs; the client-facing note is reassuring without minimizing.

BAND S2 · chat, templates, Review, export · a pull request into the public repository, use-cases folder

Studio Policy and Staff Handbook Pack

WHO IT SERVES

Fitness · studio or gym owner opening a second and third location

WHAT TO BUILD

Generate and then maintain the operating document set: member policies, staff handbook, opening and closing procedures, incident and emergency procedures, cleaning and equipment maintenance schedules, and the class cancellation policy. Location-specific facts and any jurisdictional differences live in one record. When a policy changes, the app edits every location's version, produces the staff change bulletin and the acknowledgment form, and updates the version history.

WHAT STRONG LOOKS LIKE

location differences survive a global policy change; the bulletin describes exactly what changed; version history is generated rather than hand-kept.

BAND S2 · Multi-document, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Membership Lifecycle Letter Pack

WHO IT SERVES

Fitness · membership or retention manager handling freezes, transfers and cancellations

WHAT TO BUILD

Every membership change needs paperwork that matches the signed agreement, not a generic template. The app opens the member's actual agreement, applies the requested change as a formal amendment (freeze period, transfer of holder, downgrade, cancellation with the correct notice calculated from the agreement's own terms), and produces the confirmation letter stating exactly what changed and from when. Where the agreement does not permit the request, it says so and cites the clause instead of proceeding.

WHAT STRONG LOOKS LIKE

notice periods are read from the member's own agreement rather than a default; impermissible requests are refused with a citation; the amendment attaches to the original cleanly.

BAND S1 · chat, search, templates, Review, export · a pull request into the public repository, use-cases folder

Jurisdiction Permit-Format Switcher

WHO IT SERVES

Solar · permitting coordinator at an installer working across dozens of authorities

WHAT TO BUILD

The same installation needs a differently shaped permit packet in every jurisdiction: different forms, different plan-set note requirements, different attachments, different signature blocks. From one project record and one master plan-set narrative, the app produces the packet in the target authority's required shape, maps the project's data into that authority's forms, flags requirements the project cannot yet satisfy, and produces a submission checklist. Switching target jurisdiction reshapes the packet without re-entering the project.

WHAT STRONG LOOKS LIKE

two jurisdictions produce genuinely different, correct packets from identical project data; unsatisfiable requirements are flagged rather than left blank; the checklist matches the assembled packet.

BAND S3 · Multi-document, templates, search, images, export · a pull request into the public repository, use-cases folder

Interconnection Standard to Compliance Checklist

WHO IT SERVES

Solar · engineering lead facing a utility's interconnection standard for the first time

WHAT TO BUILD

Utility interconnection standards are long, dense and different per utility. Upload the standard and the proposed design. The app produces a plain-language summary for the team, a requirement-by-requirement checklist with the standard's own section cited, a compliance assessment of the proposed design against each requirement, and a list of the documents the utility will require at application. Everything cites the source so an engineer can verify without re-reading the standard.

WHAT STRONG LOOKS LIKE

requirements are extracted comprehensively rather than skimmed; every assessment cites a real section; nothing is asserted about a requirement the standard does not contain.

BAND S2 · search, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Customer Finance and Disclosure Set

WHO IT SERVES

Solar · sales operations lead issuing consumer finance paperwork

WHAT TO BUILD

Once a customer picks a finance route (cash, loan, lease or a power-purchase structure), generate the disclosure set that route and that jurisdiction require: the finance agreement on the lender's template with the project's numbers, the mandated consumer disclosures verbatim, the savings and production estimate with its assumptions stated plainly, the cancellation rights notice, and the signature packet index. Assumptions in the savings estimate must be traceable to the proposal so the two never disagree.

WHAT STRONG LOOKS LIKE

mandated disclosure wording is verbatim per jurisdiction; savings figures reconcile to the proposal exactly; changing the finance route regenerates the correct set rather than editing the wrong one.

BAND S2 · Multi-document, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Site Survey to Design Narrative

WHO IT SERVES

Solar · surveyor or designer turning a roof visit into an engineering document

WHAT TO BUILD

A site survey comes back as 40 photographs, a few measurements and some voice notes. The app reads the imagery, drafts the site assessment document: roof structure and condition observations tied to the photographs that show them, shading assessment with the obstructions named, electrical service observations, access and safety notes, and the resulting design assumptions and constraints. Every observation links to its evidence photograph, and anything the photographs do not support is marked as requiring a second visit.

WHAT STRONG LOOKS LIKE

observations are genuinely grounded in the photographs and say so; unsupported claims are refused; the narrative is usable by a designer who did not attend the site.

BAND S3 · images, chat, search, templates, export · a pull request into the public repository, use-cases folder

Fleet Asset Report Batch

WHO IT SERVES

Solar · asset manager reporting monthly across a portfolio of operating sites

WHAT TO BUILD

Each month, produce the per-site asset report from production data, availability, alarms, work orders and inspection notes: production against forecast with a variance explanation, availability, incident summary, open work orders with age, and next-month actions, plus a portfolio-level roll-up. Reports update the same per-site document each month so a site's history reads continuously, and the roll-up ranks sites by underperformance so the asset manager reads one page first.

WHAT STRONG LOOKS LIKE

variance explanations are specific rather than generic; the monthly update preserves layout and history; the roll-up arithmetic reconciles to the site reports.

BAND S2 · API, Multi-document, images, templates, export · a pull request into the public repository, use-cases folder

Shared-mailbox document desk

WHO IT SERVES

Operations, claims, admissions, or support team working a shared inbox. Host: shared mailbox plus a small web console.

WHAT TO BUILD

Documents land in a shared mailbox all day. The desk classifies each one (claim form, invoice, transcript, certificate, contract), routes it to the right handling flow, and drafts the outgoing response document from the matching template. Handlers see a queue: source document on the left, drafted response on the right, approve or edit. Everything is logged against the case reference in the subject line.

WHAT STRONG LOOKS LIKE

classification is right often enough to be useful and honest when unsure rather than confidently wrong. The queue never loses a message. Handling twenty documents takes the time one used to.

BAND S3 · API, search, templates, Review, export, Multi-document · a pull request into the public repository, use-cases folder

Google Sheets add-on: roster to documents

WHO IT SERVES

School administrator, event organiser, HR coordinator, small-business owner. Host app: Google Sheets.

WHAT TO BUILD

One sheet, one row per record, one template. The add-on generates a document per row with the template's exact formatting, writes the resulting file link back into a column of that row, and can email each document to the address in the row. Conditional sections respond to column values. A preview runs the first three rows before committing to five hundred.

WHAT STRONG LOOKS LIKE

five hundred documents are formatting-identical to the template and to each other. Link write-back matches the right row every time. A failed row is reported, not silently skipped.

BAND S1 · API, templates, export, chat · a pull request into the public repository, extensions folder

Power Platform connector and Power Apps component

WHO IT SERVES

Citizen developer or IT pro building internal apps on Power Apps and Dataverse. Host: Microsoft Power Platform.

WHAT TO BUILD

A certified custom connector exposing the document actions (generate from template, edit by instruction, request approval, export in each format) plus a Power Apps canvas component that drops a document view and review panel straight into an app screen. A Dataverse-backed sample app ships with it, so a maker can wire a record to its document in an afternoon without writing a line of code.

WHAT STRONG LOOKS LIKE

the connector passes Microsoft's certification review. Actions are named in maker language, not developer language. The sample app runs on a fresh tenant with no manual setup beyond a key.

BAND S2 · API, templates, Review, export · a pull request into the public repository, extensions folder

Microsoft 365 Copilot agent for documents

WHO IT SERVES

Enterprise user whose employer has standardised on Copilot. Host: Microsoft 365 Copilot.

WHAT TO BUILD

A declarative agent published into the tenant so users can ask Copilot to do the things Copilot cannot do to a document: apply a precise section edit without rewriting the file, produce the branded deliverable, run the review with approvals, export with real headers, footnotes and page geometry. The agent handles the document work and returns the file and a summary of what changed, inside the Copilot conversation the user is already having.

WHAT STRONG LOOKS LIKE

it installs through normal tenant app governance. The handoff is invisible: the user asks in Copilot and gets a correct file back. It is honest about what it did rather than claiming success on a failed edit.

BAND S3 · MCP, API, chat, Review, export · a pull request into the public repository, extensions folder

Reference-manager bridge

WHO IT SERVES

Researcher, graduate student, or research-group administrator. Host: Zotero, Mendeley, or a BibTeX file, plus the document.

WHAT TO BUILD

Connect the library once. Then citations work by meaning inside the document: "cite the 2021 replication study here" inserts the correct citation and adds the entry to the bibliography. Switching target style restyles every citation and the bibliography together. An audit finds citations with no library entry, entries cited nowhere, and duplicated references. Works on documents that arrived from LaTeX as well as ones that started in Word.

WHAT STRONG LOOKS LIKE

the bibliography matches the in-text citations exactly after a style switch. Nothing is cited that does not exist in the library. LaTeX-sourced documents keep their maths intact through the whole process.

BAND S2 · search, chat, export, templates, Multi-document · a pull request into the public repository, extensions folder

Jupyter and Quarto extension: notebook to styled report

WHO IT SERVES

Data scientist or analyst who has to hand a document to a non-technical stakeholder. Host app: JupyterLab, plus a Quarto pipeline.

WHAT TO BUILD

A toolbar button turns the notebook into the report the stakeholder needs: narrative written around the outputs rather than a dump of cells, charts placed as figures with captions and numbering, tables styled, code hidden or included by choice, and the whole thing on the organisation's report template. Re-running after new data updates the numbers in the prose as well as the figures.

WHAT STRONG LOOKS LIKE

numbers stated in the narrative always match the outputs they came from. Figures land at print resolution with correct captions. A re-run after data changes updates every stated figure, not just the charts.

BAND S2 · API, chat, images, templates, export · a pull request into the public repository, extensions folder

Self-hosted chat document bot

WHO IT SERVES

Government, defence, healthcare, or bank team running Mattermost or Rocket.Chat inside its own network.
Host app: Mattermost or Rocket.Chat.

WHAT TO BUILD

The same request-and-approve document flow for organisations that cannot use public chat platforms: slash commands to generate and edit, interactive message buttons for approvals, and channel-scoped template libraries. Deployment is a container and a config file, with an outbound allow-list so a security team can see exactly what leaves the network and nothing else does.

WHAT STRONG LOOKS LIKE

it installs on a self-hosted instance in under an hour from the documentation alone. Network egress matches the documented allow-list exactly. Every feature works without any public chat dependency.

BAND S2 · MCP, API, Review, templates, export · a pull request into the public repository, extensions folder

Box workflow app with metadata write-back

WHO IT SERVES

Enterprise content manager in a regulated industry running Box as the system of record. Host app: Box.

WHAT TO BUILD

Triggered on upload or on a metadata change, the app processes the document (classify, extract key terms, generate the required companion document, normalise to the retention-approved template) and writes the results back into the file's Box metadata template as structured fields. Those fields then drive Box's own retention, search, and workflow rules, so the document work feeds the governance the enterprise already runs.

WHAT STRONG LOOKS LIKE

metadata fields are populated accurately enough to drive retention decisions, and left empty rather than guessed when the document does not say. Governance labels and legal holds are never overridden. Processing is idempotent per file version.

BAND S3 · API, search, templates, export, Review · a pull request into the public repository, extensions folder

Nextcloud app for self-hosted and sovereign deployments

WHO IT SERVES

European public body, university, or privacy-first company running Nextcloud on its own infrastructure. Host app: Nextcloud.

WHAT TO BUILD

A Nextcloud app that adds document editing, review, and export to the files a sovereignty-conscious organisation already keeps in-house: right-click a file to open an editing session, run an instruction, approve the proposed changes, and save a new version back into Nextcloud. Data-flow documentation, a configurable region, and a plain statement of exactly what is sent and what is retained, written for a data-protection officer to read.

WHAT STRONG LOOKS LIKE

it passes the Nextcloud app-store review. Versioning and sharing behave exactly like native Nextcloud. The data-flow document is accurate enough that a DPO signs it without a follow-up call.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

Dynamics 365 solution

WHO IT SERVES

Enterprise sales or service operations on Microsoft Dynamics. Host app: Dynamics 365 with Dataverse.

WHAT TO BUILD

A packaged solution that adds document generation and editing to Dynamics records: quotes and orders in sales, resolution letters in service, contracts in procurement. Uses the same Power Platform connector underneath, so an organisation gets the Dynamics experience plus the ability to build its own flows on the same actions. Includes a sample solution file, security roles, and an upgrade path.

WHAT STRONG LOOKS LIKE

the solution imports cleanly into a fresh environment and upgrades without data loss. Security roles are least-privilege by default. It works identically for a custom entity as for a built-in one.

BAND S3 · API, templates, Review, export · a pull request into the public repository, extensions folder

monday.com marketplace app

WHO IT SERVES

Agency, marketing, or project team running delivery on monday.com. Host app: monday.com.

WHAT TO BUILD

A board view and an item action: turn an item and its updates into the client deliverable, the statement of work, or the status report, on the account's brand template. The generated file attaches to the item, the status column moves, and a per-board template picker lets each client's board produce that client's branded output. An automation recipe fires the same thing when an item enters a status.

WHAT STRONG LOOKS LIKE

the brand template is genuinely per-board, so ten client boards produce ten distinct-looking deliverables. Automation recipes fire once and only once. The attached file opens correctly on a phone.

BAND S2 · API, templates, chat, export, Review · a pull request into the public repository, extensions folder

Platform-parameter axis: Trello joins ClickUp and Asana as a sibling kanban host of the monday.com build.

Odoo module

WHO IT SERVES

SMB or mid-market company running Odoo as its ERP, often in emerging markets. Host app: Odoo.

WHAT TO BUILD

A module that replaces Odoo's report templating for the documents where formatting matters: quotations, invoices, delivery notes, contracts, HR letters. It generates on the company's real branded template, supports per-company and per-language variants for multi-entity groups, and lets a user amend a generated document in plain language before sending. Installs like any other Odoo module and appears in the standard print menu.

WHAT STRONG LOOKS LIKE

it slots into the existing print menu so users need no training. Multi-company and multi-language variants resolve correctly. Amended documents keep the numbering and legal footers the jurisdiction requires.

BAND S2 · API, templates, chat, export, Review · a pull request into the public repository, extensions folder

In-page overlay for web apps with no integration

WHO IT SERVES

Recruiter in an applicant system, teacher in a learning platform, lawyer in a contract system that has no API access on their plan. Host: the browser, inside the vendor's own web app.

WHAT TO BUILD

Some of the highest-pain document work happens inside web applications that offer no useful export and no integration on the customer's plan. This overlay works inside them: it recognises the document area, lifts the content out with its formatting, edits it properly, and puts the result back through the page's own editor. Site profiles define the behaviour per application, so support for a new one is configuration rather than a rewrite.

WHAT STRONG LOOKS LIKE

the paste-back survives the host editor's own sanitisation with the formatting intact. Adding a new site is a profile, not code. It never fights the host application's autosave.

BAND S3 · chat, API, Review, export · a pull request into the public repository, extensions folder

Headless CMS publishing bridge

WHO IT SERVES

Content operations lead at a company running a headless CMS and also owing PDFs to sales and compliance. Host: Contentful, Sanity, Strapi or similar, plus a small service.

WHAT TO BUILD

Author or edit the long-form document once, then publish it two ways: mapped into the CMS as structured entries with the right fields and references, and exported as the branded PDF or Word version for the audiences that need a file. Edits flow back the same way, so the web page and the PDF never drift. A mapping screen defines which document sections become which content fields.

WHAT STRONG LOOKS LIKE

an edit in the document reaches both outputs with no manual step. Reference fields and media map correctly. Republishing does not create duplicate entries.

BAND S3 · API, chat, templates, export, Multi-document · a pull request into the public repository, extensions folder

Enterprise integration-platform connector

WHO IT SERVES

Integration engineer at a large company standardised on an enterprise integration platform such as Workato, Tray, Boomi or MuleSoft. Host: that platform.

WHAT TO BUILD

A certified connector with the document actions modelled as first-class platform operations: generate from template, edit by instruction, submit for approval, poll for completion, fetch each export format. Includes typed input and output schemas so recipes autocomplete, a documented retry and error contract, and a callback pattern for long-running work so a recipe does not sit blocked. Ships with three reference recipes covering the common shapes.

WHAT STRONG LOOKS LIKE

it passes the platform's certification. Long jobs use callbacks rather than blocking a worker. Errors are typed and retryable, so a recipe author never has to guess what went wrong.

BAND S3 · API, templates, Review, export · a pull request into the public repository, extensions folder

Component pack for developer automation registries

WHO IT SERVES

Developer building glue code fast on a hosted automation runtime such as Pipedream. Host: that runtime's component registry.

WHAT TO BUILD

A published set of small, single-purpose components (upload a document, edit a section, generate from template, request approval, export as PDF, wait for a job) that a developer composes in code with no boilerplate and no key handling. Each component ships with a working example workflow. Because the registry is public and searchable, the components become a discovery surface as well as a convenience.

WHAT STRONG LOOKS LIKE

each component does one thing and composes cleanly with the others. Authentication is handled by the platform's connected-account model, never by pasted keys. The examples run unmodified.

BAND S1 · API, chat, templates, export · a pull request into the public repository, extensions folder

Business-process engine human-task connector

WHO IT SERVES

Process engineer at a bank, insurer, or telco modelling workflows in a process engine such as Camunda or Temporal. Host: that engine.

WHAT TO BUILD

Two building blocks for a modelled process: a service task that produces or edits a document, and a human task that is a document approval, complete with the reviewer's before-and-after view, a deadline, escalation, and a decision recorded back into the process instance. A loan approval, a claim decision, or a change request can now include real document work without leaving the process model.

WHAT STRONG LOOKS LIKE

the process instance and the document never disagree about the decision, including after a restart on either side. Timers and escalation behave exactly as modelled. The recorded decision is complete enough to satisfy an audit of the process.

BAND S3 · API, Review, templates, export · a pull request into the public repository, extensions folder

Figma plugin: brand kit to document template

WHO IT SERVES

Brand designer who owns the visual system and keeps being asked for a Word template. Host app: Figma.

WHAT TO BUILD

Read the brand from the Figma file itself, the colour styles, type scale, logo assets, spacing rules, and cover-page layouts, and produce a real document template from it: heading styles mapped to the type scale, table treatments, cover page, running headers and footers, and a sample document showing every element. Update the Figma styles later and regenerate to keep the document system in step with the design system.

WHAT STRONG LOOKS LIKE

the generated template genuinely matches the design system when placed side by side. Fonts and logos are embedded correctly. Regeneration after a style change is clean rather than additive.

BAND S2 · images, templates, export, chat · a pull request into the public repository, extensions folder

Acrobat add-on: annotated PDF to applied edits

WHO IT SERVES

Compliance reviewer, auditor, or client-facing manager whose review cycle happens in PDF comments. Host app: Adobe Acrobat.

WHAT TO BUILD

Reviewers mark up the PDF the way they always have, with sticky notes, highlights, and callouts. The add-on reads every annotation with its author and date, proposes the concrete document edit each one implies, and produces a corrected document with the changes applied, plus a response log stating how each comment was addressed. The reviewer never changes how they review.

WHAT STRONG LOOKS LIKE

every annotation is accounted for, including ones it cannot act on, which are listed as needing a human. Author attribution survives into the response log. Applied edits land in the right place in the source document, not approximately near it.

BAND S3 · API, chat, Review, export, search · a pull request into the public repository, extensions folder

macOS Finder actions and Shortcuts pack

WHO IT SERVES

Mac user who wants document work to be a right-click, not an app. Host: macOS Finder, Quick Actions, and the Shortcuts app.

WHAT TO BUILD

Right-click any document in Finder for the common jobs: reformat to my template, convert to PDF or Word, summarise into a companion file, redact, translate. Each is also a Shortcuts action, so a user can compose them into their own automations, attach them to a folder action, or run them from Spotlight and the menu bar. No window opens unless a decision is needed.

WHAT STRONG LOOKS LIKE

actions appear for the right file types only and complete without a visible app launching. Shortcuts actions expose sensible inputs and outputs so they chain. Results land beside the source with predictable names.

BAND S1 · API, chat, templates, export · a pull request into the public repository, extensions folder

Mobile share-sheet app

WHO IT SERVES

Field worker, salesperson, student, or anyone whose only computer is a phone. Host apps: iOS and Android share sheets.

WHAT TO BUILD

Receive a document in email or a messaging app, share it to this app, type or dictate the instruction, and get the edited document back to share onward. The whole interaction is three taps and one sentence, with review of proposed changes on a screen designed for a thumb. Files are handled by the platform's own document picker, and the app works on a poor connection by queueing and finishing in the background.

WHAT STRONG LOOKS LIKE

a phone-only user completes a real edit on a formatted document without pinching or scrolling sideways. Background completion survives the app being closed. Formatting is not degraded because the device is small.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

macOS menu-bar document desk

WHO IT SERVES

Consultant, writer, or analyst who wants the capability always one keystroke away. Host: macOS menu bar.

WHAT TO BUILD

A global hotkey opens a small window over whatever is on screen. Drop a file on it, paste content into it, or type an instruction about the document already in progress. It keeps recent documents, shows running jobs, and notifies when approvals are waiting. Everything closes back out of the way with one key. It is the launcher pattern applied to documents rather than to search.

WHAT STRONG LOOKS LIKE

it opens instantly and never steals focus at the wrong moment. Dropped files of any supported type work. Notifications about waiting approvals are accurate and never duplicate.

BAND S2 · chat, API, Review, export, Multi-document · a pull request into the public repository, extensions folder

Windows Explorer shell extension with managed install

WHO IT SERVES

Corporate Windows user, deployed by IT. Host: Windows File Explorer.

WHAT TO BUILD

Right-click any document in Explorer, or a whole selection, for the standard jobs: reformat to the corporate template, convert, redact, summarise, send for approval. Ships as an installer IT can deploy silently with the organisation's settings baked in, including which actions appear for which groups. Results write beside the source or to a configured location, with an event-log entry per action for support.

WHAT STRONG LOOKS LIKE

the installer deploys silently and cleanly and uninstalls the same way. Multi-select on fifty files works without stalling Explorer. Group-based action visibility is enforced, not cosmetic.

BAND S2 · API, templates, export, Review · a pull request into the public repository, extensions folder

Legal document-management system integration

WHO IT SERVES

Law firm knowledge or IT lead running a legal document-management system such as iManage. Host app: that system, in Word and in the web client.

WHAT TO BUILD

Documents in a legal DMS are filed to a matter and versioned under strict rules, and lawyers will not work outside it. This integration keeps everything in place: open a document from the matter, edit or redline it with review, and save a new version back with the DMS's own version comment and metadata. The matter's prior documents become the precedent shelf for search, scoped so a matter only ever sees what its ethical wall allows.

WHAT STRONG LOOKS LIKE

versioning, check-in and check-out, and access walls are respected exactly. No document ever leaves the matter it belongs to. Search across precedents never returns a document the user is not entitled to see.

BAND S3 · API, search, Multi-document, Review, export · a pull request into the public repository, extensions folder

Cloud legal-DMS integration

WHO IT SERVES

Mid-size firm or corporate legal department on a cloud legal document system such as NetDocuments. Host app: that system.

WHAT TO BUILD

The same matter-scoped editing and review, built for the cloud DMS model: profile-driven filing, workspace templates, and document-type rules. Generating a new document picks the right profile and workspace automatically, so a drafted engagement letter is filed correctly without anyone choosing a folder. Review and approval happen in the DMS's own security context, and the audit trail lands where the firm's compliance team already looks.

WHAT STRONG LOOKS LIKE

profiles and workspace rules are honoured with no manual filing step. Nothing is created outside the security model. The audit trail is complete and matches the DMS's own record.

BAND S3 · API, templates, Review, export, search · a pull request into the public repository, extensions folder

Signature-ready packet builder for any e-sign vendor

WHO IT SERVES

Sales operations, HR, or legal operations lead whose company has already chosen an e-signature vendor. Host: a service plus a small web console.

WHAT TO BUILD

The gap between "the document is agreed" and "the document is out for signature" is manual everywhere. This builds the packet: assemble the agreement and its exhibits in the right order, generate a cover and index, and place the signature, initial, date, and named-field anchors where each party signs, in the form the chosen vendor understands. One packet definition works across the major vendors, so switching e-signature providers does not mean rebuilding the document flow.

WHAT STRONG LOOKS LIKE

anchors land in the right place across all supported vendors from a single definition. Exhibit order and pagination are correct. A packet that fails validation says why, before it reaches the signer.

BAND S3 · Multi-document, templates, export, Review, API · a pull request into the public repository, extensions folder

Drop-in document panel for any product

WHO IT SERVES

Product or engineering lead at a vertical software company that needs document editing inside their app next quarter, not next year. Host: their own web application.

WHAT TO BUILD

An embeddable panel, one script tag or one web component, that drops a complete document experience into any record page: open, edit by instruction, review proposed changes, and export, with the host application's branding, the host's users, and per-end-customer isolation so one of their customers can never see another's documents or memory. Theming, event hooks, and a permissions model are configuration. The reference implementation shows it wired into a sample vertical app end to end.

WHAT STRONG LOOKS LIKE

it embeds in an unfamiliar application in under a day using only the documentation. Per-end-customer isolation holds under a deliberate attempt to cross it. It looks like part of the host product, not a widget bolted on.

BAND S3 · API, chat, Review, export, Multi-document, memory, templates · a pull request into the public repository, extensions folder

Quote to invoice lifecycle with one document lineage

WHO IT SERVES

Sales operations, agencies, small business finance.

WHAT TO BUILD

Build the chain: quote, revised quote, accepted quote, invoice. Every stage is a real document, each carries a reference to the one before it, and the accepted quote's line items flow into the invoice with nothing retyped. A revision must be visibly a revision, showing what changed against the previous quote, rather than a fresh document with a new number. Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S2 · templates, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Dunning ladder with tone escalation

WHO IT SERVES

Accounts receivable, subscription businesses, agencies.

WHAT TO BUILD

Build the five-stage collection ladder: friendly reminder, firm reminder, formal notice, final demand, and a pre-legal letter. Each stage inherits the invoice facts and escalates the tone deliberately, with the escalation described in a written tone guide the candidate authors. Add a jurisdiction switch that changes required wording and the statutory language where it applies, for two jurisdictions. a collections letter system plus an authored tone guide, useful in any finance or revenue operations role.

WHAT STRONG LOOKS LIKE

the tone ladder is genuinely graded and would not embarrass a finance team at any rung, the jurisdiction switch changes what it must and nothing else, and the facts stay consistent across all five letters.

BAND S1 · templates, chat, export · a pull request into the public repository, use-cases folder

Expense report assembler from receipt images

WHO IT SERVES

Finance operations, travel and expense, consultants.

WHAT TO BUILD

Build the assembler: attach a folder of receipt images, read them visually, produce an itemised expense report on a company template with a category breakdown and a policy-exception section listing anything above a threshold or missing a receipt. Ambiguous receipts must be listed as needing review rather than guessed. a vision-to-document pipeline with policy checking, an appealing engineering portfolio piece.

WHAT STRONG LOOKS LIKE

the exception section is genuinely populated by policy rules, unreadable receipts are surfaced rather than hallucinated into totals, and the arithmetic is verifiable against the source images.

BAND S2 · attachments, chat, templates, export · a pull request into the public repository, use-cases folder

Input scope: receipts in languages other than English are handled; unreadable or non-English fields are surfaced honestly rather than skipped.

Purchase order and goods-received pack

WHO IT SERVES

Procurement, operations, manufacturing.

WHAT TO BUILD

Build the three-way pack: purchase order, goods received note, and the reconciliation memo that matches them against the supplier invoice. The reconciliation memo is the point: it must state, line by line, what matched, what did not, and by how much, and it must be readable by a person who will act on it. a three-way matching document pipeline, a credible finance-systems artifact.

WHAT STRONG LOOKS LIKE

seeded mismatches (quantity, price, a partial delivery, a duplicated line) are all caught and described in ordinary language, and the memo distinguishes a discrepancy from a rounding difference.

BAND S2 · Multi-document, chat, search, export · a pull request into the public repository, use-cases folder

Reconciliation output: the memo identifies WHICH of the three documents (purchase order, goods receipt, invoice) is the one that is wrong, not only that they disagree.

Credit note and correction flow

WHO IT SERVES

Billing operations, finance teams.

WHAT TO BUILD

An invoice is wrong. Build the correct response: a credit note referencing the original invoice's specific lines, a corrected invoice, and a customer-facing explanation letter, all three consistent and all three carrying the original invoice's identifiers. The original invoice must never be silently altered, and the flow must make that impossible. a correction workflow with an immutability guarantee, a strong finance-engineering artifact.

WHAT STRONG LOOKS LIKE

the original is provably untouched, the credit note references specific lines rather than a total, and a partial credit is handled as well as a full one.

BAND S1 · Multi-document, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Statement of account and ageing pack

WHO IT SERVES

Accounts receivable, credit control.

WHAT TO BUILD

Build the monthly statement: opening balance, transactions, closing balance, and an ageing table by bucket, generated for a set of fabricated customers and rendered on one branded template. Then build the part that matters to a credit controller: a one-paragraph narrative per statement summarising the account's position and the single action recommended. Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S1 · API, templates, chat, export · a pull request into the public repository, use-cases folder

Proforma invoice to shipping packet

WHO IT SERVES

Exporters, cross-border sellers, trade finance.

WHAT TO BUILD

Build the app that turns an order into the export money-documents: proforma invoice, commercial invoice, and a customs value declaration, with the currency, incoterm, and tax treatment driven by the destination. The interesting requirement is that changing the destination must correctly change several documents at once, including which fields become mandatory. a cross-border document generator with a rules layer.

WHAT STRONG LOOKS LIKE

a destination change ripples correctly and completely, mandatory-field rules are data-driven, and the candidate documents where they were unsure of a rule rather than guessing confidently.

BAND S2 · API, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Multi-currency, multi-language invoice localiser

WHO IT SERVES

Global SaaS, exporters, marketplaces.

WHAT TO BUILD

Take one invoice and localise it properly for four markets: language, currency formatting, date format, number separators, tax label, and the legally required wording, with the layout holding in every version including at least one right-to-left script. Produce a comparison sheet showing all four side by side. a localisation reference implementation for finance documents, an uncommon and valuable specialisation.

WHAT STRONG LOOKS LIKE

the right-to-left version is genuinely correct rather than mirrored badly, number and date formats are right for each locale, and the total is unambiguous in every version.

BAND S2 · chat, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Structured-data companion for the human-readable invoice

WHO IT SERVES

Fintech, e-invoicing, accounts payable automation.

WHAT TO BUILD

Build both halves and prove they agree: the human-readable invoice document and a structured field extract of the same invoice, plus a reconciliation check that fails loudly when the document and the structured data disagree. Then demonstrate the failure mode deliberately by editing the document and showing the check catching it Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S3 · API, search, chat, export · a pull request into the public repository, extensions folder

Revenue recognition memo drafter

WHO IT SERVES

Accounting, financial reporting, audit.

WHAT TO BUILD

Build a drafter that turns a contract summary into an accounting memo: the performance obligations identified, the transaction price and its allocation, the timing conclusion, and the judgements made, with each conclusion tied to the specific contract clause that drove it. Run it on three fabricated contracts with deliberately different shapes. an accounting memo drafter with clause anchoring, a credible finance or audit technology artifact.

WHAT STRONG LOOKS LIKE

conclusions quote the clause that drove them, the three contracts produce genuinely different memos, and areas of judgement are flagged as judgement rather than stated as fact.

BAND S3 · search, chat, templates, export · a pull request into the public repository, use-cases folder

Budget versus actual variance narrative

WHO IT SERVES

Financial planning and analysis, department heads.

WHAT TO BUILD

Build the monthly reporting pack: a variance table generated from a fabricated ledger, then the written commentary that finance teams actually spend their time on, with a paragraph per material variance naming the driver and the expected trajectory. Rerun with next month's numbers and update only the commentary that changed, keeping stable commentary intact Numbers that live in table rows are carried, never recomputed from prose. Pass source values through unchanged. Do not re-derive a total by reading rows back and adding them up. Where a total is needed, take it from the source data, or compute it in your own code from that source data, never from text a model has produced. Include a check that compares every figure appearing in the narrative against the source row it came from, and surfaces any figure it cannot match instead of printing it. Then tell us what you observed: where this held, where it did not, and what you did about it.

WHAT STRONG LOOKS LIKE

strong work is judged on DETECTION, not on arithmetic being right: the check catches every figure it cannot trace back to a source row and refuses to print it rather than guessing; the candidate's write-up states plainly where carrying values worked and where it did not. Correct totals are necessary but they are not the signal, because the candidate's own verification step is.

BAND S2 · API, chat, templates, export · a pull request into the public repository, use-cases folder

Payslip explainer and compensation letter pack

WHO IT SERVES

Payroll, people operations, employee experience.

WHAT TO BUILD

Build the documents employees actually read: a plain-language payslip explainer for a fabricated payslip, an annual compensation summary, and a compensation-change letter. All three must be produced from one record so that a change to the record updates all three, and the explainer must be genuinely understandable by someone with no payroll knowledge. an employee-communication document pack with a clarity method.

WHAT STRONG LOOKS LIKE

the explainer is tested for clarity, not just written; the three documents never disagree on a number; and the change letter states the effective date and the delta explicitly.

BAND S1 · templates, chat, Multi-document, export · a pull request into the public repository, use-cases folder

Tax invoice field-compliance checker

WHO IT SERVES

Finance operations, marketplaces, accountants.

WHAT TO BUILD

Different jurisdictions require different fields on a tax invoice. Build the checker: read an invoice document, determine which jurisdiction rule set applies, and report every required field that is missing, malformed, or ambiguous, with a suggested fix that can be applied in place on approval. Cover four jurisdictions and cite the public source for each rule set. a sourced multi-jurisdiction rules layer and a checker, a valuable and rare fintech artifact.

WHAT STRONG LOOKS LIKE

the rules are sourced and dated, the checker distinguishes missing from malformed, and applying a fix does not disturb the invoice layout.

BAND S2 · API, search, chat, Review, export · a pull request into the public repository, extensions folder

Freelancer engagement and billing bundle

WHO IT SERVES

Freelancers, contractor marketplaces, small agencies.

WHAT TO BUILD

Build the bundle a contractor sends and bills against: engagement letter, timesheet, invoice, and a payment reminder, all sharing one client record and one rate card. The timesheet must roll into the invoice with no retyping, and a rate change mid-engagement must be handled correctly in both. a billing bundle they can use in their own freelance work immediately.

WHAT STRONG LOOKS LIKE

the mid-engagement rate change splits the invoice correctly by period, totals reconcile to the timesheet, and the bundle is genuinely usable by a real freelancer on day one.

BAND S1 · templates, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Renewal quote and price-change notice batch

WHO IT SERVES

Subscription businesses, customer success, revenue operations.

WHAT TO BUILD

Build the batch: for a list of fabricated accounts renewing, produce a renewal quote and a price-change notice per account, where the notice explains the change in terms specific to that account's plan and usage. Include an approval gate so a human reviews anything above a defined increase before the batch is considered complete. a batch communication pipeline with an approval gate, useful in any revenue operations role.

WHAT STRONG LOOKS LIKE

each notice is genuinely account-specific rather than a merge field swap, the approval gate holds, and the batch reports which accounts it deliberately did not process and why.

BAND S2 · API, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Paper scaffolder from a research question

WHO IT SERVES

Researchers

WHAT TO BUILD

Turns a question, a method and a dataset description into a correctly structured manuscript skeleton on the target journal's template, then fills sections on request.

WHAT STRONG LOOKS LIKE

the skeleton matches the journal's actual required structure, and empty sections are honestly empty rather than filled with plausible filler.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Results narrative from tables and figures

WHO IT SERVES

Researchers

WHAT TO BUILD

Writes the results prose directly from the data, with figure references, numbering and captions that stay correct after reordering.

WHAT STRONG LOOKS LIKE

the numbers in the prose match the tables exactly, and reordering figures rennumbers every reference.

BAND S2 · images, chat, export · a pull request into the public repository, use-cases folder

Research grant application with budget justification

WHO IT SERVES

Principal investigators

WHAT TO BUILD

Assembles a funder-formatted application including budget justification, impact statement and data management plan.

WHAT STRONG LOOKS LIKE

the budget narrative reconciles with the budget table, and every mandatory annex is present.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Dataset datasheet and model card generator

WHO IT SERVES

ML and data teams

WHAT TO BUILD

Produces the documentation artefacts that publication venues and regulators now require.

WHAT STRONG LOOKS LIKE

it prompts for the uncomfortable fields, such as known limitations and collection consent, rather than skipping them.

BAND S2 · templates, chat, export · a pull request into the public repository, use-cases folder

Reproducibility checklist appendix

WHO IT SERVES

Researchers

WHAT TO BUILD

Builds the checklist appendix a venue requires and cross-checks each answer against the manuscript.

WHAT STRONG LOOKS LIKE

it flags checklist answers the manuscript does not actually support.

BAND S2 · search, chat, export · a pull request into the public repository, use-cases folder

Thesis revision tracker

WHO IT SERVES

Doctoral candidates

WHAT TO BUILD

Tracks committee feedback across drafts, showing which comments are addressed and where.

WHAT STRONG LOOKS LIKE

a comment from draft one is traceable to the paragraph that resolved it in draft four.

BAND S2 · Multi-document, Review, memory · a pull request into the public repository, use-cases folder

Abstract, poster and talk-notes trio

WHO IT SERVES

Researchers

WHAT TO BUILD

One accepted paper becomes an abstract, a poster layout and speaker notes, each correctly sized for its format.

WHAT STRONG LOOKS LIKE

the poster is genuinely readable at poster dimensions rather than a shrunken paper.

BAND S2 · chat, images, export · a pull request into the public repository, use-cases folder

Figure and caption manager

WHO IT SERVES

Researchers

WHAT TO BUILD

Generates, numbers and captions figures, and repairs numbering when figures move.

WHAT STRONG LOOKS LIKE

numbering survives insertion in the middle, and captions follow the venue's style.

BAND S2 · images, chat, export · a pull request into the public repository, use-cases folder

Supplementary materials assembler

WHO IT SERVES

Researchers

WHAT TO BUILD

Builds the supplementary document with its own numbering scheme, cross-linked to the main manuscript.

WHAT STRONG LOOKS LIKE

main-text references to supplementary items resolve correctly in both directions.

BAND S1 · Multi-document, export · a pull request into the public repository, use-cases folder

Authorship and disclosure statements

WHO IT SERVES

Researchers

WHAT TO BUILD

Generates contribution statements, conflict-of-interest declarations and funding acknowledgements in the venue's required taxonomy.

WHAT STRONG LOOKS LIKE

it maps free-text contributions onto the venue's controlled vocabulary correctly.

BAND S1 · templates, chat · a pull request into the public repository, use-cases folder

Plain-language summary generator

WHO IT SERVES

Researchers and funders

WHAT TO BUILD

Produces the lay summary funders increasingly mandate, at a target reading level, without losing accuracy.

WHAT STRONG LOOKS LIKE

reading level is measured, and a domain expert would still call the summary correct.

BAND S2 · chat, search, export · a pull request into the public repository, use-cases folder

Ethics application assembler

WHO IT SERVES

Researchers

WHAT TO BUILD

Assembles an ethics or institutional review submission with participant materials and consent forms consistent with the protocol.

WHAT STRONG LOOKS LIKE

the consent form and the protocol never contradict each other, checked automatically.

BAND S2 · Multi-document, templates, export · a pull request into the public repository, use-cases folder

Continuity and canon checker

WHO IT SERVES

Authors

WHAT TO BUILD

Finds name, timeline, point-of-view and tense inconsistencies across a long manuscript.

WHAT STRONG LOOKS LIKE

it catches a planted continuity error at chapter twenty that references chapter two.

BAND S2 · search, Multi-document · a pull request into the public repository, use-cases folder

Beta-reader feedback consolidator

WHO IT SERVES

Authors

WHAT TO BUILD

Merges many readers' comments into a prioritised revision plan, identifying where readers agree.

WHAT STRONG LOOKS LIKE

it distinguishes consensus from a single loud opinion.

BAND S2 · Multi-document, search, Review · a pull request into the public repository, use-cases folder

Full apparatus builder

WHO IT SERVES

Publishers

WHAT TO BUILD

Title page, copyright page, dedication, table of contents, acknowledgements, about-the-author and a back-matter index structure.

WHAT STRONG LOOKS LIKE

the apparatus is complete and correctly ordered for the format, and the contents list matches the actual chapters.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Book proposal with market section

WHO IT SERVES

Non-fiction authors

WHAT TO BUILD

Assembles a proposal including comparable titles, audience analysis and chapter outline.

WHAT STRONG LOOKS LIKE

the comparable-titles section is specific and real rather than generic.

BAND S2 · search, templates, export · a pull request into the public repository, use-cases folder

Serial archive to book compiler

WHO IT SERVES

Creators

WHAT TO BUILD

Compiles a run of newsletter issues or posts into a coherent book, deduplicating and bridging between pieces.

WHAT STRONG LOOKS LIKE

the bridges make it read as a book rather than a stapled archive.

BAND S2 · Multi-document, memory, export · a pull request into the public repository, use-cases folder

Manuscript translation with structure and voice

WHO IT SERVES

Publishers

WHAT TO BUILD

Translates a full manuscript preserving structure, formatting and as much voice as the target language allows, with a translator review pass.

WHAT STRONG LOOKS LIKE

a professional translator's review finds the output usable as a first pass rather than needing full retranslation.

BAND S3 · chat, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Audiobook narration script

WHO IT SERVES

Publishers

WHAT TO BUILD

Derives a narration script with pronunciation notes, speaker attribution and pacing marks.

WHAT STRONG LOOKS LIKE

it correctly identifies dialogue attribution and flags ambiguous names for pronunciation.

BAND S1 · chat, export · a pull request into the public repository, use-cases folder

Accessible and large-print edition

WHO IT SERVES

Publishers and libraries

WHAT TO BUILD

Produces a large-print or screen-reader-optimised edition from the same manuscript.

WHAT STRONG LOOKS LIKE

it meets a recognised accessibility specification and the structure is verified, not assumed.

BAND S1 · templates, export · a pull request into the public repository, use-cases folder

Multi-edition builder

WHO IT SERVES

Publishers

WHAT TO BUILD

One manuscript producing several editions from a single source of truth, each with its own front matter and formatting.

WHAT STRONG LOOKS LIKE

a correction in the source propagates to every edition without manual repetition.

BAND S2 · Multi-document, templates, export · a pull request into the public repository, use-cases folder

Interview-to-memoir pipeline

WHO IT SERVES

Ghostwriters

WHAT TO BUILD

Turns interview transcripts into memoir chapters in the subject's voice, with a fact-check pass against supplied documents.

WHAT STRONG LOOKS LIKE

the voice is the subject's, and unsupported claims are flagged rather than smoothed over.

BAND S3 · Multi-document, memory, search, export · a pull request into the public repository, use-cases folder

Evidence-backed resume rewriter

WHO IT SERVES

Job seekers

WHAT TO BUILD

Rewrites a resume against a specific job description, changing only what the evidence in the resume supports, with each change traced to its source.

WHAT STRONG LOOKS LIKE

it never invents experience. Every rewritten bullet points at the original fact it came from, and the layout survives untouched.

BAND S2 · chat, Review, export · a pull request into the public repository, use-cases folder

Keyword and requirement gap analyser

WHO IT SERVES

Job seekers

WHAT TO BUILD

Maps the posting's stated requirements against the resume, shows the genuine gaps, and edits only where there is real evidence to surface.

WHAT STRONG LOOKS LIKE

it distinguishes "you have this but did not say it" from "you do not have this", and refuses to fake the second.

BAND S2 · search, chat, Review · a pull request into the public repository, use-cases folder

Impact-quantifying bullet rewriter

WHO IT SERVES

Job seekers

WHAT TO BUILD

Turns duty statements into outcome statements, asking for the number when the resume does not contain one.

WHAT STRONG LOOKS LIKE

it asks rather than fabricating a metric.

BAND S1 · chat, Review · a pull request into the public repository, use-cases folder

Role-family resume matrix

WHO IT SERVES

Job seekers

WHAT TO BUILD

Maintains several tailored versions from one master resume, keeping them in sync when the master changes.

WHAT STRONG LOOKS LIKE

a change to the master propagates correctly to every variant without losing per-variant tailoring.

BAND S2 · Multi-document, templates, export · a pull request into the public repository, use-cases folder

Cover letter from the same evidence base

WHO IT SERVES

Job seekers

WHAT TO BUILD

Writes a letter that draws on the same verified facts as the resume, so the two never contradict.

WHAT STRONG LOOKS LIKE

no claim appears in the letter that is absent from the resume.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Interview preparation brief

WHO IT SERVES

Job seekers

WHAT TO BUILD

Combines the posting, the resume and public company material into a targeted prep document with likely questions and evidence-backed answers.

WHAT STRONG LOOKS LIKE

the questions are specific to that role and the suggested answers use the candidate's real experience.

BAND S2 · search, Multi-document, export · a pull request into the public repository, use-cases folder

Recruiter submittal with a gap memo

WHO IT SERVES

Staffing agencies

WHAT TO BUILD

Produces the client-facing submittal on the agency template plus an internal memo naming where the candidate does not match.

WHAT STRONG LOOKS LIKE

the internal memo is honest, which is what makes an agency trust the tool.

BAND S2 · templates, search, export · a pull request into the public repository, use-cases folder

Reference and recommendation letter drafter

WHO IT SERVES

Managers and academics

WHAT TO BUILD

Drafts letters from a structured intake, in the writer's own voice, with strength calibrated honestly.

WHAT STRONG LOOKS LIKE

a lukewarm reference reads as lukewarm rather than being inflated.

BAND S2 · templates, memory, export · a pull request into the public repository, use-cases folder

Promotion packet builder

WHO IT SERVES

Employees and managers

WHAT TO BUILD

Assembles a self-review or promotion case with evidence pulled from supplied artefacts.

WHAT STRONG LOOKS LIKE

every claim links to an artefact, which is exactly what promotion committees ask for.

BAND S2 · Multi-document, search, templates, export · a pull request into the public repository, use-cases folder

Academic CV and teaching portfolio

WHO IT SERVES

Academics

WHAT TO BUILD

Maintains a long-form CV plus a teaching portfolio, formatted per institution, updated from new outputs.

WHAT STRONG LOOKS LIKE

it handles the length and the institution-specific ordering rules correctly.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Career-change narrative reframer

WHO IT SERVES

Career changers

WHAT TO BUILD

Reframes existing experience for a different field without misrepresenting it.

WHAT STRONG LOOKS LIKE

the reframing is honest and a hiring manager in the target field would find it credible.

BAND S2 · chat, Review · a pull request into the public repository, use-cases folder

Negotiation brief

WHO IT SERVES

Job seekers

WHAT TO BUILD

Builds an evidence-based negotiation document from the offer, market data supplied by the user, and the candidate's leverage.

WHAT STRONG LOOKS LIKE

it is a working document the candidate can actually use in a conversation, not generic advice.

BAND S2 · search, templates, export · a pull request into the public repository, use-cases folder

Job-board side bulk normaliser

WHAT TO BUILD

Job boards and applicant tracking vendors. A server-side service that normalises inbound resumes to a house schema and format at volume, with per-file error reporting.

WHAT STRONG LOOKS LIKE

it handles the messy real distribution of resume formats and reports failures honestly instead of producing silent garbage.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Recurring invoice agent

WHO IT SERVES

SMB

WHAT TO BUILD

Issues invoices on a schedule, adjusting for changes in scope and flagging anything unusual before sending.

WHAT STRONG LOOKS LIKE

it detects and surfaces an anomaly rather than blindly repeating last month.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Freelancer invoice command-line tool

WHO IT SERVES

Developers and freelancers

WHAT TO BUILD

invoice new --client acme --hours 12 producing a branded document, with a local client registry.

WHAT STRONG LOOKS LIKE

it works offline for everything except the generation call, and the output is genuinely presentable.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

Payment plan agreement

WHO IT SERVES

SMB and healthcare

WHAT TO BUILD

Generates a payment plan agreement from terms, in plain language, ready for signature.

WHAT STRONG LOOKS LIKE

the schedule table and the prose terms agree exactly.

BAND S2 · templates, chat, export · a pull request into the public repository, use-cases folder

Financial close pack

WHO IT SERVES

Finance

WHAT TO BUILD

Assembles the close checklist, supporting memos and sign-off sheet for a period close.

WHAT STRONG LOOKS LIKE

sign-offs are attributed and the pack is complete before it can be closed.

BAND S2 · Multi-document, templates, Review, export · a pull request into the public repository, use-cases folder

Control narrative writer from evidence

WHO IT SERVES

Compliance

WHAT TO BUILD

Turns raw evidence into the control narrative an auditor will accept, one control at a time.

WHAT STRONG LOOKS LIKE

each narrative cites the specific evidence that supports it and never overstates coverage.

BAND S2 · search, Multi-document, templates · a pull request into the public repository, use-cases folder

Policy change-log agent

WHO IT SERVES

Compliance

WHAT TO BUILD

Maintains a human-readable change log across policy versions, describing what changed in substance.

WHAT STRONG LOOKS LIKE

it describes substantive change and ignores formatting churn.

BAND S2 · Multi-document, search, export · a pull request into the public repository, use-cases folder

Auditor request responder

WHO IT SERVES

Compliance

WHAT TO BUILD

Takes an auditor's request list, finds the responsive material, and assembles the response pack with an index.

WHAT STRONG LOOKS LIKE

it marks items it could not satisfy instead of quietly omitting them.

BAND S2 · search, Multi-document, export · a pull request into the public repository, use-cases folder

AI governance policy pack

WHO IT SERVES

Compliance

WHAT TO BUILD

Generates and maintains an AI management policy set, including model inventory and use-case assessments.

WHAT STRONG LOOKS LIKE

it is grounded in a real published framework rather than invented structure.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Healthcare privacy policy and procedure pack

WHO IT SERVES

Healthcare compliance

WHAT TO BUILD

A complete policy and procedure set, maintained as regulations and practices change.

WHAT STRONG LOOKS LIKE

procedures reference the specific policy sections they implement, and the links stay valid.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Payment security self-assessment narrative

WHO IT SERVES

Fintech and commerce

WHAT TO BUILD

Drafts the narrative portions of a payment-security self-assessment from the organisation's evidence.

WHAT STRONG LOOKS LIKE

it distinguishes what is implemented from what is planned.

BAND S2 · search, templates, export · a pull request into the public repository, use-cases folder

Access review attestation pack

WHO IT SERVES

IT governance

WHAT TO BUILD

Builds the periodic access review pack with reviewer sign-off and exception documentation.

WHAT STRONG LOOKS LIKE

exceptions carry a documented justification and an owner.

BAND S2 · templates, Review, export · a pull request into the public repository, use-cases folder

Vendor due diligence requester

WHO IT SERVES

Procurement

WHAT TO BUILD

The requesting side: assembles the diligence questionnaire, tracks responses, and writes the assessment memo.

WHAT STRONG LOOKS LIKE

the assessment memo takes a position rather than restating the answers.

BAND S2 · Multi-document, search, templates, export · a pull request into the public repository, use-cases folder

Regulatory change impact memo

WHO IT SERVES

Compliance

WHAT TO BUILD

Given a specific regulatory change and a policy set, produces the impact memo naming affected documents and required actions.

WHAT STRONG LOOKS LIKE

the affected-document list is precise, with paragraph-level pointers.

BAND S2 · search, Multi-document, export · a pull request into the public repository, use-cases folder

Charter party and maritime contract editor

WHO IT SERVES

Shipping

WHAT TO BUILD

Redline charter parties and bills with the clause conventions the trade uses.

WHAT STRONG LOOKS LIKE

it understands standard-form clause naming.

BAND S2 · chat, Review, templates, export · a pull request into the public repository, use-cases folder

Crew and port documentation packet

WHO IT SERVES

Shipping

WHAT TO BUILD

Assembles crew certificates, port clearance and manifest documents per port requirements.

WHAT STRONG LOOKS LIKE

per-port requirement differences are handled.

BAND S2 · Multi-document, templates, export · a pull request into the public repository, use-cases folder

Agricultural supply and certification pack

WHO IT SERVES

Agriculture

WHAT TO BUILD

Supply contracts plus organic or fair-trade certification documentation.

WHAT STRONG LOOKS LIKE

certification evidence links to the specific standard clause.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Mining permit and impact assembler

WHO IT SERVES

Mining

WHAT TO BUILD

Assembles permit applications with environmental and community impact sections.

WHAT STRONG LOOKS LIKE

the technical annexures survive with their figures intact.

BAND S2 · Multi-document, images, export · a pull request into the public repository, use-cases folder

Player contract and scouting report system

WHO IT SERVES

Sports

WHAT TO BUILD

Contracts with performance clauses plus structured scouting reports with video references.

WHAT STRONG LOOKS LIKE

performance clause structures are handled correctly.

BAND S2 · templates, images, export · a pull request into the public repository, use-cases folder

Split sheets and sync licence generator

WHO IT SERVES

Music

WHAT TO BUILD

Songwriter split sheets, sync licences and rights documentation.

WHAT STRONG LOOKS LIKE

percentage splits always total correctly and every rights holder is captured.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Dental practice document suite

WHO IT SERVES

Dentistry

WHAT TO BUILD

Treatment plans, consent forms and insurance pre-authorisations.

WHAT STRONG LOOKS LIKE

the treatment plan is patient-readable and clinically precise at once.

BAND S2 · templates, Review, export · a pull request into the public repository, use-cases folder

Collective agreement redliner

WHO IT SERVES

Labour relations

WHAT TO BUILD

Redlines collective bargaining agreements against each side's position, tracking every open item.

WHAT STRONG LOOKS LIKE

the open-items tracker stays accurate across many rounds.

BAND S3 · Review, Multi-document, export · a pull request into the public repository, use-cases folder

Public safety incident report writer

WHO IT SERVES

Police and emergency services

WHAT TO BUILD

Structured incident reports from field notes with the elements a prosecution requires.

WHAT STRONG LOOKS LIKE

it prompts for missing required elements rather than producing an incomplete report.

BAND S2 · chat, templates, Review, export · a pull request into the public repository, use-cases folder

Transcript to structured record

WHO IT SERVES

Courts and tribunals

WHAT TO BUILD

Turns a raw transcript into a structured record with speaker attribution, exhibits and an index.

WHAT STRONG LOOKS LIKE

attribution is accurate and exhibits are cross-referenced.

BAND S2 · search, Multi-document, export · a pull request into the public repository, use-cases folder

Notarial and legalisation packet

WHO IT SERVES

Notaries and legal services

WHAT TO BUILD

Assembles notarial certificates and legalisation packets per destination country requirements.

WHAT STRONG LOOKS LIKE

destination-specific requirements are researched and correct.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Case-note and care-plan writer

WHO IT SERVES

Social work and elder care

WHAT TO BUILD

Structured case notes and care plans with strict privacy handling.

WHAT STRONG LOOKS LIKE

it separates observation from interpretation, which the profession requires.

BAND S2 · templates, Review, export · a pull request into the public repository, use-cases folder

Archive finding-aid generator

WHO IT SERVES

Libraries and archives

WHAT TO BUILD

Produces structured finding aids and collection descriptions from accession records.

WHAT STRONG LOOKS LIKE

it follows a recognised descriptive standard.

BAND S2 · templates, search, export · a pull request into the public repository, use-cases folder

Standards drafting workspace

WHO IT SERVES

Standards bodies

WHAT TO BUILD

Multi-round drafting with formal comment submission, disposition and resolution tracking.

WHAT STRONG LOOKS LIKE

every comment has a recorded disposition and the audit trail is complete.

BAND S3 · Review, Multi-document, export · a pull request into the public repository, use-cases folder

Regulated-product compliance documentation

WHO IT SERVES

Highly regulated consumer products

WHAT TO BUILD

Licence applications, testing records and labelling documentation per jurisdiction.

WHAT STRONG LOOKS LIKE

jurisdiction-specific labelling requirements are handled correctly.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Protocol and audit documentation pack

WHO IT SERVES

Web3 and security

WHAT TO BUILD

Assembles audit reports, disclosures and technical documentation with code and diagram embedding.

WHAT STRONG LOOKS LIKE

code blocks and diagrams survive export intact.

BAND S2 · images, Multi-document, export · a pull request into the public repository, use-cases folder

Apparel technical pack generator

WHO IT SERVES

Fashion and manufacturing

WHAT TO BUILD

Produces the technical pack a factory needs: measurements, materials, construction notes and graded specifications with images.

WHAT STRONG LOOKS LIKE

the measurement tables grade correctly across sizes.

BAND S2 · images, templates, export · a pull request into the public repository, use-cases folder

Environmental impact assessment assembler

WHO IT SERVES

Environmental consulting

WHAT TO BUILD

Assembles a long multi-discipline assessment with figures, appendices and a non-technical summary.

WHAT STRONG LOOKS LIKE

the non-technical summary is genuinely accessible and consistent with the technical body.

BAND S3 · Multi-document, images, search, export · a pull request into the public repository, use-cases folder

Microfinance and cooperative loan documentation

WHO IT SERVES

Financial inclusion

WHAT TO BUILD

Loan documents, group guarantees and repayment schedules in local languages at low literacy levels.

WHAT STRONG LOOKS LIKE

the borrower-facing documents are understandable to the actual borrower, verified with a reading-level check in the target language.

BAND S2 · templates, chat, export · a pull request into the public repository, use-cases folder

Disclosure release pipeline with a reviewable redaction log

WHO IT SERVES

Public-records officers and litigation support teams releasing to outside parties

WHAT TO BUILD

A batch pipeline for releasing a document set to someone outside the organisation. Candidate redactions are proposed per exemption category with the reason attached, a human decides each one, redactions are applied so the removed content is genuinely gone from the released file rather than covered over, and the release package ships with a log that states, per document, what was withheld and under which exemption.

WHAT STRONG LOOKS LIKE

the released file contains no recoverable trace of redacted content, proven by extracting text from the exported file; the log is the deliverable a requester or a court can read; a document released twice under different exemption sets produces two consistent packages; anything the reviewer did not explicitly decide blocks the release rather than defaulting to release.

BAND S2 · Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Defined-term governance across a document set

WHO IT SERVES

Legal departments, standards authors and policy owners with drifting definitions

WHAT TO BUILD

One authoritative glossary as a maintained document, and every other document in the set reconciled against it: terms used but never defined, terms defined twice, terms defined differently in two documents, and definitions that have drifted from the authority. Changing a definition once proposes the matching reviewable edit everywhere the term is defined or materially relied on.

WHAT STRONG LOOKS LIKE

a genuine conflict (two documents defining the same term incompatibly) is surfaced for a decision and never merged automatically; propagation edits the definition sites and leaves ordinary uses of the word alone, which is the precision problem; the glossary cites the document each entry came from; a term retired in the authority is reported everywhere it survives.

BAND S2 · Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Bilingual authoritative pair maintainer

WHO IT SERVES

Organisations where both language versions are legally authoritative

WHAT TO BUILD

A document that exists as a pair, where neither side is subordinate. An edit to either language proposes the matching edit on the other, aligned at the clause level; the pair carries an alignment map so a reviewer can see any clause beside its counterpart; and the tool refuses to let the pair be published while any clause is unmatched or stale.

WHAT STRONG LOOKS LIKE

editing either side proposes the counterpart edit, not only the "original" side; the alignment survives a clause being inserted or moved on one side; a numbering change on one side is mirrored exactly; the publish gate genuinely blocks on a stale counterpart rather than warning about it.

BAND S2 · Multi-document, chat, Review, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Two-estate policy reconciliation for a merger

WHO IT SERVES

Post-acquisition integration teams holding two complete policy sets

WHAT TO BUILD

Two organisations' policy sets in, one target set out. The tool maps each policy in one estate to its counterpart in the other, and produces a per-policy recommendation with the evidence: keep this one, keep that one, merge them (with the merged draft written), or write a new one because neither is adequate. It also produces the two lists nobody wants to make by hand - obligations that exist in only one estate, and conflicts where the two estates require opposite things.

WHAT STRONG LOOKS LIKE

the mapping is proposed from content rather than titles, since the two estates name nothing the same way; a genuine conflict is presented as a decision with both positions quoted, never silently resolved; the coverage list proves no obligation from either estate was dropped; the merged drafts read like policy, not like two documents stapled together.

BAND S2 · Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Pre-submission validator against a regulator's published filing rules

WHO IT SERVES

Regulatory affairs, tender teams, court filers, grant offices

WHAT TO BUILD

The gate before submission. The regulator's or authority's published submission rules - mandatory sections, ordering, naming conventions, declarations, file format and size limits, signature and date requirements - are held as a checkable rule set, and a completed filing package is validated against them, producing a pass or a per-rule failure list with the location of each problem. The rule set is data, so a new authority is added without touching code.

WHAT STRONG LOOKS LIKE

each failure quotes the rule it breaks and points at where in the package it breaks, so it is fixable without reading the rulebook; the validator distinguishes a hard rejection rule from a discretionary one; adding a second authority is demonstrated and is configuration only; a package that passes is genuinely accepted by the authority's own published checker where one exists.

BAND S2 · Multi-document, templates, search, export · a pull request into the public repository, use-cases folder

Source-traceable brief where every claim resolves to the page it came from

WHO IT SERVES

Research analysts, policy teams, diligence desks

WHAT TO BUILD

A brief-builder where every factual claim in the output resolves to the exact source page it came from, and a claim that cannot be resolved is refused rather than written. It ingests a set of source documents, produces a structured brief, and attaches to each claim a verifiable pointer back into a specific page of a specific source. The judged property is refusal: when the sources do not support a claim, the tool declines to assert it and says which claim it dropped and why.

WHAT STRONG LOOKS LIKE

Every claim in a generated brief resolves to a real page in a real source, checked mechanically rather than by eye; a deliberately unsupported question produces an explicit refusal naming the gap, not a plausible sentence; the citation survives an edit to the brief; and the tool reports the count of claims it refused, so a low number is visible rather than hidden.

BAND S2 · search, Multi-document, export · a pull request into the public repository, use-cases folder

Tables out of a fixed-layout document into structured, editable output

WHO IT SERVES

Engineer. Anyone who has fought a real-world table.

WHAT TO BUILD

Take fixed-layout documents that contain tables, extract those tables faithfully, and produce a structured, editable document where the table is real structure rather than a picture of one. Handle the awkward cases on purpose: merged cells, a table split across pages, a header row that repeats, and a column of numbers that is right-aligned text.

WHAT STRONG LOOKS LIKE

merged cells and page-split tables survive, and the tool reports its own confidence per table instead of silently guessing.

BAND S2 · API, attachments, export, Review · a pull request into the public repository, use-cases folder

Guided first edit for a brand new document

WHO IT SERVES

Product-minded engineer. Cares about the first ninety seconds.

WHAT TO BUILD

Build the first-run experience for someone who has just arrived with a document and does not know what to ask for. It should read what they actually uploaded, propose two or three specific edits that make sense for that document rather than generic suggestions, and let them accept one and see it land. The proposals must be derived from the document's real content, not a fixed menu.

WHAT STRONG LOOKS LIKE

upload two very different documents and get two genuinely different sets of proposals.

BAND S2 · API, chat, Review · a pull request into the public repository, use-cases folder

Pre-signature redline desk

WHO IT SERVES

Engineer building for a contracts or deal-desk team.

WHAT TO BUILD

Build the workspace for the stage before a contract goes out for signature: pull in the counterparty's draft, mark deviations from an agreed position, propose replacement language per deviation, and produce a clean signature-ready version plus a record of what changed and why. Signature itself is always a hand-off to an external provider and is out of scope.

WHAT STRONG LOOKS LIKE

every proposed change is traceable to a stated position, and the clean version and the change record are generated from the same source rather than maintained separately.

BAND S2 · API, chat, Review, export · a pull request into the public repository, use-cases folder

Veterinary clinical note and client communication app

WHO IT SERVES

Engineer. A vertical build with a real documentation burden behind it.

WHAT TO BUILD

A small application for a veterinary practice: turn a consultation record into a structured clinical note, and from the same source produce the client-facing summary, the medication instruction sheet and the referral letter, each in the right register for its reader. The clinical note and the client letter must never contradict each other because both derive from one record.

WHAT STRONG LOOKS LIKE

one source produces four genuinely different registers, and a change to the record updates all of them consistently.

BAND S1 · API, chat, templates, export · a pull request into the public repository, use-cases folder

Past-work browser: every session and the documents it touched

WHO IT SERVES

Product-minded engineer. Cares about the moment someone cannot find what they made last week.

WHAT TO BUILD

Build the surface that answers "where is the thing I was working on". List every past working session, show the documents each one touched, let a user search across them by content rather than by filename, and open any of them back up in the state it was left in. Sessions that produced nothing should be visibly distinguishable from sessions that produced something, so the list stays useful once it is long.

WHAT STRONG LOOKS LIKE

it stays usable at several hundred sessions, search finds a document by something written inside it, and reopening genuinely restores the working state rather than just the file.

BAND S2 · API, Multi-document, search, export · a pull request into the public repository, use-cases folder

Textbook authoring agent with pedagogical structure

WHO IT SERVES

Educators and publishers producing structured textbooks with exercises and consistent pedagogy.

WHAT TO BUILD

A textbook is not a long book: every chapter carries learning objectives, worked examples, exercises with answer keys, and difficulty that must rise in order. Build an agent that authors a short textbook or course-book from a syllabus: it maintains a pedagogy ledger (objectives covered, concepts introduced, difficulty curve), generates exercises whose answer keys it verifies against its own worked examples, and refuses to use a concept in chapter N that no earlier chapter introduced.

WHAT STRONG LOOKS LIKE

a reviewer picks any exercise and finds its concepts introduced earlier plus a correct answer key; a planted out-of-order concept triggers the agent's own objection; regenerating one chapter leaves the others byte-stable.

BAND S3 · chat, Multi-document, export · a pull request into the public repository, use-cases folder

INTEGRATIONS AND PIPELINES

MCP pipeline: CRM → proposal → e-sign

WHO IT SERVES

Sales/Agents

WHAT TO BUILD

Sales teams already have a CRM record for every deal and a signature tool at the end, but nothing chains the two around SuperDocs, so proposals still get built by hand from a copy-pasted deal summary. Build an agent pipeline that pulls a deal's fields from a CRM, drops them into a SuperDocs proposal template with real branding and terms, and, once a human approves the draft, hands the finished document to an external e-sign provider's MCP server (DocuSign and PandaDoc both publish one) to route for signature. SuperDocs' role stops at generating and formatting the document; sending it out and confirming execution is entirely the e-sign provider's job, never claimed as SuperDocs' own.

WHAT STRONG LOOKS LIKE

running the pipeline against a sample CRM record produces a correctly populated proposal, a human approval step gates the handoff, and the signature leg is a real call into DocuSign's or PandaDoc's MCP server, not a stubbed step.

BAND S3 · MCP, templates, export · a pull request into the public repository, extensions folder

Coda/Notion database → doc-set generator

WHO IT SERVES

Ops

WHAT TO BUILD

teams keep repeatable document sets, say a welcome letter, a contract, and a scope-of-work, as one row per client in a Coda table or a Notion database, then assemble every document by hand for every new row. Coda's API returns a row's values keyed by immutable column IDs, with a separate columns call to map those back to human-readable names; Notion's database-query API returns typed per-property values the same way. Neither platform does anything with the document generation itself. Build a generator that reads a table or database, maps each row's fields onto a fixed set of SuperDocs templates rather than just one, and produces that whole matched document set per row through SuperDocs' multi-doc and export surfaces.

WHAT STRONG LOOKS LIKE

Point it at a table of five rows; it produces the same two-or-more distinct filled templates per row, correctly mapped from that row's own fields, with one doc tab per generated file in the multi-doc session rather than a single merged file.

BAND S2 · API, templates, Multi-document, export · a pull request into the public repository, extensions folder

Slack Canvas → exportable doc

WHO IT SERVES

Cross-industry

WHAT TO BUILD

A Slack Canvas is a persistent, section-structured document object attached to a channel, not a message, a Workflow Builder step, or an App Home tab, so none of those integration shapes can read it. Slack's `canvases.sections.lookup` API, under the `canvases:read` scope, returns a canvas back as ordered Markdown sections. Build an MCP action that looks up a channel's canvas, walks its sections in order, and exports the result as a branded DOCX or PDF through SuperDocs, so a canvas full of project notes can leave Slack as a real document.

WHAT STRONG LOOKS LIKE

Given a real multi-section canvas with headings, a checklist, and a table, the exported document keeps the sections in their original order, and a reviewer can check section-by-section that nothing from the canvas was silently dropped.

BAND S2 · MCP, export · a pull request into the public repository, extensions folder

Content brief factory for an editorial calendar

WHO IT SERVES

Content / content strategist or agency editorial lead commissioning freelance writers

WHAT TO BUILD

From a calendar of planned topics, generate a brief per piece: angle, audience, search intent, an outline with heading structure, required internal links, sources to cite, tone and word count, and what the piece must not claim. Briefs are generated against the brand's existing content library so they avoid duplicating what already ranks and explicitly name the pieces to link to. Writers receive the brief as an editable document and return the draft into that same document, where the brief becomes the review checklist.

WHAT STRONG LOOKS LIKE

briefs are genuinely differentiated against the existing library rather than interchangeable outlines; the brief-to-checklist transition is actually useful at review time; the outline structure survives into the writer's returned draft.

BAND S1 · Multi-document, search, templates, chat, export · a pull request into the public repository, use-cases folder

Inclusive-language sweep across a live requisition set

WHO IT SERVES

Talent-acquisition lead or DEI partner responsible for every open job posting

WHAT TO BUILD

Job descriptions accumulate exclusionary phrasing, inflated requirement lists and inconsistent benefit language across dozens of live requisitions. Build a tool that holds the full live set, audits each posting against the organization's own guidance document, reports the issue profile across the portfolio (which teams, which patterns), and applies the approved rewrites in place so the posting keeps its structure, its team-specific content and its ATS field boundaries. The audit reruns on a schedule and reports only new drift.

WHAT STRONG LOOKS LIKE

portfolio-level reporting, not per-document nagging; rewrites preserve team-specific content; the rerun reports only what is new since the last pass.

BAND S2 · Multi-document, search, templates, Review, API · a pull request into the public repository, use-cases folder

Negotiation playbook builder with drift audit

WHO IT SERVES

Legal operations or commercial lead turning a signed-contract portfolio into a playbook

WHAT TO BUILD

Most teams have positions in their heads and precedent in their archive but no written playbook. Build a tool that reads the signed portfolio, clusters how each key term was actually settled, and drafts the playbook document: preferred position, acceptable fallback, walk-away, with real examples cited from the portfolio. Then it audits new inbound paper against the playbook it just wrote and reports where the team is drifting from its own standard. The playbook stays a living document that is edited in place as positions change.

WHAT STRONG LOOKS LIKE

playbook positions are evidenced by real cited examples from the portfolio, and the drift audit measures deviation over time rather than judging a single contract.

BAND S3 · search, Multi-document, memory, chat, export · a pull request into the public repository, use-cases folder

Brand-Compliance Audit Across a Client's Document Estate

WHO IT SERVES

Agencies / brand teams · brand guardian auditing a client's real-world documents

WHAT TO BUILD

Point the app at a folder of a client's existing documents. It scores each file against the brand kit and style guide (typography, palette, heading structure, banned phrases, required legal boilerplate, logo presence, tone) and produces one audit document: a score per file, every violation cited to its section, and an estimated remediation effort. A second pass fixes the mechanical violations in place and re-scores, leaving judgement calls for a human.

WHAT STRONG LOOKS LIKE

the score is reproducible across runs; citations point at the right section every time; the fix pass changes presentation without rewriting meaning.

BAND S3 · API, Multi-document, search, templates, Review, export · a pull request into the public repository, extensions folder

Email-to-document service address

WHO IT SERVES

Anyone who cannot install anything: field staff, external counterparties, clients, phone-only users. Host: email itself.

WHAT TO BUILD

The organisation gets an address. Forward a document with the instruction in the email body ("reformat this to our template and remove the pricing annex") and the processed file comes back as a reply within minutes, with a link to review the changes if the instruction was ambiguous. No install, no login, works from any device and any mail client. Sender allow-lists, per-sender quotas, and a rejection reply for anything outside policy.

WHAT STRONG LOOKS LIKE

the reply arrives with a correctly named attachment and a summary a human would write for a human. Ambiguity produces a question rather than a guess. Unknown senders get a polite refusal, never processing.

BAND S2 · API, chat, export, Review · a pull request into the public repository, extensions folder

Gmail add-on: inbound attachment normaliser

WHO IT SERVES

Coordinator receiving documents from many outside parties in whatever format they have. Host app: Gmail (Google Workspace add-on, also surfaces in Drive).

WHAT TO BUILD

Open a message with an attachment and the add-on offers to normalise it: any supported input format in, the organisation's house template out, with headings mapped, tables restyled, logos placed, and a cover page added. The normalised file goes to a Drive folder and back as a reply attachment. It handles the awkward ones (legacy Word, OpenDocument, rich text, scanned PDF) rather than refusing them.

WHAT STRONG LOOKS LIKE

a scanned PDF comes back as genuinely editable text with its tables intact. A twelve-year-old .doc survives with its structure. Output is visually indistinguishable from a document authored on the template.

BAND S2 · API, templates, export, chat · a pull request into the public repository, extensions folder

Google Forms to generated document pipeline

WHO IT SERVES

Anyone running intake: clinics, schools, councils, small firms, event teams. Host: Google Forms plus a small service.

WHAT TO BUILD

A form submission triggers the document it implies: an intake form becomes a completed record on the practice template, a booking form becomes a confirmation and an agreement, an application becomes a formatted submission packet. The document is delivered to the submitter and filed to Drive. Conditional questions drive conditional document sections. Configuration is a mapping screen, not code.

WHAT STRONG LOOKS LIKE

every submission produces exactly one document, delivered, with correct values in every field. Optional questions left blank produce clean documents rather than empty placeholders. The mapping screen is usable by the person who built the form.

BAND S1 · API, templates, export, Review · a pull request into the public repository, extensions folder

Google Chat app: document requests and approvals

WHO IT SERVES

Teams running on Google Workspace who want document work to start in chat. Host app: Google Chat.

WHAT TO BUILD

Mention the app in a space with a request ("draft the renewal letter for Acme using last year's") and it replies in-thread with a link to the drafted document and a card summarising what it produced. Approval cards post to the space so a reviewer approves or rejects proposed changes from chat. A daily digest lists what the space produced and what is still waiting on someone.

WHAT STRONG LOOKS LIKE

cards render correctly on desktop and mobile Chat. Approvals from chat genuinely apply to the document. The pending list is accurate and nothing waits silently forever.

BAND S2 · MCP, chat, Review, export, Multi-document · a pull request into the public repository, extensions folder

SharePoint document-library agent

WHO IT SERVES

Records manager or intranet owner at a mid-size or enterprise organisation. Host app: SharePoint Online.

WHAT TO BUILD

A SharePoint library is a folder with columns, and the columns say what each document is. This agent uses them: when a file lands or a column changes, it applies the library's configured treatment (normalise to the department template, generate the companion summary, produce the required cover sheet) and writes the result back as a new version with the column values updated. A library-level settings page defines the rules. Every action appears in the file's version history with a readable comment.

WHAT STRONG LOOKS LIKE

versioning, check-out state, and retention labels are all respected rather than trampled. Nothing runs twice on the same version. The rules page is configurable by a SharePoint admin with no code.

BAND S3 · API, templates, export, search, Multi-document · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

OneDrive and Office file handler: open with AI editor

WHO IT SERVES

Any Microsoft 365 user with documents in OneDrive or SharePoint. Host: the M365 file context menu.

WHAT TO BUILD

Register as a file handler so any DOCX or PDF in OneDrive, SharePoint, or Teams files gets an "Open with AI editor" action in the right-click menu. Choosing it opens the document in a full editing session in the browser, with chat, proposed-change review, and export, and saves back to the same file as a new version. No download and re-upload, and no separate copy to reconcile later.

WHAT STRONG LOOKS LIKE

the save-back writes a proper new version to the original file, not a duplicate. Permissions are inherited from the file so nobody sees what they should not. A PDF opened this way comes back as an editable document that still exports as a PDF.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Teams message extension: thread to document with approval card

WHO IT SERVES

Any Teams-first organisation. Host app: Microsoft Teams.

WHAT TO BUILD

Select a message or a thread, choose the extension, and pick what it should become: minutes, a decision record, a customer-facing summary, a requirements document. It drafts on the chosen template and posts an adaptive card back into the channel with a preview and Approve and Revise buttons. Approving publishes the document to the channel's files tab. Revising takes an instruction typed into the card.

WHAT STRONG LOOKS LIKE

the adaptive card works on Teams desktop, web, and mobile. Approval genuinely writes the file to the channel. The draft attributes decisions to the people who actually made them in the thread.

BAND S2 · MCP, chat, Review, export, templates · a pull request into the public repository, extensions folder

Teams tab: document review workspace

WHO IT SERVES

Legal, compliance, or editorial team whose reviews currently happen over email. Host app: Microsoft Teams (channel tab).

WHAT TO BUILD

A channel tab that holds the team's review queue. Each item shows the document, the proposed changes as before-and-after cards, who is assigned, and what is blocked. Reviewers approve or reject in place, leave a comment, or reassign. Notifications post to the channel when something waits too long. Completed reviews are archived with the full decision record and can be exported as an evidence pack.

WHAT STRONG LOOKS LIKE

the queue survives a page reload mid-review with partly-decided batches intact. Two reviewers on the same document do not overwrite each other. The archived record is complete enough for an audit.

BAND S3 · Review, API, Multi-document, export, search · a pull request into the public repository, extensions folder

Notion integration: review-cycle round-trip

WHO IT SERVES

Startup or agency team whose source of truth is Notion but whose reviewers live in Word. Host app: Notion.

WHAT TO BUILD

Send a Notion page out for formal review: it becomes a styled document, goes to reviewers who mark it up with tracked changes and comments in Word, and comes back in. The integration reads the returned markup and proposes each change back onto the Notion blocks, one approval at a time, preserving Notion structure (toggles, callouts, databases inline). The page keeps a link to the review round it came from.

WHAT STRONG LOOKS LIKE

block structure survives the round-trip both ways. Reviewer attribution is preserved into the Notion comment. A reviewer who has never heard of Notion can do their part entirely in Word.

BAND S3 · API, chat, Review, export, Multi-document · a pull request into the public repository, extensions folder

Notion integration: controlled policy set

WHO IT SERVES

Operations or compliance lead at a company whose wiki is Notion and whose auditor wants documents.
Host app: Notion.

WHAT TO BUILD

Policies live as Notion pages for everyday reading and as approved, versioned documents for audit. The integration keeps both in step: a page edit proposes a new document revision, approval publishes it with an effective date and a change summary, and the Notion page shows a banner stating which approved version it corresponds to. Auditors get an exported evidence pack of the current set with the approval record.

WHAT STRONG LOOKS LIKE

the banner never lies about which version is current. The evidence pack is complete and internally consistent. Editing a page does not silently change what is published.

BAND S2 · templates, Review, Multi-document, export, search · a pull request into the public repository, use-cases folder

Confluence Cloud app: in-page AI edit and review

WHO IT SERVES

Technical writer, product manager, or IT policy owner on Atlassian. Host app: Confluence Cloud.

WHAT TO BUILD

An app that edits the page you are on instead of exporting it away. Select a section, give an instruction, and the proposal appears inline as an accept-or-reject highlight. A page-level review mode routes the changes to a named approver before they publish. Space-level commands normalise a whole space to the house structure, and a publish action produces the branded PDF for the audiences that will never open Confluence.

WHAT STRONG LOOKS LIKE

Confluence macros, tables, and code blocks survive an edit untouched. Page history shows a clean single revision per approved change. The space normalisation is reviewable page by page, never a bulk overwrite.

BAND S3 · chat, Review, export, search, Multi-document · a pull request into the public repository, extensions folder

GitHub App: document pull-request reviewer

WHO IT SERVES

Docs-as-code team, open-source maintainer, or platform team with docs in the repo. Host: GitHub.

WHAT TO BUILD

Install on a repo and every pull request touching documentation gets a review comment: style-guide violations, broken structure, stale references, and a rendered preview of what the change will look like as a styled PDF or Word deliverable. A command in a PR comment applies a suggested fix as a commit. Release tags trigger the branded release document, attached to the release.

WHAT STRONG LOOKS LIKE

comments are specific and few, so maintainers do not mute the app. The rendered preview matches the eventual export. Applying a fix produces a clean, attributable commit.

BAND S2 · API, chat, export, images, Review · a pull request into the public repository, extensions folder

Slack Workflow Builder custom step

WHO IT SERVES

Operations or people lead who builds Slack workflows without writing code. Host app: Slack.

WHAT TO BUILD

A custom step that any Slack workflow can drop in: "generate a document from this template with these values", "reformat the attached file", "send this for approval". Because it is a step and not a bot, it composes with the form step, the approval step, and the channel-post step that teams already use. An onboarding workflow, a request form, or a weekly report can now end in a real branded document with no developer involved.

WHAT STRONG LOOKS LIKE

the step's inputs and outputs map cleanly to other steps so a non-technical builder can wire it without documentation. Failures surface in the workflow's own error handling. The produced file lands where the workflow says.

BAND S2 · API, templates, export, Review · a pull request into the public repository, extensions folder

Slack app home: the document desk

WHO IT SERVES

Any team that wants document work visible where they already work. Host app: Slack.

WHAT TO BUILD

The app's Home tab is a personal desk: documents you have open, changes waiting on your approval, jobs still running, and recent output, all with buttons. Approve a batch of proposed edits from Slack. Open a document. Re-run a failed job. Channel messages remain for requests, but the desk is where state lives, so nothing is lost in scrollback and nobody has to remember what they asked for on Tuesday.

WHAT STRONG LOOKS LIKE

approving from Slack genuinely applies the change and the desk reflects it immediately. Nothing sits pending invisibly. The desk is accurate after a restart on either side.

BAND S2 · MCP, Review, Multi-document, chat, export · a pull request into the public repository, extensions folder

One approval relay, every chat surface

WHO IT SERVES

Platform team whose users are split across Slack, Teams, Google Chat, and plain email. Host: all of them at once.

WHAT TO BUILD

One service that renders a document approval into whichever surface each approver actually uses, and reconciles the answer back into one decision. A legal reviewer approves from Teams, an ops lead from Slack, an external counterparty from an emailed link, and the document sees one consistent set of decisions. Includes reminders, escalation after a deadline, and a single audit record across all four surfaces.

WHAT STRONG LOOKS LIKE

the same approval cannot be double-applied from two surfaces. A decided item disappears everywhere at once. The emailed link is safe to send outside the company and expires.

BAND S3 · Review, API, export · a pull request into the public repository, extensions folder

Salesforce AppExchange managed package

WHO IT SERVES

Sales operations or revenue systems admin at a Salesforce-run company. Host app: Salesforce.

WHAT TO BUILD

A managed package that puts a document panel on any record page: from an opportunity, generate the branded proposal, quote, or order form; from a case, the customer letter; from a contract record, the redline. The panel shows the drafted document, lets the rep adjust it in plain language, routes it for approval through Salesforce's own approval process, and files the final file back to the record with the stage updated. Configured by an admin, not a developer.

WHAT STRONG LOOKS LIKE

it passes AppExchange security review. Field mapping is admin-configurable across custom objects. The document that lands on the record is the one that was approved, every time.

BAND S3 · API, templates, chat, Review, export · a pull request into the public repository, extensions folder

HubSpot app marketplace CRM card

WHO IT SERVES

Marketing or sales operations lead at an SMB or mid-market company on HubSpot. Host app: HubSpot.

WHAT TO BUILD

A CRM card on the deal, contact, and company records that generates and edits the documents that stage needs: proposal, one-pager, onboarding pack, renewal letter. It reads the record's properties, drafts on the company template, and writes the finished file back to the record with an activity entry. A workflow action lets HubSpot automation trigger the same thing without a human clicking.

WHAT STRONG LOOKS LIKE

installs through HubSpot's OAuth flow with no manual key handling. Property mapping covers custom properties. The workflow action is reliable enough to run unattended on a real pipeline.

BAND S2 · API, templates, chat, export, Review · a pull request into the public repository, extensions folder

Airtable marketplace extension

WHO IT SERVES

Operations builder who runs a business process on an Airtable base. Host app: Airtable.

WHAT TO BUILD

An extension in the base plus a scripting action for automations. Pick a table, a view, and a template; produce one document per record with the file link written back to a field; or produce one combined document across a view (a catalogue, a report, a packet with a table of contents). Field mapping is visual. A preview runs on three records before the full run.

WHAT STRONG LOOKS LIKE

attachment fields, linked records, and long text with formatting all map correctly. The combined-document mode produces a coherent single document rather than a concatenation. Runs are resumable after an interruption.

BAND S1 · API, templates, Multi-document, export · a pull request into the public repository, extensions folder

Meeting-recorder ecosystem connector

WHO IT SERVES

Team already using an AI notetaker and still typing up the actual document afterwards. Host: the notetaker's webhooks, plus a small service.

WHAT TO BUILD

The notetaker already produces the transcript and the summary. This connector turns them into the deliverable the meeting was for: minutes on the board template, the client follow-up letter, the requirements document, the interview scorecard, the care note. Rules pick the output by meeting type or title pattern. The draft goes to the owner for approval before delivery, and attaches back to the meeting record.

WHAT STRONG LOOKS LIKE

the output is a real document on a real template, not a prettier transcript. Attributed decisions and owners are correct. It supports several notetakers behind one configuration rather than one hard-wired vendor.

BAND S2 · API, chat, templates, Review, export · a pull request into the public repository, extensions folder

Code review to design-review document

WHO IT SERVES

DevTools. Engineering leads who need a written design record from work that only exists as a pull request.

WHAT TO BUILD

Chains a repository server with our document server. Given a merged change set, the pipeline gathers the diff, the discussion, and the linked issue, then writes a structured design-review document: context, decision, alternatives considered, risks, and an auto-rendered architecture diagram generated from the described components. Subsequent related changes edit the same document rather than creating a new one. The final file is exported and attached back to the repository as a release artefact.

WHAT STRONG LOOKS LIKE

the pipeline is declarative and reusable rather than a script hardcoded to one repository; the diagram is generated, not pasted; a second related change updates the existing document in place; authentication for both servers handled cleanly; a failure in one server does not corrupt the document.

BAND S2 · MCP, chat, images, Multi-document, export · a pull request into the public repository, use-cases folder

Cloud-storage knowledge sync with write-back

WHO IT SERVES

Cross-industry. Teams whose documents live in a shared drive and whose AI cannot touch them.

WHAT TO BUILD

A pipeline that reads a document from a cloud storage server, edits it through our document server, and writes the edited file back to the same location with fidelity preserved, including a versioned copy so nothing is lost. A watch mode processes a folder on a schedule with a rule per subfolder. The interesting engineering is the write-back contract: what happens when someone edited the file in the meantime, which is handled by detection and an explicit choice rather than a silent overwrite.

WHAT STRONG LOOKS LIKE

a concurrent human edit is detected and never silently clobbered; formatting survives the full round trip (verified structurally, not by eye); permissions respected; the versioned copy is real; a dry-run mode that reports what it would change.

BAND S3 · MCP, chat, Review, export, Multi-document · a pull request into the public repository, use-cases folder

Database query to styled report, on a schedule

WHO IT SERVES

Analytics and operations. Data teams asked for the same report every month.

WHAT TO BUILD

Chains a database server with our document server: a saved query set runs, the results are shaped, and the narrative report is edited in place month over month, with tables refreshed and commentary revised only where the numbers moved materially. Charts render as diagrams inside the document. The report keeps its branded template. Delivery is an export plus a link. Distinct from an existing warehouse-to-board-report row because this one maintains a single living report rather than producing a new one each period.

WHAT STRONG LOOKS LIKE

the query layer is parameterised and injection-safe; a month with no material change produces a near-empty change set and says so; numbers in prose and numbers in tables never disagree, enforced by a check; the template survives twelve cycles.

BAND S2 · MCP, chat, images, templates, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Search server plus document server for a cited brief

WHO IT SERVES

Strategy and research. Analysts producing briefs under deadline.

WHAT TO BUILD

A pipeline that composes a web-search server with our document server to produce a brief where every claim carries a real footnote to its source, and a verification pass edits out anything the sources do not support. Sources are stored as reference attachments so the document can be re-checked later without re-searching. The brief is exported with genuine footnotes and a reference list.

WHAT STRONG LOOKS LIKE

footnotes are real footnotes in the exported file, not superscript text; the verification pass demonstrably removes or softens an unsupported claim; sources are captured so the brief is reproducible a month later; rate limits on the search side handled gracefully.

BAND S2 · MCP, chat, search, Review, export · a pull request into the public repository, use-cases folder

Local filesystem bridge for offline document work

WHO IT SERVES

DevTools and privacy-conscious teams. Engineers who work from a folder.

WHAT TO BUILD

Composes a filesystem server with our document server so an agent can operate on a local project folder: read the sources, produce or edit the deliverable, and write results back to disk with a manifest of what changed. Includes a guard layer that constrains the agent to an allowed path set and requires confirmation before any destructive write, because that is the part most integrations get wrong.

WHAT STRONG LOOKS LIKE

path traversal is impossible, proven by tests; nothing outside the allowed set is ever written; the manifest is accurate; large files never pass through the agent's context; a documented recovery step if a run is interrupted mid-write.

BAND S2 · MCP, chat, export, Multi-document · a pull request into the public repository, use-cases folder

Document handoff contract between agents and servers

WHO IT SERVES

Agent infrastructure. Teams building multi-agent systems where more than one agent touches one document.

WHAT TO BUILD

A small specification plus reference implementation for passing a document between agents: a handoff envelope carrying the durable document identifier, the intended scope of the next agent's work, what must not be touched, and the review state. Two demonstration agents exchange a document twice, and a validator rejects a handoff that violates the declared scope. It answers the question every multi-agent builder hits: how do two agents share one document without trampling each other.

WHAT STRONG LOOKS LIKE

the specification is written well enough that a third party implements against it; the validator catches real violations; the demonstration includes a rejected handoff; the envelope carries no secrets; versioned with a compatibility statement.

BAND S3 · MCP, chat, Review, Multi-document · a pull request into the public repository, extensions folder

Router agent that chooses between document tools honestly

WHO IT SERVES

Agent infrastructure. Teams that want the right tool per task, not a single hardcoded one.

WHAT TO BUILD

An agent given several document-capable tools (ours plus open alternatives) that decides per task which to use, records the decision and the reason, and produces a comparison report over a task set. The point is honesty: it must be free to choose against us, and the report must show where it did and why. The task set and scoring method are published so anyone can rerun it.

WHAT STRONG LOOKS LIKE

the alternatives are configured competently, not strawmen; the decision log is complete; the report states methodology, sample size, and limits; the repository includes a case where an alternative was the right choice; results reproducible from a clean checkout.

BAND S3 · MCP, chat, export · a pull request into the public repository, use-cases folder

Policy gateway in front of the document tools

WHO IT SERVES

Enterprise platform and security teams. Organisations that will not point an agent at an external tool without controls.

WHAT TO BUILD

A proxy that sits between an agent client and our tool server and adds organisational policy: an allowlist of which tools an agent may call, a pre-send scan that blocks or masks sensitive content, per-team quotas, and a full audit log of every call and its cost. Configuration is a policy file. Deploys as a container next to the team's other infrastructure.

WHAT STRONG LOOKS LIKE

the proxy is transparent (clients need only a URL change) and adds little latency; policy denials return a clear, agent-readable reason; the audit log is tamper-evident and excludes document content by default; a documented failure mode when the proxy is down; a test suite driving policy cases.

BAND S3 · MCP, API · a pull request into the public repository, extensions folder

Vertical workflow-prompt pack layered on the shipped prompts

WHO IT SERVES

Vertical software teams. Anyone whose users would rather type a slash command than a paragraph.

WHAT TO BUILD

A distributable pack of workflow prompts for one vertical, layered on top of the four that ship with the server: for example a legal pack with redline, fallback-clause, and obligation-summary commands, each capturing the right instructions, the right approval mode, and the right export settings. Installs as one entry in a client configuration. Includes a template for building a pack for any other vertical, which is the reusable part.

WHAT STRONG LOOKS LIKE

prompts work unchanged across at least three agent clients; arguments validated with helpful errors; each command documented with an example transcript; the pack-building template genuinely produces a working pack for a second vertical, demonstrated.

BAND S1 · MCP, chat, Review, templates, export · a pull request into the public repository, extensions folder

A vertical tool server that composes ours underneath

WHO IT SERVES

Vertical software teams. A company that wants its own branded agent tools without building a document engine.

WHAT TO BUILD

A thin tool server exposing a handful of domain verbs (for a contracts example: review against playbook, insert fallback, produce redline) that internally call our tools and add the domain judgement, the playbook, and the house rules. It is the reference pattern for a partner building on us: their tools, their names, their logic, our engine underneath. Includes a deployment guide and a note on where the boundary should sit.

WHAT STRONG LOOKS LIKE

the domain layer adds real judgement rather than renaming our tools; authentication of end users handled with per-customer scoping; the boundary is documented with the reasoning; error mapping from the underlying calls into domain language; a working end-to-end demonstration.

BAND S3 · MCP, chat, Review, templates, search, export · a pull request into the public repository, use-cases folder

Framework-agnostic adapter with three quickstarts

WHO IT SERVES

Agent developers on any framework.

WHAT TO BUILD

One small adapter layer that exposes our tools to any agent framework, plus three quickstarts showing the same task in three popular frameworks so a reader can compare and pick. The value over the existing per-framework rows is that this is one maintained core with thin bindings, so a new framework is an afternoon rather than a new project. Includes a written note on when to use the hosted tool server directly and when a wrapper is worth it.

WHAT STRONG LOOKS LIKE

the core has no framework dependency; each quickstart runs from a clean clone; the three produce identical output for the same task, which is the test; a documented process for adding a fourth binding; versioned releases.

BAND S2 · MCP, API, chat, export · a pull request into the public repository, extensions folder

Tracing and cost dashboard for agent document work

WHO IT SERVES

Platform engineering. Teams that need to explain what their agents spent and why.

WHAT TO BUILD

An instrumentation package that emits standard traces and metrics for every document tool call an agent makes: which tool, on which document, how long, how many operations, and whether a human approved it. A companion dashboard shows spend by agent, by document, and by outcome, and flags loops where an agent retried the same edit repeatedly. Exports a monthly summary as a styled document, which is a small joke that also works.

WHAT STRONG LOOKS LIKE

open, standard telemetry rather than a proprietary format; no document content in traces by default; the loop detector catches a deliberately looping agent; overhead measured and negligible; dashboards provided as importable definitions, not screenshots.

BAND S2 · MCP, API, chat, export · a pull request into the public repository, extensions folder

Client conformance suite for the full tool surface

WHO IT SERVES

Agent-client developers and integration teams.

WHAT TO BUILD

A test suite that drives all thirty-eight tools and all four workflow prompts through a client and reports which ones the client renders and calls correctly, including the awkward parts: streaming progress, a paused large edit that must be continued rather than approved, a review batch that is partially decided, and a large file that must go through the upload path. Produces a compatibility matrix per client, published and kept current.

WHAT STRONG LOOKS LIKE

the suite runs unattended and produces the same matrix twice in a row; failures are attributed precisely (client limitation versus configuration versus genuine defect); it covers the tools-only clients as well as the tools-and-prompts ones; the matrix is generated, never hand-edited.

BAND S3 · MCP, chat, Review, Multi-document, export · a pull request into the public repository, extensions folder

Mailbox pipeline: attachment in, edited document out

WHO IT SERVES

Operations and professional services. Teams whose intake is an email with a file attached.

WHAT TO BUILD

Chains an email server with our document server: a message arrives with a document attached and an instruction in the body, the pipeline loads the file, applies the instruction, and replies with the edited file and a summary of what changed. Sensitive requests route to a human approval step before the reply is sent. An allowlist controls who may drive it. Threading is preserved so the exchange reads as a normal conversation.

WHAT STRONG LOOKS LIKE

the allowlist is enforced and tested; instructions are treated as untrusted input with a documented handling rule; large attachments handled through the upload path; the reply states plainly what changed; a failure replies with a useful message rather than silence.

BAND S2 · MCP, chat, Review, export · a pull request into the public repository, use-cases folder

Slack workflow step with an approval card

WHO IT SERVES

Ops

WHAT TO BUILD

A native workflow step so document generation drops into Slack's own automation builder, with a Block Kit approve or reject card.

WHAT STRONG LOOKS LIKE

approval is attributed to a person and survives an app restart.

BAND S2 · MCP, Review, export · a pull request into the public repository, extensions folder

Slack message shortcut: thread to document

WHO IT SERVES

Ops

WHAT TO BUILD

Right-click any thread and turn it into a structured document with decisions and owners extracted.

WHAT STRONG LOOKS LIKE

it extracts real decisions, not a transcript summary.

BAND S2 · MCP, chat, export · a pull request into the public repository, extensions folder

Discord approval bot for community documents

WHO IT SERVES

Community

WHAT TO BUILD

Community-maintained guides where edits are proposed by anyone and approved by moderators.

WHAT STRONG LOOKS LIKE

role-gated approval, and a public change log the community can audit.

BAND S1 · MCP, Review · a pull request into the public repository, extensions folder

Mailbox label watcher

WHO IT SERVES

Ops

WHAT TO BUILD

Anything landing in a watched label or folder is processed by a standing rule and filed back.

WHAT STRONG LOOKS LIKE

idempotent, with a visible audit trail per message.

BAND S1 · API, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Outlook rules-driven auto-formatter

WHO IT SERVES

Enterprise

WHAT TO BUILD

Inbound attachments matching a rule are normalised to the house template before they reach the team.

WHAT STRONG LOOKS LIKE

it never modifies the original attachment, only produces a normalised copy.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

Object-storage pipeline

WHO IT SERVES

Platform

WHAT TO BUILD

An S3 or GCS bucket where a drop triggers ingest, edit and a written-back result, using pre-signed URLs throughout.

WHAT STRONG LOOKS LIKE

never downloads the payload to the orchestrator, and handles files far past inline limits.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Activepieces or Windmill node

WHO IT SERVES

Automation, self-hosted

WHAT TO BUILD

An open-source automation node so self-hosters get document actions natively.

WHAT STRONG LOOKS LIKE

merged upstream, or at least submitted, rather than living only in a fork.

BAND S1 · API, export · a pull request into the public repository, extensions folder

Appsmith or Budibase block

WHO IT SERVES

Internal tools

WHAT TO BUILD

A drop-in block that gives an internal tool document generation and review.

WHAT STRONG LOOKS LIKE

configurable without code, with sensible defaults.

BAND S2 · API, templates, Review · a pull request into the public repository, extensions folder

WhatsApp Business API assistant

WHO IT SERVES

SMB in high-WhatsApp markets

WHAT TO BUILD

Official Business API, template messages, a document in and a document out, in the user's language.

WHAT STRONG LOOKS LIKE

it works within template-message rules and handles a 24-hour session window correctly.

BAND S3 · API, chat, export · a pull request into the public repository, extensions folder

Browser extension for attachments in webmail and storage previews

WHO IT SERVES

Cross-industry

WHAT TO BUILD

Adds an edit action to attachments and previews across common webmail and storage sites.

WHAT STRONG LOOKS LIKE

it degrades gracefully when a site changes its DOM instead of breaking silently.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Deployable webhook receiver template

WHO IT SERVES

Platform

WHAT TO BUILD

A one-click deployable receiver that turns any product event into a generated or edited document, with signature verification and replay protection.

WHAT STRONG LOOKS LIKE

verifies signatures properly and is safe to expose publicly.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

Job-board resume tailoring extension

WHO IT SERVES

Careers

WHAT TO BUILD

On any job posting, tailor an uploaded resume to that posting with evidence-backed edits and a diff view.

WHAT STRONG LOOKS LIKE

it never invents experience, and every changed bullet traces to something already in the resume.

BAND S2 · chat, Review, export · a pull request into the public repository, extensions folder

Newsletter platform archive and repurposer

WHO IT SERVES

Marketing

WHAT TO BUILD

Publishes issues, then compiles them into a branded archive and repurposes each into a second format.

WHAT STRONG LOOKS LIKE

the compiled archive has a working table of contents and consistent styling across issues written months apart.

BAND S2 · API, Multi-document, export · a pull request into the public repository, extensions folder

Outbound document-event delivery service with replay

WHO IT SERVES

Integration engineers currently polling for document changes

WHAT TO BUILD

The emitting half of an integration. A service that turns document lifecycle events (a change applied, an approval decided, a job finished, a document archived) into signed, ordered, retried webhooks for a customer's own systems, with a replay endpoint so a consumer that was down for an hour can request exactly the window it missed.

WHAT STRONG LOOKS LIKE

at-least-once delivery with idempotency keys the consumer can deduplicate on; per-document ordering preserved under retry; a consumer down for an hour replays exactly the missed window, no more and no less; signature verification documented with a working example in two languages; a dead-letter path that a human can drain.

BAND S2 · API, Multi-document, export · a pull request into the public repository, extensions folder

Signature round trip, where the executed copy becomes the record

WHO IT SERVES

Legal ops, sales operations, HR and procurement whose workflow ends at signature

WHAT TO BUILD

The catalog is full of documents that get negotiated and none that get signed. This closes the loop with an electronic-signature provider: the approved document is sent for signature exactly as approved, signer routing and field placement are derived from the document rather than positioned by hand, the executed copy and its certificate return and are filed as an immutable version, and the record afterwards shows what was signed rather than what was last edited.

WHAT STRONG LOOKS LIKE

the file sent for signature is byte-identical to the approved version, proven by hash; signature and initial fields land on the right anchors in a document whose pagination changed since the template was made; the executed copy is stored immutably and a later edit forks a new version rather than mutating the signed one; a request cancelled mid-flight leaves no orphan version.

BAND S2 · API, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Print and mail fulfilment pipeline with proof reconciliation

WHO IT SERVES

Operations teams with statutory obligations to post paper

WHAT TO BUILD

Generated documents to paper through a print-and-post vendor, treated as a pipeline with a reconciliation obligation. Address validation and correction happen before spend, each piece gets a proof that is checked against what was generated, the batch carries a per-piece audit record, and returned-undeliverable outcomes flow back and mark the record.

WHAT STRONG LOOKS LIKE

the vendor's returned proof is compared to the generated file automatically rather than by eye, and a mismatch stops the batch; address failures surface before any postage is spent; the audit record reconciles piece by piece with the vendor's dispatch report; a partial batch failure never posts a duplicate.

BAND S2 · API, templates, export · a pull request into the public repository, use-cases folder

Device-side scan intake with cover-sheet routing

WHO IT SERVES

Back offices that still receive paper: clinics, law firms, schools, local government.

WHAT TO BUILD

Paper in at a multifunction device, a filed document out. A scan profile posts to the pipeline, a printed cover sheet carries the routing instruction so the sender chooses the treatment without touching a screen, low-quality scans escalate to a stronger extraction path rather than producing confident nonsense, and the result is filed with the derived metadata.

WHAT STRONG LOOKS LIKE

a skewed, stapled, double-sided scan is handled or honestly refused, never half-read; the cover sheet is the only training a sender needs; the escalation decision is made on measured page quality rather than on a guess; a re-scan of the same document does not create a second record.

BAND S2 · API, attachments, search, export · a pull request into the public repository, use-cases folder

Long-run editing job that survives a restart without duplicating work

WHO IT SERVES

Ops teams running document jobs that take longer than one sitting

WHAT TO BUILD

A durable runner for an editing job that is too long to finish in one sitting. It plans a multi-step edit over a document set, commits progress after every step, and on restart resumes from the last committed step rather than replaying from the beginning. The judged property is that a restart is invisible in the output: no step runs twice, no step is skipped, and the document after an interrupted-and-resumed run is byte-identical to the same run completed without interruption.

WHAT STRONG LOOKS LIKE

A run killed at a random step and restarted produces a document byte-identical to the uninterrupted run, proven by diff across at least twenty kill points; no edit is applied twice; the resume point is recoverable from durable state alone with no in-memory assumption; and the runner reports honestly which steps it skipped as already-done rather than silently re-issuing them.

BAND S3 · Multi-document, chat, export · a pull request into the public repository, use-cases folder

Published integration for the visual automation platforms

WHO IT SERVES

Integrations engineer. Comfortable shipping inside someone else's review process.

WHAT TO BUILD

Build and submit a real integration to the visual automation platforms where non-developers assemble workflows, covering at least two of them from one shared core. The integration exposes a small set of actions: create a document from a template, edit a section, request approval, export. It must use each platform's own connected-account model for authentication rather than asking a user to paste a key, and it must pass that platform's published submission review.

WHAT STRONG LOOKS LIKE

it is actually submitted, the authentication follows the host platform's model, and the same core serves both platforms without a fork.

BAND S2 · API, templates, export, Review · a pull request into the public repository, extensions folder

Issue tracker action: a ticket becomes a styled document

WHO IT SERVES

Integrations engineer who has lived in an issue tracker.

WHAT TO BUILD

Build an action inside a popular issue tracker that turns a ticket into a finished, styled document: a specification, a customer-facing summary, or a handover note, chosen by template. The document links back to the ticket, and re-running the action updates the existing document rather than creating a second one.

WHAT STRONG LOOKS LIKE

re-running is genuinely idempotent, and the link back survives a ticket being renamed or moved.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

The missing document-edit verb for agent tool suites

WHO IT SERVES

Teams running agent tool suites that can read, file, share and track documents but cannot change what is inside them.

WHAT TO BUILD

Many agent tool suites can read, create, move and share documents, but not edit them: an agent can do everything to a document except change what is inside it. Build the edit capability that closes that hole: a tool layer an existing agent suite can call, with SuperDocs as the editing engine behind it, so an instruction like 'update the renewal date in this agreement' becomes a reviewed, applied change and an exported file.

WHAT STRONG LOOKS LIKE

the edit verb feels native to the suite it extends, changes land only where intended with the review step intact, and the demo shows a document the agent could previously only read being changed end to end.

BAND S3 · API, MCP, chat, export · a pull request into the public repository, extensions folder

EXTENSIONS, DEV TOOLS AND SDK SURFACES

Raycast/Alfred quick-doc command

WHO IT SERVES

Power user or developer living in a launcher

WHAT TO BUILD

Build an extension for a launcher the user already runs: a single command takes a typed instruction ("offer letter for X, start date Y"), drafts the document through the API, and drops the exported file in the user's chosen folder, with progress states surfaced in the launcher's own UI. Package it the way that launcher's store expects, with configuration for the API key kept out of plaintext.

WHAT STRONG LOOKS LIKE

installable from a clean checkout by the launcher's normal mechanism; the command degrades honestly when the API is unreachable; a second invocation with the same phrasing produces a fresh document, never a cached one.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Dify/Flowise SuperDocs block

WHO IT SERVES

No-code AI-platform builder wiring agent flows on a visual canvas

WHAT TO BUILD

Build a drag-in canvas node for a visual agent-building platform: the node exposes document generate, targeted-edit, and export as configurable block inputs and outputs inside the platform's own flow model, packaged to that platform's plugin standard and submitted through its review process. This is the no-code sibling of the code adapters, which cover developers only.

WHAT STRONG LOOKS LIKE

a non-programmer can wire source-to-document-to-export on the canvas alone; node failures surface in the platform's own error model instead of silent nulls; the packaged node installs from the platform's standard mechanism.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Vibe-coder quickstart CLI (npx create-superdocs-app)

WHO IT SERVES

DevTools. A builder who wants a working, wired-up project from one terminal command rather than a checklist.

WHAT TO BUILD

Wiring a new project to SuperDocs today means reading the API docs and hand-assembling key setup, a first call, and project scaffolding yourself. Build a one-command CLI, the kind a builder runs as `npx create-superdocs-app`, that scaffolds a new project already wired to the public API and MCP server: key setup prompts, a working first call, and a choice of a couple of common stack templates. This is a scaffold generator, not another API client library: the deliverable is the zero-to-running-project command itself.

WHAT STRONG LOOKS LIKE

Running the published command on a clean machine with no prior setup produces a project that makes a real successful API call within the first few minutes, for every stack template it offers, not only the first one tested.

BAND S2 · API, MCP · a pull request into the public repository, extensions folder

Public changelog document maintainer

WHO IT SERVES

Open-source maintainer or devrel keeping a public changelog current release after release

WHAT TO BUILD

Build the living-changelog maintainer: a standing document that gets a targeted update per release - new entries added in the house style, older entries untouched, the document re-exported and published on each cycle. The one-time release-notes generators already exist; this is the maintainer that keeps one document alive across releases.

WHAT STRONG LOOKS LIKE

release N's update leaves releases 1 through N-1 byte-identical; entry style stays consistent across ten releases without manual reformatting; the published artifact and the maintained document never drift.

BAND S2 · chat, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Content refresh engine for a published library

WHO IT SERVES

Content / content or SEO lead at a company with hundreds of published articles, or an agency running client libraries

WHAT TO BUILD

Point the tool at an existing content library (exported pages or documents). It detects staleness signals per article: outdated years and version numbers, deprecated product names, dead references, statistics older than a threshold, and claims contradicted by a newer source you supply. For each article it proposes surgical edits with the reason attached, leaving structure, internal links and formatting untouched, and produces a refresh queue ranked by traffic value. The editor approves per article, and a change log documents what was updated and when. Runs as a command-line pass over a directory or inside CI.

WHAT STRONG LOOKS LIKE

edits are genuinely surgical so an article's structure, headings and internal links survive; staleness detection is specific and cites the reason rather than warning that something "may be outdated"; the change log is complete enough to defend the refresh to a client who pays for it.

BAND S2 · API, Multi-document, search, chat, Review, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Obsidian plugin: in-vault surgical editor

WHO IT SERVES

Writer, researcher, or developer with a large personal vault. Host app: Obsidian.

WHAT TO BUILD

Not an export plugin: an editing one. A command palette entry and a side panel let the user edit the current note, or a selection, by instruction, and see the proposed change as a diff inside Obsidian before accepting. Vault-wide commands work across many notes at once ("update the project name across every note in this folder, leave the daily notes alone"), still with a per-note diff to accept. Links, tags, front matter, and callouts survive every edit.

WHAT STRONG LOOKS LIKE

wiki links and block references never break. Front matter is preserved exactly. A vault-wide change is reviewable note by note and reversible as a whole. Nothing writes to disk without acceptance.

BAND S2 · chat, Review, Multi-document, search, API · a pull request into the public repository, extensions folder

Obsidian plugin: long-form manuscript compiler

WHO IT SERVES

Author, thesis writer, or documentation lead building a book out of notes. Host app: Obsidian.

WHAT TO BUILD

Define the manuscript's spine (an index note listing chapters in order) and the plugin compiles the vault into a real long-form document: chapters in order, heading levels normalised, front and back matter generated, footnotes converted from inline notes, figures placed with captions and numbering, citations resolved from a bibliography file, and a table of contents built. Output is a print-ready PDF and a DOCX an editor can mark up.

WHAT STRONG LOOKS LIKE

a 300-page compile is stable and repeatable, with figure and footnote numbering continuous across chapters. Citations resolve to the right entries. Recompiling after editing one chapter changes only what should change.

BAND S3 · Multi-document, templates, export, images, chat · a pull request into the public repository, extensions folder

VS Code extension: open .docx and .pdf inside the IDE

WHO IT SERVES

Developer, technical writer, or solo builder who resents leaving the editor. Host app: Visual Studio Code.

WHAT TO BUILD

Register a custom editor so double-clicking a .docx or .pdf in the file explorer opens it as a readable, editable document inside VS Code rather than launching Word. A side panel takes instructions, shows proposed changes as a familiar diff, and saves back to the same file. Multi-file commands work across a docs folder. The whole thing works over remote and container workspaces where Office is not installed at all.

WHAT STRONG LOOKS LIKE

the round-trip is lossless: open, edit one paragraph, save, and Word still shows the original styling everywhere else. Diffs are readable. It works in a remote container with no local Office.

BAND S3 · MCP, API, chat, Review, export · a pull request into the public repository, extensions folder

JetBrains IDE plugin

WHO IT SERVES

Java, Kotlin, Python, or PHP developer on IntelliJ, PyCharm, or WebStorm. Host app: JetBrains IDEs.

WHAT TO BUILD

The same document capability delivered as a JetBrains plugin: a tool window that edits documentation, specifications, and release material by instruction, previews changes with the IDE's own diff viewer, and exports the styled deliverable. Integrates with the IDE's version-control view so a document change lands in the same commit as the code change it describes.

WHAT STRONG LOOKS LIKE

it uses the IDE's native diff and notification surfaces rather than reinventing them. It respects project scope and never touches files outside it. Installs from the marketplace with no external setup beyond a key.

BAND S2 · MCP, API, chat, Review, export, images · a pull request into the public repository, extensions folder

LaTeX bridge for VS Code and Overleaf users

WHO IT SERVES

Academic, physicist, or economist whose source of truth is LaTeX but whose collaborators and journals want Word. Host app: VS Code with a LaTeX workflow, plus an Overleaf-compatible sync.

WHAT TO BUILD

Take a LaTeX project (root file, chapter includes, figures, bibliography, custom class) and bring it into an editable document with the maths lossless, footnotes real, citations resolved against the bibliography, cross-references numbered, and the source's page geometry honoured. Edit collaboratively or hand a Word file to a co-author, then bring their tracked changes back as reviewable edits. Anything that genuinely cannot be represented is flagged visibly, never dropped silently.

WHAT STRONG LOOKS LIKE

every equation survives both directions and exports as native Word maths. The co-author round-trip does not scramble the bibliography. Unsupported constructs produce a visible placeholder and a named warning, never a silent loss.

BAND S3 · API, chat, Review, export, Multi-document · a pull request into the public repository, extensions folder

Agent document workspace for coding tools

WHO IT SERVES

Developer running a coding agent that produces documents as part of its work. Host app: VS Code (panel), usable next to any agent.

WHAT TO BUILD

When an agent edits a document, the developer currently gets a wall of text and no way to check it. This panel gives them one: every document change an agent proposes appears as a reviewable before-and-after, grouped by turn, with approve, reject, and revert per change and a single log of what the agent actually did versus what it claimed. It is the human control surface for agent-driven document work, not another way to run the agent.

WHAT STRONG LOOKS LIKE

the claimed-versus-actual log catches a turn where the agent said it edited something and did not. Reverting one change does not disturb the others. It works with more than one agent client without special-casing any of them.

BAND S3 · MCP, Review, chat, export, Multi-document · a pull request into the public repository, extensions folder

Chrome side-panel document copilot

WHO IT SERVES

Anyone whose work document lives inside a web application. Host app: Chrome (side panel).

WHAT TO BUILD

A side panel that stays open next to whatever web app the user is in, holding the document they are working on. Select text in the page, send it to the panel, edit it there with full formatting fidelity, and paste the result back. Or work on an uploaded document alongside the web app entirely. It is the missing document workspace for the hundreds of web tools that have a text box and nothing more.

WHAT STRONG LOOKS LIKE

the panel survives navigation without losing state. Paste-back keeps the formatting the destination editor can accept and degrades predictably where it cannot. It never reads page content the user did not send.

BAND S2 · chat, API, Review, export, Multi-document · a pull request into the public repository, extensions folder

Enterprise browser extension with managed deployment

WHO IT SERVES

IT administrator rolling a browser tool out to a workforce. Host apps: Microsoft Edge and Firefox, alongside the Chrome build.

WHAT TO BUILD

The same extension, packaged for managed deployment: force-installed by policy, pre-configured with the organisation's key so no user ever sees one, allowed on an approved list of sites only, with data-handling settings the administrator controls and users cannot change. Includes the policy templates, the enterprise deployment guide, and the privacy summary a security review will ask for.

WHAT STRONG LOOKS LIKE

the policy templates apply cleanly and the settings genuinely cannot be overridden by a user. The site allow-list is enforced, not advisory. The privacy summary matches observable behaviour.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

Browser extension: clean up what the chat just gave you

WHO IT SERVES

Everyone who has pasted an AI answer into a document and watched the formatting collapse. Host: the browser, on any AI chat interface.

WHAT TO BUILD

A button next to any AI chat response that turns it into a real document: headings that are headings, tables that are tables, lists that nest correctly, code that stays monospaced, maths that renders, and no stray markdown characters. Choose a template and get the branded version. Keep appending answers from the same conversation into one growing document rather than pasting five times and fixing it five times.

WHAT STRONG LOOKS LIKE

a response with a nested table, code, and maths survives completely. The growing-document mode keeps consistent styling across appends. It works across several chat interfaces without a per-site rewrite.

BAND S1 · chat, export, templates · a pull request into the public repository, extensions folder

Zoom app: draft the deliverable during the call

WHO IT SERVES

Consultant, recruiter, solicitor, or sales engineer whose call produces a document. Host app: Zoom (in-meeting side panel).

WHAT TO BUILD

Not a post-call recap. During the meeting, the side panel builds the deliverable live: the scope of work, the offer, the advice note, the requirements list. Participants see it take shape, correct it in the moment ("no, the start date is the fifteenth"), and agree the content before anyone leaves. The document is finished and sent while the call is still open, on the firm's template.

WHAT STRONG LOOKS LIKE

the panel keeps up with the conversation without becoming noise. Corrections land immediately and visibly. The document sent at the end is one participants recognise, so nobody has to reconcile it later.

BAND S3 · chat, Review, templates, export, Multi-document · a pull request into the public repository, extensions folder

sdx, a pipe-friendly document CLI for Node

WHO IT SERVES

DevTools. Platform engineer who lives in a terminal and pipes JSON between tools.

WHAT TO BUILD

One npm-installed binary that treats documents like text streams. `sdx edit report.docx "tighten the executive summary"` prints a section-level diff; `--json` emits machine-readable results for `jq`; `cat notes.md | sdx new --title "Q3 review" --export pdf` builds a styled file from stdin. Exit codes distinguish "nothing needed changing" from a real failure. Profiles live in a dotfile, `--session` keeps one conversation across invocations, and `--approve` walks each proposed change in the terminal before it lands. Ships bash, zsh, and fish completions.

WHAT STRONG LOOKS LIKE

a real 120-page styled .docx handled through the pre-signed upload path so file bytes never bloat the process; documented exit codes; a stable, tested `--json` schema; backoff on rate limits; API keys never printed to logs or shell history; integration tests against a recorded fixture server so CI needs no live key.

BAND S2 · chat, API, export, Review · a pull request into the public repository, extensions folder

superdocs-cli, the Python-native command line client

WHO IT SERVES

DevTools. Backend Python teams and data engineers who already have a requirements.txt.

WHAT TO BUILD

A pip install-able CLI with named profiles, environment-file loading, and --dry-run that shows exactly which sections a request would touch before spending an operation. Subcommands cover draft, edit, search, export, and files. Rich terminal rendering shows a per-section before/after. A --plan flag writes a reviewable YAML plan that a later apply consumes, which is what makes it safe to run against a client's real documents. Config precedence (flag, environment, profile, default) is documented and tested.

WHAT STRONG LOOKS LIKE

typed with full annotations and shipped with a py.typed marker; --dry-run provably spends zero operations; unicode and right-to-left filenames handled; a clean uninstall; docs generated from the code, not hand-maintained; tested on Windows as well as macOS and Linux.

BAND S2 · chat, API, search, export, Multi-document · a pull request into the public repository, extensions folder

Single-binary Go CLI with real release engineering

WHO IT SERVES

DevTools. Site-reliability and infrastructure engineers who refuse to install a language runtime in a container.

WHAT TO BUILD

A statically linked, zero-dependency binary cross-compiled for macOS, Linux, and Windows on amd64 and arm64. Same verbs as the Node and Python clients, but the value is distribution: signed release archives, checksums, a Homebrew tap, a Scoop manifest, and a go install path. Long operations use the asynchronous job flow with a resumable job id, so a dropped SSH session does not lose work. --output supports table, JSON, and a compact CI format.

WHAT STRONG LOOKS LIKE

reproducible builds, a real release workflow (tag to signed artefacts), a version command that reports build metadata, graceful handling of a job that outlives the process, and a documented minimum Go version. Bonus for a --verify command that checks connectivity, auth, and quota in one call.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

sdoc, a Rust batch-editing CLI

WHO IT SERVES

DevTools. Performance-minded engineers and Rust-shop platform teams.

WHAT TO BUILD

A crates.io CLI built for volume: point it at a folder or a glob, give one instruction, and it fans the work out across a bounded worker pool while a live progress display shows per-file state. Every run writes a manifest so an interrupted batch resumes exactly where it stopped, and never re-spends an operation on a file already finished. Structured, typed errors distinguish a file-level failure from a run-level one. Optional --budget caps total operations for the run.

WHAT STRONG LOOKS LIKE

concurrency that respects rate limits rather than hammering through them; a resumable manifest proven by killing the process mid-run; zero panics on malformed input files; memory that stays flat across a thousand-file run; cargo test plus a soak test in CI.

BAND S3 · API, chat, export · a pull request into the public repository, extensions folder

POSIX shell client and curl cookbook, zero runtime

WHO IT SERVES

DevTools. Engineers scripting inside minimal CI containers, and anyone who cannot install packages.

WHAT TO BUILD

One portable shell file plus a companion cookbook of copy-paste curl commands covering the whole flow: create a session, upload a large file through the pre-signed URL path, run an edit, poll an asynchronous job, approve changes, and export. Works with only curl and a JSON tool present, and degrades to pure curl with documented parsing when it is not. Each recipe states which operation count it spends. Doubles as the fallback documentation any agent can read and execute.

WHAT STRONG LOOKS LIKE

genuinely POSIX (verified under dash and busybox, not just bash); secrets passed by environment, never as visible arguments; every recipe runnable end to end from a clean container; failure paths shown, not only happy paths; a matching one-page cheat sheet.

BAND S1 · API, export, Review · a pull request into the public repository, extensions folder

superdocs diff, git-style review in the terminal

WHO IT SERVES

DevTools. Engineers who already trust a diff-and-approve workflow for code and want it for documents.

WHAT TO BUILD

A review-first CLI. Ask for a change and nothing lands until you say so: the tool renders a coloured, section-scoped before/after in the terminal, exactly the way git add -p walks a patch, with keys for approve, deny, skip, and send back with feedback. It reads the proposed batch from the human-in-the-loop flow, so this is the real approval path, not a local imitation. --stat prints a summary of sections touched; --patch-file saves the review for someone else to run.

WHAT STRONG LOOKS LIKE

approvals survive a terminal reload because pending state is fetched from the job rather than held in memory; denied changes provably leave the document untouched and bill nothing; wide tables and non-Latin scripts render without breaking the column layout; keyboard model documented and consistent.

BAND S2 · chat, Review, API · a pull request into the public repository, extensions folder

Batch-edit pipeline runner with a result ledger

WHO IT SERVES

Operations engineering. Teams that have a folder of documents and a policy change to apply to all of them.

WHAT TO BUILD

Point the runner at a source folder and a rules file: each rule is a selector (filename pattern, or a semantic search hit inside the document) plus an instruction. It processes files concurrently, writes one row per file to a ledger (input, sections touched, operations spent, output path, outcome), and produces a summary report as a styled document at the end. Re-running skips completed rows unless --force. A --sample N mode proves the rules on a subset before committing to the full run.

WHAT STRONG LOOKS LIKE

the ledger is the source of truth and is crash-safe; partial failures never corrupt an output file; the sample mode genuinely predicts full-run behaviour; per-file cost reported honestly including failed turns that billed nothing; a real corpus (mixed .docx, PDF, Markdown) in the fixtures.

BAND S3 · API, chat, search, export · a pull request into the public repository, extensions folder

docbuild, a declarative document build graph

WHO IT SERVES

DevTools. Documentation engineers who want documents built the way software is built.

WHAT TO BUILD

A build tool for documents. A docbuild.toml declares targets: sources (Markdown, data files, an existing .docx template), the instruction or transformation to apply, and outputs in any of the five export formats. The tool computes a dependency graph, rebuilds only what changed since the last run using content hashes, and caches results. docbuild watch rebuilds on save; docbuild graph renders the dependency graph as a diagram inside the generated documentation itself.

WHAT STRONG LOOKS LIKE

incremental rebuilds provably skip unchanged targets (and therefore spend zero operations); the cache is correct across machine and clock changes; cycles detected with a clear error; parallel target execution; the tool's own documentation built by itself, which is the demo.

BAND S3 · API, chat, templates, export, images · a pull request into the public repository, extensions folder

Terminal user-interface document editor

WHO IT SERVES

DevTools. Developers who never leave the terminal, plus anyone on a remote server.

WHAT TO BUILD

A full-screen terminal application: left pane lists the document's sections, right pane shows the selected section, bottom pane is the chat. Type an instruction, watch progress stream in, and approve or deny each proposed change with a keystroke. Tabs across multiple open documents in one session. / runs a semantic search and jumps to the hit. Export from a menu. It is the web app's core loop, rebuilt for people who work over SSH.

WHAT STRONG LOOKS LIKE

live progress rendered from the stream rather than a spinner; the interface stays responsive while a long edit runs; resize, mouse-less operation, and screen-reader-friendly output; state recovered cleanly after a disconnect; a demo recording that shows a real multi-section edit under a minute.

BAND S3 · chat, Review, Multi-document, search, export · a pull request into the public repository, extensions folder

docs-lint, a house-style gate for documents in a repository

WHO IT SERVES

DevTools and Marketing operations. Teams with a style guide nobody follows.

WHAT TO BUILD

A linter for prose and formatting that runs as a pre-commit hook and in CI. Point it at a style guide (uploaded once as a reference the AI consults) and it reviews changed documents against it, reporting findings with section anchors. --fix proposes the corrections as a reviewable change set rather than silently rewriting, so a human still approves. Configurable severity per rule; exit non-zero only on the severities you choose.

WHAT STRONG LOOKS LIKE

fast on unchanged files (only touched documents are checked); findings map to precise locations; --fix never applies anything without approval; the style guide is data, not code, so any team can swap theirs in; a sample repository showing a failing then passing run.

BAND S2 · chat, search, templates, Review · a pull request into the public repository, extensions folder

superdocs watch, the save-triggered document dev loop

WHO IT SERVES

DevTools. Anyone who wants documents to feel like hot-reloading code.

WHAT TO BUILD

A watcher that monitors a source file or folder and re-runs a stored instruction on every save, then opens or refreshes the exported preview. Debounced so a fast typist does not trigger a run per keystroke, with a per-run cost line printed so you always know what a save just spent. A --budget flag stops the watcher after N operations. Includes a "preview only" mode that renders without spending, for layout iteration.

WHAT STRONG LOOKS LIKE

debounce and coalescing that provably prevent runaway spend; watcher survives editors that write files atomically (rename-over); clean shutdown; the cost line is accurate including zero-cost turns; works on all three operating systems.

BAND S1 · API, chat, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

superdocs shell, an interactive session REPL

WHO IT SERVES

DevTools. Developers exploring the API before writing integration code.

WHAT TO BUILD

A stateful read-evaluate-print loop. open contract.docx loads a document, then plain English lines are edits, and slash commands cover the structural operations: /tabs lists open documents, /focus switches, /search finds a clause, /history shows the turns, /revert 3 rewinds to a message, /export pdf writes a file. Command history persists. /replay turns the session into a runnable script in the language of your choice, which is how exploration becomes integration code.

WHAT STRONG LOOKS LIKE

/replay emits code that actually runs unchanged; revert and redo both work and are demonstrated; the loop recovers from a network drop mid-turn; tab completion for document and session names; a transcript in the README that a reader can follow verbatim.

BAND S2 · chat, Multi-document, search, export, Review · a pull request into the public repository, extensions folder

Container image of the CLI, one-line document jobs anywhere

WHO IT SERVES

DevOps. Teams that run everything as a container and want no local installs.

WHAT TO BUILD

A small, multi-architecture image published to a public registry that wraps the CLI: docker run --rm -v "\$PWD:/work" superdocs/cli edit /work/report.docx "add a risks section". Volume conventions, a non-root user, and an entrypoint that passes through exit codes. Documented tags for pinning. A companion Compose file shows a scheduled job and a queue-worker shape for teams that want batch processing without writing a service.

WHAT STRONG LOOKS LIKE

image under a sensible size with a documented base; builds for amd64 and arm64; runs as non-root and handles file ownership on the mounted volume correctly; no secret ever baked into a layer; a published software bill of materials and a vulnerability scan in CI.

BAND S1 · API, chat, export · a pull request into the public repository, extensions folder

PowerShell module for Windows and enterprise shops

WHO IT SERVES

Enterprise IT. Windows administrators automating document workflows in organisations where Word is the standard.

WHAT TO BUILD

A module of idiomatic cmdlets: New-SuperDoc, Edit-SuperDoc, Export-SuperDoc, Get-SuperDocSession, with pipeline support so `Get-ChildItem *.docx | Edit-SuperDoc -Instruction "apply the 2026 letterhead"` works the way a Windows administrator expects. Credentials stored through the platform's secret management rather than plaintext. Verbose and what-if support on every state-changing cmdlet.

WHAT STRONG LOOKS LIKE

genuinely idiomatic (approved verbs, pipeline input by property name, `-WhatIf` honoured), works on both Windows PowerShell and cross-platform PowerShell, help written as comment-based help so `Get-Help` is complete, module signed, and published with a clear licence.

BAND S2 · API, chat, templates, export · a pull request into the public repository, extensions folder

mcp-doctor, a setup diagnostic for agent clients

WHO IT SERVES

DevTools. Anyone whose agent client says "server failed to connect" with no further detail.

WHAT TO BUILD

A diagnostic command that finds the configuration files of the major agent clients, checks them for the usual defects (wrong transport key, trailing slash, a header written in the wrong shape, a key that is present but revoked), performs a real handshake against the server, lists the tools and prompts it can see, and prints a numbered fix for each problem it found. `--fix` writes the corrected configuration after showing a diff. `--report` produces a shareable, redacted summary for a support thread.

WHAT STRONG LOOKS LIKE

it never prints or transmits a key; it distinguishes a network problem from an authentication problem from a configuration-shape problem; the fixes are correct for each client's actual configuration format; a test matrix covering the client configurations, driven from fixtures.

BAND S2 · MCP, API · a pull request into the public repository, extensions folder

Fuzzy-finder document picker for the terminal

WHO IT SERVES

DevTools. Terminal power users with a growing document library.

WHAT TO BUILD

A small integration that pipes the durable Files library into a fuzzy finder: type a few characters, preview the document's outline and section count in the preview pane, hit enter to open it in an editing session, and pick an action (edit, search, export, rename, archive). Shell functions and key bindings for bash, zsh, and fish. A companion widget inserts the chosen document's identifier into the current command line, which makes it composable with every other tool here.

WHAT STRONG LOOKS LIKE

the listing is fetched cheaply (metadata only, no document bodies pulled into the terminal); the preview is instant and cached; key bindings do not clash with common configurations; archived documents handled distinctly; works with the two most common fuzzy finders, not just one.

BAND S1 · API, Multi-document, search, export · a pull request into the public repository, extensions folder

sd-migrate, port a regenerate-everything pipeline to section edits

WHO IT SERVES

DevTools. An engineer at a company that already built its own document pipeline and is now paying for it every month.

WHAT TO BUILD

A tool aimed at the exact prospect our business brief calls prime: a team whose current pipeline rebuilds a whole document to change one paragraph. Point it at their pipeline's inputs and outputs, and it produces a migration report: which of their operations are section-scoped in principle, a side-by-side of the same change made both ways, the measured token and time difference on their own documents, and generated adapter code for the language they use. It changes nothing automatically; it produces evidence and a starting patch.

WHAT STRONG LOOKS LIKE

the measurement is honest, reproducible, and run on the customer's documents rather than a rigged sample; the report states its methodology and its limits; the generated adapter compiles and passes a smoke test; no customer content leaves the machine except the documents the operator explicitly submits.

BAND S3 · API, chat, export · a pull request into the public repository, extensions folder

Python SDK, typed and both synchronous and asynchronous

WHO IT SERVES

DevTools. Backend engineer at a legal-tech or fintech startup wiring document editing into an existing service.

WHAT TO BUILD

A first-class client library: one class with typed methods for sessions, documents, chat, approvals, files, and export, offered in both blocking and async forms from the same source. Uploads pick the right path automatically by file size so the caller never thinks about it. Long operations expose both a blocking wait() and an iterator over progress events. Rich exception types carry the machine-readable error code and the suggested action. A context manager cleans up sessions.

WHAT STRONG LOOKS LIKE

complete type coverage checked in CI; async and sync behaviour verified identical by a shared test suite; retries idempotent so a retried edit cannot apply twice; timeouts configurable per call; a documented pagination helper; published with source and wheel distributions and a changelog.

BAND S2 · API, chat, Review, Multi-document, export, search · a pull request into the public repository, extensions folder

TypeScript SDK, isomorphic and fully typed

WHO IT SERVES

DevTools. Full-stack engineer shipping a document feature in a web product.

WHAT TO BUILD

One package that works in Node, in the browser through a server proxy, and in edge runtimes. Types generated from the published specification so they never drift. First-class streaming: a typed async iterator over progress events with automatic reconnection that resumes from the last event received rather than replaying the whole turn. Helpers for the review flow return the proposed change batch already shaped for rendering a diff component. Tree-shakeable, dual module formats, no runtime dependency on a specific HTTP client.

WHAT STRONG LOOKS LIKE

types compile under strict mode and are tested with type-level assertions; the browser path never requires shipping a secret key to the client and the README says so loudly; reconnection tested by killing the connection mid-stream; bundle size published and watched in CI.

BAND S2 · API, chat, Review, export, Multi-document · a pull request into the public repository, extensions folder

Go SDK with context, options, and streaming readers

WHO IT SERVES

DevTools. Backend teams running Go services in regulated or high-throughput environments.

WHAT TO BUILD

An idiomatic client: every call takes a context, configuration uses functional options, uploads accept an `io.Reader` so a large file streams from disk without being held in memory, and exports return a reader the caller can pipe straight to storage. Errors are typed values that support the standard errors-as helpers. A separate sub-package provides the review flow as a small state machine. Zero third-party dependencies beyond the standard library.

WHAT STRONG LOOKS LIKE

cancellation actually cancels in-flight work; no goroutine leaks under a race-detector test; a documented, stable module path with semantic versioning; example tests that appear in generated documentation; benchmark tests for the streaming upload path.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

Java SDK plus a Spring Boot starter

WHO IT SERVES

Enterprise. Java platform teams in banking, insurance, and large corporates.

WHAT TO BUILD

A dependency-light Java client with both blocking and future-returning methods, plus a companion auto-configuration starter so a Spring application gets a ready client from properties alone. Immutable request builders, a pluggable HTTP layer, and a retry policy that is configuration, not code. The starter registers health and metrics for the document client so it appears in the existing observability stack. Java 17 baseline with a module descriptor.

WHAT STRONG LOOKS LIKE

published to a central artefact repository with signed artefacts, sources, and Javadoc; no dependency version conflicts in a realistic enterprise application; thread-safe client documented as such; an integration test using a mock server; a migration note for the next major version.

BAND S3 · API, chat, Review, export, templates · a pull request into the public repository, extensions folder

.NET SDK with dependency-injection integration

WHO IT SERVES

Enterprise. Microsoft-stack teams, which is most corporate document work.

WHAT TO BUILD

A modern .NET client built on the platform's HTTP factory so pooling, resilience, and telemetry come from the host application rather than being reinvented. Everything asynchronous with cancellation tokens, nullable annotations enabled, and an extension method that registers the client in one line. A source-generated serialisation path keeps it trim-friendly. Includes a small worker-service sample that watches a folder and produces branded documents.

WHAT STRONG LOOKS LIKE

works on the current long-term-support framework and the previous one; no synchronous-over-asynchronous deadlock traps; package published with symbols and deterministic builds; analyser-clean; the worker sample runs from a clean clone with only a key set.

BAND S2 · API, chat, templates, export, Review · a pull request into the public repository, extensions folder

Ruby gem with a background-job adapter

WHO IT SERVES

DevTools. Rails shops, which are heavily represented in legal-tech and human-resources software.

WHAT TO BUILD

An idiomatic gem: a configuration block, plain-old-Ruby result objects, and a documented error hierarchy. The differentiator is the job adapter: long document work is enqueued as a background job with the polling, retry, and review-state handling already written, so a Rails developer adds one job class and a callback rather than building an asynchronous pipeline from scratch. Generators scaffold a controller and view for the approve-or-deny screen.

WHAT STRONG LOOKS LIKE

the adapter is queue-agnostic; retries are idempotent so a re-enqueued job cannot double-apply an edit; the generated review screen renders a real before-and-after diff; tested against the two most recent Rails versions; gem signed and released with a changelog.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

PHP SDK with a Laravel service provider

WHO IT SERVES

DevTools. PHP agencies and the enormous long tail of business software built in PHP.

WHAT TO BUILD

A standards-based client (interoperable HTTP client and message interfaces so it drops into any framework) plus a Laravel package: service provider, facade, configuration file, queued job class, and an artisan command for one-off edits. Uploads use the streaming path. The queued job handles the review flow, exposing events an application can listen to when changes are proposed and when they are applied.

WHAT STRONG LOOKS LIKE

framework-agnostic core with the framework glue in a separate namespace; supported on current PHP versions with static analysis at a high level; no global state; the artisan command handles a large file without memory limits being hit; published with a clear versioning policy.

BAND S2 · API, chat, Review, export, templates · a pull request into the public repository, extensions folder

Kotlin multiplatform client for server and Android

WHO IT SERVES

Mobile and backend. Teams building a document feature into an Android application and its backend with one library.

WHAT TO BUILD

A coroutine-first client compiled for the Java virtual machine and Android from one source set, with flows for progress events and suspending functions everywhere else. Android specifics handled properly: lifecycle-aware collection, work-manager integration for long jobs that must survive the app being backgrounded, and file access through the platform's document provider so a user picks a file from their device and gets an edited one back.

WHAT STRONG LOOKS LIKE

no blocking calls on the main thread, proven by a strict-mode test; work survives process death; a sample application in the repository that a reviewer can build and run; artefacts published for both targets; a documented minimum Android level.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Swift package for Apple platforms

WHO IT SERVES

Mobile. iOS and macOS developers building document features, including the growing set of AI note-taking apps.

WHAT TO BUILD

A native Swift package using structured concurrency throughout, with typed errors and an actor-isolated client. A sample application shows the full loop on iOS: pick a file, ask for a change in plain English, review the proposed edits in a native diff view, and export back to the system file provider. On macOS the same package powers a small menu-bar utility. Async sequences carry progress events.

WHAT STRONG LOOKS LIKE

strict concurrency checking enabled and clean; no key stored in the application bundle, with the correct pattern documented; the sample builds from a clean checkout; accessibility labels on the review interface; package published with platform and version requirements stated.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Rust SDK crate with typed errors

WHO IT SERVES

DevTools. Rust services and anyone building the batch CLI above on a shared foundation.

WHAT TO BUILD

An async crate with a typed error enumeration mapping every documented failure to a variant, builder-pattern requests, and streaming upload and download that never buffers a whole file. Feature flags keep the dependency tree small for callers who only need one part. Provides the shared foundation the Rust CLI card uses, so the two ship as one workspace: library plus binary.

WHAT STRONG LOOKS LIKE

documented public API with examples that compile as doc tests; no unsafe code; minimum supported Rust version stated and enforced in CI; the error enumeration is exhaustive against the published specification; published with full documentation on the crate host.

BAND S3 · API, chat, export, Review · a pull request into the public repository, extensions folder

Elixir client with supervised job polling

WHO IT SERVES

DevTools. Elixir teams, common in messaging, healthcare, and high-concurrency back offices.

WHAT TO BUILD

An idiomatic package where the interesting part is the concurrency model: each long-running document job becomes a supervised process that polls, emits telemetry events, and reports completion to a caller process, so a thousand concurrent document jobs is an ordinary supervision tree rather than a bespoke worker pool. Includes telemetry events for cost and duration, and a live-view sample showing real-time progress for many documents at once.

WHAT STRONG LOOKS LIKE

back-pressure and rate-limit handling at the supervisor level; crash of one job never affects another; telemetry events documented; the live-view sample genuinely runs many concurrent jobs; published with typespecs and generated documentation.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Specification-driven client generation pipeline

WHO IT SERVES

DevTools. Platform engineers who would rather generate a client than hand-write one.

WHAT TO BUILD

A repository that consumes the published API specification and produces typed clients for several languages on a schedule, publishing each to its registry when the specification changes. Includes the generator configuration, the post-generation patches that make each output idiomatic, a compatibility report between versions, and a smoke test per language that runs against a mock server. The point is not one more client; it is the machinery that keeps every client honest.

WHAT STRONG LOOKS LIKE

regeneration is fully reproducible from a clean checkout; breaking changes detected and reported rather than silently published; each generated client passes its smoke test before release; the patches are minimal and explained; the whole thing runs in CI with no manual step.

BAND S3 · API · a pull request into the public repository, extensions folder

Idempotency and safe-retry middleware

WHO IT SERVES

DevTools. Anyone who has double-charged a customer or double-applied an edit because of a retry.

WHAT TO BUILD

A small library, implemented in two languages from one written specification, that makes document operations safe to retry: caller-supplied operation keys, a local record of in-flight and completed operations, a decision table for which failures are retryable, and jittered backoff. It refuses to retry anything whose outcome is unknown without an explicit policy choice, and it exposes a reconciliation call that checks the server's job state before deciding. Ships with a fault-injection test harness.

WHAT STRONG LOOKS LIKE

the fault-injection harness proves no double-apply under network partition, duplicated request, and process restart; the policy is data, not conditionals scattered through the code; the two implementations are verified against the same behavioural test vectors; every decision documented with the reasoning.

BAND S3 · API, chat, Review · a pull request into the public repository, extensions folder

Streaming client library with resume

WHO IT SERVES

DevTools. Front-end and backend engineers rendering live progress.

WHAT TO BUILD

A focused library for the progress stream: typed events, automatic reconnection that resumes from the last received sequence instead of replaying the turn, a heartbeat-based staleness detector, and back-pressure so a slow consumer cannot be overwhelmed. Adapters for browser, Node, and a server-side language. Includes a small reference interface component that renders the standard event set (progress, proposed changes, document changes, model notice, completion, usage, error) so an integrator sees the whole vocabulary working.

WHAT STRONG LOOKS LIKE

resume verified by cutting the connection at several points and proving no event is lost or duplicated; unknown future event types ignored safely rather than crashing; memory bounded on a long stream; the reference component is genuinely reusable, not a demo page.

BAND S2 · API, chat, Review, Multi-document · a pull request into the public repository, extensions folder

Asynchronous job and review state machine as a library

WHO IT SERVES

DevTools. Teams building an approval interface, which is most regulated-industry integrations.

WHAT TO BUILD

The review flow is a state machine: submitted, working, awaiting approval, partially decided, resuming, finished, failed, plus the pause-and-continue branch for very large edits. This library implements it once, correctly, with persistence hooks so the state survives a page reload or a server restart, and a reconciliation call that rebuilds an interface from the server's durable job state after a client comes back. Framework-agnostic core plus one framework binding.

WHAT STRONG LOOKS LIKE

every state transition covered by tests including the awkward ones (partially decided batch, reload mid-approval, a continue prompt mistaken for an approval); reconciliation proven by killing and restoring a client; the persistence hook is an interface so any store works; the diagram in the README is generated from the code.

BAND S3 · API, chat, Review · a pull request into the public repository, extensions folder

Python web-service reference integration

WHO IT SERVES

DevTools. A team that needs a working service, not a snippet.

WHAT TO BUILD

A production-shaped reference service in a modern Python web framework: authentication, per-end-customer scoping so one tenant's documents and memory never touch another's, a queue for long jobs, webhook or polling delivery of results, an approval endpoint pair, and structured logging with cost per request. Deployable with one command to a container host. It is written to be read: every architectural decision has a short comment explaining why, including the ones about scoping and idempotency.

WHAT STRONG LOOKS LIKE

tenant isolation demonstrated by a test that tries to cross the boundary and fails; secrets never in the repository; load-tested with a report; migrations and configuration handled properly; a genuine README that gets a stranger running in ten minutes.

BAND S3 · API, chat, Review, Multi-document, memory, export · a pull request into the public repository, extensions folder

Next.js integration package with React hooks

WHO IT SERVES

DevTools. Product engineers adding an AI document feature to a web application this quarter.

WHAT TO BUILD

A package plus a sample application: server route handlers that keep the key server-side, React hooks for a document session, a streaming chat hook, and a review hook that returns proposed changes ready to render. Includes an unstyled diff component and a headless approval component so any design system can adopt them. The sample is a complete small product: upload, edit, review, export, and a files list.

WHAT STRONG LOOKS LIKE

the key provably never reaches the browser; hooks handle unmount and rapid re-render without leaking streams; the sample deploys from a clean fork with one environment variable; accessible components; a bundle-size budget.

BAND S2 · API, chat, Review, Multi-document, export · a pull request into the public repository, extensions folder

Serverless adapter pack for long-running document jobs

WHO IT SERVES

DevTools. Teams on function platforms who hit the execution-time wall the first day they try this.

WHAT TO BUILD

Three deployable templates (a major cloud's functions, an edge worker platform, and a hosting platform's functions) that solve the one hard problem: a document job can outlive a function invocation. Each template shows the correct shape: start the job, return immediately, and finish through a durable path (queue, scheduled poller, or a stateful worker), with the review flow surviving across invocations. Includes cost notes and cold-start measurements.

WHAT STRONG LOOKS LIKE

each template actually deploys and completes a long job end to end; no busy-waiting inside a function; state is durable, not in memory; the failure mode when a platform limit is hit is explained and handled; the measurements are real and reproducible.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Typed error taxonomy and remediation library

WHO IT SERVES

DevTools. Support-conscious engineers and anyone writing an agent that must recover on its own.

WHAT TO BUILD

One canonical, machine-readable catalogue of every error a caller can receive, each with its code, its cause in plain language, whether it is retryable, and the concrete remediation. Published as data (a schema file) plus thin libraries in two languages that turn a raw failure into a typed object with a suggested_action. Includes the size and capacity refusals, which are expandable limits rather than hard caps, so an agent knows to ask for an increase rather than give up.

WHAT STRONG LOOKS LIKE

the catalogue is generated or verified against real responses rather than hand-written from memory; unknown codes degrade to a safe default; remediation text is written for both a human and an agent; a test that asserts the catalogue covers every documented status.

BAND S2 · API · a pull request into the public repository, extensions folder

Mock server and contract-test fixture package

WHO IT SERVES

DevTools. Every team that wants continuous integration to run without a live key or a monthly quota.

WHAT TO BUILD

A runnable mock built from the published specification, with recorded realistic responses (a styled document with sections, a proposed change batch, a paused large edit, a capacity refusal, a rate limit) and a record-and-replay mode for capturing new scenarios. Ships as a container and a library, plus a contract-test suite an integrator runs against both the mock and the live service to prove they agree. Scenario switching by header so one mock covers many test cases.

WHAT STRONG LOOKS LIKE

the contract suite genuinely catches a drift between mock and live; recorded fixtures carry no real customer content; deterministic responses so tests do not flake; the mock starts in under a second; documented scenario catalogue.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

Notion two-way sync agent

WHO IT SERVES

Content ops

WHAT TO BUILD

Keeps a Notion page and a SuperDocs document in sync, with conflict prompts when both moved.

WHAT STRONG LOOKS LIKE

a real conflict policy, not last-writer-wins, and block-level mapping that survives a round trip.

BAND S3 · API, Multi-document · a pull request into the public repository, extensions folder

Notion database to contract pipeline

WHO IT SERVES

Sales

WHAT TO BUILD

Each database row becomes a real contract or proposal with the layout the counterparty expects.

WHAT STRONG LOOKS LIKE

the produced file looks like a lawyer made it, not like a Notion export.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

Obsidian agent-maintained living note

WHO IT SERVES

Personal knowledge

WHAT TO BUILD

A plugin where an agent keeps a designated note current from the rest of the vault, editing only the sections it owns.

WHAT STRONG LOOKS LIKE

the user's own paragraphs are never rewritten, and the ownership boundary is explicit in the file.

BAND S2 · API, search, memory · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Obsidian vault to manuscript compiler

WHO IT SERVES

Authors

WHAT TO BUILD

Compiles ordered notes into a single manuscript with front matter, consistent voice and a print-ready output.

WHAT STRONG LOOKS LIKE

chapter order and cross-references are preserved, and a second compile after edits changes only what moved.

BAND S2 · Multi-document, chat, export · a pull request into the public repository, extensions folder

Logseq or Roam block-level editor

WHO IT SERVES

Networked thought

WHAT TO BUILD

Edits at block granularity so the graph structure survives.

WHAT STRONG LOOKS LIKE

block refs still resolve after the edit.

BAND S2 · API, chat · a pull request into the public repository, extensions folder

Apple Pages and iWork bridge

WHO IT SERVES

Mac-first users

WHAT TO BUILD

Reads a .pages bundle, edits the content, and returns a document the user can keep working in.

WHAT STRONG LOOKS LIKE

honest about what round-trips and what does not, with a documented capability matrix.

BAND S2 · API, export · a pull request into the public repository, extensions folder

LibreOffice extension

WHO IT SERVES

Open-source and public sector

WHAT TO BUILD

A native LibreOffice Writer extension with a sidebar editing panel.

WHAT STRONG LOOKS LIKE

installs as a single .oxt, works on Linux, and respects the document's existing styles.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

OnlyOffice or Collabora plugin

WHO IT SERVES

Self-hosted collaboration

WHAT TO BUILD

Adds AI editing to a self-hosted document server.

WHAT STRONG LOOKS LIKE

works fully inside the customer's network boundary except for the API call, and documents exactly what leaves.

BAND S2 · API, Review, export · a pull request into the public repository, extensions folder

Zoho Writer extension

WHO IT SERVES

SMB, India and global

WHAT TO BUILD

Sidebar editing inside Zoho's editor.

WHAT STRONG LOOKS LIKE

uses Zoho's own review and track-change primitives rather than reinventing them.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

WPS Office add-in

WHO IT SERVES

Asia-Pacific

WHAT TO BUILD

An add-in for the office suite with an enormous installed base outside the Microsoft ecosystem.

WHAT STRONG LOOKS LIKE

the UI is localised properly and CJK text is handled correctly end to end.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Live-doc agent for Quip, Slite or Nuclino

WHO IT SERVES

Team wikis

WHAT TO BUILD

An agent that maintains a team's canonical pages as decisions land.

WHAT STRONG LOOKS LIKE

it edits the specific stale paragraph rather than rewriting the page.

BAND S2 · API, memory, search · a pull request into the public repository, extensions folder

SharePoint and Teams document approval app

WHO IT SERVES

Enterprise

WHAT TO BUILD

Documents living in SharePoint get an AI edit plus a Teams approval card before the version is committed.

WHAT STRONG LOOKS LIKE

version history in SharePoint shows a clean lineage, and rejection leaves no trace in the library.

BAND S3 · API, Review, export · a pull request into the public repository, extensions folder

Dropbox Paper or Box Notes finaliser

WHO IT SERVES

Collaboration

WHAT TO BUILD

Turns a collaborative draft into the finished, formatted deliverable.

WHAT STRONG LOOKS LIKE

the informal structure of the draft is correctly interpreted into real headings and sections.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

Local-first editor plugin (AFFiNE, Anytype or similar)

WHO IT SERVES

Privacy-first users

WHAT TO BUILD

Adds document AI to a local-first editor, with an explicit data-boundary explainer.

WHAT STRONG LOOKS LIKE

nothing leaves the device except the specific section being edited, and that is demonstrated with a network log.

BAND S2 · API, chat · a pull request into the public repository, extensions folder

VS Code ADR keeper

WHO IT SERVES

Engineering

WHAT TO BUILD

Watches the repository and keeps architecture decision records current as code changes, proposing edits as a diff in the editor.

WHAT STRONG LOOKS LIKE

it proposes an ADR only when a real architectural change happened, and never rewrites the human's rationale.

BAND S2 · MCP, API, Review · a pull request into the public repository, extensions folder

VS Code docs-scoped chat sidebar

WHO IT SERVES

Engineering

WHAT TO BUILD

A sidebar chat whose context is the repository's docs folder, editing files in place with a diff view.

WHAT STRONG LOOKS LIKE

edits land as reviewable diffs in the working tree, never as a silent write.

BAND S2 · MCP, search, Review · a pull request into the public repository, extensions folder

Neovim plugin

WHO IT SERVES

Engineering

WHAT TO BUILD

Async Lua plugin exposing edit, review and export over the API with a diff-in-buffer review flow.

WHAT STRONG LOOKS LIKE

never blocks the editor, and the diff review uses native buffers rather than a floating hack.

BAND S2 · API, Review, export · a pull request into the public repository, extensions folder

Emacs package

WHO IT SERVES

Engineering

WHAT TO BUILD

An Emacs package with an org-mode aware editing command.

WHAT STRONG LOOKS LIKE

org structure survives, and the package is MELPA-ready.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Zed extension

WHO IT SERVES

Engineering

WHAT TO BUILD

An extension for the newer editor, with the MCP surface wired as a native tool.

WHAT STRONG LOOKS LIKE

uses Zed's extension API properly rather than shelling out.

BAND S1 · MCP · a pull request into the public repository, extensions folder

Sublime Text plugin

WHO IT SERVES

Engineering

WHAT TO BUILD

Package Control-ready plugin for edit and export commands.

WHAT STRONG LOOKS LIKE

no UI freeze, and a sensible keymap that does not fight defaults.

BAND S1 · API, export · a pull request into the public repository, extensions folder

Jupyter notebook to paper extension

WHO IT SERVES

Research

WHAT TO BUILD

Turns an analysis notebook into a written paper section with figures, captions and correctly typeset math.

WHAT STRONG LOOKS LIKE

figures carry real captions with numbering, and re-running the notebook updates only the affected prose.

BAND S2 · API, images, export · a pull request into the public repository, extensions folder

RStudio and Quarto addin

WHO IT SERVES

Statistics and research

WHAT TO BUILD

Renders a Quarto or R Markdown analysis into a formatted manuscript and edits the narrative sections by instruction.

WHAT STRONG LOOKS LIKE

code chunks and their outputs stay intact while only prose is edited.

BAND S2 · API, images, export · a pull request into the public repository, extensions folder

Overleaf and LaTeX companion

WHO IT SERVES

Academic

WHAT TO BUILD

A browser extension plus a small service that edits a LaTeX project section by section and keeps math and citations intact.

WHAT STRONG LOOKS LIKE

it edits one section without recompiling meaning elsewhere, and equations survive verbatim.

BAND S3 · API, chat, images, export · a pull request into the public repository, extensions folder

Observable or Marimo report export

WHO IT SERVES

Data

WHAT TO BUILD

Turns a reactive notebook into a written report with live figures baked in.

WHAT STRONG LOOKS LIKE

figures render at print quality rather than as screenshots.

BAND S2 · API, images, export · a pull request into the public repository, extensions folder

Coding-agent tool integration

WHO IT SERVES

Agents

WHAT TO BUILD

Wires SuperDocs into a coding agent's tool surface so the agent that writes the code also maintains its documentation in the same run.

WHAT STRONG LOOKS LIKE

the agent chooses to update docs at the right moments rather than on every commit.

BAND S1 · MCP, export · a pull request into the public repository, extensions folder

Cursor workspace pack

WHO IT SERVES

Agents

WHAT TO BUILD

Rules file, MCP configuration and a scaffolded docs workspace so a Cursor user is productive against SuperDocs in one minute.

WHAT STRONG LOOKS LIKE

a new user with no context gets a correct first edit without reading documentation.

BAND S1 · MCP · a pull request into the public repository, extensions folder

Devcontainer and Codespaces template

WHO IT SERVES

Developer onboarding

WHAT TO BUILD

A one-click cloud dev environment with the MCP server preconfigured and a sample document set.

WHAT STRONG LOOKS LIKE

opens in a browser and the first tool call succeeds with zero local install.

BAND S1 · MCP, API · a pull request into the public repository, extensions folder

Terminal HITL review TUI

WHO IT SERVES

Developer

WHAT TO BUILD

Renders proposed changes as a coloured diff in the terminal with keyboard accept, reject and edit, then commits the batch.

WHAT STRONG LOOKS LIKE

200 proposed changes are reviewable without the terminal becoming unusable, and partial approval works.

BAND S2 · Review, API · a pull request into the public repository, extensions folder

Pre-commit hook for documentation

WHO IT SERVES

Engineering

WHAT TO BUILD

Blocks a commit when the docs contradict the change, and offers to fix them.

WHAT STRONG LOOKS LIKE

low false-positive rate. A hook that cries wolf gets uninstalled, and the submission should show the tuning that avoided it.

BAND S2 · MCP, search, Review · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

GitHub Action that opens a documentation pull request

WHO IT SERVES

Engineering

WHAT TO BUILD

On merge, the action drafts the documentation update and opens a PR with the diff for human review.

WHAT STRONG LOOKS LIKE

the PR is small and scoped, and it is skipped entirely when nothing needs saying.

BAND S2 · MCP, API, Review · a pull request into the public repository, extensions folder

Documentation build pipeline recipes

WHO IT SERVES

Engineering

WHAT TO BUILD

A Makefile and Just recipe pack that turns a repository's markdown into a versioned, styled release deliverable.

WHAT STRONG LOOKS LIKE

reproducible output, byte-stable across two runs on the same input.

BAND S2 · API, images, export · a pull request into the public repository, extensions folder

gh CLI extension

WHO IT SERVES

Engineering

WHAT TO BUILD

gh superdocs commands that operate on a repository's documents from inside the GitHub CLI.

WHAT STRONG LOOKS LIKE

respects gh auth and conventions rather than inventing its own.

BAND S1 · MCP, API · a pull request into the public repository, extensions folder

Document diff tool

WHO IT SERVES

Legal and engineering

WHAT TO BUILD

docdiff a.docx b.docx producing a real redline document plus a summary of substantive versus cosmetic changes.

WHAT STRONG LOOKS LIKE

it separates meaning changes from formatting noise, which naive diffs cannot do.

BAND S2 · Multi-document, chat, export · a pull request into the public repository, extensions folder

Dart and Flutter package

WHAT TO BUILD

One client for mobile, web and desktop, with a review widget.

WHAT STRONG LOOKS LIKE

the same code path works on all three targets, proven by a running build for each.

BAND S2 · API, Review · a pull request into the public repository, extensions folder

R package

WHAT TO BUILD

For researchers: ingest, edit, export and a knitr integration.

WHAT STRONG LOOKS LIKE

passes R CMD check cleanly with a vignette.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Julia package

WHAT TO BUILD

Scientific computing client with a Pluto notebook example.

WHAT STRONG LOOKS LIKE

registered package, typed responses, math handled correctly.

BAND S2 · API, images, export · a pull request into the public repository, extensions folder

Deno and Bun native client

WHAT TO BUILD

Zero-dependency, standard-library-only client for modern runtimes.

WHAT STRONG LOOKS LIKE

no npm compatibility shim, and it runs under both runtimes unmodified.

BAND S1 · API · a pull request into the public repository, extensions folder

Generated clients plus a conformance suite

WHAT TO BUILD

Clients generated from the public OpenAPI spec in several languages, plus a test suite that proves each behaves identically.

WHAT STRONG LOOKS LIKE

it finds and reports at least one real spec ambiguity.

BAND S3 · API · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Salesforce Apex client

WHAT TO BUILD

Named credentials, governor-limit-safe batching, and a Lightning component.

WHAT STRONG LOOKS LIKE

stays inside Apex callout and heap limits on a large document.

BAND S3 · API, templates, export · a pull request into the public repository, extensions folder

Browser-only client with signed uploads

WHAT TO BUILD

A client-side library that uploads large files through pre-signed URLs without a backend of its own.

WHAT STRONG LOOKS LIKE

a 60 MB file uploads from a static site with no server, and the key is never exposed in the bundle.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Continuous-integration document-jobs action pack - CI system: GitHub Actions

WHO IT SERVES

DevTools. Release engineers and technical writers who want documents produced by the pipeline, not by hand.

WHAT TO BUILD

One reusable job, published for three CI systems: a GitHub Action, a GitLab CI template, and a Bitbucket pipe. It takes a document source or session, an instruction, and an output format, then attaches the result as a build artefact and optionally comments the section-level summary back on the pull request. Handles long runs through the asynchronous job flow with a resumable id so a runner timeout does not lose the work. Secrets are read from the CI secret store only. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

an action that is actually reusable (versioned tag, documented inputs and outputs, no hardcoded paths); artefact upload verified byte-for-byte; failure output that says which step failed and what to do; caching so a rerun of an unchanged build spends nothing; published to the GitHub Marketplace.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Continuous-integration document-jobs action pack - CI system: GitLab CI

WHO IT SERVES

DevTools. Release engineers and technical writers who want documents produced by the pipeline, not by hand.

WHAT TO BUILD

One reusable job, published for three CI systems: a GitHub Action, a GitLab CI template, and a Bitbucket pipe. It takes a document source or session, an instruction, and an output format, then attaches the result as a build artefact and optionally comments the section-level summary back on the pull request. Handles long runs through the asynchronous job flow with a resumable id so a runner timeout does not lose the work. Secrets are read from the CI secret store only. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

an action that is actually reusable (versioned tag, documented inputs and outputs, no hardcoded paths); artefact upload verified byte-for-byte; failure output that says which step failed and what to do; caching so a rerun of an unchanged build spends nothing; published to the GitHub Marketplace.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Continuous-integration document-jobs action pack - CI system: Bitbucket Pipelines

WHO IT SERVES

DevTools. Release engineers and technical writers who want documents produced by the pipeline, not by hand.

WHAT TO BUILD

One reusable job, published for three CI systems: a GitHub Action, a GitLab CI template, and a Bitbucket pipe. It takes a document source or session, an instruction, and an output format, then attaches the result as a build artefact and optionally comments the section-level summary back on the pull request. Handles long runs through the asynchronous job flow with a resumable id so a runner timeout does not lose the work. Secrets are read from the CI secret store only. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

an action that is actually reusable (versioned tag, documented inputs and outputs, no hardcoded paths); artefact upload verified byte-for-byte; failure output that says which step failed and what to do; caching so a rerun of an unchanged build spends nothing; published to the GitHub Marketplace.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Document language server

WHO IT SERVES

Teams keeping long structured documents in a repository

WHAT TO BUILD

A server speaking the editor language-server protocol so that any editor that already speaks it gains document intelligence with no per-editor work: go to the definition of a defined term, find every reference to a section, rename a defined term across a document set, diagnostics for cross-references that no longer resolve, and code actions that apply a section-scoped edit in place.

WHAT STRONG LOOKS LIKE

it runs unmodified in at least three editors that speak the protocol, demonstrated; rename is complete and reversible, including inside notes and tables; diagnostics are tuned for a low false-positive rate with the tuning method written up; a very long document stays responsive because the index is incremental, not rebuilt per keystroke.

BAND S3 · API, search, chat · a pull request into the public repository, extensions folder

bisect, find the turn that broke the document

WHO IT SERVES

Developers and support engineers with a wrong document and a long history

WHAT TO BUILD

Given a document whose current state is wrong, a session history, and a predicate that says whether a given state is good, walk the history the way a version-control bisect walks commits: restore candidate states into a scratch copy, halve the search space, and name the turn that introduced the defect along with the sections it touched and the instruction that caused it.

WHAT STRONG LOOKS LIKE

a planted regression in a sixty-turn history is found in a logarithmic number of restores, not by walking every turn; the live document is provably never mutated; the predicate can be a shell command so any check can drive it; the report explains the guilty turn in terms of sections and instruction, not internal identifiers.

BAND S2 · API, search · a pull request into the public repository, extensions folder

Structural contract linter for generated documents

WHO IT SERVES

Teams generating documents in a pipeline who need a shippable/not-shippable call

WHAT TO BUILD

A declared contract for a document - required sections present and in order, tables carrying the declared columns, no placeholder text surviving, numeric totals internally consistent, page count within a stated bound, required parts (headers, footnotes) present - checked as a gate in a build. The contract is a file that lives beside the template, and a failure names the clause of the contract that broke.

WHAT STRONG LOOKS LIKE

the contract file is readable by the non-engineer who owns the template; every failure names the contract clause and the location, so nobody has to open the document to know what happened; the placeholder check catches a leftover token that reads like real prose; the gate is fast enough to run on every generated document, not on a sample.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

Offline-capable review and approval client for a phone

WHO IT SERVES

The approver who is the bottleneck and is never at a desk: partners, clinicians, site managers, executives.

WHAT TO BUILD

A mobile client for the approval half of the workflow only. Pending change batches sync down with enough context to decide (the before, the after, the reason), decisions are made and queued offline, and they apply in order when the connection returns - including the case where the underlying document moved on while the reviewer was in a tunnel, which must be surfaced rather than resolved by force.

WHAT STRONG LOOKS LIKE

a decision made offline never applies to a document state the reviewer did not see; a batch that became stale while queued is returned to the reviewer with what changed, not silently applied or silently dropped; the diff is legible on a small screen including tables; the queue survives the application being killed.

BAND S2 · API, Review, Multi-document · a pull request into the public repository, extensions folder

Numbering-scheme migration across a document estate

WHO IT SERVES

Standards bodies, legal departments and regulated manufacturers who inherited two numbering conventions and must end with one.

WHAT TO BUILD

Move an entire library from one clause or section numbering convention to another in a single reviewable migration: renumber consistently, repair every internal cross-reference, rewrite every reference between documents in the set, and publish an old-to-new mapping table so downstream systems, spreadsheets and external citations can be updated rather than broken.

WHAT STRONG LOOKS LIKE

the mapping table is complete and machine-readable, and is the deliverable a reviewer checks first; no reference in the set is left pointing at a number that no longer exists, proven by a post-migration audit; documents outside the migration that cite the old numbers are listed rather than ignored; the migration is reversible from the mapping table alone.

BAND S2 · API, Multi-document, search, export · a pull request into the public repository, extensions folder

Reviewer console for a document that is being revised while it is read

WHO IT SERVES

Editorial and review teams working on a document in flight

WHAT TO BUILD

A reading surface for a document that another party is actively revising, where the reader is never shown a torn or stale state. It subscribes to the document's change stream, applies committed changes in order, and guarantees that what the reader sees is always some consistent committed version rather than a half-applied mixture. The judged property is behaviour under adversarial conditions: dropped connections, out-of-order delivery, duplicate events and reconnect after a long gap.

WHAT STRONG LOOKS LIKE

Under a scripted stream containing dropped, duplicated and out-of-order events, the reader never renders a state that never existed; a reconnect after an arbitrary gap converges to the correct current version without a full reload; every rejected or replayed event is logged with the reason; and the failure modes are demonstrated deliberately rather than described.

BAND S2 · Multi-document, chat, Review · a pull request into the public repository, use-cases folder

Finish and publish the abandoned ProseMirror integration

WHO IT SERVES

DevTools. An engineer who has shipped and maintained an open-source package, not only written one.

WHAT TO BUILD

There is a partly built ProseMirror integration sitting unfinished in a public demo repository. Take it to publishable quality: complete the editing loop against the public API, add the tests it lacks, write the README with a working quickstart, package it under a real name and publish it to the language's package registry. Then write a short listing note covering what the package does and does not do.

WHAT STRONG LOOKS LIKE

a package a stranger can install and get an AI edit working inside ten minutes, with the failure modes documented honestly rather than hidden.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Lexical plugin with a working review layer

WHO IT SERVES

DevTools. Comfortable inside an editor framework's own node and state model.

WHAT TO BUILD

Build an AI editing plugin for Lexical, the one major rich-text framework whose ecosystem has no serious AI story. Load a document, send a selection or the whole document to the public API, and bring the result back into the editor as a reviewable change rather than a silent overwrite. The interesting part is the review layer: rendering a proposed change inside that framework's node model, and letting a user accept or reject it per change.

WHAT STRONG LOOKS LIKE

the diff renders as native nodes, survives an undo, and does not corrupt document state when a change is rejected.

BAND S3 · API, chat, Review · a pull request into the public repository, extensions folder

TipTap extension package, built on the ProseMirror core

WHO IT SERVES

DevTools. Someone who prefers the smallest change that works.

WHAT TO BUILD

TipTap is a wrapper over ProseMirror, which the first card targets, so a TipTap extension is a thin shell that mounts the same underlying plugin. Build and publish that shell. Prove the claim in the README: show the same AI editing working in TipTap and in at least one downstream editor that itself wraps TipTap, such as BlockNote or Novel, from the same core code. one small package reaches several editor communities at once.

WHAT STRONG LOOKS LIKE

the shell is genuinely thin, the shared core is not copy-pasted, and the downstream editors work without extra per-editor code.

BAND S2 · API, chat, Review · a pull request into the public repository, extensions folder

Quill module for drop-in AI editing

WHO IT SERVES

DevTools. Willing to fight a data model that does not want to cooperate.

WHAT TO BUILD

Build a drop-in AI editing module for Quill, whose document model is a flat stream of operations rather than a block tree. The constraint is the point: our editing model addresses stable blocks, and this editor does not have them natively. Design and document a mapping between the two, then implement editing on top of it, and be explicit in the README about what round-trips cleanly and what does not.

WHAT STRONG LOOKS LIKE

the mapping is written down and defended, and the known-lossy cases are listed rather than quietly broken.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

CKEditor 5 plugin using native suggestions

WHO IT SERVES

DevTools. Comfortable with a plugin architecture that has strong opinions.

WHAT TO BUILD

Build an AI editing plugin for CKEditor 5, which ships its own plugin system and its own change-tracking feature. Route edits through the public API and surface them using the host editor's native suggestion or track-changes mechanism rather than inventing a parallel review UI. Publish it with a quickstart.

WHAT STRONG LOOKS LIKE

proposed edits appear as native suggestions a reviewer already knows how to accept, not as a bolted-on panel.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

Slate plugin that is schema-agnostic

WHO IT SERVES

DevTools. React engineer who has built custom editor behaviour before.

WHAT TO BUILD

Build AI editing for Slate, the React-first framework whose documents are plain JSON trees the host application owns. Because the host owns the document shape, the plugin has to be configurable about where text lives rather than assuming a schema. Ship it with two example schemas so the configurability is demonstrated rather than claimed.

WHAT STRONG LOOKS LIKE

the same plugin works against two genuinely different document shapes with configuration only, no forking.

BAND S2 · API, chat, Review · a pull request into the public repository, extensions folder

One adapter package covering every editor integration

WHO IT SERVES

DevTools. Library design taste; has thought about public API surfaces.

WHAT TO BUILD

Design and publish a single package that holds the shared editing core once and exposes a thin adapter per supported editor. Define the adapter interface, implement at least three adapters, and document how someone adds a fourth without touching the core. Include a conformance test suite an adapter must pass.

WHAT STRONG LOOKS LIKE

the conformance suite is real, a fourth adapter is genuinely addable from the README alone, and the core has no editor-specific branches.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Integration starter kit, downloadable and runnable

WHO IT SERVES

DevTools. Cares about the first five minutes of someone else's experience.

WHAT TO BUILD

Produce a downloadable starter that a developer can unzip, add a key to, and run to see an AI edit inside an editor. One per supported editor, generated from a shared template rather than hand-maintained. Include the smallest possible dependency set and a single command to start.

WHAT STRONG LOOKS LIKE

it runs on a clean machine with no undocumented step, and the generator means adding an editor does not mean maintaining another starter by hand.

BAND S1 · API, chat, export · a pull request into the public repository, extensions folder

Monaco and Draft.js adapters

WHO IT SERVES

DevTools. Comfortable in two unrelated APIs in one build.

WHAT TO BUILD

Two smaller integration targets in one card: Monaco, the embeddable component used inside many web IDEs, and Draft.js, which a large installed base has still not migrated off. Build an adapter for each against the shared core, and write a short note on whether the older library is worth continuing to support. completes the editor matrix and gives us a defensible answer on the legacy target.

WHAT STRONG LOOKS LIKE

both adapters work, and the recommendation on Draft.js is argued from its actual install base rather than taste.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

OAuth 2.1 authorization-code flow for the agent server

WHO IT SERVES

DevTools or platform. Has implemented an auth flow, not only consumed one.

WHAT TO BUILD

Our agent server currently authenticates with a bearer token only, which several agent registries treat as insufficient. Build a working authorization-code flow in front of it: the authorization and token endpoints, the metadata document a client discovers, scope handling, refresh, and revocation. Ship a reference client that completes the flow end to end, and a written note on which registry requirements the implementation now satisfies and which it still does not.

WHAT STRONG LOOKS LIKE

a real client completes the round trip without hand-editing tokens, and the requirements note is checked against each registry's current published rules rather than assumed.

BAND S3 · MCP, API · a pull request into the public repository, extensions folder

Review-first editing where the pending diff survives

WHO IT SERVES

Engineer comfortable with state that outlives a request.

WHAT TO BUILD

Build an editing surface where every AI change arrives as a pending, reviewable diff rather than an applied edit, and where that pending state persists: close the tab, come back, and the review is still there and still approvable. Support handing a pending review to a second person. Approvals are per change, not all-or-nothing.

WHAT STRONG LOOKS LIKE

the review survives a reload and a different browser, partial approval works, and a rejected change leaves no trace in the document.

BAND S3 · API, chat, Review, Multi-document · a pull request into the public repository, extensions folder

A runtime guard for lost block identity

WHO IT SERVES

DevTools. Enjoys building the thing that catches the bug.

WHAT TO BUILD

Our editing model depends on each block carrying a stable identifier through a round trip. Integrations break this quietly by stripping attributes, and the failure shows up later as an edit landing in the wrong place. Build a small library that validates a document round trip and reports precisely which blocks lost identity, where, and what the likely cause was. Ship it as a package plus an optional check that can run in a build pipeline. every editor integration in this block will hit exactly this failure, and this is the tool that makes it a five-minute problem.

WHAT STRONG LOOKS LIKE

it names the specific block and the specific transformation that dropped it, rather than reporting that something went wrong.

BAND S1 · API, export · a pull request into the public repository, extensions folder

Switch between document variants without losing either

WHO IT SERVES

Front-end engineer who can design a small amount of state well.

WHAT TO BUILD

Give a document more than one working variant and a way to move between them: create a variant from the current state, edit it independently, compare it to another, and promote one. The storage side already preserves variants, so this is a build about the interface and the state model rather than about persistence.

WHAT STRONG LOOKS LIKE

switching variants never loses unsaved work, and it is always obvious which variant is being edited.

BAND S2 · API, Multi-document, Review, export · a pull request into the public repository, extensions folder

DOCX round-trip through the open-source SuperDoc editor

WHO IT SERVES

DevTools. Someone who treats format fidelity as a measurable property rather than a vibe.

WHAT TO BUILD

SuperDoc is an open-source, browser-based document editor with a current V2 package and public documentation. Build a working demo that loads a real DOCX file into it, sends the document through our public API or agent server for an AI edit, and lands the edited result back in that editor with formatting intact. Then prove it: run a corpus of documents through the full round trip and publish a fidelity report that measures what survived and what did not, direction by direction. It also puts a working integration in front of that editor's users.

WHAT STRONG LOOKS LIKE

the fidelity report is quantitative and unflattering. It names the constructs that break, with a reproducible case for each, rather than showing one document that happened to work.

BAND S3 · API, MCP, attachments, export, Review · a pull request into the public repository, extensions folder

Miro board to structured document

WHO IT SERVES

Design and workshop teams that end every session with a board full of stickies and no document.

WHAT TO BUILD

Workshop output dies on the board: sticky notes, frames and votes never become the brief or decision document anyone can circulate. Build a Miro app that reads a selected frame or the whole board through the Miro REST API, groups stickies by cluster and votes, and produces a structured, styled document (workshop summary, decision record or project brief) with the board's images embedded, then pushes a link back onto the board.

WHAT STRONG LOOKS LIKE

a messy 60-sticky demo board becomes a coherent sectioned document where every section traces to a named cluster or frame; re-running after the board changes updates only the affected sections.

BAND S2 · API, images, templates, export · a pull request into the public repository, extensions folder

Basecamp project to status document

WHO IT SERVES

Agencies and project teams that run delivery in Basecamp and owe clients a written status.

WHAT TO BUILD

Basecamp deliberately keeps things simple, which means no reporting layer: message boards, to-do sets and check-in answers never roll up into the status document a stakeholder outside the tool expects. Build an integration against the Basecamp API that assembles a chosen project's recent activity into a styled weekly or monthly status document with sections for progress, decisions, blockers and next steps, each traceable to the source thread or to-do.

WHAT STRONG LOOKS LIKE

a reviewer can click from any sentence in the status document to the Basecamp item it came from; two consecutive runs produce a document that updates rather than repeats.

BAND S2 · API, templates, export · a pull request into the public repository, extensions folder

Replit workspace document panel

WHO IT SERVES

Builders working entirely in Replit who need finished documents beside their code.

WHAT TO BUILD

Projects built and shared inside Replit's browser IDE ship without real documents: the README, spec and user guide stay unwritten because leaving the browser breaks the flow. Build a Replit extension or template that adds a document panel to the workspace: it reads the project's files, drafts and edits the README, spec or user-facing guide through the document API, and exports a styled PDF or DOCX artifact into the project.

WHAT STRONG LOOKS LIKE

from a fresh fork of the template, a first-time user gets a real generated-and-edited document exported inside the workspace in one sitting; the panel demonstrably reflects code changes in the next document revision.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Xcode source editor companion for release documents

WHO IT SERVES

iOS and Mac teams that ship release notes and App Store documents with every build.

WHAT TO BUILD

Apple-platform developers write release notes, App Store descriptions and technical documents in the one IDE that has no document tooling. Build an Xcode source editor extension plus a small companion app: select a changelog, commit list or code region, send it to the document API, and get back edited release notes or App Store copy as a styled document, with the companion app handling review and export since editor extensions are UI-constrained.

WHAT STRONG LOOKS LIKE

the extension installs and runs under current macOS signing rules; a raw commit list becomes reviewable release notes and an exported document without leaving the Xcode-plus-companion pair; the App Store copy variant respects the store's own length fields.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

AI-terminal integration for document work

WHO IT SERVES

Developers who live in AI terminals and want documents created and edited without leaving them.

WHAT TO BUILD

AI-native terminals like Warp treat commands and their output as structured blocks and expose their own agent surface, which the existing shell and TUI tooling does not touch. Build a first-class integration for one AI-forward terminal: a workflow or agent-mode tool that captures selected command blocks (a deploy log, a test run, a migration transcript) and turns them into a styled runbook, incident note or change record through the document API, reviewable before export.

WHAT STRONG LOOKS LIKE

a captured real command sequence becomes a runbook whose steps match the blocks one-to-one; the integration uses the terminal's own workflow or agent surface, not a wrapper script pretending to be one.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

AGENT-LOOP DEMOS AND LIVING DOCUMENTS

MCP pipeline: data warehouse → board report

WHO IT SERVES

Finance/Ops

WHAT TO BUILD

Finance and ops teams pull the same board-report numbers from a data warehouse every cycle and rewrite the surrounding narrative by hand each time. Build a pipeline that queries a data source for a set of metrics, drafts a narrative board report around those numbers with supporting charts, and hands the result off through a wrapper-side scheduler for recurring delivery.

WHAT STRONG LOOKS LIKE

Reviewer supplies one fictional metrics payload and gets back one styled board report where every number in the narrative matches the input exactly, plus one chart image built from the same numbers.

BAND S3 · API, images, export · a pull request into the public repository, use-cases folder

Legal-hold notice and custodian acknowledgment kit

WHO IT SERVES

Legal / legal-ops manager or compliance counsel at a company facing litigation

WHAT TO BUILD

From a matter description and a custodian list, generate the legal-hold notice on company or firm letterhead, personalized per custodian with their actual systems and data sources named, plus an acknowledgment form and a scheduled reminder series. A tracking document records who acknowledged and when, and regenerates follow-ups only for those who have not. Release notices come from the same matter record when the hold lifts. Every version is retained so the company can show a court exactly what was sent, to whom, and on what date.

WHAT STRONG LOOKS LIKE

Personalization goes past the name to the custodian's systems, role and scope; the version trail is defensible under challenge; re-running for newly added custodians never reissues to people who already acknowledged.

BAND S1 · API, templates, memory, export · a pull request into the public repository, use-cases folder

Dual student and instructor edition generator

WHO IT SERVES

Education / lab coordinator, course author or textbook publisher

WHAT TO BUILD

One authored source produces two linked documents: the student edition (procedure, data blanks, questions) and the instructor edition (the same content plus answers, marking guidance, safety notes and timing). Authors write once and mark blocks as instructor-only. Editing either edition keeps shared content in sync and edition-specific content independent. Answer keys regenerate to a different question order for exam security, and both editions export print-ready with consistent numbering.

WHAT STRONG LOOKS LIKE

shared and edition-specific content are modelled so an edit to shared prose updates both editions while an instructor note can never leak into the student edition, with a leak test shipped in the repo; numbering stays aligned after a reorder.

BAND S2 · Multi-document, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Employee-relations letter engine with case file

WHO IT SERVES

HR / HR business partner or employee-relations specialist handling conduct and performance

WHAT TO BUILD

Improvement plans, absence letters, invitation-to-hearing letters, outcome letters and appeal responses follow a process with strict sequencing and content requirements. The app generates each letter from the case record on the approved template, refuses to skip a required step in the sequence, and maintains the case-file document showing the chronology, what was sent and what was said at each meeting. Meeting notes convert into the outcome letter's summary section with the employee's own response included verbatim.

WHAT STRONG LOOKS LIKE

process sequencing is enforced with a readable explanation whenever a step is refused; the employee's words are always verbatim and never paraphrased; the case file is complete enough to be disclosed in a tribunal without embarrassment.

BAND S2 · chat, templates, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Transfer pricing documentation assembler

WHO IT SERVES

Accounting / international tax manager at a multinational or an advisory firm

WHAT TO BUILD

Documentation is a master file plus a local file per country, each with functional analysis, industry analysis, financial data, benchmarking extracts and intercompany agreement summaries. The app maintains one master narrative and generates each local file by inserting jurisdiction-required sections and local financial data while keeping the group narrative consistent everywhere. A change to the group functional analysis proposes the same change in every local file. Local-language versions generate where the jurisdiction demands them.

WHAT STRONG LOOKS LIKE

group-narrative consistency across twenty local files is provable with a test; jurisdiction-specific required sections are correct per country; a master change propagates without disturbing legitimate local additions.

BAND S3 · Multi-document, templates, search, Review, export · a pull request into the public repository, use-cases folder

Expert-interview to insight memo pipeline

WHO IT SERVES

Consulting / research or strategy consultant running an expert-interview programme

WHAT TO BUILD

Each expert call produces a transcript. The app generates the interview memo (background, key points, quotes, implications) on the firm template with identifying details handled per the compliance rules, then maintains a rolling synthesis document across all interviews: recurring themes with the interviews supporting them, disagreements between experts stated as disagreements, and open questions to carry into the next calls. Adding interview eleven updates the synthesis surgically rather than regenerating the whole document.

WHAT STRONG LOOKS LIKE

disagreements are preserved rather than averaged into a consensus nobody actually stated; the post-interview synthesis update is a targeted edit with a visible diff; compliance handling of expert identity is consistent and configurable.

BAND S2 · chat, Multi-document, search, memory, Review, export · a pull request into the public repository, use-cases folder

Analyst relations briefing and RFI response pack

WHO IT SERVES

Marketing / analyst relations or product marketing lead at a B2B company

WHAT TO BUILD

Analyst evaluations demand a long structured response: capability by capability with evidence, plus a briefing document and a demo script mapped to the evaluation criteria. The app maintains an answer library of previously approved responses, drafts each criterion's answer from it, flags criteria where the library has nothing approved, enforces the per-question word limits, and assembles the briefing narrative and one-page summary. Cycle over cycle it shows what changed since last time so the team can explain a movement in position.

WHAT STRONG LOOKS LIKE

word limits are respected exactly, because analyst forms enforce them mechanically; unanswered criteria are surfaced rather than padded with adjacent content; the year-over-year change view is derived from the documents rather than remembered.

BAND S2 · search, templates, Multi-document, chat, Review, export · a pull request into the public repository, use-cases folder

Risk register and treatment plan maintenance

WHO IT SERVES

Compliance / risk manager or GRC lead

WHAT TO BUILD

The risk register is a living document: risks with cause and consequence, inherent and residual ratings, controls, treatment actions, owners and review dates, plus a heat-map diagram. The app maintains it, drafts new entries from incident reports or audit findings in the register's existing wording style, recalculates and redraws the heat map when ratings change, and generates the board risk report with the movement since last quarter derived rather than retyped. Overdue treatments produce a chase document per owner.

WHAT STRONG LOOKS LIKE

new entries match the register's established wording conventions instead of reading as foreign text; the heat map regenerates correctly after every rating change; the board report's movement narrative is derived from register history, not written from memory.

BAND S2 · chat, images, Multi-document, templates, memory, export · a pull request into the public repository, use-cases folder

ISO 9001 clause-to-document traceability maintainer

WHO IT SERVES

Quality manager preparing for a certification or surveillance audit

WHAT TO BUILD

The auditor asks which document satisfies each clause of the standard. Build an app that holds the quality manual and procedure set, maintains a traceability matrix mapping every clause to the document sections that address it, and re-checks the mapping whenever a procedure is edited. Clauses that lost their evidence because a section was rewritten or deleted are raised as gaps with the specific edit that caused the drift. The matrix and a short audit-readiness summary export as one branded pack.

WHAT STRONG LOOKS LIKE

deleting a covering section immediately shows as a clause gap with the causing edit named, and the matrix survives a document rename without breaking its links.

BAND S3 · search, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Supplier corrective-action request and closure record

WHO IT SERVES

Supplier quality engineer managing nonconforming supplier performance

WHAT TO BUILD

Build a workflow that issues the corrective-action request from the nonconformance evidence (photos, inspection data, the affected purchase order), receives the supplier's response, drafts the assessment of whether the root cause and containment are adequate, and maintains the closure record with the verification evidence attached. Requests that go past due generate escalation letters at the contractually specified intervals, referencing the clause that permits escalation. The full history stays in one document rather than a chain of emails.

WHAT STRONG LOOKS LIKE

the adequacy assessment cites specific weaknesses in the supplier's stated root cause instead of rubber-stamping it, and the escalation intervals come from the contract rather than a hardcoded default.

BAND S2 · chat, images, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Series bible and continuity checker

WHO IT SERVES

Novelist or series showrunner keeping a long-running world consistent

WHAT TO BUILD

Build a tool that reads the manuscripts of a series and maintains a living bible: characters with physical details and relationships, places, timeline, and established rules of the world, each entry citing the chapter it came from. When a new chapter is written, the tool checks it against the bible and reports contradictions with both passages quoted (an eye colour that changed, a character in two places, a season that does not follow). Resolving a contradiction updates either the chapter or the bible entry, never silently both.

WHAT STRONG LOOKS LIKE

contradictions are found across hundreds of pages with both sources quoted, the bible entries carry real chapter citations, and resolution is an explicit choice rather than an automatic overwrite.

BAND S3 · Multi-document, search, memory, chat, export · a pull request into the public repository, use-cases folder

Fleet Inspection and Compliance File

WHO IT SERVES

Automotive / fleet · compliance manager keeping a commercial fleet audit-ready

WHAT TO BUILD

Per vehicle, assemble and maintain the compliance file: driver inspection reports, scheduled maintenance records, defect rectification evidence with photographs, roadworthiness certificates and their expiry dates. The app produces a fleet-level readiness report showing which vehicles have gaps and which certificates expire inside the next window, then generates the per-vehicle file as a single indexed document an auditor can read end to end.

WHAT STRONG LOOKS LIKE

expiry detection reads dates out of the certificate documents themselves; the per-vehicle file is genuinely audit-ready with a working index; gaps are listed rather than papered over.

BAND S2 · API, Multi-document, images, search, export · a pull request into the public repository, use-cases folder

Warranty Policy Manual and Bulletin Push

WHO IT SERVES

Automotive · warranty policy owner maintaining the manual the whole dealer network works from

WHAT TO BUILD

Maintain the dealer warranty policy and procedure manual as a living document. When a policy changes, the app edits the affected sections in place, generates the change bulletin describing what changed and from when, produces the dealer-facing summary and acknowledgment form, and updates the manual's revision history and effective date. Dealers in different regions get the variant that applies to them, generated from the same master with regional differences preserved.

WHAT STRONG LOOKS LIKE

regional variants stay correct through a global change; the bulletin describes exactly what changed with no drift from the manual edit; revision history is generated rather than hand-maintained.

BAND S2 · Multi-document, chat, templates, Review, export · a pull request into the public repository, use-cases folder

Art Direction Bible with Generated Reference

WHO IT SERVES

Gaming / creative · art director setting the visual rules for an outsourced art pipeline

WHAT TO BUILD

Build the art direction document that external studios actually work from: palette with swatches, silhouette and proportion rules, material and lighting language, do and do-not example pairs, and per-asset-class specification tables. Reference images are generated inside the document from described intent and iterated by chat ("make the do-not example more obviously wrong"), so the bible ships complete without a separate concept pass. Sections carry the asset classes they govern so an outsourcer can be sent only their slice.

WHAT STRONG LOOKS LIKE

generated reference images actually illustrate the stated rule; iterating one example leaves the rest of the document alone; the per-outsourcer slice exports cleanly.

BAND S2 · chat, images, Multi-document, templates, export · a pull request into the public repository, use-cases folder

Google Drive folder watcher

WHO IT SERVES

Operations or delivery team with a shared Drive that everyone dumps into. Host: Google Drive plus a small console.

WHAT TO BUILD

Watch a Drive folder. Every new file gets a configured treatment: normalise to the house template, extract a summary into a companion document, generate the standard response document, or flag it for a human. Results are written back to Drive beside the source with a naming convention and a manifest sheet recording what was done, when, and by which rule. Rules are per folder, so different teams get different behaviour.

WHAT STRONG LOOKS LIKE

no file is processed twice and none is missed, including files that arrive during a restart. The manifest reconciles exactly with the folder. Nothing overwrites the source file.

BAND S2 · API, templates, export, search, Multi-document · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

SharePoint controlled-document set

WHO IT SERVES

Quality or policy manager running a controlled document library under ISO, SOC 2, or a regulator. Host app: SharePoint Online.

WHAT TO BUILD

Policies, SOPs, and work instructions with a real lifecycle: draft, review, approve, publish, scheduled re-review. Each document carries owner, effective date, and review date. The app drafts revisions by instruction, routes them through the approval chain, publishes the approved version with an auto-generated change summary and revision history table, and chases overdue re-reviews. Superseded versions are retained and marked, never quietly replaced.

WHAT STRONG LOOKS LIKE

the revision history table in the document matches the system's record exactly. Nothing publishes without the named approvals. An auditor can trace any published clause back to who approved it and when.

BAND S3 · templates, Review, Multi-document, export, search · a pull request into the public repository, use-cases folder

Slack Connect cross-company redline exchange

WHO IT SERVES

Deal desk, procurement, or partnerships lead negotiating with a counterparty in a shared Slack Connect channel. Host app: Slack.

WHAT TO BUILD

Negotiation now happens in shared channels, but the document still bounces over email. This app keeps it in the channel: drop the contract, each side's proposed changes appear as cards visible to both, each side approves its own positions, and the app maintains one authoritative current version with a visible history of who changed what. Internal-only commentary stays internal and never posts to the shared channel.

WHAT STRONG LOOKS LIKE

the internal and external boundary is airtight, verified by test. Both sides always see the same current version. The exported file carries the full negotiation trail as track changes for the lawyers.

BAND S3 · Multi-document, Review, chat, export, search · a pull request into the public repository, extensions folder

Dropbox folder watcher with write-back

WHO IT SERVES

Small firm, studio, or freelancer collective running its whole business out of Dropbox. Host: Dropbox.

WHAT TO BUILD

Nominate folders and their treatments. A client drops a file into the shared folder and it comes back normalised to the studio template, or summarised into a companion brief, or turned into the standard response document, filed beside the original with a consistent name. A small web console shows what ran, what it produced, and what needs a human. Shared-folder permissions are respected so a client only ever sees their own outputs.

WHAT STRONG LOOKS LIKE

it survives Dropbox's sync behaviour without processing half-written files or looping on its own output. Client folders never cross. The console's record matches the folder exactly.

BAND S2 · API, templates, export, search · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Local folder watcher tray app

WHO IT SERVES

Small business or professional practice with no cloud storage, no IT, and a shared drive. Host: Windows and macOS desktop.

WHAT TO BUILD

A tray application with an in folder and an out folder. Drop a document into "in", it comes back into "out" processed according to the rule set the owner picked (reformat to our template, produce the client letter, extract the summary). Rules are chosen from a plain-language list, not written. Everything is visible on disk, so a non-technical owner can see it working and can stop it by quitting the app.

WHAT STRONG LOOKS LIKE

it never loses or duplicates a file, including across a laptop sleep or a network-drive hiccup. Errors produce a readable note beside the file, not a silent failure. A non-technical owner can install and configure it alone.

BAND S2 · API, templates, chat, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Batch intake worker with manifests and receipts

WHO IT SERVES

Enterprise operations in insurance, banking, logistics, or payroll that still exchanges files on a schedule. Host: a secure file endpoint plus a scheduled worker.

WHAT TO BUILD

Partners drop a nightly batch. The worker validates the manifest, processes every document under the agreed rule set, writes outputs to the return path, and produces a receipt file per batch: counts, per-document outcomes, failures with reasons, and timings. Re-runs are safe and idempotent. Exceptions route to a human queue rather than blocking the batch. Retention and deletion follow the contract.

WHAT STRONG LOOKS LIKE

the receipt reconciles to the manifest exactly, every night, including on failure nights. A re-run produces no duplicates. Ten thousand documents complete inside the window with a documented throughput number.

BAND S3 · API, templates, export, search, Review · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Data-orchestration operator for scheduled document production

WHO IT SERVES

Data engineer or analytics lead whose pipelines end in a report somebody reads. Host: Apache Airflow, Dagster, or a similar scheduler.

WHAT TO BUILD

An operator and an asset definition so a document is a first-class output of a data pipeline: the pipeline finishes, the operator produces the branded report with the numbers written into the prose and the charts placed as figures, and delivery happens on the same schedule with the same retry and alerting the rest of the pipeline uses. Failed report runs surface in the existing monitoring rather than in somebody's inbox on Monday.

WHAT STRONG LOOKS LIKE

it is idempotent and backfillable, so re-running last Tuesday produces last Tuesday's report. Figures and stated numbers always agree. Failures appear in the pipeline's own alerting.

BAND S2 · API, images, templates, export, chat · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Research-paper writing agent, outline to submission-ready manuscript

WHO IT SERVES

Academic and research tooling. Founder or engineer at a company building an AI research-writing or manuscript product.

WHAT TO BUILD

An agent that writes a real paper end to end and never regenerates the manuscript to change one part. It drafts the outline, gathers sources, writes each section into the document, inserts equations as real mathematics and citations against a bibliography, builds the table of contents, then runs revision passes that edit only the sections under review. A supervisor pass checks claim-to-citation grounding and rewrites just the unsupported sentences. Export is a submission-shaped file with real footnotes, headers, and page geometry. Every turn logs sections touched and cost.

WHAT STRONG LOOKS LIKE

a genuine paper produced end to end and opened in Word without a single broken equation, footnote, or reference; revision passes provably touch only the named sections (show the before and after section identifiers); the citation checker catches a planted unsupported claim; cost per paper reported; a second run on a different topic proving it is not hand-tuned to one example.

BAND S3 · chat, MCP, Review, search, images, export, Multi-document · a pull request into the public repository, use-cases folder

Living document that updates itself from a data source

WHO IT SERVES

Operations and analytics. A team whose weekly report is rebuilt by hand every week.

WHAT TO BUILD

A scheduled agent owns one document forever. On each run it reads a data source, compares it with what the document currently says, and edits only the parts that changed: a number in a table cell, a sentence in the commentary, a chart caption. The document keeps its identity, its formatting, and its history, so the change log is a list of surgical edits rather than a stack of regenerated files. A summary of what changed is appended to a revision note at the end of the document.

WHAT STRONG LOOKS LIKE

an unchanged input produces zero edits and spends nothing; only the affected cells and sentences are touched, proven across at least ten scheduled runs; formatting after fifty runs is identical to run one; a failure to reach the data source leaves the document untouched and says so.

BAND S2 · API, chat, search, Multi-document, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Multi-agent editorial review loop

WHO IT SERVES

Publishing and content operations. Editorial teams and anyone shipping high-volume written work.

WHAT TO BUILD

Three agents share one document: a writer that drafts, a reviewer that critiques against a style guide, and a fact-checker that verifies claims against attached sources. The reviewer and checker do not rewrite; they propose changes into the review queue, and a human referee approves or denies each one. The loop terminates on a quality bar or a round limit. The final document carries reviewer comments as real comments, so the audit trail survives export.

WHAT STRONG LOOKS LIKE

agents never fight over the same section (a documented coordination rule, not luck); denied proposals leave no residue; the loop provably terminates; comments survive the round trip to Word; a transcript showing a bad draft becoming a good one, with the human decisions visible.

BAND S3 · chat, Review, Multi-document, search, templates, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Agent that files and then edits its own reports

WHO IT SERVES

DevTools and agent infrastructure. Teams running agents whose output nobody reads because it is a wall of appended text.

WHAT TO BUILD

An agent keeps a running report document for each long task. Instead of appending a new block per step, it edits the report in place: it updates the status line, revises the findings section as understanding improves, corrects an earlier wrong conclusion rather than contradicting it later, and keeps a compact decision log. The result is a document that reads like a finished report at any moment during the run, not a transcript.

WHAT STRONG LOOKS LIKE

a mid-run snapshot of the document is readable and coherent; corrections replace the wrong statement rather than stacking a contradiction; the document survives hundreds of edits with structure intact; cost per run is bounded and reported; a side-by-side against an append-only version of the same run.

BAND S2 · chat, API, search, export, memory · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Scheduled agent job runner with an idempotent change log

WHO IT SERVES

Operations engineering. Anyone running document maintenance on a timer.

WHAT TO BUILD

A small scheduler service where each job is a document, an instruction, and a cadence. It guarantees one run at a time per document, records every run with what changed and what it cost, and skips work when nothing needs changing. Failures retry with backoff and never leave a half-applied edit. A dashboard document (itself maintained by the runner) lists every job, its last run, and its spend, which is a neat demonstration in its own right.

WHAT STRONG LOOKS LIKE

two runners racing the same document cannot both edit it; a run that changes nothing costs nothing and says so; the change log is durable; the self-maintained dashboard is genuinely produced by the system, not hand-written; clean shutdown mid-job.

BAND S2 · API, chat, export, Multi-document · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Long-horizon book-writing agent with consistency checks

WHO IT SERVES

Authoring and publishing. Writers, and product teams building writing tools.

WHAT TO BUILD

An agent that writes a full-length manuscript over many sessions: it plans the structure, drafts chapter by chapter into one document, and maintains a separate continuity record (characters, timeline, terminology) that it consults and updates. Before each new chapter it checks the record for contradictions and, when it finds one, fixes the earlier passage surgically rather than rewriting the book. Front and back matter and a live table of contents are generated at the end.

WHAT STRONG LOOKS LIKE

a real manuscript of meaningful length produced across multiple sessions with state recovered each time; a deliberately planted continuity error found and corrected in place; the export opens cleanly with correct heading hierarchy; cost per chapter reported and stable as the book grows, which is the point.

BAND S3 · chat, Multi-document, memory, search, export · a pull request into the public repository, use-cases folder

Agent that maintains a decision record from repository activity

WHO IT SERVES

DevTools. Engineering organisations that lose their architectural reasoning in pull-request comments.

WHAT TO BUILD

An agent watches a repository's merged changes and discussions, decides which ones represent a durable decision, and maintains a single decision-record document: adding a new record, marking an old one superseded by editing its status line, and updating the index. It never regenerates the document. Each record cites the change that produced it. A weekly digest section summarises what moved.

WHAT STRONG LOOKS LIKE

noise filtered sensibly (a documented rule for what counts as a decision); superseding edits the existing record rather than duplicating it; citations resolve; the document stays navigable after a hundred records; a dry-run mode for the first week. Distinct from the existing chat-thread-to-record row: the trigger is repository activity and the record is maintained continuously.

BAND S2 · MCP, chat, search, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Incident agent: live incident document to postmortem

WHO IT SERVES

DevTools and site reliability. On-call teams.

WHAT TO BUILD

When an incident opens, the agent creates a live incident document with the standard sections and then keeps it current as the incident evolves: timeline entries appended, impact statement revised as understanding changes, current status rewritten in place, workstreams tracked. When the incident closes it converts the same document into a postmortem by restructuring the sections and drafting the analysis, keeping the timeline intact. Exports a branded file for distribution.

WHAT STRONG LOOKS LIKE

the document is genuinely usable during a live incident (fast, never locked, never regenerated under the reader); status revisions replace rather than accumulate; the postmortem conversion preserves every timeline entry verbatim; a rehearsal run against a simulated incident with a written record of the result.

BAND S3 · chat, API, search, images, export · a pull request into the public repository, use-cases folder

Agent that keeps reference documentation in sync with a specification

WHO IT SERVES

DevTools. Teams whose published documentation drifts from their API within a week of every release.

WHAT TO BUILD

The agent diffs the current interface specification against the previous one, classifies each change, and applies targeted edits to the human-written documentation: updating the affected parameter description, adding a new section for a new capability, marking something deprecated, and leaving every untouched page alone. Changes go through review so a writer approves the wording. It reports what it could not decide instead of guessing.

WHAT STRONG LOOKS LIKE

untouched pages are byte-identical afterwards; a renamed field updates every place it appears, proven by search; the uncertain cases are escalated rather than silently written; the run is reproducible from two specification snapshots committed in the repository.

BAND S3 · API, chat, Review, search, Multi-document, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Compliance-checklist agent loop with a verifier

WHO IT SERVES

Proposals and compliance. Teams answering requirements documents under a deadline.

WHAT TO BUILD

Given a requirements document and a draft response, the agent builds a checklist of every requirement, then loops: for each unmet item it edits the relevant response section, and a separate verifier pass re-reads the edited section and decides whether the requirement is now met. Unresolved items stay on the checklist with a reason. The checklist lives in its own document alongside the response, updated in place each round.

WHAT STRONG LOOKS LIKE

the verifier is genuinely independent (it can and does reject the writer's work, with examples in the repository); the loop terminates; a planted impossible requirement is reported honestly rather than faked; the final checklist reconciles exactly with the response document.

BAND S3 · chat, Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Literature-review agent with grounded drafting and a verification pass

WHO IT SERVES

Research and analysis. Research teams, and product teams building research assistants.

WHAT TO BUILD

The agent gathers a source set, indexes it, and drafts a structured review where every claim is written together with the source that supports it. A second pass then re-checks each claim against the indexed sources and edits only the sentences it cannot support, either softening them or removing them, and records what it removed in an appendix. The result is a document plus an honest provenance appendix, exported with real footnotes.

WHAT STRONG LOOKS LIKE

the verification pass demonstrably changes the document (a planted unsupported claim is caught); footnotes survive export as real footnotes; the provenance appendix is complete and machine-checkable; sources with conflicting findings are represented rather than averaged away.

BAND S3 · chat, search, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Parallel section drafting with a merge referee

WHO IT SERVES

Agent infrastructure. Teams building multi-agent systems that produce one artefact.

WHAT TO BUILD

Several agents draft assigned sections of one long document at the same time. A referee agent owns integration: it resolves overlapping content, enforces one voice, fixes cross-references between sections written independently, and produces a merge report. The demonstration measures wall-clock time and cost against a single agent writing the same document sequentially, and reports both honestly, including the cases where parallelism did not pay.

WHAT STRONG LOOKS LIKE

no section is silently overwritten by a peer; the referee's decisions are logged and reviewable; the comparison is fair (same prompt, same model settings, repeated runs, variance reported); a failure case is documented rather than hidden.

BAND S3 · chat, Multi-document, Review, search, export · a pull request into the public repository, use-cases folder

Editor agent that learns a house style across sessions

WHO IT SERVES

Content and agencies. Teams with a voice that new writers take months to absorb.

WHAT TO BUILD

An agent that edits documents and, with the operator's consent, keeps a durable note of the corrections it received: preferences accepted, phrasings rejected, formatting conventions. Over a series of sessions it needs fewer corrections on the same class of document. The demonstration is a measured curve: correction count per document across twenty documents, with the memory turned off and turned on. A single command wipes the memory, and that is shown too.

WHAT STRONG LOOKS LIKE

the measurement is real and repeatable, with the control arm run properly; memory is opt-in and scoped per end customer so one client's voice never leaks into another's document, proven by a test; the wipe genuinely wipes; the note stays small rather than growing without bound.

BAND S3 · chat, memory, Review, templates, export · a pull request into the public repository, use-cases folder

Self-healing document agent for structure and numbering

WHO IT SERVES

Legal, pharmaceutical, and standards-bound documentation.

WHAT TO BUILD

A scheduled agent that inspects long documents for structural defects (broken heading hierarchy, clause numbering that no longer runs in sequence, cross-references pointing at a renumbered clause, an out-of-date table of contents, orphaned footnotes) and repairs only the defective parts, routing every repair through review. It produces a health report per document and a trend across the library.

WHAT STRONG LOOKS LIKE

it detects defects on a corpus of deliberately damaged documents and repairs them without touching anything else; a clean document produces zero proposed changes; cross-reference repair is verified by following every reference after the fix; the health report is honest about what it could not repair.

BAND S3 · chat, search, Review, Multi-document, export · a pull request into the public repository, use-cases folder

Data-watcher agent that updates only the affected table rows

WHO IT SERVES

Finance and operations. Anyone with a document containing a table that mirrors a data set.

WHAT TO BUILD

The agent watches a data set and, when a row changes, edits exactly that row in the document's table, leaving the table's borders, shading, merged cells, and every other row untouched. Added rows are inserted in the right position; removed rows are deleted; a change note is written beneath the table. The demonstration runs fifty updates over a heavily styled table and shows the table is visually identical at the end except for the intended cells.

WHAT STRONG LOOKS LIKE

cell-level precision proven with a visual comparison before and after; merged cells and alternating shading survive; the round trip to Word keeps the table intact; an unchanged data set produces no operations; row insertion position is correct, not appended.

BAND S2 · API, chat, search, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Two-agent document negotiation with a human referee

WHO IT SERVES

Legal and sales operations. Teams where two sides mark up the same paper.

WHAT TO BUILD

Two agents, each with its own playbook, negotiate one document: one proposes edits, the other counters, and every round lands as reviewable proposed changes rather than silent rewrites. A human referee approves what lands. The transcript, the playbooks, and the final marked-up document are the deliverable, exported with tracked changes so the trail is real. Includes a stopping rule and an escalation path when the agents cannot converge.

WHAT STRONG LOOKS LIKE

every round is auditable and reversible; the export carries genuine tracked changes readable in Word; a non-converging case is handled by escalation rather than an infinite loop; the playbooks are data files a reader can edit and re-run.

BAND S3 · chat, Review, templates, Multi-document, export · a pull request into the public repository, use-cases folder

Budget-aware autonomous agent with self-service onboarding

WHO IT SERVES

Agent infrastructure. Fully autonomous agents with no human operator in the loop.

WHAT TO BUILD

An agent that handles its own lifecycle: it signs itself up, connects, checks its account status before starting work, plans the run against its remaining monthly operations, chooses cheaper strategies as the budget tightens, and when it genuinely needs more capacity it asks for a limit increase through the supported path rather than failing silently. It also knows how to hand the account to a human when one appears. The deliverable is the agent plus a written record of a real cold start.

WHAT STRONG LOOKS LIKE

a genuine cold start from nothing to a finished document with no human step, recorded; budget planning that measurably changes behaviour near the cap; the handover path exercised end to end; no credential ever written to a log. Distinct from the existing self-signup row because the subject here is budget planning and graceful degradation.

BAND S3 · MCP, chat, export, memory · a pull request into the public repository, use-cases folder

Weekly digest that is edited, never regenerated

WHO IT SERVES

Cross-industry. Any team with a recurring written update.

WHAT TO BUILD

One document, fifty-two weeks. Each week the agent revises the standing sections (what changed, what is at risk, what is next), moves last week's content into a compact archive section, and preserves formatting, so the reader always opens the same familiar document rather than a new file. A companion measurement shows the cumulative cost against regenerating the digest weekly, on the same content, with the method stated.

WHAT STRONG LOOKS LIKE

run for enough weeks (real or simulated) that drift would have appeared, with the final document still clean; archive growth handled so the document stays navigable; the cost comparison is honest and reproducible; a reader can diff week over week and see only intended changes.

BAND S1 · chat, API, search, export, memory · a pull request into the public repository, use-cases folder

Document as agent working memory

WHO IT SERVES

Agent infrastructure. Researchers and builders of long-running agents.

WHAT TO BUILD

A reference implementation in which the agent's working state lives in a structured document instead of a context window: a plan section, an evidence section, an open-questions section, and a conclusions section, each edited as the agent learns. The agent reads back only the sections it needs, so its context stays small while its memory stays complete and human-readable. Includes a comparison against an in-context approach on the same task.

WHAT STRONG LOOKS LIKE

context usage measured and clearly lower; the document is legible to a human at any point in the run; the agent recovers correctly after a restart with no in-memory state; the comparison names its limits rather than overclaiming.

BAND S3 · chat, search, Multi-document, export · a pull request into the public repository, use-cases folder

Evidence-binder agent for a control framework

WHO IT SERVES

Compliance and governance. Security and compliance teams preparing for audit.

WHAT TO BUILD

An agent maintains one evidence binder document across a whole audit period: as a control's evidence changes it updates that control's narrative and its evidence index entry in place, dates the change, and leaves everything else alone. It flags controls whose evidence is stale and drafts the request. Human approval is required for every narrative change. The binder exports as a single navigable file with a live table of contents.

WHAT STRONG LOOKS LIKE

control identifiers and cross-references stay consistent over many updates; stale detection is rule-driven and explained; nothing lands without approval; the exported binder is genuinely navigable; a simulated three-month period run in the repository as fixtures. Distinct from the existing evidence-index row because this is a continuous agent loop rather than a one-time builder.

BAND S2 · chat, Review, Multi-document, templates, search, export · a pull request into the public repository, use-cases folder

Living-document agent

WHO IT SERVES

Cross-industry

WHAT TO BUILD

An agent owns a named section of a standing document and keeps it current from a changing source, editing only inside its boundary and leaving human paragraphs alone.

WHAT STRONG LOOKS LIKE

the ownership boundary is explicit and never crossed, and a week of runs produces a readable history rather than churn.

BAND S2 · MCP, memory, search, Review · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Nightly document QA agent

WHO IT SERVES

Compliance and technical writing

WHAT TO BUILD

Runs overnight across a document set looking for contradictions, stale facts and broken references, then files a review queue for the morning.

WHAT STRONG LOOKS LIKE

precision over volume. A clean corpus produces an empty queue, and every finding is reproducible.

BAND S2 · MCP, search, Multi-document · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Cold-start agent onboarding run

WHO IT SERVES

Agents

WHAT TO BUILD

A fully autonomous agent with no prior account signs itself up, connects, and completes a real edit and export, with the whole transcript published.

WHAT STRONG LOOKS LIKE

zero human steps, and the transcript shows where the agent got confused so we can fix the surface.

BAND S2 · MCP, API, export · a pull request into the public repository, use-cases folder

Quota-aware agent with a budget guard

WHO IT SERVES

Agents

WHAT TO BUILD

The agent checks its remaining allowance before planning, sizes the work to fit, and degrades gracefully instead of failing halfway.

WHAT STRONG LOOKS LIKE

it never starts work it cannot finish, and it reports the trade-off it made in plain language.

BAND S1 · MCP · a pull request into the public repository, use-cases folder

Exhaustion-to-human handoff

WHO IT SERVES

Agents

WHAT TO BUILD

The agent runs until it hits its allowance, then hands off to a named human operator who adopts the account with the one-time takeover code and continues the work.

WHAT STRONG LOOKS LIKE

the handoff email is recognisable to the human, the account keeps its documents and history, and an intercepted link alone is not enough to take over.

BAND S2 · MCP, API · a pull request into the public repository, use-cases folder

Self-correcting agent

WHO IT SERVES

Agents

WHAT TO BUILD

The agent evaluates its own edit against the instruction, detects that it overreached, and reverts that specific turn rather than patching over it.

WHAT STRONG LOOKS LIKE

it reverts precisely one turn and the document returns byte-exact, with the reasoning for the revert recorded.

BAND S2 · MCP, Review · a pull request into the public repository, use-cases folder

Approval-gated autonomous run

WHO IT SERVES

Regulated industries

WHAT TO BUILD

The agent works autonomously but pauses at defined risk boundaries, waits for a human decision, and resumes from where it stopped, possibly hours later.

WHAT STRONG LOOKS LIKE

the pause survives a process restart, partial approval is honoured exactly, and denial produces no partial write.

BAND S3 · Review, MCP · a pull request into the public repository, use-cases folder

Bounded-run agent with an honest report

WHO IT SERVES

Agents

WHAT TO BUILD

Given a fixed budget and deadline, the agent does what it can and then reports precisely what it did not do and why.

WHAT STRONG LOOKS LIKE

the report is accurate. Nothing is claimed that was not done, which is harder than it sounds.

BAND S1 · MCP · a pull request into the public repository, use-cases folder

Confidence-gated escalation

WHO IT SERVES

Regulated

WHAT TO BUILD

The agent scores its own confidence per edit and routes anything below threshold to a human queue instead of applying it.

WHAT STRONG LOOKS LIKE

the confidence signal correlates with actual correctness on a labelled set. Show the correlation, not the claim.

BAND S2 · Review, MCP · a pull request into the public repository, use-cases folder

Crash-resume durability demo

WHO IT SERVES

Platform

WHAT TO BUILD

Kill the process mid-approval and restart it. The pending decisions, the document state and the job all come back intact.

WHAT STRONG LOOKS LIKE

a scripted, reproducible kill test with the recovery shown, not described.

BAND S2 · Review, API · a pull request into the public repository, use-cases folder

Very long document across turns

WHO IT SERVES

Cross-industry

WHAT TO BUILD

A several-hundred-page document processed in a chain of turns that respects the wall-clock boundary, with state carried forward and a single coherent result.

WHAT STRONG LOOKS LIKE

the seams are invisible in the output, and the run is resumable from any turn.

BAND S3 · chat, Multi-document, search, export · a pull request into the public repository, use-cases folder

Voice-driven editing loop

WHO IT SERVES

Accessibility and field work

WHAT TO BUILD

Speak an instruction, the agent edits, and it reads back what changed so the user can approve by voice.

WHAT STRONG LOOKS LIKE

the spoken diff summary is genuinely understandable without seeing the screen.

BAND S2 · chat, Review, export · a pull request into the public repository, use-cases folder

Ticket-queue documentation agent

WHO IT SERVES

Engineering and support

WHAT TO BUILD

Pulls from an issue tracker, updates the affected documents, links the change back to the ticket and closes it.

WHAT STRONG LOOKS LIKE

it closes only tickets it actually resolved, and links are bidirectional.

BAND S2 · MCP, search, Review · a pull request into the public repository, use-cases folder

Pre-release compliance gate

WHO IT SERVES

Regulated publishing

WHAT TO BUILD

A document cannot be exported until it passes a configurable gate, with each failure explained and fixable in place.

WHAT STRONG LOOKS LIKE

the gate is configurable per organisation and its verdicts are explainable and appealable.

BAND S2 · Review, search, export · a pull request into the public repository, use-cases folder

Agent-to-agent document handoff

WHO IT SERVES

Agents

WHAT TO BUILD

A research agent produces findings and hands the same live document to a writing agent, which never re-reads the source material from scratch.

WHAT STRONG LOOKS LIKE

the handoff is by document reference, not by copying text through the context window. Show the token difference.

BAND S2 · MCP, Multi-document, memory · a pull request into the public repository, use-cases folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Calendar and mail to a meeting brief pack

WHO IT SERVES

Executive support

WHAT TO BUILD

Before each meeting, an agent assembles a brief from the invitation, prior threads and past documents.

WHAT STRONG LOOKS LIKE

the brief is short and specific to that meeting rather than a generic dossier.

BAND S2 · MCP, search, memory, export · a pull request into the public repository, use-cases folder

Local repository documentation maintainer

WHO IT SERVES

Engineering

WHAT TO BUILD

Filesystem server plus document server, so an agent maintains a repository's documentation set directly on disk.

WHAT STRONG LOOKS LIKE

it works on a real open-source repository and its output is mergeable as a pull request.

BAND S2 · MCP, search, Review · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Browser research to a teardown document

WHO IT SERVES

Strategy

WHAT TO BUILD

A browsing server gathers evidence and the document server assembles a structured teardown with screenshots.

WHAT STRONG LOOKS LIKE

screenshots are captioned and sourced, and the structure is consistent across three different subjects.

BAND S2 · MCP, images, export · a pull request into the public repository, use-cases folder

Scope: The browsing and the screenshots are done by a separate MCP server that you choose and run. We never browse the web and we never take screenshots. We receive the evidence you gathered and assemble the document from it. Wiring two servers together, one that gathers and ours that writes, is the point of this card.

Issue tracker to linked release notes

WHO IT SERVES

Engineering

WHAT TO BUILD

Release notes generated from resolved issues, with each entry linked back and grouped by audience.

WHAT STRONG LOOKS LIKE

it distinguishes user-facing from internal changes correctly.

BAND S1 · MCP, export · a pull request into the public repository, use-cases folder

Restricted-network deployment kit

WHO IT SERVES

Enterprise IT

WHAT TO BUILD

A guide and configuration for running the agent surface behind a corporate proxy, with an explicit egress list.

WHAT STRONG LOOKS LIKE

it names exactly what must be allowed and why, and it has been verified behind a real proxy.

BAND S1 · MCP, API · a pull request into the public repository, extensions folder

Multi-registry listing sweep

WHO IT SERVES

Distribution

WHAT TO BUILD

Correct listings across every agent and MCP registry that exists, each with the metadata that registry actually ranks on.

WHAT STRONG LOOKS LIKE

a maintained table of registries, submission status and requirements, not a one-off submission.

BAND S1 · MCP · publish it on your own channel and put the link in the submission form

Skill packs for non-Claude agent clients

WHO IT SERVES

Agents

WHAT TO BUILD

Equivalent capability packs for other agent clients, each idiomatic to that client rather than a copy.

WHAT STRONG LOOKS LIKE

each pack is tested in its own client and the differences between clients are documented.

BAND S1 · MCP · a pull request into the public repository, extensions folder

Authenticated agent connection walkthrough

WHO IT SERVES

Agents

WHAT TO BUILD

A working demonstration of the standards-based discovery and authorisation flow an agent client follows to connect.

WHAT STRONG LOOKS LIKE

a real client completes the flow end to end and the walkthrough explains each step for someone implementing it.

BAND S2 · MCP · a pull request into the public repository, extensions folder

Soft-collaboration reconciler for two independent writers on one document

WHO IT SERVES

Platform engineers with two humans, or a human and an agent, on one document

WHAT TO BUILD

A reusable client that makes concurrent authorship survivable without a real-time server. It polls the per-document event stream, applies other writers' committed changes into the local view, rebases a locally pending edit onto them, and - when a rebase is genuinely ambiguous - surfaces the conflict for a human decision rather than picking a winner. Ships with a reference interface showing incoming changes, pending local edits, and the conflict queue.

WHAT STRONG LOOKS LIKE

no lost update across a scripted 200-edit interleaving of two writers; every stale write is surfaced, never silently discarded; the reconciliation is deterministic and replayable from the event log alone, so a disputed outcome can be reconstructed; a writer offline for an hour rejoins without a full reload.

BAND S2 · API, Multi-document, Review · a pull request into the public repository, extensions folder

Learned model-tier routing policy, measured and then shipped

WHO IT SERVES

Platform teams paying for document AI at volume on one fixed model setting

WHAT TO BUILD

An agent that chooses, per sub-task, the cheapest capability tier and reasoning depth that still clears a stated quality gate, records the outcome, and learns a routing table from its own history. The deliverable is not the agent: it is the published routing table plus the harness that produced it, so another team can adopt the policy and re-derive it on their own workload.

WHAT STRONG LOOKS LIKE

the learned policy beats both always-cheapest and always-strongest on cost at equal measured quality, shown with the runs; the quality gate is defined before the runs and not adjusted after; the cases where the cheap tier was wrong are published, not hidden; re-deriving the table on a second workload is one documented command.

BAND S2 · API, chat, search · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Rollback-and-replan agent, where undo is the control structure

WHO IT SERVES

Teams running unattended document agents that compound their own mistakes

WHAT TO BUILD

An agent whose failure path is not "fix it forward". After each turn it verifies its own work against the turn's stated intent; when verification fails it reverts the entire turn to the pre-turn state and re-plans from scratch with the failure as an input, rather than issuing a corrective edit on top of a bad one. The reverted attempt stays visible in the history.

WHAT STRONG LOOKS LIKE

a deliberately bad instruction leaves the document byte-identical to its pre-turn state, proven by diff; the second plan is genuinely different from the first rather than a retry of it; the history shows the reverted turn honestly instead of hiding it; a bounded number of attempts, after which it stops and says so.

BAND S2 · API, chat, Review · a pull request into the public repository, use-cases folder

Public-repository disclosure audit against a supplied policy

WHO IT SERVES

Engineering teams that publish code and need to know it discloses nothing it should not.

WHAT TO BUILD

Write a short disclosure policy for a fictional company: the categories that must never appear in its public repositories (model providers, framework names, database technology, internal file and class names, credential-shaped strings, customer names). Then audit a public repository of your own or a synthetic one against that policy, using a scripted sweep plus your own reading, and open a pull request with minimal, targeted edits that remove every violation. Anything that looks sensitive but is not on the list gets flagged in the pull request without describing what you think it means.

WHAT STRONG LOOKS LIKE

the sweep is reproducible, the edits are surgical and complete, and the judgement call shows: unlisted-but-suspicious findings are flagged with restraint, never ignored and never speculated about.

BAND S2 · chat, export · a pull request into the public repository, use-cases folder

RAG, TEMPLATES AND BENCHMARK

Legal template pack

WHO IT SERVES

Legal / a solo counsel, small firm, or startup legal-ops lead who needs a clean baseline template set, not a negotiation-fallback engine.

WHAT TO BUILD

Small legal teams and startups without in-house counsel need a clean baseline set of contract templates to start from, not a negotiation-ladder engine with position-by-position fallbacks. Build a starter template pack: NDA, MSA, SOW, and employment-agreement templates plus a basic supporting clause library, delivered as ready-to-use content, no negotiation logic attached.

WHAT STRONG LOOKS LIKE

A reviewer opens each of the four templates and confirms every one has clearly labeled placeholder fields and no orphaned cross-reference to a clause that isn't in the pack.

BAND S1 · templates · a pull request into the public repository, use-cases folder

Template-gallery microsite

WHO IT SERVES

Cross-industry. A content-minded builder who writes real category copy and wires a working deep link, not placeholder text.

WHAT TO BUILD

Someone searching for a contract template or an onboarding-packet template should land on a page that actually shows it, not a locked product screenshot. Build a small public microsite that browses SuperDocs' own template packs by category, with real preview content per template and a one-click "open this template" link that lands the visitor directly inside the product with that template loaded. This is distinct from an in-app template pane and from contributing to a third party's gallery: it is our own public showcase, built to be found by search.

WHAT STRONG LOOKS LIKE

Every listed template's one-click link opens that exact template in the product on the first click, and at least three real industry categories each carry genuinely distinguishing copy rather than the same paragraph with a noun swapped.

BAND S1 · templates · a pull request into the public repository, use-cases folder

Benchmarking study report builder

WHO IT SERVES

Consulting / research lead at a consulting firm or an industry association

WHAT TO BUILD

A benchmarking study collects responses from many participants, then delivers a public report plus a confidential personalized report per participant showing their position against the cohort. The app builds the cohort analysis into charts, drafts the report narrative, then generates each participant's version by inserting their own figures and tailored commentary while keeping every other participant's data anonymous. A disclosure check confirms no participant's identity is inferable from a small-cohort cut before anything is exported.

WHAT STRONG LOOKS LIKE

the anonymity check is real and blocks any report where a small cell would identify a participant; personalized commentary is accurate to that participant's data; public and personalized versions share one narrative source so they cannot contradict each other.

BAND S3 · Multi-document, images, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Carrier and shipper contract redline desk

WHO IT SERVES

Transportation procurement or legal ops lead negotiating carrier agreements at volume

WHAT TO BUILD

Carrier agreements arrive as the carrier's paper, and the shipper has a playbook of required positions on liability limits, detention, fuel surcharge, insurance and indemnity. Build an app that loads the playbook as a knowledge base, redlines inbound paper clause by clause, inserts the fallback wording where the carrier's position is unacceptable, leaves acceptable clauses untouched, and produces a negotiation summary listing every position taken with the playbook rule that drove it. Export a tracked-changes Word file the carrier's counsel can work with directly.

WHAT STRONG LOOKS LIKE

the redline touches only the clauses the playbook cares about, the negotiation summary and the redline never disagree, and the export opens in Word as genuine tracked changes.

BAND S2 · chat, templates, Review, export · a pull request into the public repository, use-cases folder

Newsroom house-style copy desk

WHO IT SERVES

Copy chief or section editor at a publication with a written stylebook

WHAT TO BUILD

Build a copy desk where the publication's stylebook and past corrected articles are the knowledge base, and an incoming draft is edited to house style as tracked changes the editor accepts or rejects one at a time. Style calls are explained inline ("stylebook says figures under 10 are spelled out"), the writer's voice is preserved, and quotes inside quotation marks are never altered. The accepted version exports for the CMS, and rejected calls teach the desk not to repeat that suggestion for that writer.

WHAT STRONG LOOKS LIKE

direct quotes come through untouched, every suggestion cites the rule it enforces, and rejections measurably reduce repeat suggestions.

BAND S2 · chat, templates, Review, memory, export · a pull request into the public repository, use-cases folder

Style-guide linter that fixes in place

WHO IT SERVES

Technical writer or docs lead enforcing a writing style guide across a large docs set

WHAT TO BUILD

Build a linter whose rules are the team's actual style guide, attached as the knowledge base rather than hardcoded. It reads a page, flags each violation with the rule it breaks and the suggested rewrite, and applies the accepted fixes in place without touching code samples, tables or headings that are already compliant. Batch mode runs the whole docs set and reports the violation profile by rule so the team can see which guidance is being ignored most, and a rule change reruns without any code change.

WHAT STRONG LOOKS LIKE

rules live entirely in the attached style guide (swapping in a different guide changes behaviour with no code edit), and applied fixes never alter code samples or meaning.

BAND S2 · templates, chat, Review, API, export · a pull request into the public repository, extensions folder

Preprint to journal submission converter

WHO IT SERVES

Author moving a manuscript from a preprint to a specific journal's requirements

WHAT TO BUILD

Each journal wants its own structure, section order, reference style, figure and table placement, word limits, blinding requirements and declarations. Build a tool that holds the target journal's guidelines as the knowledge base, restructures the manuscript to match without rewriting the science, reformats the references, moves figures and captions to the required placement, strips identifying information when blinding is required, and reports every guideline it could not satisfy. Switching target journals re-runs the transformation from the original.

WHAT STRONG LOOKS LIKE

the scientific content is provably unchanged while structure and style move, blinding is thorough (acknowledgements and self-citations included), and unmet guidelines are reported honestly.

BAND S3 · templates, chat, Multi-document, export, search · a pull request into the public repository, use-cases folder

Flow-Down Clause Checker

WHO IT SERVES

GovCon · contracts administrator issuing a subcontract under a prime award

WHAT TO BUILD

Upload the prime contract and the draft subcontract. The app identifies the clauses the prime requires to be flowed down, checks whether each one appears in the subcontract, and inserts the missing clauses from a clause library at the correct place, repairing surrounding numbering and cross-references. Output is the updated subcontract plus a flow-down matrix marking each clause required, present, inserted, or deliberately omitted with a stated reason.

WHAT STRONG LOOKS LIKE

insertions land in the right section with numbering and cross-references repaired around them; the matrix is complete and honest about omissions; running it twice inserts nothing the second time.

BAND S3 · Multi-document, search, templates, Review, export · a pull request into the public repository, use-cases folder

Grounded drafting engine with sentence-level source anchors

WHO IT SERVES

Research, legal, and analyst tooling. Teams whose buyers ask "how do I know it did not make this up".

WHAT TO BUILD

A drafting service where generation and verification are separate stages. Each drafted sentence is written together with the retrieved passage that supports it, then an independent pass re-checks each sentence against the source set and edits only the ones it cannot support. The final document carries a provenance appendix mapping sentences to sources, exportable as a real appendix with footnotes. Unsupported material is removed and listed, never quietly kept.

WHAT STRONG LOOKS LIKE

planted unsupported claims are caught at a measured rate on a labelled test set; the appendix is machine-checkable; verification is genuinely independent of drafting; the trade-off between coverage and caution is exposed as a setting and documented.

BAND S3 · chat, search, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Freshness-tracked knowledge-base to document generator

WHO IT SERVES

Documentation and support operations. Teams whose generated documents go stale silently.

WHAT TO BUILD

Generates a document from a knowledge base, then keeps it fresh by tracking which source each section came from. When a source changes, only the sections derived from it are regenerated, and the rest of the document is left untouched, with a freshness stamp per section. A staleness report lists what is out of date and by how long. This is the maintenance half that generation-only tools never build.

WHAT STRONG LOOKS LIKE

the section-to-source mapping survives edits and reordering; an unchanged knowledge base triggers zero work; a changed source updates exactly its sections, proven; the staleness report is accurate after many cycles; a human edit to a section is respected rather than overwritten on the next refresh.

BAND S3 · API, chat, search, Multi-document, export · a pull request into the public repository, use-cases folder

Bring-your-own-retrieval adapter layer

WHO IT SERVES

DevTools. Teams that already run a vector store and are not replacing it.

WHAT TO BUILD

An adapter that lets an existing retrieval stack feed our editing surface: retrieved passages are supplied as reference material for the edit, and the resulting change set is written back with the retrieval identifiers attached. Adapters for three common store shapes plus a written interface so a fourth is straightforward. It removes the false choice between "use their retrieval" and "use ours" by making both work together.

WHAT STRONG LOOKS LIKE

the interface is small and documented; each adapter has an integration test against a real store instance in CI; passage size and count handled sensibly for cost; the identifiers survive into the change record; a written comparison of when each option is the better choice.

BAND S2 · API, chat, search, Review · a pull request into the public repository, extensions folder

Entity-graph retrieval to a structured report with a diagram

WHO IT SERVES

Diligence, investigations, and research teams.

WHAT TO BUILD

Builds an entity and relationship graph across a document set (parties, obligations, dates, references), then uses the graph to plan the report: the outline follows the graph, each section is written from the passages attached to those entities, and an auto-rendered relationship diagram is embedded in the document. Clicking through the report, a reader can trace any statement back to the entity and the source that produced it.

WHAT STRONG LOOKS LIKE

the graph is built from the documents rather than hand-seeded; the diagram is generated and stays in sync with the graph; the report is genuinely better than a flat retrieval baseline, shown side by side; the whole run is reproducible on the included corpus.

BAND S3 · chat, search, Multi-document, images, export · a pull request into the public repository, use-cases folder

Retrieval evaluation harness for document work

WHO IT SERVES

DevTools and machine-learning engineering. Teams tuning a document assistant.

WHAT TO BUILD

A harness that scores a document retrieval-and-edit pipeline on three axes rather than one: did it retrieve the right passage, was the answer grounded, and did the resulting edit touch only what it should. Ships with a labelled corpus, a scoring implementation, a command-line runner, and a report generator. Configurations are files, so a team can compare their variants and publish the result.

WHAT STRONG LOOKS LIKE

the labelled corpus is genuinely labelled by a human and documented; scores are stable across repeated runs and variance is reported; the edit-precision axis is the differentiator and is implemented rigorously; results reproducible from a clean clone; a written note on what the harness cannot measure.

BAND S3 · API, chat, search, Review · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Provenance ledger for AI edits

WHO IT SERVES

Regulated industries. Legal, audit, clinical, and financial teams that must explain any change.

WHAT TO BUILD

A layer that records, for every AI edit applied to a document, which retrieved source justified it, which instruction requested it, who approved it, and when. The ledger exports as an appendix or a separate audit document, and a verification command re-checks that the document's current content still matches the ledger. It is the answer to "prove why this sentence changed" nine months later.

WHAT STRONG LOOKS LIKE

the ledger is append-only and tamper-evident; verification catches an out-of-band edit; it survives revert and redo correctly; export is readable by a non-technical auditor; storage of the ledger is pluggable.

BAND S2 · API, chat, Review, search, export · a pull request into the public repository, extensions folder

Table-aware retrieval with write-back

WHO IT SERVES

Finance and operations. Anyone whose source of truth is a table inside a document.

WHAT TO BUILD

Retrieval that understands tabular structure: it indexes rows and headers as records rather than flattened prose, answers questions about them, and, when asked to update, edits the specific cells in the document while keeping the table's formatting intact. Handles merged cells, multi-page tables, and tables inside uploaded files. Includes a question set demonstrating that row-level retrieval beats flat text retrieval on the same corpus.

WHAT STRONG LOOKS LIKE

the structural comparison against flat retrieval is measured, not asserted; merged and multi-page tables handled; cell edits preserve borders, shading, and alignment through export; a corpus of genuinely awkward real-world tables in the fixtures.

BAND S3 · search, chat, API, export · a pull request into the public repository, use-cases folder

Multi-tenant retrieval isolation reference

WHO IT SERVES

Platform engineering. Any company serving document AI to their own customers.

WHAT TO BUILD

A reference implementation of correct isolation when one platform serves many end customers: per-end-customer scoping of documents, retrieval, and long-term memory, with keys chosen so nothing crosses a boundary, plus a deliberately adversarial test suite that tries to make one tenant see another's content and fails every time. Includes a written threat model and the failure cases the authors considered.

WHAT STRONG LOOKS LIKE

the adversarial suite is genuinely adversarial and is the centrepiece; the threat model names what is out of scope; scoping is applied consistently across every surface used, not just the obvious one; a checklist a reader can apply to their own integration.

BAND S3 · API, search, memory, Multi-document, chat · a pull request into the public repository, extensions folder

Event-driven incremental re-index pipeline

WHO IT SERVES

Platform engineering. Teams whose library grows and changes constantly.

WHAT TO BUILD

A pipeline that keeps a document library's retrieval index fresh without periodic full rebuilds: it listens for document change events, re-indexes only the affected parts, handles deletions and restorations, and reports lag between a change and its visibility in search. Includes a backfill mode for the initial load and a consistency checker that detects and repairs drift between the library and the index.

WHAT STRONG LOOKS LIKE

measured lag with a target and an alert; the consistency checker finds injected drift; deletion and restoration handled correctly; the backfill is resumable; cost of steady-state operation reported.

BAND S3 · API, search, Multi-document · a pull request into the public repository, extensions folder

Knowledge-gap detector that authors the missing article

WHO IT SERVES

Customer support and success. Teams whose help centre is always one incident behind.

WHAT TO BUILD

A loop that watches incoming questions, retrieves against the existing help library, and identifies clusters that retrieval cannot answer well. For each cluster it drafts the missing article, links it from the related existing articles by editing them in place, and routes both for approval. It closes the loop by re-testing the original questions against the updated library and reporting the improvement.

WHAT STRONG LOOKS LIKE

the gap detection is measured rather than vibes-based, with a stated threshold; existing articles are edited surgically to add links, not regenerated; the before-and-after retrieval improvement is reported honestly including cases with no gain; approval required before anything publishes. Distinct from the existing ticket-to-article rows because the trigger is a measured retrieval gap across a library.

BAND S3 · search, chat, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Token-cost benchmark: whole-document regeneration against section-level editing

WHO IT SERVES

Positioning and developer evangelism. Every engineer who has been asked "why not just send the whole document to a model".

WHAT TO BUILD

The flagship measurement. A corpus of documents at several sizes (one page to several hundred), a set of realistic edit tasks per document, and two arms: regenerate the whole document per edit, or edit the target section only. It records tokens, wall-clock time, and cost per edit and cumulatively across a multi-turn session, then publishes tables, charts, and the raw data. A runner reproduces everything with one command. Distinct from the existing interactive calculator row, which estimates for marketing; this measures and publishes.

WHAT STRONG LOOKS LIKE

both arms implemented competently and reviewed for fairness; repeated runs with variance reported; document sizes and edit types varied enough that the curve, not one number, is the finding; raw data committed; the write-up states exactly what was measured and what was not; results published whatever they show.

BAND S3 · chat, API, export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Round-trip fidelity harness with structural assertions

WHO IT SERVES

Quality engineering. Anyone whose pipeline ingests Word and must emit Word.

WHAT TO BUILD

Upload, edit, export, then assert structurally on the result: footnotes are still real footnotes, headers and footers are still real parts with working page fields, tracked changes and comments survived with their authors and dates, equations are still native mathematics rather than pictures, bookmarks and internal cross-references still resolve, page geometry per section is unchanged. Failures report the exact structural element that changed. Runs as a suite anyone can point at their own documents.

WHAT STRONG LOOKS LIKE

assertions inspect the real file structure rather than converting to text and comparing; the suite runs on a stranger's document without modification; a genuine regression is caught when one is deliberately introduced; the report names the element and the section; runtime short enough to belong in continuous integration.

BAND S3 · API, chat, export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Public leaderboard repository for document-editing systems

WHO IT SERVES

Positioning and community. The whole category.

WHAT TO BUILD

A benchmark repository with a defined task set, a scoring harness, a submission process, and a published table. Anyone can submit a system by opening a pull request with a runner; the results are produced by continuous integration, not by the submitter, so the table is trustworthy. Categories cover edit precision, format fidelity, cost, and latency. The maintainers commit publicly to publishing results that place us below a competitor.

WHAT STRONG LOOKS LIKE

submissions genuinely run in continuous integration with pinned versions and no submitter-controlled scoring; the task set is public and versioned; a written governance note covers conflicts of interest, since we author it; at least two non-our systems present at launch; results reproducible by a third party.

BAND S3 · API, chat, export, Review · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Edit-precision evaluation: collateral damage measurement

WHO IT SERVES

Quality engineering and product. Teams burned by an AI that changed more than it was asked.

WHAT TO BUILD

Measures the thing users actually complain about: when asked to change one thing, how much else changed. A labelled corpus pairs each instruction with the sections that should change and the sections that must not, and the scorer reports intended-change coverage alongside a collateral-damage rate, per instruction type (single sentence, one section, a table row, a global rename). Publishes a breakdown by instruction type rather than one aggregate.

WHAT STRONG LOOKS LIKE

labels created carefully with a documented process and a second reviewer; the metric handles legitimately ambiguous instructions rather than punishing them; results broken out by instruction type; failure examples included in the write-up; the corpus is reusable by anyone.

BAND S3 · chat, API, search · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Measured scope: competing approaches are measured on the same corpus, and losses are published.

Latency and throughput benchmark across document sizes

WHO IT SERVES

Platform engineering. Teams sizing an integration before they commit.

WHAT TO BUILD

Measures time to first progress event, time to completion, and sustained throughput across document sizes and concurrency levels, reporting median and tail latencies rather than an average. Separates the phases so a reader can see where time goes: upload, processing, editing, export. Publishes a sizing guide derived from the results, so an integrator can predict what their workload will feel like.

WHAT STRONG LOOKS LIKE

tail latencies reported, never just an average; the load generator does not become the bottleneck, and that is verified; rate-limit responses are excluded from latency statistics and reported separately rather than being counted as slow successes; enough runs for the variance to be meaningful; the sizing guide follows from the data.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Synthetic long-document corpus generator

WHO IT SERVES

Quality engineering. Everyone building the other benchmark cards.

WHAT TO BUILD

A generator that produces documents of controlled size and structure with known ground truth: a stated number of sections, tables with known contents, footnotes at known positions, deliberate structural defects on request, and multiple languages and scripts. Because the ground truth is known, any scorer can check an edit exactly. Output in several formats so the same logical document can be tested through different ingest paths. Deterministic from a seed.

WHAT STRONG LOOKS LIKE

deterministic and seedable; the ground-truth manifest is complete and machine-readable; generated documents are realistic enough to be meaningful rather than obviously synthetic; scales to very large documents without exhausting memory; licence-clean by construction.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Human review-cost evaluation

WHO IT SERVES

Product and operations. Buyers deciding whether review-every-change is affordable.

WHAT TO BUILD

Measures the human side: with a small panel of reviewers and a set of AI-proposed change batches, it records time to decide per change, agreement between reviewers, and how often a reviewer's decision was later judged wrong. Compares presenting changes as one batch against one at a time, and section diffs against whole-document diffs. The finding is a practical recommendation about how to present AI changes to humans.

WHAT STRONG LOOKS LIKE

a real panel, however small, with participants informed and consenting; the protocol is written before the data is collected; inter-reviewer agreement reported; conclusions proportionate to the sample size, explicitly; the raw timings published with identities removed.

BAND S2 · Review, chat, export · a pull request into the public repository, use-cases folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Multilingual and right-to-left fidelity evaluation

WHO IT SERVES

International product teams. Buyers outside English-first markets.

WHAT TO BUILD

A reproducible test of editing and export quality across many languages and scripts, including right-to-left and East Asian text, mixed-direction paragraphs, and bilingual documents. It checks that the instruction language and the document language can differ, that direction and shaping survive edit and export, that fonts do not fall back to unreadable substitutes in the exported file, and that line breaking is correct. Publishes a per-language profile.

WHAT STRONG LOOKS LIKE

native-speaker review for at least a few languages, credited; the export is inspected as a file rather than assumed correct from the editor view; mixed-direction cases included because they are where things break; the profile is honest per language rather than a single pass mark.

BAND S3 · chat, export, search · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Agent tool-selection accuracy harness

WHO IT SERVES

Agent infrastructure. Anyone shipping a tool surface to an agent.

WHAT TO BUILD

A harness that measures whether an agent picks the right tool from a large tool surface: a labelled set of user requests with the correct tool sequence, scored on selection accuracy, unnecessary calls, and recovery after a wrong first choice. Because it also measures cost per completed task, it turns tool-description quality into a number that can be improved deliberately rather than by intuition.

WHAT STRONG LOOKS LIKE

the labelled set covers ambiguous requests, not just easy ones; a change to a tool description produces a measurable score change, demonstrated; results reported per tool so the weakest descriptions are identified; the harness works against any tool server, not only ours.

BAND S3 · MCP, chat, export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Cost regression gate for continuous integration

WHO IT SERVES

Platform engineering. Teams that have to keep a document feature inside a budget.

WHAT TO BUILD

A small library and continuous-integration step that runs a fixed scenario set on every change and fails the build when token cost, operation count, or latency regresses beyond a threshold. Baselines are committed, so the history of a project's document costs is visible in version control. Produces a trend chart and, of course, a monthly summary document. Works for any integration, not only ours.

WHAT STRONG LOOKS LIKE

scenarios are stable enough that the gate does not flake, with the flake rate measured; thresholds configurable with sensible defaults; a real regression is caught when one is deliberately introduced; baseline updates require an explicit, reviewable commit; runtime acceptable for every-change use.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Runnable cookbook, fifty recipes tested in continuous integration

WHO IT SERVES

DevTools. Every developer in the first hour of an evaluation.

WHAT TO BUILD

A cookbook site where each recipe is a complete, runnable program solving one real task (upload a large file without blowing up memory, stream progress into a web interface, approve changes from a server, target one document among several open ones, export five formats in one pass, recover from a capacity refusal), in several languages, and every recipe is executed in continuous integration against a mock so a broken example fails the build. Each recipe states the operations it spends.

WHAT STRONG LOOKS LIKE

recipes are executable files in the repository, not snippets pasted into prose; the continuous-integration run genuinely executes them; each one is minimal, complete, and copy-pasteable; a contribution guide so recipes can be added; a machine-readable index an AI assistant can consume.

BAND S2 · API, chat, Review, Multi-document, search, templates, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

API client collections with a sixty-second key path

WHO IT SERVES

DevTools. Evaluators who click before they code, and API-first teams.

WHAT TO BUILD

Maintained collections for the popular API clients (a request collection, an environment file, and a saved example response per request), generated from the published specification so they never drift. The differentiator is the onboarding script bundled with the collection: it walks a new evaluator from nothing to a working key and a first successful edit in under a minute, then leaves them with a folder of working requests covering the whole surface, including the review flow and the large-file path.

WHAT STRONG LOOKS LIKE

generated, not hand-maintained, with the regeneration step in continuous integration; every request has a saved example; secrets live in the environment file and are never committed; the collection is published where those clients are discovered; a stranger completes the first edit without reading anything else.

BAND S1 · API, chat, Review, export · a pull request into the public repository, extensions folder

One-command local sandbox

WHO IT SERVES

DevTools. Teams evaluating without wanting to sign anything yet, and teams building offline.

WHAT TO BUILD

A single compose file that brings up a local development environment: the mock server from card 863, a small sample web application, a queue worker, and seeded sample documents including a heavily styled one, a scanned one, and a multi-language one. One command, no key required, everything works offline. A flag switches the same environment to the live service when the evaluator is ready, which is the moment of conversion.

WHAT STRONG LOOKS LIKE

truly one command on a clean machine, tested on all three operating systems; no key needed for the offline path; the switch to live requires only setting one value; the seeded documents are realistic; teardown leaves nothing behind.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

One feature, six stacks

WHO IT SERVES

DevTools. Engineering leads choosing how to integrate.

WHAT TO BUILD

A monorepo where the same small feature (upload a document, ask for a change, review the proposed edits, export the result) is implemented six times in six stacks, each idiomatic to its ecosystem, each deployable, each with the same test suite proving identical behaviour. A comparison document written from the exercise explains the trade-offs honestly: where each stack made the work easier or harder.

WHAT STRONG LOOKS LIKE

all six genuinely run and pass the shared behavioural suite; each is idiomatic rather than a transliteration of the first one; the comparison names real trade-offs including where our surface was awkward; each stack deployable independently with its own instructions.

BAND S3 · API, chat, Review, export · a pull request into the public repository, extensions folder

Terminal tour that teaches the surgical-edit model

WHO IT SERVES

DevTools. Developers who learn by doing and skip documentation.

WHAT TO BUILD

A guided interactive tour run with one command. It walks a developer through the mental model in about ten minutes: here is a document with sections, here is a change that touches one of them, here is what the whole-document approach would have cost by comparison, here is the review flow, here is multiple documents in one session, here is export. Each step is a real call against a sample document, with the cost shown, ending with a generated starter project in the language of choice.

WHAT STRONG LOOKS LIKE

the tour is resumable and skippable; every step is a real call, not a simulation; the cost comparison shown is measured, consistent with card 908, and not overstated; it works on a free-tier key inside the monthly allowance; the generated starter project actually runs.

BAND S2 · chat, Review, Multi-document, export · a pull request into the public repository, extensions folder

Failure playbook with a runnable simulator

WHO IT SERVES

DevTools and support. Teams preparing an integration for production.

WHAT TO BUILD

A written playbook covering every failure an integration will meet (too large, too complex, rate limited, a job already running, a turn that exceeded the time cap, an approval that arrived twice, a network drop mid-stream) paired with a simulator that produces each one on demand so a team can test their handling before a customer does. Each entry states the cause, the correct response, and whether it was billable, and the last part matters because a failed turn costs nothing.

WHAT STRONG LOOKS LIKE

the simulator reproduces each failure faithfully and deterministically; every playbook entry is verified against the real behaviour, not written from assumption; the capacity refusals are correctly described as expandable limits with the request path shown; the playbook is machine-readable so an agent can consult it.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

Published schema and type package across registries

WHO IT SERVES

DevTools. Teams that want types without adopting a client library.

WHAT TO BUILD

A single source of truth published as data: the request and response schemas, the event vocabulary, and the error catalogue from card 862, released as versioned schema files plus generated type packages for several languages. A consumer gets editor completion and compile-time checking without taking a runtime dependency on anyone's client. Versioned with a compatibility policy and a changelog that names every breaking change.

WHAT STRONG LOOKS LIKE

generated from the published specification in continuous integration rather than hand-written; a compatibility check blocks a release that would silently break consumers; every published package carries the same version; the schema files are usable directly for validation; documented policy for how long a version is supported.

BAND S2 · API, chat, Review, export · a pull request into the public repository, extensions folder

Fallback-ladder clause library

WHO IT SERVES

Legal, contract managers, deal desk. (a growth candidate can build the content; an engineer can build the retrieval and insertion loop).

WHAT TO BUILD

Build a clause library where every clause carries three graded positions: preferred, acceptable, and walk-away, plus a one-line rationale for each. Load the library as a knowledge base, then drive a session that reads an inbound contract, finds the matching clause, and proposes the correct rung of the ladder given a stated risk posture. Cover 20 clauses that appear in almost every commercial agreement. Ship the library, the assembly session, and a recorded run against a public sample contract. a published clause library plus a working playbook engine under their name. This is a legal-ops portfolio piece that reads as domain expertise, not as a tutorial.

WHAT STRONG LOOKS LIKE

the rationale lines are written by someone who understands why a cap on liability moves, not paraphrased from a generic web page. The retrieval picks the right clause when the contract uses different wording than the library does. The proposed change lands as a reviewable diff, not a rewrite of the surrounding section.

BAND S2 · templates, search, chat, Review, export · a pull request into the public repository, use-cases folder

Multi-jurisdiction employment document pack

WHO IT SERVES

HR, people ops, employer of record.

WHAT TO BUILD

Build offer letter, employment agreement, and policy acknowledgement templates for four jurisdictions with genuinely different requirements, plus a variables sheet naming what changes per jurisdiction and why. Drive a session that takes one role and one jurisdiction and produces the correct set on a single branded template. Every jurisdiction-specific difference must be annotated in a companion notes document, sourced. a jurisdiction-annotated HR document pack, which is a credible people-ops or HR-tech portfolio artifact.

WHAT STRONG LOOKS LIKE

the four jurisdictions are chosen because they differ in interesting ways, not because they were easy. The annotations cite public sources. The branded template survives all four fills with identical layout, proven by exporting all four and diffing the styling.

BAND S2 · templates, chat, Multi-document, export · a pull request into the public repository, use-cases folder

Plain-language rewrite pack for consumer documents

WHO IT SERVES

Consumer finance, insurance, utilities, public sector.

WHAT TO BUILD

Take five real consumer-facing documents in the public domain (terms, a policy summary, a notice, a bill explainer, a benefits letter) and build a plain-language rewrite pack: a reading-level target, a substitution guide for the jargon that actually appears, and a session that rewrites in place while preserving every legally load-bearing sentence. Produce a before and after readability score for each document and state the method used to score. a measured before-and-after readability study, which is a strong writing and content-design portfolio piece.

WHAT STRONG LOOKS LIKE

the rewrite preserves meaning exactly where meaning is load-bearing and simplifies only where it is safe, and the candidate can point at which sentences they refused to touch and say why. Readability scores are computed, not asserted.

BAND S2 · chat, templates, search, export · a pull request into the public repository, use-cases folder

Open-source project governance document pack

WHO IT SERVES

Developer tools, open-source maintainers, foundations.

WHAT TO BUILD

Build the governance document set a serious open-source project needs: code of conduct, contributing guide, security policy, maintainer handbook, contributor licence agreement, and a release policy. Then build the part that matters: a session that adapts the whole set to a new project's name, licence, governance model, and enforcement contact in one pass, and exports both Markdown for the repo and a styled PDF for a foundation submission. a governance pack other projects will actually adopt, with their name on it. Reusable forever.

WHAT STRONG LOOKS LIKE

the set is opinionated and internally consistent, the adaptation pass changes every dependent mention (not just the title), and the Markdown output is clean enough to commit directly.

BAND S1 · templates, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Machine-readable brand style kit

WHO IT SERVES

Agencies, marketing ops, brand teams.

WHAT TO BUILD

Turn a brand style guide into something an AI can enforce: a structured style contract naming fonts, sizes, spacing, heading behaviour, table treatment, tone rules, and forbidden constructions, plus a reference document that demonstrates each rule. Then build the enforcement session: load the kit, hand it any off-brand document, and get back the same document brought to standard with a change report saying which rule triggered each edit. a brand enforcement kit and a written method, which is a genuinely rare marketing-ops portfolio artifact.

WHAT STRONG LOOKS LIKE

the change report maps every edit to a named rule, so a brand manager can audit it. Nothing outside the named rules is touched, and the candidate proves that with a diff on an already-compliant document that should come back with zero changes.

BAND S2 · templates, chat, Review, export, memory · a pull request into the public repository, use-cases folder

Newsroom house-style pack and style-drift report

WHO IT SERVES

Media, publishing, content teams.

WHAT TO BUILD

Build a house-style pack for a publication (real public style guides exist and can be cited) and a drift report: run 20 published articles through a style check, produce a report naming the rules broken most often, and apply the fixes to three of them in place with tracked review. The report is the deliverable that makes the pack interesting. an editorial study with real numbers, publishable under their byline.

WHAT STRONG LOOKS LIKE

the drift report is quantified and ranked, the fixes preserve the writer's voice, and the candidate can show an article where the checker was wrong and explain why.

BAND S2 · templates, chat, search, Review, export · a pull request into the public repository, use-cases folder

Citation and reference style kit

WHO IT SERVES

Academia, research, publishing.

WHAT TO BUILD

Build a kit that converts a document's references between four citation styles without touching anything else, including in-text markers, footnote numbering, and the bibliography order. Cover the awkward cases on purpose: a work with no author, a translated work, a preprint, a standard, and a dataset. Ship a test document containing every awkward case and a matrix showing the output in all four styles. a citation conversion test corpus and matrix, useful to any research-tools team they ever join.

WHAT STRONG LOOKS LIKE

footnote numbering survives, the bibliography reorders correctly when the style demands it, and the awkward cases are handled rather than quietly dropped. The matrix is the proof.

BAND S2 · chat, search, export · a pull request into the public repository, use-cases folder

Accessible-document standard kit

WHO IT SERVES

Government, education, enterprise, anyone under accessibility procurement rules.

WHAT TO BUILD

Build a kit that takes a structurally messy document and brings it to an accessible standard: real heading hierarchy, list semantics instead of manual bullets, table headers, meaningful link text, image descriptions, and language attributes. Produce a before-and-after structure report, and export a tagged PDF plus a Word file that both carry the structure. State clearly what is fixed automatically and what still needs a human. an accessibility remediation method with measured before and after structure, which is a hireable specialisation on its own.

WHAT STRONG LOOKS LIKE

the structure report is generated from the document, not written by hand. The candidate is honest about the residue, and the honesty is specific rather than a disclaimer.

BAND S3 · chat, export, images, Review · a pull request into the public repository, use-cases folder

Board and corporate-governance document pack

WHO IT SERVES

Startups, nonprofits, company secretaries.

WHAT TO BUILD

Build the governance set: board charter, committee charters, written resolutions, notice of meeting, minutes template, and a conflict-of-interest declaration. Then build the cycle: a session that takes an agenda and a transcript and produces the notice, the minutes, and the resolutions as three linked documents with consistent numbering and dates across all three. a governance pack plus a multi-document consistency demo, which is a strong ops or legal-ops artifact.

WHAT STRONG LOOKS LIKE

the three documents stay consistent when the agenda changes, resolution numbering does not drift, and the minutes read like minutes rather than a transcript summary.

BAND S2 · templates, Multi-document, chat, export · a pull request into the public repository, use-cases folder

RFP answer library with provenance tags

WHO IT SERVES

Procurement, bid teams, public-sector sales.

WHAT TO BUILD

Build a starter answer library of 60 reusable RFP answers across security, support, commercials, and implementation, where every answer carries a provenance tag: who owns it, when it was last verified, and whether it may be sent without review. Then build the responder: given a requirements list, draft the response document pulling from the library, and mark every answer that is stale or review-gated so a human sees it before it ships. a provenance-tagged answer library and a responder, which is directly reusable in any bid or presales role.

WHAT STRONG LOOKS LIKE

the staleness and review gates actually block, the requirements-to-answer mapping is shown rather than assumed, and unanswerable requirements are flagged as gaps instead of being filled with waffle.

BAND S2 · templates, search, chat, Review, export · a pull request into the public repository, use-cases folder

Independent consultant contract pack

WHO IT SERVES

Freelancers, agencies, small professional-services firms.

WHAT TO BUILD

Build the working set an independent consultant needs: engagement letter, statement of work, rate card, change order, and a project closure note, all on one branded system. The interesting part is the change order: it must reference the original statement of work section by section and restate only what changed, with the original left untouched as a linked document. a contract pack they can use in their own consulting work immediately.

WHAT STRONG LOOKS LIKE

the change order genuinely references sections of the original rather than restating the whole scope, and the pair reads correctly when opened together months later.

BAND S1 · templates, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Non-English legal template pack, one language done properly

WHO IT SERVES

International legal, cross-border services.

WHAT TO BUILD

Pick one language the candidate is genuinely fluent in and build a template pack in it: a service agreement, an NDA, and a notice letter, correct in register and convention rather than translated from English. Then build the bilingual pairing: the same agreement produced side by side in both languages with the layout preserved across scripts. Document every convention that differs from the English original. a bilingual legal template pack in their own language, which is a rare and defensible specialisation.

WHAT STRONG LOOKS LIKE

a native reader would not identify it as translated. The bilingual layout holds, including in a right-to-left script if that is the language chosen. The conventions document is specific.

BAND S2 · templates, chat, Multi-document, export · a pull request into the public repository, use-cases folder

Trade and export documentation pack

WHO IT SERVES

Logistics, exporters, trade finance.

WHAT TO BUILD

Build the export document set that travels with a shipment: proforma invoice, commercial invoice, packing list, certificate of origin, and a letter-of-credit document checklist. The pack must enforce internal consistency, because in trade documents a mismatch between the invoice and the packing list is what gets a shipment held. Build a consistency checker that reads all five and reports every field that disagrees. a trade document pack and a cross-document consistency checker, an unusual and concrete supply-chain portfolio piece.

WHAT STRONG LOOKS LIKE

the consistency checker catches seeded mismatches in a test set, including subtle ones like a weight unit change, and the report names the document and the field rather than saying documents disagree.

BAND S2 · templates, Multi-document, search, export · a pull request into the public repository, use-cases folder

Grant post-award compliance reporting pack

WHO IT SERVES

Nonprofits, research administration, grant managers.

WHAT TO BUILD

Existing rows cover proposals. This covers what happens after the money lands: interim and final narrative reports, budget narratives with variance explanations, a no-cost extension request, and a change-of-scope notification. Build the pack plus a session that takes the original proposal as a reference document and drafts the interim report against the commitments actually made in it, quoting them. a post-award reporting pack, immediately useful to any nonprofit or research office.

WHAT STRONG LOOKS LIKE

the report is anchored to real commitments pulled from the proposal rather than invented, and unmet commitments are surfaced rather than smoothed over.

BAND S2 · templates, Multi-document, search, chat, export · a pull request into the public repository, use-cases folder

Negotiation playbook pack

WHO IT SERVES

Deal desk, commercial legal, sales operations.

WHAT TO BUILD

A playbook is not a clause library. Build the document that governs a negotiation: for each of 15 negotiation points, the opening position, the concession sequence, the trade it may be exchanged for, the approver required to go past a line, and the words to use. Then produce the artifact a deal desk actually wants: a one-page negotiation brief generated for a specific inbound contract, listing which of the 15 points that contract touches and where it sits versus the playbook. a negotiation playbook and brief generator, a commercially valuable artifact in any revenue-operations or legal-operations career.

WHAT STRONG LOOKS LIKE

the brief is genuinely derived from the inbound contract, points that the contract is silent on are called out as silent (not as compliant), and the approver routing is correct.

BAND S3 · templates, search, chat, export · a pull request into the public repository, use-cases folder

Technical-writing document standard kit

WHO IT SERVES

Engineering organisations, developer relations, technical writers.

WHAT TO BUILD

Build the standard set an engineering org writes against: design document, architecture decision record, runbook, incident postmortem, and release note, each with a required-sections contract and a worked example. Then build the check: hand the kit any draft of one of those types and get a completeness report naming missing required sections and thin ones, plus the option to fill the gaps in place. a document standard kit engineers will actually use, with their name on it.

WHAT STRONG LOOKS LIKE

the worked examples are good enough to be adopted as-is, the completeness check distinguishes missing from thin, and the runbook example includes a real diagram generated in the flow.

BAND S1 · templates, chat, images, export · a pull request into the public repository, use-cases folder

Teacher feedback comment bank and rubric pack

WHO IT SERVES

Education, tutoring platforms, teachers.

WHAT TO BUILD

Build a rubric set for three assignment types plus a comment bank of 100 feedback comments organised by rubric criterion and severity. Then build the loop: given a student document, apply feedback as tracked review comments drawn from the bank, adapted to the specific sentence rather than pasted, and produce a summary sheet showing the rubric score with the evidence for each criterion. a rubric and comment-bank pack plus a marking demo, useful to anyone working in education technology.

WHAT STRONG LOOKS LIKE

the comments are anchored to specific sentences and adapted, the rubric evidence quotes the student's own text, and the student's document is otherwise untouched.

BAND S2 · templates, chat, Review, export · a pull request into the public repository, use-cases folder

Template pack fidelity QA harness

WHO IT SERVES

Anyone shipping template packs, internal tooling.

WHAT TO BUILD

Every pack in this group needs proof it survives contact. Build a harness that takes any template pack, fills it with a generated data set, applies a standard battery of edits, exports to Word and PDF, and reports what drifted: fonts, table borders, list numbering, header and footer content, page geometry. Output a pass or fail scorecard per template with the specific drift named. Run it against three of the packs in this group and publish the scorecard. a fidelity test harness and a published scorecard method, a strong quality-engineering portfolio piece.

WHAT STRONG LOOKS LIKE

the harness detects seeded drift reliably, the scorecard is specific enough to act on, and the candidate reports failures found in real packs rather than a clean sheet.

BAND S3 · API, templates, export, Multi-document · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Clause-library extraction from a firm's own past documents

WHO IT SERVES

Legal operations, in-house counsel.

WHAT TO BUILD

Most firms do not have a clause library. They have 500 signed contracts. Build the tool that closes that gap: ingest a folder of past agreements, cluster the recurring clauses, pick the most common formulation of each, and emit a starter clause library with a frequency count and the source documents each formulation came from. Run it on a public contract corpus and publish the resulting library plus the frequency table. a clause-mining tool and a published corpus study, which is a serious legal-tech engineering credential.

WHAT STRONG LOOKS LIKE

clustering groups genuinely equivalent clauses that use different wording, the frequency table is honest about long-tail variants, and every library entry links back to its sources.

BAND S3 · API, search, Multi-document, export · a pull request into the public repository, extensions folder

Template pack versioning and changelog kit

WHO IT SERVES

Enterprise template owners, brand and compliance teams.

WHAT TO BUILD

Templates change and nobody knows what changed. Build the kit: a version scheme, a diff report between two versions of the same template rendered as a readable change summary rather than markup, a changelog document that accumulates those summaries, and an impact list naming which downstream documents were generated from the superseded version. a template governance kit and a written versioning method.

WHAT STRONG LOOKS LIKE

the change summary is readable by a non-technical template owner, the impact list is correct, and the changelog reads like a document rather than a log dump.

BAND S2 · templates, Multi-document, chat, export · a pull request into the public repository, use-cases folder

AI system technical-documentation pack

WHO IT SERVES

AI product teams, compliance leads, anyone selling AI into Europe.

WHAT TO BUILD

Build the documentation set an AI product is increasingly asked to produce: an intended-purpose statement, a data-provenance description, a risk assessment, a human-oversight description, a performance and limitation statement, and a change log. Then build the maintenance loop, which is the real problem: when the system changes, take a change description and update only the affected sections across all six documents, leaving the rest byte-identical. an AI documentation pack and a change-propagation demo. This is a credential in a field that barely has practitioners yet.

WHAT STRONG LOOKS LIKE

the change propagation touches exactly the affected sections and the candidate proves it with a diff over the whole set. The documents read as though a compliance lead wrote them, not as though a checklist was expanded.

BAND S3 · templates, Multi-document, chat, Review, export · a pull request into the public repository, use-cases folder

Model card and system card generator

WHO IT SERVES

ML teams, research groups, platform providers.

WHAT TO BUILD

Build a generator that turns a structured description of a model or system into a published card: purpose, training data summary, evaluation results as real tables, limitations, out-of-scope uses, and a maintenance contact. Include the awkward requirement that makes it interesting: evaluation tables must be regenerated from a results file so the card can be refreshed after every evaluation run without a human retyping numbers. a card generator repo under their name, which sits well on any ML engineering profile.

WHAT STRONG LOOKS LIKE

rerunning with new results updates only the tables and the dated fields. Markdown and PDF outputs both look right. Nothing is invented when a field is absent from the input; it is marked as not provided.

BAND S2 · API, templates, chat, export · a pull request into the public repository, extensions folder

Data protection impact assessment builder

WHO IT SERVES

Privacy teams, data protection officers, SaaS compliance.

WHAT TO BUILD

Build a guided assessment that turns answers about a processing activity into a full assessment document: necessity and proportionality analysis, risk table with likelihood and severity, mitigations mapped to each risk, and a residual-risk conclusion. The document must be editable afterwards without the structure collapsing, because these get reviewed and revised for weeks. a working privacy assessment builder, a direct credential for privacy or governance roles.

WHAT STRONG LOOKS LIKE

the risk table stays a real table through several rounds of edits, mitigations stay linked to their risks when rows are added or removed, and the conclusion updates when the residual risk changes.

BAND S2 · templates, chat, Review, export · a pull request into the public repository, use-cases folder

Processing-activity register maintainer

WHO IT SERVES

Privacy operations, compliance managers.

WHAT TO BUILD

A register is a living table that nobody keeps current. Build the maintainer: ingest an existing register document, accept a change feed of new or retired processing activities, and update the register in place, preserving the surrounding formatting and appending a dated revision note. Produce a quarterly delta report saying what was added, changed, and removed. a register maintenance tool and a long-table stability study.

WHAT STRONG LOOKS LIKE

a 200-row register survives repeated updates with no drift in column widths, ordering, or styling, and the delta report is generated from the document rather than from the input feed.

BAND S2 · API, chat, search, export · a pull request into the public repository, extensions folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Subject access request response assembler

WHO IT SERVES

Privacy operations, customer support in regulated firms.

WHAT TO BUILD

Build the response packet: a cover letter explaining what was found and what was withheld and why, an index of the disclosed items, and the disclosure bundle itself with third-party details redacted. Redaction must be proposed for human confirmation rather than applied silently. Use entirely fabricated records. a response assembly workflow with a documented redaction review step.

WHAT STRONG LOOKS LIKE

redaction candidates are surfaced with their reason and confirmed by a human before anything is applied, the index matches the bundle exactly, and the cover letter states withholding grounds specifically rather than generically.

BAND S2 · Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Statement of applicability and control-narrative maintainer

WHO IT SERVES

Information security, compliance engineers.

WHAT TO BUILD

Build the maintainer for the two documents that carry an information security programme: the statement of applicability naming which controls apply and why, and the control narratives describing how each is implemented. Then the loop: given a control that changed, update its narrative and its applicability row together, and generate a reviewer summary of every change made this cycle. an information security documentation maintainer, a direct credential for security or compliance engineering roles.

WHAT STRONG LOOKS LIKE

applicability and narrative never fall out of sync, exclusions carry a written justification, and the reviewer summary is derived from the documents.

BAND S2 · templates, Multi-document, search, chat, export · a pull request into the public repository, use-cases folder

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Annual policy-set refresh with a change report

WHO IT SERVES

Compliance managers, security leads.

WHAT TO BUILD

Distinct from generating a policy set: this refreshes one that already exists. Take a policy set of 12 documents, apply a change brief (new tooling, a changed retention period, a renamed team, a new obligation), and produce the refreshed set plus a board-ready change report saying what changed in each policy and why. The unchanged policies must come back untouched, and the candidate must prove it. a policy refresh method with a measured no-drift proof.

WHAT STRONG LOOKS LIKE

a structural proof that untouched policies are untouched, a change report a board would accept, and correct handling of a change that affects several policies inconsistently.

BAND S2 · Multi-document, chat, Review, export · a pull request into the public repository, use-cases folder

Audit-trail evidence pack

WHO IT SERVES

Auditors, regulated operations, quality management.

WHAT TO BUILD

Auditors ask who changed what, when, and who approved it. Build the demonstration: run a document through a realistic review cycle using proposed changes, approvals, denials, and a rollback, then assemble an evidence pack showing the full trail as a readable document with the before and after of each approved change and the reason each denial was recorded. Explain in the pack which parts of the trail are produced by the platform and which the integrator must record themselves. an audit-evidence method and a worked example, valuable in any regulated-industry role.

WHAT STRONG LOOKS LIKE

the trail is complete and readable by a non-technical auditor, the rollback is represented honestly, and the boundary between platform-provided and integrator-provided evidence is stated plainly rather than blurred.

BAND S2 · Review, chat, export, API · a pull request into the public repository, use-cases folder

Accessibility conformance report builder

WHO IT SERVES

Enterprise sales engineering, public-sector suppliers.

WHAT TO BUILD

Build a generator for the conformance report enterprise and government procurement asks for: criterion-by-criterion tables with a conformance level and remarks per row. Feed it a findings file and produce the report; feed it an updated findings file and update only the changed rows, keeping the remarks that still apply. Ship it with a fabricated product's findings as the example. a conformance report generator, directly useful in sales engineering and compliance roles.

WHAT STRONG LOOKS LIKE

the criterion table is complete and correctly structured, incremental updates preserve untouched remarks, and the output is presentable to a procurement officer without hand editing.

BAND S2 · API, templates, chat, export · a pull request into the public repository, use-cases folder

Records retention schedule and disposition certificates

WHO IT SERVES

Records management, legal operations, regulated industries.

WHAT TO BUILD

Build the retention schedule as a maintained document, category by category, with the retention period, the trigger event, the legal basis, and the disposition method. Then build the certificate flow: when a category is disposed of, generate a dated disposition certificate that quotes the schedule row it acted on, and append a line to the schedule's disposal history without disturbing the rest. a records governance kit, useful in any regulated operations role.

WHAT STRONG LOOKS LIKE

the certificate quotes the live schedule rather than a copy, the history append does not reflow the schedule, and legal-hold exceptions are represented rather than ignored.

BAND S1 · templates, chat, Multi-document, export · a pull request into the public repository, use-cases folder

Clinical correspondence pack, synthetic data only

WHO IT SERVES

Healthcare, clinical software, medical scribing.

WHAT TO BUILD

Build the correspondence set a clinician produces daily: referral letter, discharge summary, results letter to a patient in plain language, and a prior-authorisation letter, all on one letterhead. Drive them from a fabricated encounter record. Every letter must pass through explicit human sign-off before it is considered final, and the pack must include a written statement that this is a demonstration on fabricated records and not a clinical or certified system. a clinical documentation demo built to a safety standard, which is a strong health-tech portfolio piece.

WHAT STRONG LOOKS LIKE

the patient letter is genuinely plain language while the referral is genuinely clinical, sign-off is enforced rather than decorative, and no real or realistic-looking identifiers appear anywhere.

BAND S2 · templates, chat, Review, export · a pull request into the public repository, use-cases folder

Informed-consent multi-language kit

WHO IT SERVES

Clinical research, ethics committees, research operations.

WHAT TO BUILD

Build a consent form kit in three languages where the versions must stay structurally identical section by section, because that is what an ethics committee checks. Then build the amendment flow: when one section changes in the source language, the other versions surface as out of date and are updated with the structure preserved. Include a structural comparison report across all three. a multi-language consent kit with a parity checker, unusual and credible in research operations.

WHAT STRONG LOOKS LIKE

the structural comparison is generated and catches a deliberately desynchronised section, reading level is controlled in every language, and layout survives across scripts.

BAND S3 · templates, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Device technical-file section assembler

WHO IT SERVES

Medical device, hardware regulatory affairs.

WHAT TO BUILD

Build an assembler for a technical file: device description, intended purpose, risk management summary, verification and validation summary, labelling, and a post-market surveillance plan, assembled as one navigable document with a table of contents and internal cross-references that survive export. Then prove the hard part: change one section and confirm that cross-references and numbering elsewhere stay correct. a regulatory assembly demo with a cross-reference integrity proof.

WHAT STRONG LOOKS LIKE

cross-references genuinely resolve after edits and after export, the table of contents regenerates, and the candidate shows a case where an edit would break numbering and demonstrates it does not.

BAND S3 · Multi-document, chat, search, export · a pull request into the public repository, use-cases folder

Safety narrative drafter, fabricated cases only

WHO IT SERVES

Life sciences, pharmacovigilance, safety operations.

WHAT TO BUILD

Build a drafter that turns a structured fabricated adverse-event record into a written case narrative in the conventional order: patient background, event onset, treatment, outcome, and assessment. Then build the correction loop, which is where the real work is: when a field in the record is corrected, the narrative must be amended in place with the previous wording retained as a tracked change rather than silently replaced. a regulated-writing demo with a tracked-correction workflow.

WHAT STRONG LOOKS LIKE

narratives read as though written by a safety writer, corrections appear as tracked amendments, and any field absent from the record is stated as unknown rather than invented. All records visibly fabricated.

BAND S3 · API, chat, Review, export · a pull request into the public repository, use-cases folder

Cross-reference integrity checker for long filings

WHO IT SERVES

Regulatory affairs, legal filings, standards bodies.

WHAT TO BUILD

Build a checker that reads a long filing and reports every broken internal reference: a section referenced but absent, a numbering gap, a defined term used before its definition, an appendix referenced but not attached, and a cross-reference pointing at the wrong section after a renumber. Run it over a public long document and publish the findings with the fix applied to three of them. an integrity checker plus published findings on a real document, a strong engineering artifact.

WHAT STRONG LOOKS LIKE

the checker finds real problems in a real public document, not only seeded ones, and the fixes applied do not disturb anything else.

BAND S3 · API, search, chat, Multi-document, export · a pull request into the public repository, extensions folder

Operational resilience documentation pack

WHO IT SERVES

Financial services technology, critical infrastructure suppliers.

WHAT TO BUILD

Build the documentation an operationally critical supplier is asked for: a service inventory with criticality ratings, dependency mapping including third parties, tolerance thresholds for disruption, a testing plan, and an exit and substitutability plan. The dependency map must be a generated diagram inside the document, regenerated when the inventory changes. a resilience documentation pack with a live diagram pipeline.

WHAT STRONG LOOKS LIKE

the diagram is generated from the inventory rather than drawn once, changing the inventory updates both the table and the diagram, and the exit plan is specific rather than aspirational.

BAND S3 · templates, Multi-document, images, chat, export · a pull request into the public repository, use-cases folder

Suspicious activity narrative drafter, fabricated cases only

WHO IT SERVES

Financial crime compliance, fintech.

WHAT TO BUILD

Build a drafter that turns a fabricated alert record into a written narrative covering who, what, when, where, why, and how, in the order investigators are trained to write. Then the important part: a completeness check that refuses to finalise a narrative missing any of the six, naming what is missing rather than filling it in. Every record and entity fabricated and visibly labelled as such. a regulated-narrative drafter with an enforced completeness gate.

WHAT STRONG LOOKS LIKE

the refusal behaviour is real and specific, narratives read like investigator writing, and nothing about the demo could be mistaken for real case data.

BAND S2 · API, templates, chat, Review, export · a pull request into the public repository, use-cases folder

Cross-clause contradiction detector

WHO IT SERVES

Legal and compliance

WHAT TO BUILD

Finds places where clause seven contradicts clause twelve, explains the conflict, and proposes a resolution as an edit.

WHAT STRONG LOOKS LIKE

it finds genuine contradictions in a real contract and does not flag ordinary tension as conflict. Precision is the whole test.

BAND S3 · search, chat, Review · a pull request into the public repository, use-cases folder

Weakest-section critique engine

WHO IT SERVES

Writing and proposals

WHAT TO BUILD

Ranks the sections of a document by how weak they are against a stated goal, with reasons, then rewrites the worst on request.

WHAT STRONG LOOKS LIKE

the ranking is defensible and stable across runs, and the rewrite measurably addresses the stated weakness.

BAND S2 · search, chat, Review · a pull request into the public repository, use-cases folder

Document question answering with span citations

WHO IT SERVES

Cross-industry

WHAT TO BUILD

Answers questions about a document with the exact supporting span highlighted, and refuses when the document does not say.

WHAT STRONG LOOKS LIKE

it says "not stated in this document" when that is the truth, rather than inferring plausibly.

BAND S2 · search, chat · a pull request into the public repository, use-cases folder

Brief-to-section matcher

WHO IT SERVES

Careers, sales, grants

WHAT TO BUILD

Given a brief and a set of candidate sections, picks the best fit with an explanation and adapts it.

WHAT STRONG LOOKS LIKE

the choice is explained in terms of the brief's actual requirements, and a poor fit is admitted.

BAND S2 · search, Multi-document, chat · a pull request into the public repository, use-cases folder

Obligation graph builder

WHO IT SERVES

Legal and procurement

WHAT TO BUILD

Extracts who owes what to whom by when, builds a graph, and writes a summary document that stays linked to the source.

WHAT STRONG LOOKS LIKE

the graph is correct on a real contract, including conditional obligations, and it renders as a diagram.

BAND S3 · search, Multi-document, images, export · a pull request into the public repository, use-cases folder

Evidence-grounded fact checker

WHO IT SERVES

Research and journalism

WHAT TO BUILD

Every factual claim in a document is traced to an attached source, and unsupported claims are marked in place.

WHAT STRONG LOOKS LIKE

it distinguishes unsupported from contradicted, which are different problems.

BAND S3 · search, Multi-document, Review · a pull request into the public repository, use-cases folder

Terminology consistency auditor

WHO IT SERVES

Technical writing and legal

WHAT TO BUILD

Finds terms defined once and used inconsistently across a document set, then normalises them with approval.

WHAT STRONG LOOKS LIKE

it respects deliberate distinctions instead of flattening them.

BAND S2 · search, Multi-document, Review · a pull request into the public repository, use-cases folder

Numeric and date consistency auditor

WHO IT SERVES

Finance and legal

WHAT TO BUILD

Checks that totals add up, dates are consistent, and figures repeated in prose match the tables.

WHAT STRONG LOOKS LIKE

it catches a deliberately planted inconsistency in a long financial document and explains the arithmetic.

BAND S2 · search, chat, Review · a pull request into the public repository, use-cases folder

Redundancy and bloat finder

WHO IT SERVES

Long-form

WHAT TO BUILD

Identifies sections that repeat each other and proposes a merge that keeps every distinct point.

WHAT STRONG LOOKS LIKE

nothing unique is lost in the merge. Prove it with a before-and-after point inventory.

BAND S2 · search, chat, Review · a pull request into the public repository, use-cases folder

Argument structure mapper

WHO IT SERVES

Academic and legal

WHAT TO BUILD

Maps claims, supporting evidence and gaps, renders the structure as a diagram, and proposes edits to close the gaps.

WHAT STRONG LOOKS LIKE

the map is a genuine analysis, and the identified gaps are ones a reviewer would also raise.

BAND S3 · search, images, chat · a pull request into the public repository, use-cases folder

Weighted precedent retriever

WHO IT SERVES

Legal and consulting

WHAT TO BUILD

Finds the closest prior document weighted by recency and context, then adapts it to the new matter.

WHAT STRONG LOOKS LIKE

weighting is explainable and the adaptation catches every detail that must change.

BAND S2 · search, templates, memory · a pull request into the public repository, use-cases folder

Multi-hop question answering across a set

WHO IT SERVES

Research and diligence

WHAT TO BUILD

Answers questions that require combining facts from several documents, showing the chain.

WHAT STRONG LOOKS LIKE

the reasoning chain is shown and each hop is verifiable.

BAND S3 · search, Multi-document · a pull request into the public repository, use-cases folder

Citation verification pass

WHO IT SERVES

Academic and journalism

WHAT TO BUILD

Checks that each cited source actually supports the claim it is attached to.

WHAT STRONG LOOKS LIKE

it catches a citation that exists but does not say what the sentence claims, which is the common real failure.

BAND S3 · search, Multi-document, Review · a pull request into the public repository, use-cases folder

Material-change analysis between versions

WHO IT SERVES

Legal and procurement

WHAT TO BUILD

Compares two versions and reports what materially changed, who it favours, and what to push back on.

WHAT STRONG LOOKS LIKE

it separates substance from formatting and takes a defensible position on impact.

BAND S3 · Multi-document, search, export · a pull request into the public repository, use-cases folder

Tracked-changes preservation suite

WHAT TO BUILD

Documents carrying existing revision marks are edited, and the suite verifies that prior marks, authorship and the new change all survive.

WHAT STRONG LOOKS LIKE

covers nested and conflicting revisions, which is where implementations break.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Notes, citations and cross-reference suite

WHAT TO BUILD

Footnotes, endnotes, numbering and internal references are checked for survival and correctness after edits.

WHAT STRONG LOOKS LIKE

it tests an insertion in the middle, which forces renumbering, not only appends.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Long-document scaling study

WHAT TO BUILD

Same task at ten, fifty, two hundred and five hundred pages, measuring quality, time and cost.

WHAT STRONG LOOKS LIKE

it reports where quality degrades honestly rather than stopping at the flattering size.

BAND S3 · API, search · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Table fidelity benchmark

WHAT TO BUILD

Merged cells, nested tables, spanning headers and numeric formatting through an edit and export cycle.

WHAT STRONG LOOKS LIKE

it includes the pathological cases real documents contain.

BAND S2 · API, export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Visual content survival benchmark

WHAT TO BUILD

Images, diagrams and typeset math through ingest, edit and export.

WHAT STRONG LOOKS LIKE

it verifies that a diagram remains regenerable and math remains editable, not merely that a picture is present.

BAND S2 · images, export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Section-targeting benchmark

WHAT TO BUILD

Given an instruction, measures whether the correct section is identified, on documents with repetitive or ambiguous headings.

WHAT STRONG LOOKS LIKE

ambiguous cases are included deliberately and the failure taxonomy is published.

BAND S2 · search, chat · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Approval decision fidelity benchmark

WHAT TO BUILD

Verifies that approving a subset applies exactly that subset and denying applies nothing.

WHAT STRONG LOOKS LIKE

includes partial approvals, repeated decisions and out-of-order decisions.

BAND S2 · Review, API · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Revert correctness benchmark

WHAT TO BUILD

Verifies that reverting a turn restores the prior state exactly, including formatting, across many turn shapes.

WHAT STRONG LOOKS LIKE

byte-level verification, and it covers reverting a turn that was itself a revert.

BAND S2 · Review, API · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Concurrent editing benchmark

WHAT TO BUILD

Two writers, one document, overlapping edits. Measures whether both survive and how conflicts are surfaced.

WHAT STRONG LOOKS LIKE

a reproducible harness, and it reports any case where a write is silently lost.

BAND S3 · Multi-document, API · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Cost-per-outcome across model tiers

WHAT TO BUILD

For a fixed set of tasks, measures quality and cost at each selectable tier and recommends a default per task type.

WHAT STRONG LOOKS LIKE

the recommendation is evidence-backed and includes tasks where the cheapest tier is genuinely sufficient.

BAND S2 · chat, API · a pull request into the public repository, extensions folder

Public golden corpus

WHAT TO BUILD

A licence-clean, redistributable document set spanning formats, languages, lengths and structural features, with per-document annotations.

WHAT STRONG LOOKS LIKE

every document's licence is verified and recorded. This is the hard part and the reason it is valuable.

BAND S3 · export, attachments · a pull request into the public repository, extensions folder

Judging rubric with agreement study

WHAT TO BUILD

A rubric for scoring document edit quality, plus an inter-rater agreement study between human and automated judges.

WHAT STRONG LOOKS LIKE

reports agreement honestly, including where automated judging is unreliable.

BAND S3 · export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Reproducibility kit

WHAT TO BUILD

One command reproduces every published number, with pinned inputs and a cost estimate before it runs.

WHAT STRONG LOOKS LIKE

an independent person reproduces the results without help. That is the only test that counts.

BAND S2 · API · a pull request into the public repository, extensions folder

India employment and HR pack

WHAT TO BUILD

Offer letters, appointment letters, statutory policies and separation documents shaped to Indian practice.

WHAT STRONG LOOKS LIKE

the statutory documents that Indian employers actually need, correctly named.

BAND S1 · templates · a pull request into the public repository, use-cases folder

United States employment pack with state variants

WHAT TO BUILD

Offer letters and handbooks with state-specific required notices.

WHAT STRONG LOOKS LIKE

state variants differ where the law differs and are identical where it does not.

BAND S2 · templates, Multi-document · a pull request into the public repository, use-cases folder

Multi-language contract pack

WHAT TO BUILD

The same contract set produced in several languages with layout preserved and legal meaning checked.

WHAT STRONG LOOKS LIKE

a native speaker finds each version usable, and the structure matches across languages.

BAND S2 · templates, chat, export · a pull request into the public repository, use-cases folder

Public tender response pack

WHAT TO BUILD

The document set a government tender response requires, with the compliance annexures.

WHAT STRONG LOOKS LIKE

it covers the annexures that get bids disqualified.

BAND S2 · templates, Multi-document, export · a pull request into the public repository, use-cases folder

Clinical forms pack

WHAT TO BUILD

Intake, consent, discharge, referral and instruction documents in plain language.

WHAT STRONG LOOKS LIKE

the patient-facing documents are genuinely readable at a low reading level.

BAND S1 · templates, export · a pull request into the public repository, use-cases folder

Startup incorporation and governance pack

WHAT TO BUILD

Founder agreements, board resolutions, cap table memos, shareholder communications.

WHAT STRONG LOOKS LIKE

the documents cross-reference each other correctly.

BAND S1 · templates, Multi-document · a pull request into the public repository, use-cases folder

Real estate pack by jurisdiction

WHAT TO BUILD

Leases and disclosures where the required disclosures differ meaningfully by market.

WHAT STRONG LOOKS LIKE

the jurisdiction-specific disclosures are correct and dated.

BAND S2 · templates, Multi-document · a pull request into the public repository, use-cases folder

Academic style pack

WHAT TO BUILD

Major citation styles plus a set of journal and thesis templates, switchable on one manuscript.

WHAT STRONG LOOKS LIKE

switching style rewrites every reference correctly including in notes.

BAND S2 · templates, export · a pull request into the public repository, use-cases folder

Accessible document template pack

WHAT TO BUILD

Templates built to produce properly structured, accessible output by default.

WHAT STRONG LOOKS LIKE

documents produced from these templates pass an accessibility check without remediation.

BAND S2 · templates, export · a pull request into the public repository, use-cases folder

Insurance document pack

WHAT TO BUILD

Policy wordings, endorsements, certificates and claim correspondence with variable structures.

WHAT STRONG LOOKS LIKE

the variable substitution handles nested conditions rather than only flat fields.

BAND S1 · templates, export · a pull request into the public repository, use-cases folder

Construction contract pack

WHAT TO BUILD

Standard-form-shaped agreements, change orders, requests for information and payment applications.

WHAT STRONG LOOKS LIKE

the change order and payment application reference the base agreement's clause numbering correctly.

BAND S2 · templates, Multi-document · a pull request into the public repository, use-cases folder

Multi-jurisdiction employment document pack - jurisdiction: Europe

WHO IT SERVES

HR, people ops, employer of record.

WHAT TO BUILD

Build offer letter, employment agreement, and policy acknowledgement templates for four jurisdictions with genuinely different requirements, plus a variables sheet naming what changes per jurisdiction and why. Drive a session that takes one role and one jurisdiction and produces the correct set on a single branded template. Every jurisdiction-specific difference must be annotated in a companion notes document, sourced. a jurisdiction-annotated HR document pack, which is a credible people-ops or HR-tech portfolio artifact. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

the four jurisdictions are chosen because they differ in interesting ways, not because they were easy. The annotations cite public sources. The branded template survives all four fills with identical layout, proven by exporting all four and diffing the styling.

BAND S2 · templates, chat, Multi-document, export · a pull request into the public repository, use-cases folder

Multi-jurisdiction employment document pack - jurisdiction: United States, with state variants

WHO IT SERVES

HR, people ops, employer of record.

WHAT TO BUILD

Build offer letter, employment agreement, and policy acknowledgement templates for four jurisdictions with genuinely different requirements, plus a variables sheet naming what changes per jurisdiction and why. Drive a session that takes one role and one jurisdiction and produces the correct set on a single branded template. Every jurisdiction-specific difference must be annotated in a companion notes document, sourced. a jurisdiction-annotated HR document pack, which is a credible people-ops or HR-tech portfolio artifact. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

the four jurisdictions are chosen because they differ in interesting ways, not because they were easy. The annotations cite public sources. The branded template survives all four fills with identical layout, proven by exporting all four and diffing the styling.

BAND S2 · templates, chat, Multi-document, export · a pull request into the public repository, use-cases folder

Board and corporate-governance document pack - entity type: startup incorporation and founder governance

WHO IT SERVES

Startups, nonprofits, company secretaries.

WHAT TO BUILD

Build the governance set: board charter, committee charters, written resolutions, notice of meeting, minutes template, and a conflict-of-interest declaration. Then build the cycle: a session that takes an agenda and a transcript and produces the notice, the minutes, and the resolutions as three linked documents with consistent numbering and dates across all three. a governance pack plus a multi-document consistency demo, which is a strong ops or legal-ops artifact. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

the three documents stay consistent when the agenda changes, resolution numbering does not drift, and the minutes read like minutes rather than a transcript summary.

BAND S2 · templates, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Board and corporate-governance document pack - entity type: nonprofit and charity governance, by country

WHO IT SERVES

Startups, nonprofits, company secretaries.

WHAT TO BUILD

Build the governance set: board charter, committee charters, written resolutions, notice of meeting, minutes template, and a conflict-of-interest declaration. Then build the cycle: a session that takes an agenda and a transcript and produces the notice, the minutes, and the resolutions as three linked documents with consistent numbering and dates across all three. a governance pack plus a multi-document consistency demo, which is a strong ops or legal-ops artifact. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

the three documents stay consistent when the agenda changes, resolution numbering does not drift, and the minutes read like minutes rather than a transcript summary.

BAND S2 · templates, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Refusal-honesty evaluation

WHO IT SERVES

Buyers deciding whether a system can be trusted unattended, and the teams who have to answer them.

WHAT TO BUILD

Most evaluations measure what a system does when it can succeed. This one measures what it says when it cannot: a document too large to process, a format that is not supported, an instruction that is already satisfied, a partial success where some sections changed and some did not, a run that hit a limit. Each case is scored on whether the system refused at all, whether the stated reason was accurate, and whether the refusal was actionable.

WHAT STRONG LOOKS LIKE

the corpus provokes each failure deliberately and reproducibly; the score separates an honest refusal from a silent partial success, which is the failure that actually damages trust; expandable capacity limits are scored as expandable rather than as hard failures; our own dishonest cases are published alongside anyone else's.

BAND S2 · API, chat, export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Document-borne prompt-injection resistance evaluation

WHO IT SERVES

Security reviewers signing off on pointing an agent at outside documents

WHAT TO BUILD

A corpus of documents carrying instructions aimed at the system rather than the reader, planted everywhere a document can hide text: body prose, tables, comments, footnotes and endnotes, headers and footers, tracked-change deletions, image alternative text, and metadata. A scorer runs an ordinary editing task over each and reports whether the document's embedded instruction was obeyed, mentioned, or correctly ignored.

WHAT STRONG LOOKS LIKE

every part of a document is covered, not only body text, because the non-body parts are where this is never tested; the score distinguishes obeying an injection from noticing and reporting it; the corpus is licence-clean and redistributable; results published per part and per system, including our own failures.

BAND S2 · API, chat, search, export · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Template inference from finished documents

WHO IT SERVES

Anyone with forty documents made from a template nobody can find

WHAT TO BUILD

Given several finished documents of the same kind, infer the template: the invariant structure, the wording that never changes, the slots that vary and what type each slot holds, the sections that are conditional and the condition that seems to govern them. Output is a usable template plus a confidence report naming what the sample was too small to decide.

WHAT STRONG LOOKS LIKE

a slot is only proposed when the evidence supports it, and the uncertain ones are listed rather than guessed; regenerating each source document from the inferred template and its own values reproduces it closely, and the differences are reported honestly; the conditional sections are correct on both branches; the template carries the source documents' formatting, not a clean-room approximation.

BAND S2 · templates, Multi-document, search, export · a pull request into the public repository, use-cases folder

Derivation-lineage map for a copied document estate

WHO IT SERVES

Knowledge managers and auditors facing hundreds of copied-from-each-other files

WHAT TO BUILD

Reconstruct the family tree. Which document was copied from which, in what order, and - the part that matters - which descendants never received a change that was made upstream. The output is a rendered lineage diagram plus a divergence report: per upstream change, the descendants that took it and the descendants that did not.

WHAT STRONG LOOKS LIKE

the tree is inferred from the documents alone, with no filenames, timestamps or folder structure trusted as evidence; a deliberately planted late fork is placed correctly; the divergence report distinguishes a descendant that deliberately diverged from one that was simply missed; the diagram stays legible at a hundred documents.

BAND S2 · search, Multi-document, images, export · a pull request into the public repository, use-cases folder

GROWTH NETWORKING

Segment map: AI products that can reason but cannot change the file

WHO IT SERVES

Growth, market research.

WHAT TO BUILD

Build the map of the primary segment: AI products whose model produces an answer that then has to become a document a human edits. Cover at least eight verticals, and for each one name the product category, the document that gets produced, who edits it afterwards, and the observable public evidence that the gap exists (their own documentation, changelog, or help centre). Assemble it as a document with a summary table and a one-page brief per vertical. No individual people, no contact details, company level only. a market segmentation study with a stated evidence standard, a strong strategy and research portfolio piece.

WHAT STRONG LOOKS LIKE

every gap claim is backed by something publicly observable and cited, not inferred from a marketing page. The candidate names the verticals where the gap is weak or absent, which is more useful than eight confident yeses.

BAND S2 · chat, Multi-document, search, export · the submission form, into Drive

Ideal-customer scoring rubric and a scored account list

WHO IT SERVES

Growth, revenue operations.

WHAT TO BUILD

Build a scoring rubric with weighted, observable criteria, then apply it to 80 companies and publish the scored list with the score decomposition visible per row. Criteria must be checkable by a stranger from public sources. Include a section on the criteria the candidate tried and discarded and why. Companies only, no named individuals, no contact information anywhere in the deliverable. a scoring methodology and a worked application of it, directly transferable to any revenue-operations role.

WHAT STRONG LOOKS LIKE

two different people applying the rubric to the same company would land within a point of each other. The discarded criteria section shows real iteration. The bottom of the list is as carefully scored as the top.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Teardown of one competitor's content and discovery machinery

WHO IT SERVES

Growth, competitive intelligence.

WHAT TO BUILD

Pick one competitor that currently outranks us on the measured queries. Take their discovery machinery apart: how many pages, of what types, on what publishing cadence, targeting which query shapes, with what internal structure, what evidence they put on a page, where they get cited, and which of it looks like it is working. Assemble it as a document with a summary, a page-type inventory, and a list of the specific plays worth copying and the ones worth refusing. a competitive teardown document, the standard proof-of-work artifact for a growth role.

WHAT STRONG LOOKS LIKE

the inventory is counted rather than described, the "worth refusing" list exists and has reasons, and the candidate distinguishes what is working from what merely exists.

BAND S2 · chat, search, Multi-document, export · the submission form, into Drive

Roundup and listicle inclusion research

WHO IT SERVES

Growth, partnerships research.

WHAT TO BUILD

Several third-party roundups rank on the "best tool" query shapes where we are absent. Map them: which roundups rank for which queries, who publishes each, how entries appear to be selected, whether inclusion appears paid or editorial, what the stated submission mechanism is if any, how often each is updated, and what an entry contains. Produce a ranked table of which are worth pursuing and why. Research only. The candidate contacts nobody and submits nothing; we handle every approach. a placement-research method and a worked map, a directly reusable growth skill.

WHAT STRONG LOOKS LIKE

paid placement is distinguished from editorial with evidence rather than assumption, the ranking is justified by reach and update frequency, and the candidate flags any roundup that would damage us to appear in.

BAND S1 · chat, search, export · the submission form, into Drive

Agent tool registry and directory landscape map

WHO IT SERVES

Growth working with a developer product. (an engineer will find it easy).

WHAT TO BUILD

Map every registry, directory, and marketplace where agent tool servers are listed: what each requires to list, what a listing contains, how entries are ranked or sorted, how often the directory is refreshed, and how much traffic or visibility it plausibly carries. Note which ones currently list competitors and how those entries are written. Produce a prioritised table plus a template listing draft for the top five. We submit; the candidate drafts. a channel map and listing drafts, evidence of distribution thinking.

WHAT STRONG LOOKS LIKE

requirements are captured precisely enough to act on without revisiting the site, competitor entries are quoted, and the drafts are written to each directory's actual conventions rather than one generic blurb.

BAND S1 · chat, templates, export · the submission form, into Drive

Developer community map

WHO IT SERVES

Growth, community, developer relations.

WHAT TO BUILD

Map where developers ask document-editing questions: forums, question sites, chat communities, subreddits, and discussion boards. For each, capture the size, the activity level, the rules on vendor participation stated in their own words, the moderation posture, and the shape of the questions that come up. Rank by fit and note which communities would react badly to any vendor presence at all. Observation only. The candidate joins nothing on our behalf and posts nowhere. a community landscape map with a participation-risk assessment, a genuinely senior community-strategy artifact.

WHAT STRONG LOOKS LIKE

the vendor rules are quoted from each community's own policy, not summarised, and the candidate is willing to write "stay out of this one" and explain why.

BAND S1 · chat, export · the submission form, into Drive

Hiring-signal mining

WHO IT SERVES

Growth, market intelligence.

WHAT TO BUILD

Companies advertise their gaps in their job posts. Mine public postings for roles that imply a document-automation problem: document engineers, template specialists, contract operations, report automation, and product roles that name document generation. Build a table of companies with the quoted phrase that qualified them, the date, and a one-line read on what it implies about their build-versus-buy position. Companies only, never the named hiring manager or recruiter. a signal-mining method and a worked dataset, a strong market-intelligence portfolio piece.

WHAT STRONG LOOKS LIKE

every row carries the quote that qualified it, false positives are visibly filtered, and the candidate reports the search patterns that produced good yield so the exercise is repeatable.

BAND S1 · chat, search, export · a pull request into the public repository, use-cases folder

Funding and launch signal mining

WHO IT SERVES

Growth, market intelligence.

WHAT TO BUILD

Map recently funded or recently launched AI products whose stated purpose implies a document output: their category, what their own materials say the product produces, whether public evidence suggests they solved the document step or stopped short, and the timing read on when a document problem would become urgent for them. Company level only. Produce a table plus five one-page briefs on the clearest fits. a market-timing research method and worked briefs.

WHAT STRONG LOOKS LIKE

the "did they solve it already" question is answered from evidence, the timing argument is reasoned rather than asserted, and companies that clearly built their own good solution are marked as such rather than padded into the list.

BAND S2 · chat, Multi-document, search, export · a pull request into the public repository, use-cases folder

Integration marketplace and plugin ecosystem map

WHO IT SERVES

Growth, partnerships research.

WHAT TO BUILD

Map the platforms whose marketplaces would carry a document-editing integration: automation platforms, no-code builders, workspace app directories, and developer tool marketplaces. Capture the listing requirements, the review process and its typical duration, technical prerequisites, revenue-share terms where published, and how discoverable a new listing actually is. Rank by effort against plausible reach. We apply; the candidate maps. a partnership channel analysis with effort-to-reach reasoning.

WHAT STRONG LOOKS LIKE

requirements are captured to the level of what a build would actually have to satisfy, the discoverability assessment is evidenced rather than assumed, and marketplaces that are effectively dead are called dead.

BAND S2 · chat, templates, export · the submission form, into Drive

Review and comparison site landscape

WHO IT SERVES

Growth.

WHAT TO BUILD

Map the software review and comparison sites that rank on our category's queries: which categories we would fit, what a profile contains, how reviews are gathered, what the free versus paid tiers change, how heavily each site is cited by answer engines, and how competitors' profiles are written. Produce a prioritised recommendation with a draft profile for the top two. No profile is created and nobody is asked for a review; drafts only. a channel evaluation with drafted assets.

WHAT STRONG LOOKS LIKE

the citation analysis is checked rather than assumed, the category-fit question is answered specifically (which category, and what we would be compared against there), and the drafts read like our voice rather than boilerplate.

BAND S1 · chat, templates, export · the submission form, into Drive

Creator and independent media landscape map

WHO IT SERVES

Growth, partnerships research.

WHAT TO BUILD

Map the independent creators, video channels, newsletters, and podcasts that cover AI developer tooling and document workflows: audience shape, format, publishing cadence, how they appear to source what they cover, whether they take sponsorship and on what published terms, and which of our competitors they have already covered and how. Rank by audience fit rather than raw size. Mapping only. Names of public creators and their public channels are fine; no contact, no messages, no comments. a media landscape map with a fit methodology, a standard and valuable growth artifact.

WHAT STRONG LOOKS LIKE

the audience-fit judgement is argued rather than assumed from subscriber counts, sponsorship terms are quoted where published, and prior competitor coverage is summarised so we know what has already been said.

BAND S2 · chat, Multi-document, export · the submission form, into Drive

Event, hackathon and meetup calendar map

WHO IT SERVES

Growth, developer relations research.

WHAT TO BUILD

Map the next twelve months of events where a document-tooling talk, workshop, or sponsorship would land: audience, format, call-for-papers deadlines and requirements, sponsorship tiers where published, cost, and the practical logistics. Include online-only events and hackathons. Rank by fit against cost. We submit every proposal and make every approach; the candidate produces the calendar and, optionally, a draft abstract. an event strategy calendar with a fit rubric.

WHAT STRONG LOOKS LIKE

deadlines are exact and sourced, the fit argument is specific to the audience of each event, and events that look prestigious but are a poor fit are recommended against.

BAND S1 · chat, templates, export · the submission form, into Drive

Twelve-week content plan tied to the measured query corpus

WHO IT SERVES

Growth, content strategy.

WHAT TO BUILD

Take the short-query corpus and the zero-appearance baseline and build the plan: for each query family, what asset type would plausibly rank, who it is for, what evidence it must carry, what already ranks and what would have to be beaten, and the honest expected difficulty. Sequence twelve weeks of work with dependencies marked and a measurement plan stating what would count as movement. Argue explicitly for what not to make. a full content strategy document tied to measured baseline data, the single strongest artifact a content-strategy candidate can show.

WHAT STRONG LOOKS LIKE

the plan is grounded in what currently ranks per query rather than in general content advice, the difficulty estimates are honest including the queries judged unwinnable in twelve weeks, and the measurement plan would actually detect failure.

BAND S2 · chat, Multi-document, templates, export · a pull request into the public repository, use-cases folder

One non-English market, mapped deeply

WHO IT SERVES

Growth, international expansion.

WHAT TO BUILD

Pick one language market the candidate knows from the inside. Map it properly: how developers and document-heavy professionals there search, in which language and on which engines, what the local competitive set is including local products invisible from English, which communities matter, what content formats work locally, and what document conventions differ. Include a translated query corpus with the local phrasing people actually type rather than translations of our English queries. an international market study nobody else could have written, and a demonstrated specialisation.

WHAT STRONG LOOKS LIKE

the query corpus is idiomatic and locally sourced, the local competitors named are ones we could not have found from English, and the document conventions section contains something genuinely surprising.

BAND S2 · chat, search, export · the submission form, into Drive

Pricing and packaging teardown across the competitive set

WHO IT SERVES

Growth, product marketing.

WHAT TO BUILD

Map how the competitive set prices: model (seat, usage, volume, enterprise-only), what a free tier includes, where the paywalls sit, what is gated to enterprise, whether the model is published or hidden, and what each pricing choice implies about who they are built for. Twelve companies minimum. Assemble a comparison document with a written analysis of where the gaps and the crowded spots are. Public pricing pages and published material only. No sales calls, no quote requests, no trials taken under a false pretext. a pricing and packaging analysis, a core product-marketing portfolio piece.

WHAT STRONG LOOKS LIKE

hidden pricing is recorded as hidden rather than estimated, the analysis reasons from pricing to intended buyer, and the gaps identified are argued rather than asserted.

BAND S2 · chat, Multi-document, templates, export · the submission form, into Drive

Objection library built from public evidence

WHO IT SERVES

Growth, product marketing, presales.

WHAT TO BUILD

Build the objection library from what people already say in public: forum threads, review-site complaints, community discussions, and competitor pages that pre-empt objections. For each objection, capture the wording people actually use, how common it appears to be, whether it is fair, and what an honest answer would be, marking clearly where the honest answer is "we do not do that yet". Twenty objections minimum, sourced. a sourced objection library and method, immediately useful in any product-marketing or presales role.

WHAT STRONG LOOKS LIKE

the objections are quoted in real people's words rather than rewritten into marketing phrasing, fairness is assessed honestly, and at least a few honest answers are concessions.

BAND S2 · chat, search, templates, export · a pull request into the public repository, use-cases folder

Agency and implementation partner landscape

WHO IT SERVES

Growth, partnerships research.

WHAT TO BUILD

Map the agencies, consultancies, and implementation partners who build document workflows for clients: what they build, which platforms they already implement, whether they publish a partner programme and on what terms, and what would make a document-editing capability commercially interesting to them. Rank by fit. Mapping only; every approach is ours. a partner landscape analysis with per-partner commercial reasoning.

WHAT STRONG LOOKS LIKE

the commercial logic is worked out from the partner's side, including how they would make money, and partners who are effectively competitors are identified as such.

BAND S2 · chat, Multi-document, export · the submission form, into Drive

Community question map and drafted answers, for our review

WHO IT SERVES

Growth, community.

WHAT TO BUILD

Find the public threads where people ask exactly the questions our product answers, across forums, question sites, and discussion boards. For each: the question, its age, its visibility, the existing answers and their quality, whether a vendor answer would be welcome under that community's stated rules, and a drafted honest answer written to that community's conventions. Nothing is posted. The candidate drafts; we decide and we post, if at all. A candidate who posts has failed the card. a community engagement plan with drafted answers, demonstrating both research skill and judgement about when not to speak.

WHAT STRONG LOOKS LIKE

the drafts answer the question first and mention us last, or not at all where that is the right call. Threads where any vendor answer would be unwelcome are marked do-not-answer. The candidate visibly resisted the temptation to post.

BAND S2 · chat, templates, export · the submission form, into Drive

Message testing document: six positioning variants and a test design**WHO IT SERVES**

Growth, product marketing.

WHAT TO BUILD

Write six genuinely different positioning statements for the same product, each with the audience it targets, the belief it depends on, and the objection it invites. Then design the test that would separate them: what to run, on which surfaces, what sample would be needed, what would count as a result, and what confound would invalidate it. The test is designed, not run, and nothing is sent to anyone. a positioning and experiment-design document, strong evidence of rigorous marketing thinking.

WHAT STRONG LOOKS LIKE

the six variants are genuinely different bets rather than six phrasings of one idea, and the test design is honest about what a small sample can and cannot establish.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

Discoverability scoreboard: a repeatable runbook and its first run**WHO IT SERVES**

Growth, measurement. (an engineer could automate parts).

WHAT TO BUILD

The zero-appearance baseline was measured once. Turn it into something repeatable: a runbook naming the query set, the engines and their logged-in and logged-out states, the geography, the capture method, and the recording format, precise enough that a different person gets comparable results next month. Then run it once, record the full scoreboard, and note every place the method was ambiguous and how it was resolved. Report the coverage honestly, including cells not run. a measurement runbook and a dated scoreboard under their name, which is the rarest and most valued skill in growth: making the work measurable.

WHAT STRONG LOOKS LIKE

the runbook is specific enough to hand to a stranger, the results state coverage honestly rather than implying completeness, and the ambiguities encountered are documented rather than silently resolved.

BAND S2 · chat, templates, Multi-document, export · the submission form, into Drive

Editor-stack signal map

WHO IT SERVES

Growth

WHAT TO BUILD

Identify products whose applications already run a rich-text editor and therefore have somewhere to plug document AI in. Record the detection evidence.

WHAT STRONG LOOKS LIKE

the detection method is reproducible and the candidate states its false-positive rate honestly. Output: table plus the detection script.

BAND S1 · chat, search, export · the submission form, into Drive

Shipped-feature signal map

WHO IT SERVES

Growth

WHAT TO BUILD

Track AI products that recently shipped a document generation or export feature, from public changelogs and release notes. A product that just shipped a weak version of this has proven demand and has felt the pain.

WHAT STRONG LOOKS LIKE

the entry quotes the actual changelog line and assesses what is still missing. Value: the sharpest timing signal there is. Nobody is watching changelogs systematically. Output: dated table with quotes.

BAND S1 · chat, search, export · the submission form, into Drive

Public pain map

WHO IT SERVES

Growth

WHAT TO BUILD

Collect real people describing document generation or formatting pain in public: forums, issue trackers, question sites. Quote them verbatim with a link, and classify the pain.

WHAT STRONG LOOKS LIKE

verbatim quotes, correctly classified, with the frequency of each pain type counted. Value: the raw language real buyers use, which is what good positioning is built from. Output: quote bank with a taxonomy.

BAND S1 · chat, search, templates, export · the submission form, into Drive

Dependency map

WHO IT SERVES

Growth and engineering

WHAT TO BUILD

Find open-source projects and products that depend on document manipulation libraries. Heavy dependency usually means a homegrown pipeline that is expensive to maintain.

WHAT STRONG LOOKS LIKE

uses public package and code search properly, and separates casual use from a real pipeline. Value: the business brief says a company that built its own editing pipeline is often a prime target. This finds them mechanically. Output: ranked table plus the query method.

BAND S2 · chat, search, export · the submission form, into Drive

Marketplace listing map

WHO IT SERVES

Growth

WHAT TO BUILD

Catalogue listings in major software marketplaces whose function involves producing documents, and note which are thin or poorly rated.

WHAT STRONG LOOKS LIKE

the assessment of "thin" is evidenced by review content, not asserted. Value: shows where a document capability is already being paid for, and where the incumbent is weak. Output: marketplace-by-marketplace table.

BAND S1 · chat, search, export · the submission form, into Drive

Review-mining map

WHO IT SERVES

Growth

WHAT TO BUILD

Mine public software reviews for complaints that name document formatting, generation or export pain, attributed to the product being reviewed.

WHAT STRONG LOOKS LIKE

the complaint is quoted and the product's category is recorded, so patterns emerge by category. Value: complaints about a competitor's document capability are the highest-intent signal available. Output: quote bank grouped by category.

BAND S1 · chat, search, templates, export · the submission form, into Drive

Event and session room map

WHO IT SERVES

Growth

WHAT TO BUILD

For one buyer persona, map where they physically gather: which conferences, which specific sessions and tracks, which dates, and what the attendee profile actually is.

WHAT STRONG LOOKS LIKE

session-level detail, not conference-level. The output tells us which room, not which city. Value: the founder's "get us into rooms" ask, made concrete. Output: calendar-shaped map with cost and access notes.

BAND S1 · chat, export · the submission form, into Drive

Audience-channel map

WHO IT SERVES

Growth

WHAT TO BUILD

For one buyer persona, map the podcasts, newsletters and publications they actually consume, with audience size and how each accepts contributions.

WHAT STRONG LOOKS LIKE

it distinguishes what the persona reads from what people who sell to the persona read. Those are different lists. Value: earned-media routing that does not require paid distribution. Output: channel table with submission mechanics.

BAND S1 · chat, export · the submission form, into Drive

Cohort map

WHO IT SERVES

Growth

WHAT TO BUILD

Map recent accelerator and program cohorts for companies whose product produces documents, with what each says it does.

WHAT STRONG LOOKS LIKE

correctly separates genuine fits from superficial keyword matches. Value: dense, well-documented, publicly listed clusters of exactly our target profile. Output: cohort-by-cohort table.

BAND S1 · chat, search, export · the submission form, into Drive

Competitor-customer map

WHO IT SERVES

Growth

WHAT TO BUILD

From public case studies, logo walls and testimonials, map who is already paying for adjacent document tooling, and what they said they needed.

WHAT STRONG LOOKS LIKE

each row cites the public source, and the candidate notes what the customer said the problem was in their own words. Value: proven budget, proven problem, proven willingness to be referenced. Output: table with source links and quoted needs.

BAND S1 · chat, search, export · the submission form, into Drive

Buyer-persona interview kit

WHO IT SERVES

Growth

WHAT TO BUILD

For one persona, build the interview guide that would learn what they actually do today: open questions, no pitching, structured to avoid leading the witness.

WHAT STRONG LOOKS LIKE

the questions ask about past behaviour rather than future intent, which is the difference between useful and useless research. Value: reusable asset for every future conversation. Also a strong signal of who understands customer discovery. Output: interview guide plus a synthesis template.

BAND S1 · chat, templates, export · the submission form, into Drive

Warm-path map

WHO IT SERVES

Growth

WHAT TO BUILD

For a set of target companies, map the publicly visible connection paths: shared alumni networks, prior employers, open-source collaboration, public co-authorship. No private data, no scraping of anything gated.

WHAT STRONG LOOKS LIKE

every path is verifiable from a public source, and the candidate states clearly where they stopped for privacy reasons. Value: the founder's "who would you introduce us to and why", answered with evidence rather than guesswork. Output: connection map with sources.

BAND S2 · chat, search, export · the submission form, into Drive

Directory placement map

WHO IT SERVES

Growth

WHAT TO BUILD

Every directory, index, registry and listing where a product like ours should appear, with requirements, cost, review time and expected traffic per entry.

WHAT STRONG LOOKS LIKE

it verifies requirements against the directory's current rules rather than assuming, and ranks by effort against likely return. Value: directory placement is a repeatedly proven channel for us and the list is not written down anywhere. Output: ranked directory table.

BAND S1 · chat, search, export · the submission form, into Drive

Vertical search-demand map

WHO IT SERVES

Growth

WHAT TO BUILD

For one industry, map what practitioners actually search when they have a document problem, using public keyword and question sources.

WHAT STRONG LOOKS LIKE

the queries are the practitioner's language, not the vendor's, and volume estimates cite their source. Value: feeds intent pages directly. Each map becomes several landing pages. Output: query corpus with intent classification.

BAND S1 · chat, search, export · the submission form, into Drive

Agent-ecosystem map

WHO IT SERVES

Growth

WHAT TO BUILD

Map agent products and frameworks whose users will need a document capability, with how each one adds tools and what its distribution surface is.

WHAT STRONG LOOKS LIKE

it identifies the specific extension mechanism for each, which is what makes the map actionable. Value: our most agent-native channel, mapped properly for the first time. Output: ecosystem table with extension mechanics.

BAND S1 · chat, export · the submission form, into Drive

Marketplace-gap map

WHO IT SERVES

Growth

WHAT TO BUILD

For each major platform marketplace, determine whether a document AI application exists, and if not, why the gap persists.

WHAT STRONG LOOKS LIKE

the "why" is investigated rather than assumed. Sometimes the gap exists because the platform blocks it. Value: turns directly into a build backlog with distribution attached. Output: gap table with feasibility notes.

BAND S1 · chat, search, export · the submission form, into Drive

Public procurement map

WHO IT SERVES

Growth

WHAT TO BUILD

Find published tenders and procurement notices that name document automation, generation or management requirements.

WHAT STRONG LOOKS LIKE

it extracts the actual stated requirement, which tells us what public buyers ask for in their own words. Value: government buying is public, structured and almost never mined by early startups. Output: tender table with requirement extracts.

BAND S1 · chat, search, export · the submission form, into Drive

Archetype-to-demo worksheet

WHO IT SERVES

Growth and product

WHAT TO BUILD

For one buyer archetype, define precisely what that buyer would need to SEE working before they believed the problem was solved, then specify that demo well enough for an engineer to build it.

WHAT STRONG LOOKS LIKE

the demo is specific, buildable on the shipped surface, and the belief it creates is named explicitly. A vague "show them a demo" fails. Value: this is the highest-leverage GTM thinking on the whole list. The best worksheets become build cards in the next catalog wave. Output: worksheet per archetype, feeding the prospect-to-demo map.

BAND S2 · chat, templates, Multi-document, export · the submission form, into Drive

THE CONTENT WAVE

Social-to-doc repurposer

WHO IT SERVES

Marketing / a social or content marketer repurposing a high-performing thread or post into a polished, shareable article.

WHAT TO BUILD

A thread or post that performed well is stuck in the format it was written in, and turning it into a polished article or handout means reformatting by hand. Build an app where a user pastes in the text of a thread or post and the app restyles it into a structured, styled article, headings, pull quotes, sectioned, exported as a document or PDF. This starts from pasted post text, not a connected newsletter archive.

WHAT STRONG LOOKS LIKE

A reviewer pastes in a ten-post thread and confirms the exported article reads as sectioned prose with headings, not the original post-by-post fragments.

BAND S2 · chat, export · publish it on your own channel and put the link in the submission form

Dev.to / Hashnode cross-post formatter

WHO IT SERVES

Marketing/DevTools

WHAT TO BUILD

A technical post written once should not need manual reformatting for every place it gets published. Build a cross-post formatter: take one written draft and produce correctly formatted, platform-specific exports for at least two developer publishing platforms such as Dev.to and Hashnode, each with that platform's own front-matter fields, code-block syntax, and a canonical-URL tag pointing back to the source. This runs outbound, one draft fanning out to many platforms, the opposite direction from a tool that normalizes many inbound submissions into one house style.

WHAT STRONG LOOKS LIKE

Each platform-specific export pastes into that platform without a manual fix-up pass, code blocks keep correct syntax highlighting on both, and the canonical-URL tag actually points back to the source on every export.

BAND S2 · chat, export · publish it on your own channel and put the link in the submission form

Specification section editor with numbering integrity

WHO IT SERVES

Construction / specification writer or project architect at an AEC firm

WHAT TO BUILD

Project specifications run to hundreds of sections under a strict numbering system. The app opens a section, edits it by instruction ("add a submittal requirement for the sealant warranty", "swap the referenced standard to the current edition"), keeps the three-part structure and article numbering intact, and repairs cross-references in related sections when a change breaks them. A project-wide check reports orphan references, superseded standards, and articles that contradict the front-end documents. Output is the section DOCX plus the assembled project manual with a generated table of contents.

WHAT STRONG LOOKS LIKE

article numbering never drifts across dozens of edits; cross-section reference repair is real and demonstrated on a broken project; the assembled manual's table of contents matches the section set exactly.

BAND S3 · chat, Multi-document, search, Review, export · a pull request into the public repository, use-cases folder

Project closeout and O&M manual assembler

WHO IT SERVES

Construction / project engineer or document controller at handover

WHAT TO BUILD

Closeout means assembling warranties, as-built references, equipment cut sheets, maintenance schedules, spare-parts lists, training records and certificates into an owner manual that satisfies the specification's closeout requirements. The app reads the closeout specification section to derive the required structure, matches received documents to each requirement, generates the index, drafts the maintenance-schedule summary from the equipment sheets, and produces a deficiency list naming the responsible subcontractor for each missing item. Output is a bookmarked PDF manual plus a per-system folder.

WHAT STRONG LOOKS LIKE

the required structure is derived from the actual project specification rather than a fixed template; document-to-requirement matching survives inconsistent subcontractor naming; the deficiency list is accurate enough to justify withholding retention.

BAND S3 · Multi-document, search, templates, export · a pull request into the public repository, use-cases folder

Recorded-session to article and show-notes pipeline

WHO IT SERVES

Content / producer running a podcast, webinar or video series

WHAT TO BUILD

One recording becomes a document set: a structured article with headings and pull quotes, timestamped show notes, a quote sheet for social, a cleaned transcript with speaker labels and filler removed, and a summary card. Diagrams described verbally are drawn as diagrams where the description supports it. Speakers review the article, and any edit to a quote flows back into the quote sheet automatically. Everything exports branded, with the article also available as clean Markdown for a content system.

WHAT STRONG LOOKS LIKE

quotes stay verbatim and remain synchronized across artefacts after an edit; the article reads as written prose rather than a lightly tidied transcript; the Markdown export pastes into a content system without repair.

BAND S2 · chat, Multi-document, images, templates, export · a pull request into the public repository, use-cases folder

Corrections and clarifications manager

WHO IT SERVES

Standards editor responsible for the correction record on published articles

WHAT TO BUILD

When a published piece is wrong, the archive copy has to be amended precisely, a dated correction note appended, and the corrections page updated, all without disturbing the rest of the article. Build a tool that takes the correction request, makes the minimal edit to the archived article, appends the standard correction note in the house format, adds the entry to the corrections index, and keeps a before-and-after record for the archive. The tool refuses to make an edit larger than the correction describes.

WHAT STRONG LOOKS LIKE

the diff on the amended article is limited to the corrected words plus the appended note, and an over-broad instruction is challenged instead of executed.

BAND S1 · chat, Multi-document, Review, export · a pull request into the public repository, use-cases folder

Transcript to show notes, chapters and press one-sheet

WHO IT SERVES

Podcast or video producer shipping episode collateral on a schedule

WHAT TO BUILD

One recording has to become four artifacts: show notes with timestamped chapters, a quotable pull-quote sheet, a press one-sheet with the guest bio and key art, and a blog-length write-up. Build a pipeline that takes the transcript and produces all four as separate documents in one session, keeps guest names, titles and timestamps consistent across them, generates the key art for the one-sheet, and lets the producer revise one artifact without regenerating the others.

WHAT STRONG LOOKS LIKE

a correction to the guest's title updates every artifact that carries it in one pass, and revising the blog post leaves the other three files untouched.

BAND S1 · Multi-document, chat, images, export · a pull request into the public repository, use-cases folder

Syndication reformatter at volume

WHO IT SERVES

Digital production editor at a publisher taking wire and partner content

WHAT TO BUILD

Incoming wire and partner articles arrive in the sender's structure and style and have to land in the publication's template: house headline case, dateline format, subheading rhythm, image credit lines, boilerplate footer and localized spellings. Build a server-side worker that normalizes each incoming piece, preserves the sender's required attribution exactly, flags anything that needs a human read (legal-sensitive terms, unverified claims, missing credit), and pushes the finished file to the CMS. Throughput and the flag rate are reported per run.

WHAT STRONG LOOKS LIKE

required attribution is provably untouched across the whole batch, the flag rate is measured rather than assumed, and one instruction change restyles the next run without code changes.

BAND S2 · API, templates, chat, export · a pull request into the public repository, extensions folder

Country-conditions evidence brief assembler

WHO IT SERVES

Asylum or humanitarian-protection attorney building the country-conditions record

WHAT TO BUILD

Build a tool that takes reports, articles and expert statements as attachments, extracts the passages that bear on the specific claim elements, assembles the brief with each proposition supported by a quoted excerpt and a full citation, and produces the accompanying indexed source packet. Sources that contradict each other are surfaced rather than cherry-picked, and passages the attorney marks as central are tracked so a source swap shows every proposition that depended on it.

WHAT STRONG LOOKS LIKE

every quoted excerpt is verbatim and correctly cited, contradicting sources are surfaced, and removing a source reports exactly which propositions lose support.

BAND S3 · search, Multi-document, chat, export · a pull request into the public repository, use-cases folder

Knowledge-base freshness sweeper

WHO IT SERVES

Support-content or knowledge manager owning hundreds of help articles

WHAT TO BUILD

Product changes silently invalidate help articles, and nobody knows which ones. Build a tool that holds the article set, takes a change feed (release notes, renamed screens, changed limits, retired plans), finds every article affected including ones that describe the old behaviour without naming it, proposes the surgical edit per article, and routes the batch for review. Screenshots that no longer match are flagged for replacement. The tool reports a portfolio freshness score and the articles it could not assess.

WHAT STRONG LOOKS LIKE

measured precision and recall against a seeded stale-article corpus, an honest could-not-assess bucket, and approved edits that change only the affected sentences and leave the article's structure intact.

BAND S3 · search, Multi-document, images, Review, API · a pull request into the public repository, extensions folder

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Help-centre localization with screenshot swaps

WHO IT SERVES

Localization or support-content owner shipping help content in several languages

WHAT TO BUILD

Translated help articles break when the screenshots still show the English interface. Build a pipeline that localizes each article while protecting product names, keyboard shortcuts and code from translation, then swaps every screenshot for the locale-specific capture supplied by the team, keeping the image position, caption and callout numbering intact. Articles whose locale screenshot is missing are published with a flag rather than the wrong image. Source updates re-localize only the paragraphs that changed.

WHAT STRONG LOOKS LIKE

screenshot swaps preserve layout, captions and numbering; protected terms come through untouched in every language; incremental updates re-translate only the changed paragraphs.

BAND S2 · images, Multi-document, API, templates, export · a pull request into the public repository, extensions folder

Coverage and PR Report Builder

WHO IT SERVES

PR agencies · account executive reporting earned media to a client

WHAT TO BUILD

Feed clippings (article links, screenshots, saved PDFs) plus reach numbers. The app builds the coverage report: summary narrative, a coverage table with outlet, date, sentiment and reach, verbatim quoted highlights pulled from each article, and screenshots placed inline next to their row. Non-English coverage is summarized in the client's language while the original quote is preserved in its own script.

WHAT STRONG LOOKS LIKE

quotes are verbatim and correctly attributed; screenshots land beside the right row; multilingual coverage renders without mangling scripts, including right-to-left ones.

BAND S2 · images, search, chat, templates, export · a pull request into the public repository, use-cases folder

Playtest to Design Change Register

WHO IT SERVES

Gaming · design lead turning playtest findings into decisions on a living design document

WHAT TO BUILD

Playtest output (session notes, survey responses, telemetry summaries, video timestamps) is messy and repetitive. The app clusters findings into distinct issues, ranks them by frequency and severity, and for each proposes a concrete design change staged directly against the relevant section of the design document. Accepted changes edit the design document in place and append to a decision register recording what changed, why, and which playtest evidence drove it. Rejected findings are recorded too, with the reason.

WHAT STRONG LOOKS LIKE

clustering merges the same complaint phrased five ways; changes land in the right design section; the register makes a six-month-old decision reconstructable.

BAND S3 · Multi-document, search, Review, chat, memory, export · a pull request into the public repository, use-cases folder

Commissioning and As-Built Pack

WHO IT SERVES

Solar · project manager closing out an installation

WHAT TO BUILD

At handover, assemble the closeout pack from the field record: commissioning test results with pass and fail per test, as-built single-line and layout diagrams reflecting what was actually installed rather than what was designed, equipment serial and warranty registration list, inspection sign-offs with photographs, the customer operations and maintenance guide written for a non-technical owner, and the warranty summary. Deviations from the design are listed explicitly with their reason.

WHAT STRONG LOOKS LIKE

as-built genuinely reflects field deviations rather than restating the design; photographs attach to their test or inspection; the owner guide is readable by a homeowner.

BAND S2 · Multi-document, images, chat, templates, export · a pull request into the public repository, use-cases folder

Canva app: design to long-form document

WHO IT SERVES

Marketer, small-business owner, or educator who designs in Canva and needs the long version. Host app: Canva.

WHAT TO BUILD

A Canva app that turns a design into the document that accompanies it: a deck becomes a leave-behind report, a one-page flyer becomes the full brochure with the same look, a social set becomes a styled article. It reads the design's copy and visual system and produces a document that carries the brand into a format Canva does not do well, exported as Word or a print-ready PDF. Designs stay linked so a copy change can be pulled through.

WHAT STRONG LOOKS LIKE

the document looks like it belongs to the design rather than to a generic template. Images arrive at print resolution. Long documents keep the visual system consistent from page one to page forty.

BAND S2 · images, chat, templates, export · a pull request into the public repository, extensions folder

Zendesk app: the document sidebar in the ticket

WHO IT SERVES

Support agent, claims handler, or student-services officer working a ticket queue. Host app: Zendesk.

WHAT TO BUILD

A sidebar app on the ticket that produces the document the ticket needs without leaving it: the formal response letter, the credit note, the incident report, the knowledge-base article that stops the next ten tickets. It reads the ticket thread and the customer record, drafts on the right template, lets the agent adjust in plain language, and attaches the finished file to the ticket while filing the article to the help centre.

WHAT STRONG LOOKS LIKE

the drafted document reflects what the ticket actually says, including the customer's own wording where it matters. The template is right for the ticket type without the agent choosing. Attachment and help-centre publishing both land first time.

BAND S2 · API, chat, templates, export, Review · a pull request into the public repository, extensions folder

Build the do-it-yourself version, measure it, publish the result

WHO IT SERVES

Developers evaluating build versus buy. (a growth candidate with real coding ability may also take it).

WHAT TO BUILD

The logged-out conversational engines currently answer "AI document editing API" by recommending a do-it-yourself route: a document library plus a language model plus your own diffing. Build that route properly, in good faith, over a weekend. Then build the same three tasks on the hosted editing surface. Publish a bylined write-up with both repositories, the code, the failure cases each hit, the wall-clock, and the token counts the candidate measured. No claim in the article may come from anywhere except their own runs. two public repositories and a bylined engineering study with original measurements, the kind of piece that gets a developer hired.

WHAT STRONG LOOKS LIKE

the do-it-yourself arm is built in genuine good faith and is allowed to win where it wins, for example on a trivially small document. The failure modes are shown with real output, not described. The conclusion is defensible by a hostile reader.

BAND S3 · chat, API, export · publish it on your own channel and put the link in the submission form

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Format-fidelity round-trip benchmark and scorecard

WHO IT SERVES

Technical buyers, document engineers.

WHAT TO BUILD

Assemble a corpus of 25 documents that stress formatting: merged table cells, nested lists, footnotes, headers that change by section, tracked changes, multi-column layout, equations, embedded images, right-to-left text. Define a scoring rubric for what survives an upload, an edit, and an export. Score SuperDocs and at least two openly available alternatives, publish the corpus, the rubric, the scores, and the screenshots of every failure including ours. a public benchmark corpus and rubric under their name, a genuine contribution to the field.

WHAT STRONG LOOKS LIKE

the rubric is defined before the runs and not adjusted afterwards. Our failures are shown as prominently as anyone else's. The corpus is redistributable.

BAND S3 · API, chat, export · publish it on your own channel and put the link in the submission form

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

The long-document stress test, written as a lab report

WHO IT SERVES

Anyone with documents that broke their last tool.

WHAT TO BUILD

Take a genuinely long public document, several hundred pages, and run a realistic 30-edit session against it: targeted changes, insertions, deletions, a renumber, a revert, and an export. Record what happened at every step including the slow steps and anything that failed. Publish it as a lab report, with the document named so a reader can repeat it. a published lab report and a reusable stress-test protocol.

WHAT STRONG LOOKS LIKE

the report includes the boring middle and the failures, timings are recorded rather than estimated, and the final document is published so a reader can inspect the end state themselves.

BAND S2 · chat, search, Multi-document, Review, export · publish it on your own channel and put the link in the submission form

Three-part tutorial series in one stack, done properly

WHO IT SERVES

Developers in one named ecosystem.

WHAT TO BUILD

Pick one stack the candidate genuinely knows. Write three connected tutorials: connect and make the first edit; add human review before changes apply; add file handling and export for real-world document sizes. Each part ships runnable code, states its prerequisites, and ends somewhere useful rather than at a toy. Part three must handle at least one thing that goes wrong. a bylined tutorial series and a public example repository, the standard portfolio evidence of teaching ability.

WHAT STRONG LOOKS LIKE

a stranger with the named stack can follow part one to a working edit in under 20 minutes without asking a question. The error handling in part three is real, not a try block with a comment.

BAND S2 · API, chat, Review, export · publish it on your own channel and put the link in the submission form

Video walkthrough: an agent edits a real document end to end

WHO IT SERVES

Developers and AI-tool users who learn by watching.

WHAT TO BUILD

Record a single unbroken walkthrough, 8 to 15 minutes, of connecting an AI coding assistant to the document tools and then completing a real task: open a messy document, find a section, change it, review the proposed change, apply it, export it, and open the export to prove it survived. Show the moment something does not work and how it is handled. Real screen, real latency, no time-lapse over the hard part. a video on their own channel, which is portfolio, proof of communication skill, and an audience asset all at once.

WHAT STRONG LOOKS LIKE

the task is one a viewer recognises from their own week. The awkward moments are left in. The description carries timestamps and the exact configuration used.

BAND S1 · MCP, chat, Review, export · publish it on your own channel and put the link in the submission form

"Parsing is not editing", the corrective explainer

WHO IT SERVES

Developers and technical buyers.

WHAT TO BUILD

One measured engine answered "AI document editing API" by recommending document parsers and extraction vendors. That is a category confusion worth fixing in public. Write the explainer that separates the three jobs cleanly: extracting data out of documents, generating a new document, and changing an existing document while leaving the rest intact. Use one worked example carried through all three so the reader sees why the third is a different problem. Name the tool categories fairly, including the ones that do the other two jobs well. a bylined category explainer, the kind of writing that establishes someone as a voice in a space.

WHAT STRONG LOOKS LIKE

a reader who arrived believing they wanted a parser leaves knowing which of the three they actually need, even when the answer is not us. Fairness to the other two categories is what makes it citable.

BAND S1 · chat, export · publish it on your own channel and put the link in the submission form

The category essay: what a coding-assistant equivalent for documents has to get right

WHO IT SERVES

Developers, founders, product people.

WHAT TO BUILD

A widely read community thread asks why nobody has built the document equivalent of a coding assistant, and it ranks. Write the serious answer: what makes documents harder than code (no compiler, no test suite, formatting carries meaning, no diff a normal person can read), what an honest attempt has to solve, and what is still unsolved. Ground every claim in something the candidate built or measured while writing it. a bylined essay on a live industry question, the strongest single writing sample a growth candidate can produce.

WHAT STRONG LOOKS LIKE

it reads as a thinking piece rather than a product post, it concedes the genuinely unsolved parts, and someone who disagrees would still link to it.

BAND S2 · chat, Review, export · publish it on your own channel and put the link in the submission form

Editor toolkit plus your own model versus a hosted editing harness

WHO IT SERVES

Product engineers adding AI to an existing editor.

WHAT TO BUILD

Many teams already run a rich-text editor and are told to add AI to it. Build the same feature twice: once by wiring an editor toolkit to a model the candidate supplies and prompts themselves, once against the hosted editing surface. Publish both repositories and an honest account of what each costs in code, in tokens, in latency, and in the ongoing maintenance the candidate can foresee. two working implementations and a comparison study, a strong senior-engineering portfolio piece.

WHAT STRONG LOOKS LIKE

the toolkit arm is built the way its documentation recommends, not strawmanned. The maintenance argument is concrete, naming what breaks when a model changes. The candidate says plainly where the toolkit route is the right answer.

BAND S3 · API, chat, Review · publish it on your own channel and put the link in the submission form

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Fill a real gap in the public developer documentation

WHO IT SERVES

Developers, technical writers.

WHAT TO BUILD

Work through the public developer documentation as a newcomer, keep a friction log, then pick the single biggest genuine gap and write the missing guide properly: prose, runnable examples, and the errors a newcomer will actually hit. Open a pull request to the public documentation repository. The friction log is submitted alongside the pull request and is judged as heavily as the guide. a merged pull request to a public repository, which is durable, verifiable open-source contribution history.

WHAT STRONG LOOKS LIKE

the gap chosen is the one a newcomer actually hits first, not the one that was easiest to write. The examples run as written. The friction log is specific and unflattering where it should be.

BAND S1 · export · a pull request into our public developer documentation

Notebook walkthrough for the data and research audience

WHO IT SERVES

Data scientists, researchers, analysts.

WHAT TO BUILD

Build an executable notebook that takes a research-shaped job end to end: load a paper or long report, find and summarise the relevant sections, edit a section with tracked review, regenerate a table from a dataset, and export to both a styled document and Markdown. Every cell runs top to bottom on a fresh environment. Publish it with the outputs saved so a reader sees the result before running anything. a published, executable notebook, a well-understood credential in the data community.

WHAT STRONG LOOKS LIKE

it runs top to bottom without a hidden state cell, the saved outputs match what a fresh run produces, and the narrative between cells teaches rather than labels.

BAND S2 · API, chat, search, Review, export · the submission form, into Drive

Write the tutorial in a language you actually speak

WHO IT SERVES

Developers outside the English-first internet.

WHAT TO BUILD

Pick a language the candidate is fluent in and write a full getting-started tutorial in it, natively, not translated: idiomatic phrasing, locally relevant examples, and a worked document in that language including correct script handling. Include a section on editing documents in that language, since that is the part an English tutorial never covers. a bylined technical tutorial in their own language, which is both a portfolio piece and a rare skill signal.

WHAT STRONG LOOKS LIKE

a native speaker would not guess it started in English. The worked example uses a document type common in that market. Script and typography issues are addressed rather than avoided.

BAND S1 · chat, export · publish it on your own channel and put the link in the submission form

A curated, honest resource list for document AI

WHO IT SERVES

Developers surveying the space.

WHAT TO BUILD

Build a community-style curated list repository covering the whole space: parsing and extraction, editor toolkits, document generation, conversion libraries, format specifications, agent tooling, and hosted editing. Include competitors generously and describe each accurately, including where they are stronger than us. Add a short "how to choose" section organised by the job to be done. a curated list repository under their name, which accrues stars and contributors over time and belongs to them.

WHAT STRONG LOOKS LIKE

a competitor could read their own entry and find it fair. The "how to choose" section sometimes points away from us. Every entry is checked to be alive, not copied from an older list.

BAND S1 · export · publish it on your own channel and put the link in the submission form

Lightning-talk deck and speaker notes

WHO IT SERVES

Conference and meetup audiences, developer relations.

WHAT TO BUILD

Build a 10-minute talk: the problem (an agent that can reason but cannot change the file), one live demonstration, one honest limitation, and one takeaway the audience can act on tomorrow. Deliver the deck, full speaker notes, and a recorded run-through. The demonstration must be scripted so it survives bad conference network conditions. a delivered talk with a recording, which is speaking-history evidence they can use for years.

WHAT STRONG LOOKS LIKE

the recording is watchable at normal speed, the limitation named is a real one, and the demonstration has an offline fallback that is prepared rather than promised.

BAND S1 · chat, Review, export · publish it on your own channel and put the link in the submission form

The integration postmortem

WHO IT SERVES

Developers who have been burned before.

WHAT TO BUILD

While completing any build card in this catalog, keep a real defect log: what was misunderstood, what error message was unhelpful, what took three attempts, what the documentation implied but did not say. Then write the postmortem: five real problems, the root cause of each, and the fix. Send the same log to us as a product report. a bylined engineering postmortem, which demonstrates debugging judgement better than any successful tutorial.

WHAT STRONG LOOKS LIKE

it is specific and unflattering where warranted, every problem has a root cause rather than a workaround, and the tone is that of an engineer helping the next engineer rather than complaining or flattering.

BAND S1 · chat, export · publish it on your own channel and put the link in the submission form

One exact-query demo page, plus the build log

WHO IT SERVES

Whoever types that query.

WHAT TO BUILD

Pick one short query from the measured corpus, for example editing a Word document with tracked changes from an agent. Build the page that genuinely deserves to be the best answer to it: a working demonstration, the code, the honest limits, and the alternatives. Then write the build log explaining why the page is structured the way it is and what the query's searcher actually wants. Page and log are both deliverables; we handle any placement or promotion. a shipped page with a written rationale, which is portfolio evidence of search and conversion thinking, not just writing.

WHAT STRONG LOOKS LIKE

the page answers the query in the first screen without a signup wall, the demonstration works on a phone, and the alternatives section is real. The build log shows the searcher's intent was thought about rather than assumed.

BAND S2 · chat, Review, export · publish it on your own channel and put the link in the submission form

The definitive guide to changing a Word document programmatically

WHO IT SERVES

Developers hitting this problem for the first time.

WHAT TO BUILD

On this query today, the results are a scattering of forum threads, vendor documentation, and partial answers. Write the complete one: the file format's actual structure at the level a developer needs, the library route with working code and its real limits, why formatting breaks when you take the naive path, and the hosted route as one option among several. Vendor-neutral in structure, with the affiliation disclosed near the top. a definitive technical guide under their byline, the kind of piece that becomes the thing people link to for years.

WHAT STRONG LOOKS LIKE

the format explanation is correct and would satisfy someone who works on document formats. The library code runs. Our option is presented as one route with stated trade-offs, not as the conclusion.

BAND S3 · API, chat, export · publish it on your own channel and put the link in the submission form

Honesty note: write the guide from what the shipped product and public docs actually do today, tested by you, never from what any roadmap or ticket describes. A guide that marks the boundaries of today's capability honestly is the definitive one.

Case study from your own build, with numbers**WHO IT SERVES**

Buyers who want proof, not description.

WHAT TO BUILD

Take any build the candidate completed from this catalog and write it up as a case study: the problem, the constraints, what was built, what it cost to run, what it does not do, and what the candidate would change. Include at least three numbers they measured themselves. Written for a reader deciding whether to attempt the same thing. a bylined case study of their own work, standard and valuable portfolio material.

WHAT STRONG LOOKS LIKE

the numbers are measured and the method for each is stated. The limitations section is longer than one line. A reader could decide against the approach based on it, which is what makes it trustworthy.

BAND S1 · chat, export · publish it on your own channel and put the link in the submission form

Generate a typed client from the public specification, then write about it**WHO IT SERVES**

Developers in typed languages.

WHAT TO BUILD

Take the published interface specification, generate a typed client in a language of the candidate's choice, use it to complete a real editing task, and publish both the client and a short post on what generation got right, what it got wrong, and what a human had to fix. Include the exact commands so a reader can regenerate it. a published client library in their language of choice, a concrete open-source artifact.

WHAT STRONG LOOKS LIKE

the generated client actually completes a real task, the post is honest about the rough edges of generation, and the fixes are shown as a diff.

BAND S2 · API, chat, export · publish it on your own channel and put the link in the submission form

Agent client compatibility matrix, actually tested

WHO IT SERVES

Developers choosing an agent client.

WHAT TO BUILD

Install the document tool server in every major agent client the candidate can access, run the same five tasks in each, and publish a matrix: what connected first time, what needed a workaround, what silently failed, what the error text was, and the exact configuration that worked. Include the clients where the experience was poor. Date it, because these change. a maintained compatibility matrix with their name on it, a recognised contribution in the agent-tooling community.

WHAT STRONG LOOKS LIKE

every cell is from a real run, the workarounds are copy-pasteable, and the poor results are published rather than trimmed.

BAND S2 · MCP, chat, export · publish it on your own channel and put the link in the submission form

Where a document harness sits in an agent stack, drawn properly

WHO IT SERVES

Architects, technical decision makers.

WHAT TO BUILD

Write the architecture explainer with real diagrams, generated inside the document itself so the piece demonstrates the capability it describes: the naive stack (agent plus model plus file library), where it breaks, and the shape that works, including where human review sits and where state lives. Ship the diagram sources so readers can adapt them. a bylined architecture piece with reusable diagram sources, strong evidence of systems thinking.

WHAT STRONG LOOKS LIKE

the diagrams are genuinely explanatory rather than decorative, the naive stack is described accurately enough that its builders recognise themselves, and the piece is useful to someone who chooses a different vendor.

BAND S2 · chat, images, export · publish it on your own channel and put the link in the submission form

What a cited page contains: an anatomy study, then one page built to it

WHO IT SERVES

Content and growth practitioners. (an engineer may take it).

WHAT TO BUILD

Take 15 pages that answer engines currently cite on document-tooling questions, and take them apart: structure, question phrasing, how directly the answer appears, evidence density, freshness signals, schema, and length. Write up the pattern with the evidence. Then build one page to the pattern, on a question we can answer honestly, and hand it to us. No contact with any site owner. an original research study on answer-engine citation patterns, a genuinely differentiating growth portfolio piece.

WHAT STRONG LOOKS LIKE

the study is evidence-driven rather than a recitation of general advice, the pattern is stated as testable claims, and the page built at the end visibly follows the pattern.

BAND S2 · chat, export · publish it on your own channel and put the link in the submission form

Build in public, five posts over ten days

WHO IT SERVES

Developer and founder audiences.

WHAT TO BUILD

Take one build card from this catalog and document it live in five posts: the plan, the first thing that broke, a decision the candidate changed their mind about, a measured result, and the finished thing. Each post stands alone and shows real work, including screenshots of failures. Affiliation disclosed in the first post. Five posts total, deliberately capped: this is a depth exercise, not a volume exercise. an audience and a public record of how they work, which belongs entirely to them.

WHAT STRONG LOOKS LIKE

the failure post is the best one. Nothing reads as a scheduled marketing cadence. A reader could follow along and build the same thing.

BAND S1 · chat, export · publish it on your own channel and put the link in the submission form

Build log

WHAT TO BUILD

Write what it was actually like to build a thing on the API in a weekend, including what broke.

WHAT STRONG LOOKS LIKE

it is honest about friction. A build log with no friction reads as an advertisement and performs worse. Value: developer-credible, and the friction becomes a product backlog. Output: post on the candidate's own channel.

BAND S1 · chat, API, export · publish it on your own channel and put the link in the submission form

Benchmark write-up

WHAT TO BUILD

Take one benchmark from the benchmark repo, run it, and write it up for a general technical audience.

WHAT STRONG LOOKS LIKE

the numbers are reproducible and the limitations section is real. Value: turns a data artefact into a citable story. Output: post plus results.

BAND S1 · chat, API, export · publish it on your own channel and put the link in the submission form

Budget guard: build and prove this in small-sample mode first; state the cap a full run may spend before you run it, and report actual spend alongside your results.

Talk proposal and deck

WHAT TO BUILD

A conference talk proposal on a document-AI topic, with the deck.

WHAT STRONG LOOKS LIKE

the proposal fits a specific named call for papers and its stated criteria. Value: earned speaking slots, which we have none of. Output: proposal plus deck.

BAND S1 · chat, templates, export · publish it on your own channel and put the link in the submission form

Podcast pitch kit

WHAT TO BUILD

For one show, a pitch built around what that show's audience cares about, with three concrete episode angles.

WHAT STRONG LOOKS LIKE

it demonstrates the candidate has actually listened to the show. Value: earned media routing. Output: pitch document. Note: the candidate drafts it. Any sending is ours.

BAND S1 · chat, templates, export · publish it on your own channel and put the link in the submission form

Repository with an exemplary README

WHAT TO BUILD

The README is the deliverable: what it does, why, how to run it in one minute, and what it does not do.

WHAT STRONG LOOKS LIKE

someone runs it successfully from the README alone. Value: repositories are discovery surfaces and most READMEs are bad. Output: public repo.

BAND S1 · API · publish it on your own channel and put the link in the submission form

Public answer with working code

WHAT TO BUILD

Find a real unanswered public question about a document problem and answer it properly with working code.

WHAT STRONG LOOKS LIKE

the answer is genuinely useful and would stand even if it mentioned no product at all. Value: durable long-tail discovery. Output: link to the answer. Note: answering a real question is the candidate's own contribution, not promotion.

BAND S1 · API · publish it on your own channel and put the link in the submission form

Audience newsletter issue

WHAT TO BUILD

Write one issue for a specific defined audience about a document workflow problem.

WHAT STRONG LOOKS LIKE

the audience is narrow and the issue is specific to it. Value: newsletter formats travel. Output: issue draft published on the candidate's channel.

BAND S1 · export · publish it on your own channel and put the link in the submission form

Community gallery contribution

WHAT TO BUILD

Publish a template or starter into a community gallery where such things are shared.

WHAT STRONG LOOKS LIKE

it follows that gallery's conventions and is genuinely useful standalone. Value: distribution inside an existing community. Output: published entry.

BAND S1 · templates, export · publish it on your own channel and put the link in the submission form

Workshop or livestream

WHAT TO BUILD

Run a live session building something on the API.

WHAT STRONG LOOKS LIKE

live, unedited, with real questions taken. Value: highest-trust format, and the recording has a long tail. Output: recording.

BAND S2 · API · publish it on your own channel and put the link in the submission form

Real-user case study

WHAT TO BUILD

Help someone the candidate already knows with a real document problem, then write up what happened with that person's permission.

WHAT STRONG LOOKS LIKE

a real person with a real problem and an honest account of how well it worked. Value: proof from outside the company. Output: case study. Note: this is the candidate helping someone in their own network, with consent. It is not outreach, not a pitch, and not conducted on our behalf.

BAND S1 · chat, export · publish it on your own channel and put the link in the submission form

Five short clips, one workflow each - workflow: redline a clause

WHO IT SERVES

Practitioners who search for a single task.

WHAT TO BUILD

Record five clips of 60 to 120 seconds, each showing exactly one workflow start to finish with no preamble: redline a clause, fix a table, add a footnote, review and approve a change, export with the formatting intact. Each clip is titled with the task in the words someone would search for. No intro, no outro, no music bed over the important part. a small video portfolio on their own channel demonstrating product-communication skill. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

each clip answers its title within the first ten seconds. The screen is readable at small sizes. Five distinct workflows, not one workflow filmed five ways, which would be exactly the near-duplicate pattern this shard forbids.

BAND S1 · chat, Review, export · publish it on your own channel and put the link in the submission form

Five short clips, one workflow each - workflow: fix a table**WHO IT SERVES**

Practitioners who search for a single task.

WHAT TO BUILD

Record five clips of 60 to 120 seconds, each showing exactly one workflow start to finish with no preamble: redline a clause, fix a table, add a footnote, review and approve a change, export with the formatting intact. Each clip is titled with the task in the words someone would search for. No intro, no outro, no music bed over the important part. a small video portfolio on their own channel demonstrating product-communication skill. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

each clip answers its title within the first ten seconds. The screen is readable at small sizes. Five distinct workflows, not one workflow filmed five ways, which would be exactly the near-duplicate pattern this shard forbids.

BAND S1 · chat, Review, export · publish it on your own channel and put the link in the submission form

Five short clips, one workflow each - workflow: add a footnote**WHO IT SERVES**

Practitioners who search for a single task.

WHAT TO BUILD

Record five clips of 60 to 120 seconds, each showing exactly one workflow start to finish with no preamble: redline a clause, fix a table, add a footnote, review and approve a change, export with the formatting intact. Each clip is titled with the task in the words someone would search for. No intro, no outro, no music bed over the important part. a small video portfolio on their own channel demonstrating product-communication skill. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

each clip answers its title within the first ten seconds. The screen is readable at small sizes. Five distinct workflows, not one workflow filmed five ways, which would be exactly the near-duplicate pattern this shard forbids.

BAND S1 · chat, Review, export · publish it on your own channel and put the link in the submission form

Five short clips, one workflow each - workflow: review and approve a change

WHO IT SERVES

Practitioners who search for a single task.

WHAT TO BUILD

Record five clips of 60 to 120 seconds, each showing exactly one workflow start to finish with no preamble: redline a clause, fix a table, add a footnote, review and approve a change, export with the formatting intact. Each clip is titled with the task in the words someone would search for. No intro, no outro, no music bed over the important part. a small video portfolio on their own channel demonstrating product-communication skill. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

each clip answers its title within the first ten seconds. The screen is readable at small sizes. Five distinct workflows, not one workflow filmed five ways, which would be exactly the near-duplicate pattern this shard forbids.

BAND S1 · chat, Review, export · publish it on your own channel and put the link in the submission form

Five short clips, one workflow each - workflow: export with the formatting intact

WHO IT SERVES

Practitioners who search for a single task.

WHAT TO BUILD

Record five clips of 60 to 120 seconds, each showing exactly one workflow start to finish with no preamble: redline a clause, fix a table, add a footnote, review and approve a change, export with the formatting intact. Each clip is titled with the task in the words someone would search for. No intro, no outro, no music bed over the important part. a small video portfolio on their own channel demonstrating product-communication skill. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

each clip answers its title within the first ten seconds. The screen is readable at small sizes. Five distinct workflows, not one workflow filmed five ways, which would be exactly the near-duplicate pattern this shard forbids.

BAND S1 · chat, Review, export · publish it on your own channel and put the link in the submission form

Zendesk app: the document sidebar in the ticket - help-desk host: Freshdesk

WHO IT SERVES

Support agent, claims handler, or student-services officer working a ticket queue. Host app: Freshdesk.

WHAT TO BUILD

A sidebar app on the ticket that produces the document the ticket needs without leaving it: the formal response letter, the credit note, the incident report, the knowledge-base article that stops the next ten tickets. It reads the ticket thread and the customer record, drafts on the right template, lets the agent adjust in plain language, and attaches the finished file to the ticket while filing the article to the help centre. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

the drafted document reflects what the ticket actually says, including the customer's own wording where it matters. The template is right for the ticket type without the agent choosing. Attachment and help-centre publishing both land first time.

BAND S2 · API, chat, templates, export, Review · a pull request into the public repository, extensions folder

Zendesk app: the document sidebar in the ticket - help-desk host: Help Scout

WHO IT SERVES

Support agent, claims handler, or student-services officer working a ticket queue. Host app: Help Scout.

WHAT TO BUILD

A sidebar app on the ticket that produces the document the ticket needs without leaving it: the formal response letter, the credit note, the incident report, the knowledge-base article that stops the next ten tickets. It reads the ticket thread and the customer record, drafts on the right template, lets the agent adjust in plain language, and attaches the finished file to the ticket while filing the article to the help centre. This card is the specific variant named in its title: build exactly that variant.

WHAT STRONG LOOKS LIKE

the drafted document reflects what the ticket actually says, including the customer's own wording where it matters. The template is right for the ticket type without the agent choosing. Attachment and help-centre publishing both land first time.

BAND S2 · API, chat, templates, export, Review · a pull request into the public repository, extensions folder

Landing page per editor integration, with a working snippet

WHO IT SERVES

Growth or engineer. Can write for a developer audience without marketing voice.

WHAT TO BUILD

Write one short page per editor integration, each aimed at the developer already using that editor and asking whether AI editing is possible. Each page: the problem in one paragraph, a snippet that actually runs, what it costs to adopt, and an honest limits section. Not a brochure and not a tutorial series, one page each that a search engine and an AI assistant can both answer from. our public documentation already covers these editors with working snippets and nobody can find them.

WHAT STRONG LOOKS LIKE

the snippets run as written, and the limits section is specific enough to be useful rather than defensive.

BAND S1 · API, export · the submission form, into Drive

GROWTH EXPERIMENTS AND PLACEMENTS

Zero-signup micro-tool

WHAT TO BUILD

Build one tool that solves exactly one document job with no account required.

WHAT STRONG LOOKS LIKE

time from landing to result is under thirty seconds. Value: the lowest-friction discovery surface we can have.

Output: hosted tool plus repo.

BAND S1 · API, export · a pull request into the public repository, use-cases folder

Browser store listing

WHAT TO BUILD

Ship a browser extension through a public store, with the listing copy, screenshots and privacy disclosure done properly.

WHAT STRONG LOOKS LIKE

it passes review first time, and the privacy disclosure is accurate. Value: a directory with real organic traffic where we have nothing. Output: published listing.

BAND S1 · API · publish it on your own channel and put the link in the submission form

Directory submission, done properly

WHAT TO BUILD

Take one directory from the placement map and complete the submission to that directory's standard, with assets.

WHAT STRONG LOOKS LIKE

the entry ranks or is accepted, and the requirements are documented for the next person. Value: repeatable channel work. Output: submission package. Note: the candidate prepares it. Submission is ours.

BAND S1 · export · publish it on your own channel and put the link in the submission form

Template gallery contribution

WHAT TO BUILD

Publish a usable template set into a gallery where the target audience already browses.

WHAT STRONG LOOKS LIKE

the templates work standalone and are formatted to the gallery's conventions. Value: discovery plus a genuinely useful asset. Output: gallery entry.

BAND S1 · templates, export · publish it on your own channel and put the link in the submission form

Launch-ready repository

WHAT TO BUILD

Prepare a repository to the standard a technical community launch requires: clear scope, honest limitations, reproducible demo, no overclaiming.

WHAT STRONG LOOKS LIKE

it would survive a hostile technical audience's first ten questions. Value: launch surfaces punish overclaiming, and this is the artefact that avoids it. Output: repo plus a launch note draft. Note: the founder decides if and when anything launches.

BAND S1 · API · a pull request into the public repository, extensions folder

Community bot with real users

WHAT TO BUILD

Build a bot for a community the candidate is genuinely part of, and get it used.

WHAT STRONG LOOKS LIKE

real recurring usage by people who are not the candidate, reported honestly. Value: genuine distribution inside a community, with consent from that community. Output: bot plus a usage note.

BAND S2 · API · a pull request into the public repository, extensions folder

Embeddable artefact

WHAT TO BUILD

Build something other sites can embed, for example a widget or a badge that does something useful.

WHAT STRONG LOOKS LIKE

it embeds in three different sites without breaking. Value: distribution that compounds. Output: embeddable component.

BAND S1 · API · a pull request into the public repository, use-cases folder

Index others want to be listed in

WHAT TO BUILD

Build a public index of something the target audience cares about, maintained honestly.

WHAT STRONG LOOKS LIKE

the inclusion criteria are public and applied consistently, including where they exclude us. Value: indexes attract inbound and citations. Output: hosted index.

BAND S2 · export · the submission form, into Drive

Showcase entry

WHAT TO BUILD

Build something and submit it to the community showcase with a written account of how it was built.

WHAT STRONG LOOKS LIKE

the account is detailed enough that someone else could rebuild it. Value: populates the showcase, which is currently empty. Output: showcase entry.

BAND S1 · export, images · publish it on your own channel and put the link in the submission form

Localised landing demo

WHAT TO BUILD

Build a working demo and its landing page natively in a non-English market's language, sized to what that market actually needs.

WHAT STRONG LOOKS LIKE

written natively with locally relevant examples. Value: markets where we have users and no presence. Output: hosted demo.

BAND S1 · API · a pull request into the public repository, use-cases folder

Forkable hackathon starter

WHAT TO BUILD

A starter template designed to be forked at a hackathon, working within thirty seconds of clone.

WHAT STRONG LOOKS LIKE

it is actually forked and used by someone else. Value: activates builders, which is a stated goal. Output: template repo.

BAND S1 · API · a pull request into the public repository, extensions folder

Partner marketplace publication

WHAT TO BUILD

Take one integration through a partner marketplace's full publication process.

WHAT STRONG LOOKS LIKE

it clears the partner's review requirements, and the process is documented for the next one. Value: partner marketplaces are distribution we have not touched. Output: published listing plus a process write-up.

BAND S2 · API · publish it on your own channel and put the link in the submission form

Instrumented demo with honest measurement

WHAT TO BUILD

Instrument any demo so real usage is visible, then report what actually happened including if nothing did.

WHAT STRONG LOOKS LIKE

honest reporting of a flat result is a pass. Inflated numbers are a fail. Value: teaches us which artefacts actually get used, which we currently guess at. Output: demo plus a measurement report.

BAND S1 · API · a pull request into the public repository, use-cases folder

Organic traction sprint on one channel you already own

WHO IT SERVES

Growth generalist who already has a small audience of their own.

WHAT TO BUILD

Pick one channel where you already have standing: your own newsletter, your own social account, or a community whose written rules explicitly permit self promotion. Run one organic attempt to get real people to try the free tier. State the hypothesis before you publish, then measure. The deliverable is the experiment log, not the content: what you published, what the funnel did at each step, which step lost the most people, what you would change next, and what the sample cannot support. No paid promotion, no messages to strangers, no company contact, nothing on our accounts. Done means a log another person could rerun and compare against.

WHAT STRONG LOOKS LIKE

The hypothesis was written before publishing and is judged honestly against what happened. The candidate names the worst step in the funnel and refuses to claim significance a small sample cannot carry.

BAND S1 · chat, export · publish it on your own channel and put the link in the submission form

Watch along: five first runs, recorded as friction

WHO IT SERVES

Growth with a user research instinct, observing non technical professionals.

WHAT TO BUILD

Ask five people you already know personally to try SuperDocs while you watch, screen share or in person, on a task they genuinely have, for example redlining a long agreement the way a Pune legal practice would. Do not coach. Record where each person hesitated, the exact words they used when confused, what they expected a control to do, and where they gave up. Deliver a friction report ordered by how much time each friction cost, with a proposed fix per item. Participants consent and are anonymised. Never a company, never a prospect, never anyone approached on our behalf.

WHAT STRONG LOOKS LIKE

Confusion is quoted verbatim rather than paraphrased into product language, and the ranking is by cost to the user rather than by how easy each fix would be.

BAND S1 · chat, attachments, Review, export · the submission form, into Drive

One demo, five framings, tested in public

WHO IT SERVES

Content led growth, positioning.

WHAT TO BUILD

Build one short demonstration of a single capability, then write five genuinely different framings of it: different audience, different promise, different first line. Publish each to a channel you already own, spaced so they do not compete with each other, and record attention per framing with the counting method stated. Deliver the five framings, the raw numbers, a read on which bet won and why, and an explicit list of the confounds that make the result weak. Five different bets, not five rewrites of one sentence. Nothing is boosted, nothing is bought, no company is contacted, and the deliverable is the test rather than the posts.

WHAT STRONG LOOKS LIKE

The five framings depend on genuinely different beliefs about the reader, and the framing that flopped is reported as prominently as the one that worked.

BAND S2 · chat, export · publish it on your own channel and put the link in the submission form

The share loop that offers nothing

WHO IT SERVES

Product marketing, lifecycle.

WHAT TO BUILD

Design an organic sharing loop that works with no incentive at all: no cash, no rewards, no discounts, no referral prize, because we offer none and will not. Find the moment in a real workflow where a person would want to show someone else what they just did, write the in product copy for that moment in three variants, and write what the receiving person sees when they arrive. Then test the variants for preference with your own audience, stating sample and method. Deliver the loop map, the copy, the result, and the honest case against shipping it if that is where the evidence lands.

WHAT STRONG LOOKS LIKE

The chosen moment is one that actually occurs in the product rather than an invented one, and the candidate is willing to recommend against their own design.

BAND S2 · chat, templates, export · a pull request into the public repository, use-cases folder

The living experiment ledger

WHO IT SERVES

Growth operations., with light scripting or a no-code automation.

WHAT TO BUILD

Growth work dies in scattered notes. Build one document that every experiment writes itself into. Use a small script or a no-code automation that appends an entry through the API or an agent client, on a fixed schema: hypothesis, what changed, the measure, the expected result recorded before the run, the actual result, and the keep or kill decision. Seed it with at least six real entries from your own experiments, including one you killed. Keep the ledger as a single durable file across sessions rather than a pile of copies. Deliver the ledger, the automation, and the operating rules that keep it honest.

WHAT STRONG LOOKS LIKE

The schema forces an expected result to be written before the actual one, and the killed entry records its reason plainly instead of being quietly dropped.

BAND S3 · API, MCP, chat, export · the submission form, into Drive

Stopping rule: the loop needs an explicit stopping condition and a no-spend preview mode; it must never re-trigger on its own output.

Curated list placement research, entries drafted to each list's own format

WHO IT SERVES

Growth working a developer adjacent channel.

WHAT TO BUILD

Community maintained curated lists are a placement surface whose rules are written down in public. Map them: which lists cover the relevant categories, what the contribution guidance requires, the exact entry format including ordering and punctuation conventions, how active the maintainers are, what gets entries rejected, and whether a vendor entry is acceptable under the stated policy. Then draft the entry text for every list worth pursuing, matching that list's conventions exactly. Rank by reach against the chance of acceptance, and name any list we should stay out of. Research and drafting only. We open every submission, the candidate opens none.

WHAT STRONG LOOKS LIKE

The drafts match each target list's conventions closely enough to be pasted without editing, and at least one list is recommended against with the policy line quoted.

BAND S1 · chat, search, export · the submission form, into Drive

Education and professional body resource page placement

WHO IT SERVES

Growth working an academic and professional association channel.

WHAT TO BUILD

Universities, libraries, training providers and professional bodies publish resource pages and tool lists that quietly rank and stay cited for years. Map them for document heavy disciplines: who owns each page, the inclusion criteria in their own words, whether a vendor may propose an entry at all, what an entry contains, and how often the page is refreshed. Then draft a one page description written for that audience, plus one classroom ready exercise a lecturer could set without ever contacting us. Mark every page where a vendor proposal would be unwelcome. Mapping and drafting only, every approach is ours.

WHAT STRONG LOOKS LIKE

Inclusion criteria are quoted rather than summarised, and the exercise is good enough that a lecturer would use it whether or not they mention us.

BAND S1 · chat, search, templates, export · publish it on your own channel and put the link in the submission form

The placement pitch pack

WHO IT SERVES

Product marketing.

WHAT TO BUILD

Every placement asks for the same things in different lengths, and the versions have to agree. Build the pack once: descriptions at roughly 25, 50 and 150 words, a feature list at three depths, category fit statements for five taxonomies, a captioned screenshot set, an honest limitations paragraph, and a claims safety sheet. The claims sheet records what is public and true, that SuperDocs is used in 40 countries and 19 languages and was invited to Y Combinator Startup School, then lists what nobody may write: any wording implying we are funded or backed, any certification claim, real time co editing, spreadsheet or slide output, guaranteed redaction. E signature is always an external hand off.

WHAT STRONG LOOKS LIKE

The three lengths are the same claim at three resolutions rather than three different claims, and the safety sheet is written so a stranger could not break a rule by accident.

BAND S2 · chat, templates, Multi-document, export · the submission form, into Drive

Placement portfolio and submission sequencing

WHO IT SERVES

Growth lead, channel strategy.

WHAT TO BUILD

Take the placement surfaces already mapped and turn them into a plan we can actually work. Build an effort to reach model with the assumptions stated per surface, sequence the queue including what unblocks what, define what would count as a result for each surface and how we would detect it, and argue explicitly for the surfaces to skip and why. Add a refresh policy for the ones that decay, and one honest paragraph naming which parts of the model are guesses. Light scripting is welcome for the scoring, not required. Deliver a plan a founder could execute without a follow up question. The candidate submits nothing.

WHAT STRONG LOOKS LIKE

The skip list exists and is argued, and the assumptions behind every effort estimate are exposed clearly enough for someone to disagree with a specific number.

BAND S3 · chat, Multi-document, templates, export · the submission form, into Drive

Maker community launch runbook, fully drafted

WHO IT SERVES

Growth, launches.

WHAT TO BUILD

Draft a complete maker community launch, ready to hand over. That means the listing copy, the gallery assets and their captions, the opening comment, a reply bank covering the twenty questions most likely to arrive, the sequencing measured in elapsed hours from the moment the post goes live, who does what at each step, and the community's own rules quoted wherever they constrain us. Name the assets we do not have yet instead of assuming them. Deliver a runbook someone else could execute cold. Nothing is scheduled, nothing is submitted, and the candidate posts nowhere.

WHAT STRONG LOOKS LIKE

The reply bank answers the hard questions honestly, including the ones whose honest answer is a limitation, and the missing assets are listed rather than glossed over.

BAND S1 · chat, templates, Multi-document, export · the submission form, into Drive

The hostile room: a technical news community post and its objection bank

WHO IT SERVES

Growth with real technical literacy.

WHAT TO BUILD

Technical news communities punish marketing language and reward candour. Draft the post that would survive one: what it is, what it does not do, what is genuinely hard about the problem, and why a reader should care, with the community's own guidelines quoted wherever they shape the draft. Then predict the twenty hardest comments the post would attract and write the honest reply to each, marking every one where the honest reply is that we do not do that yet. Finish with a go or no go recommendation and the reasoning behind it. Drafting only, the candidate posts nothing.

WHAT STRONG LOOKS LIKE

The draft would not embarrass us if a sceptical engineer read it line by line, and the objection bank concedes wherever conceding is the correct answer.

BAND S1 · chat, export · publish it on your own channel and put the link in the submission form

Resource first launch into a practitioner community

WHO IT SERVES

Growth for a non developer practitioner audience, for example an academic writing group or a document heavy professional community.

WHAT TO BUILD

Some communities never accept a product launch but welcome a genuinely useful free resource. Build the resource in SuperDocs, for example a reviewer ready manuscript checklist, or the project handover pack a Kochi architecture studio would use, then plan the launch around giving it away: which community, what its written rules permit, how the resource gets introduced, what is said about who made it, and what is deliberately left unsaid. Include the follow up plan for the questions it will attract. Deliver the resource, exported, plus the plan. The candidate posts nothing and joins nothing on our behalf.

WHAT STRONG LOOKS LIKE

The resource is good enough that the community would keep it even with the vendor line removed, and the affiliation disclosure sits near the top rather than buried at the end.

BAND S1 · chat, templates, Multi-document, export · publish it on your own channel and put the link in the submission form

Launch pre mortem and the incident plan

WHO IT SERVES

Growth and communications.

WHAT TO BUILD

Assume the launch went badly, then write backwards from there. List the ways it fails: a factual claim challenged in public, the service struggling during the busiest hour, a competitor comparison we got wrong, a moderator removing the post, and the quiet launch nobody sees at all. For each one give the early signal, who decides, the prepared response drafted in full rather than described, and the line we will not cross to save it. Include the correction procedure for a published claim that turns out to be wrong, because that failure is reputational rather than technical. Deliver a plan short enough to read while something is on fire.

WHAT STRONG LOOKS LIKE

The responses are written out ready to send rather than summarised, and at least one branch concludes with taking the loss quietly instead of fighting it.

BAND S2 · chat, Multi-document, export · the submission form, into Drive

Attribution badge and a public gallery of what was built

WHO IT SERVES

Growth. Design sense plus enough code to ship an embeddable component.

WHAT TO BUILD

Design and build the mechanism that makes a project publicly visible as built with SuperDocs: a small embeddable badge a builder can add to their own project, and a gallery page that collects projects that carry it. Include the submission and review path for getting into the gallery, and the rules for what qualifies.

WHAT STRONG LOOKS LIKE

the badge is something a developer would actually put on their own README, and the gallery has a real curation rule rather than accepting everything.

BAND S2 · API, export · the submission form, into Drive

The deep growth plan: next users, first paying companies

WHO IT SERVES

Growth. Someone with a point of view who can defend it under a word limit.

WHAT TO BUILD

Write the real version of the growth plan. Where do the next users come from, and where do the first paying companies come from, and why those and not the obvious alternatives. Name the channel you would bet on and the one you would refuse, and say what would have to be true for you to be wrong. Ground it in the product as it actually exists: use it first, then write. Hard rules. Under 1,500 words. End with the one concrete thing you would do in week one, specific enough that someone could start it the next morning. We are reading for judgement and instincts, not for a business plan. Concrete beats comprehensive. the growth application form asked a 500-word version of this question, and every growth applicant answered it. That answer is the seed for this card and not a substitute for it. This is the long form: more room, a harder cap on waffle, and the same week-one ending. we are pre-revenue with the first paying companies still ahead of us, so a well-argued outside read on channel choice is directly useful whatever the hiring outcome. It is also the cheapest way to see how a growth candidate thinks before they ever touch our funnel.

WHAT STRONG LOOKS LIKE

a real bet with a named channel, a stated cost, an honest failure condition, and a week-one action that is a task rather than a theme. A plan that recommends everything recommends nothing.

BAND S3 · export, search · the submission form, into Drive

Pricing page and upgrade-moment teardown of our own product

WHO IT SERVES

Growth. Someone who can criticise a product they want to work on.

WHAT TO BUILD

Use the product on the free plan until you hit a real limit, then write the teardown. Where exactly does a user meet the ceiling, what were they in the middle of, and what does the product say to them at that moment. Then: which published plan would a user of that kind actually choose, what makes the choice hard, and what single change to the page or to the moment would convert more of them. Ground every claim in something you personally did, not in a heuristic.

WHAT STRONG LOOKS LIKE

the writer clearly hit the ceiling themselves and can say what they were doing when it happened. One prioritised change, argued, rather than a list of twelve.

BAND S2 · chat, export · the submission form, into Drive